

ECE 550

Fundamentals of Computer Systems and Engineering

Transistors -> Gates

Last time....

(Almost) every class will start with the same question:

- Who can remind us what we talked about last time?
(besides course policies)

Last time....

(Almost) every class will start with the same question:

- Who can remind us what we talked about last time?
(besides course policies)

- Abstraction
 - Interface vs Implementation
- Tools
 - VHDL, Quartus
- Transistors and Gates
 - More on this today

Power (Vcc) and Ground (Gnd)

Vcc




Gnd



- Two supply rails:
 - Power (aka Vcc, sometimes called Vdd) , e.g., +1.0 V
 - Logically, 1
 - Ground (Gnd, or Vss), e.g., 0 V
 - Logically, 0
 - I'm going to use Vcc/Gnd because that's what Quartus uses

Power (Vcc) and Ground (Gnd)



- Water analogy
 - Power: think of this as a pump, pushing water in
 - Ground: think of this as a pump sucking water out

Wires

Vcc




Gnd




- A wire (or other conductor) causes current to flow
 - Attempts to equalize voltage
 - Water analogy: think of a pipe

Short circuit




- Short circuit: direct connection from power to ground
 - Very high current (think of fast, continuous flow of water)
 - Generates a lot of heat
 - Destroys your chip
 - Very bad!

Switching





- Suppose instead we had some sort of switch
 - Think “valve”
 - Here, top connection conducts, connecting A to V_{cc}
 - Think of A as a pipe, pumped full of water
 - The bottom half resists (think closed valve) insulating A from Gnd
 - The water in A cannot get sucked down

Switching



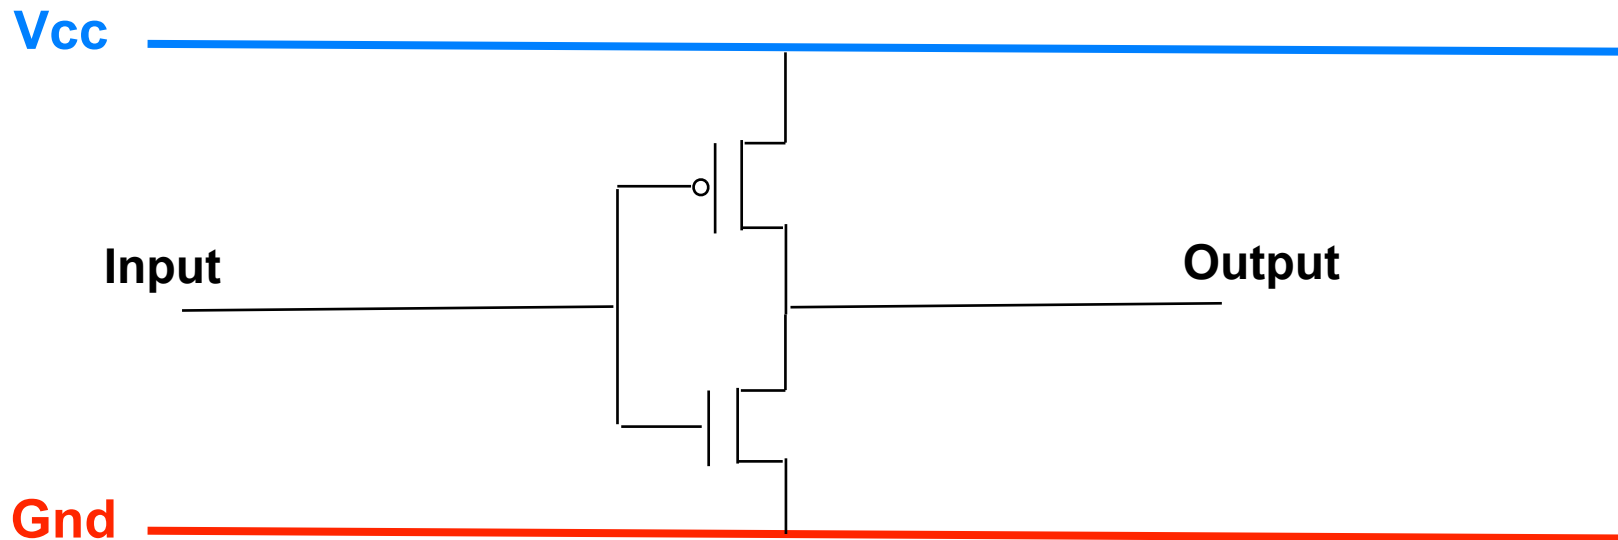
- If we switch our connection...
 - Current flows as A changes voltage levels
 - Think of a pipe draining out as its connected to suction
 - Connection to power is closed, so no short circuit

Transistors: Electrically controlled




- Two types:
 - **NMOS** (left: no circle):
 - Conducts when gate is 1, resists when gate is 0
 - Connect source to either Ground or (Drain of another NMOS)
 - **PMOS** (right: circle):
 - Conducts when gate is 0, resists when gate is 1
 - Connect source to either Vcc or (Drain of another PMOS)

CMOS: Complementary MOS




- CMOS (most common, all we care about):
 - Put PMOS and NMOS in complementary fashion
 - Either PMOS conducts or NMOS conducts, but not both
- Form a logic gate
 - Input (s): connected to gates of transistors
 - Output: connected to drains

CMOS: Complementary MOS



- Let's see how this works.
 - Suppose Input = 1 (circuitry to control input, not shown)
 - PMOS transistor resists
 - No connection between Output and V_{cc}
 - NMOS transistor conducts
 - Connection between Output and Gnd
 - Output is 0

CMOS: Complementary MOS

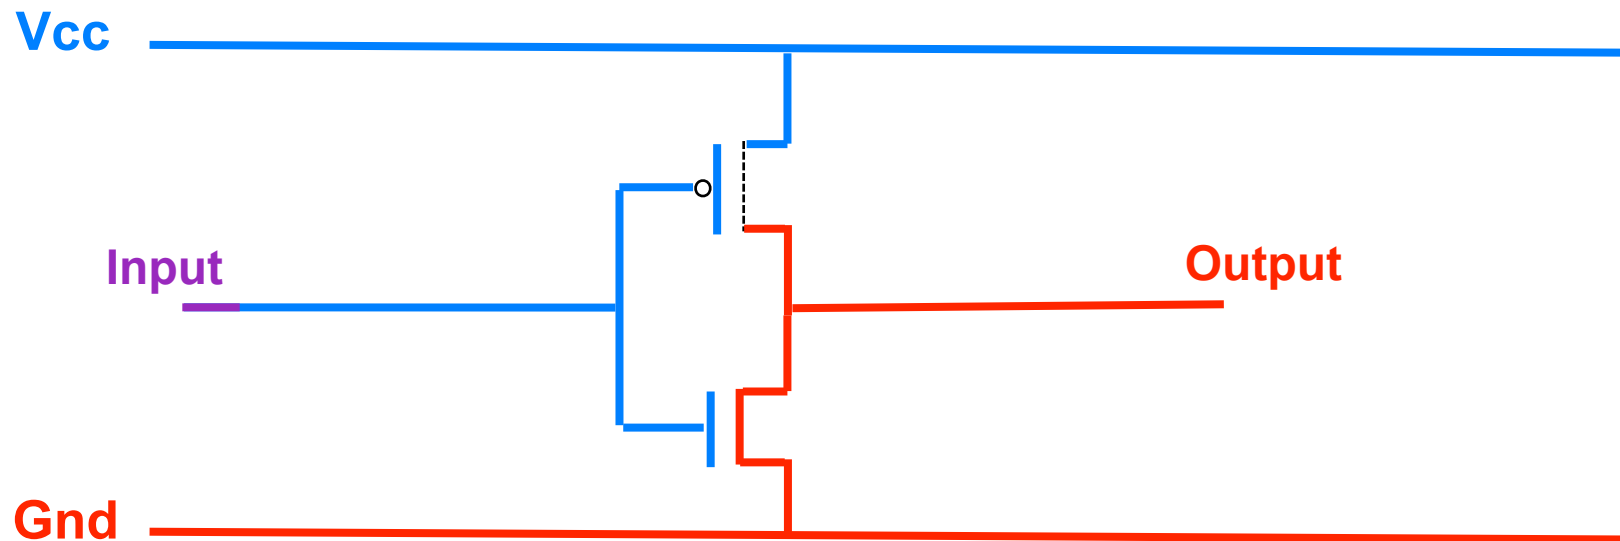


- Now suppose Input changes to 0
 - PMOS transistor conducts
 - Connection between Output and V_{cc}
 - NMOS transistor conducts
 - No Connection between Output and Gnd
 - Output is 1

Switching delays

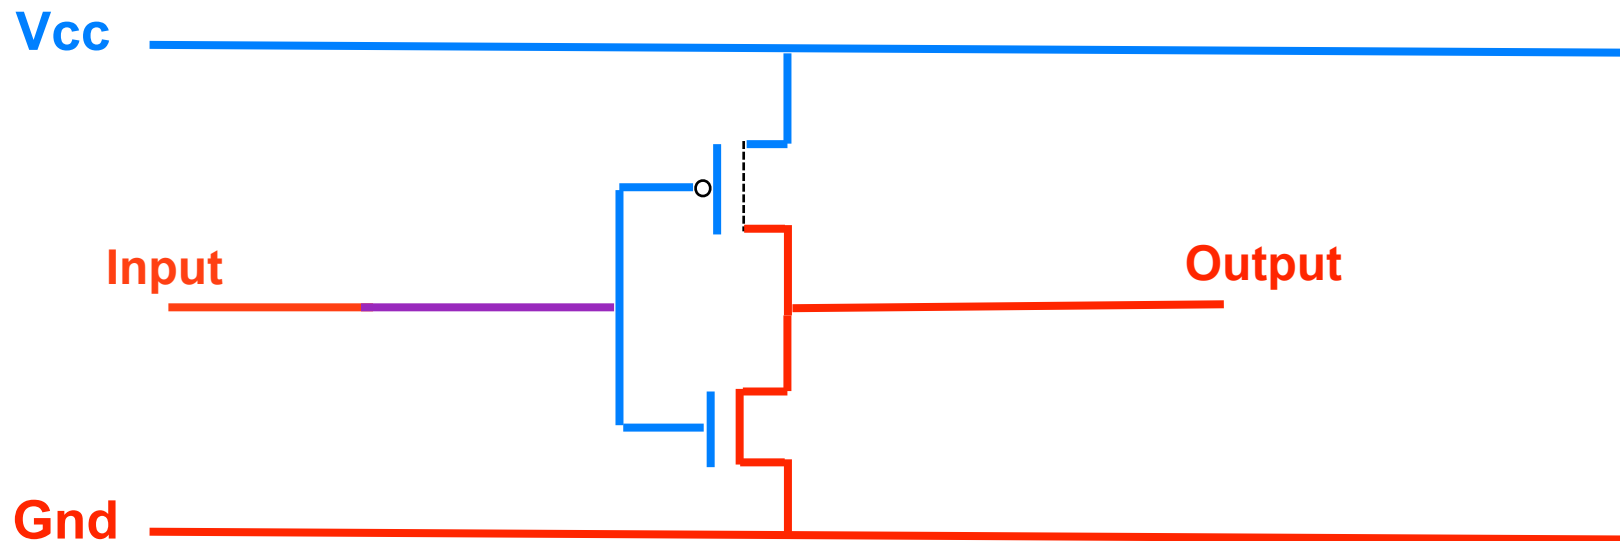
- Note: this doesn't happen instantly
 - There is some delay as these change
 - Imagine again, pipes full of water
 - Draining out input pipe takes time...
 - Once its drained enough the valves start to change...
 - Filling the output pipe takes time
 - Factors the affect the delay
 - Voltage: analogous to water pressure
 - Higher voltage = faster switching, but more power/energy
 - Resistance: analogous to pipe narrowness
 - Lower resistance = faster switching
 - Capacitance: analogous to pipe volume (how much to fill)
 - Lower capacitance = faster switching
 - Calculating delay = hard, so we let our tools do it

CMOS: Complementary MOS




- Slightly more accurate with respect to time
 - Input starts to swing from 1 to 0 (not instant)

CMOS: Complementary MOS




- Slightly more accurate with respect to time
 - Input starts to swing from 1 to 0 (not instant)
 - Change propagates along wires (also takes time)

CMOS: Complementary MOS




- Slightly more accurate with respect to time
 - Input starts to swing from 1 to 0 (not instant)
 - Change propagates along wires (also takes time)
 - Transistors start to switch (partially conductive)

CMOS: Complementary MOS



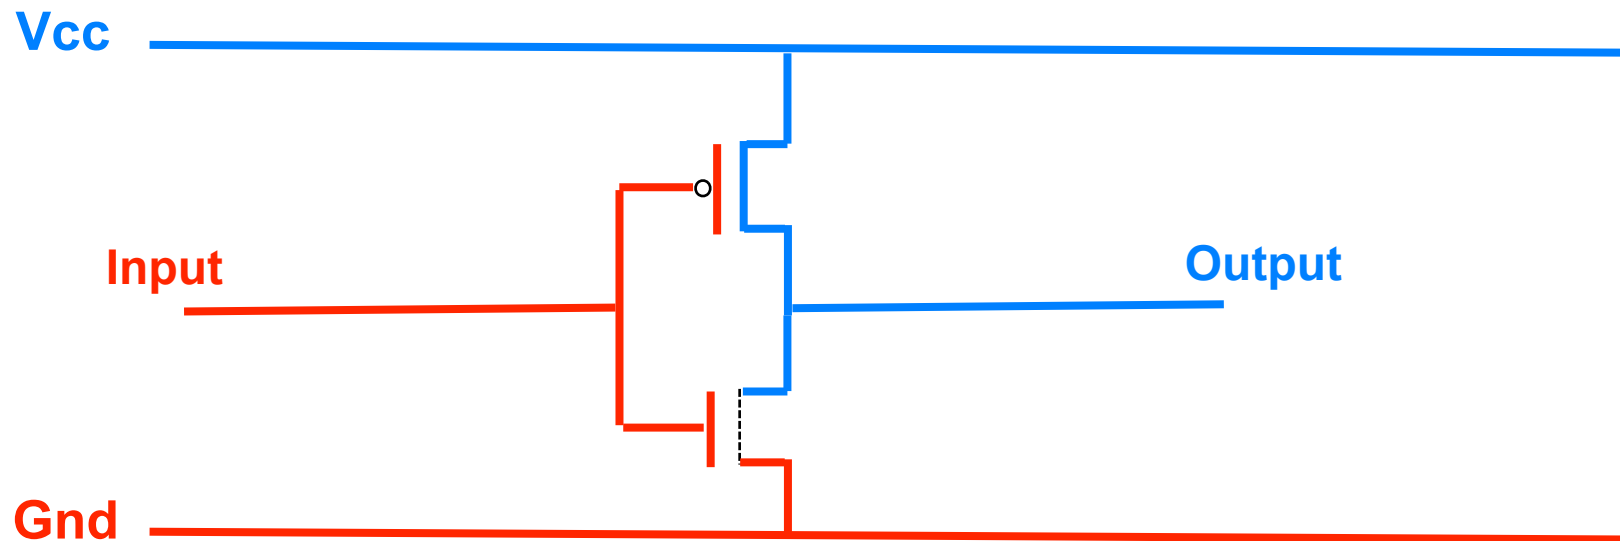
- Slightly more accurate with respect to time
 - Input starts to swing from 1 to 0 (not instant)
 - Change propagates along wires (also takes time)
 - Transistors start to switch (partially conductive)
 - Inputs reach 0 (sometime)
 - Transistors fully open/closed
 - Output may take time to transition

CMOS: Complementary MOS




- Slightly more accurate with respect to time
 - Input starts to swing from 1 to 0 (not instant)
 - Change propagates along wires (also takes time)
 - Transistors start to switch (partially conductive)
 - Inputs reach 0 (sometime)
 - Transistors fully open/closed
 - Output may take time to transition


Our first logic gate: The inverter




- This circuit is a logic gate: inverter or “NOT gate”
 - Gives logical negation of its input
 - Input = 0, Output = 1
 - Input = 1, Output = 0
 - Typically, just draw the gate, instead of the transistors:



Our first logic gate: The inverter



- (Small) Example of abstraction
 - Interface: “do logical negation”
 - Implementation: how to hook up the transistors



Let's build a more interesting gate

- Next, let us build a 2-input NOR gate

- Here is a **truth table** for NOR

- Shows output values for all possible inputs

- Output = 1 when A and B = 0

- Connect PMOS in **series**

- (Not A) and (Not B)

- Output = 0 when A or B = 1

- Connect NMOS in **parallel**

- Not (A or B)

A	B	Output
0	0	1
0	1	0
1	0	0
1	1	0


Note: two formulas are logically equivalent (DeMorgan's Laws)

- PMOS formula has NOTs on inputs
 - NMOS formula has NOT on output

(why?)

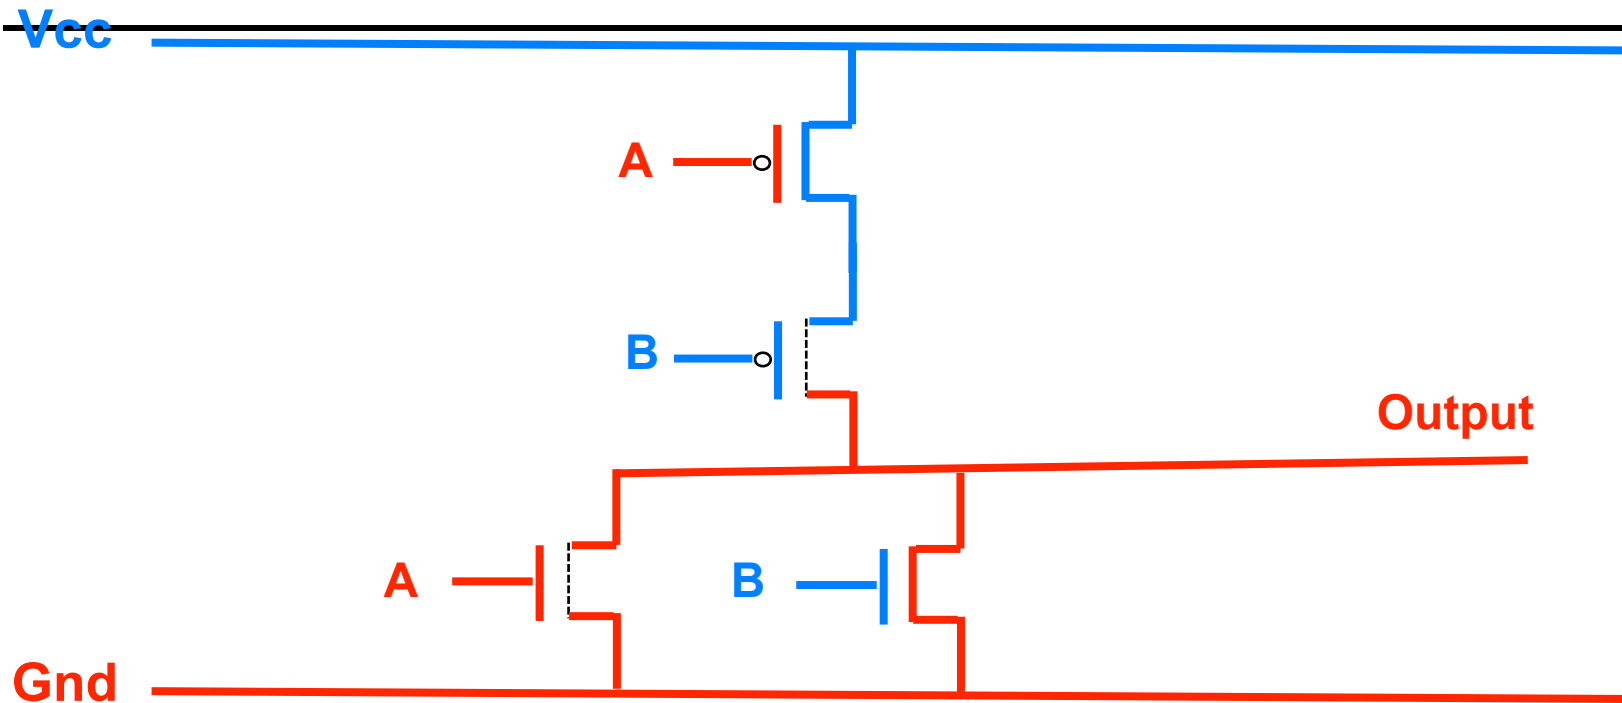
ECE 550 (Hilton): Transistors -> Gates

The NOR Gate



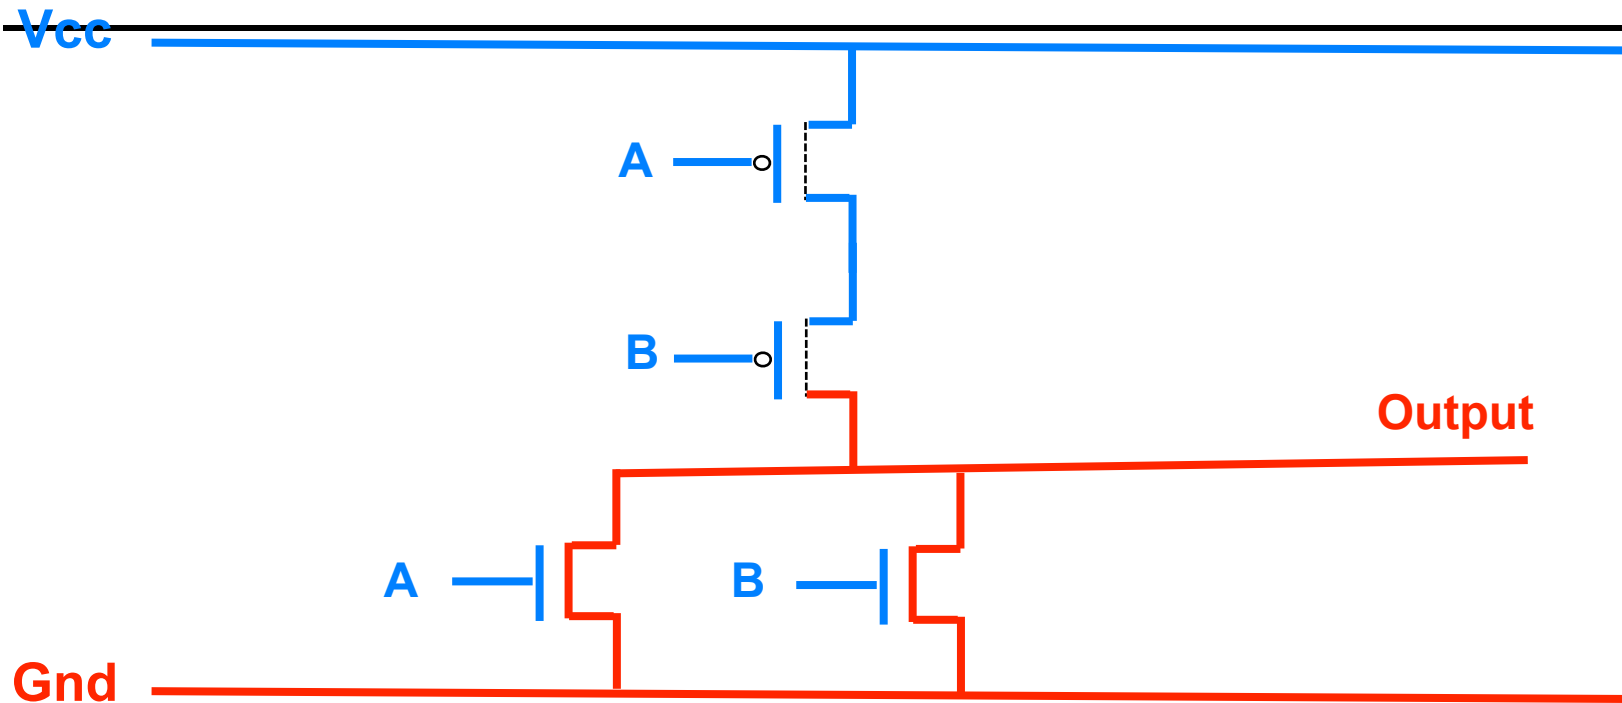
- NOR Gate
 - PMOS in series (both A and B must be 0 to get 1)
 - NMOS in parallel (either A or B at 1 results in 0)
- Side note: real chips have several layers to route wires
 - 3D drawing is hard, just label inputs

The NOR gate




- NOR Gate
 - Same gate, just changed B's value to 1
 - Now output = 0
 - PMOS connected to B resists, blocking connection to V_{cc}
 - NMOS connected to B conducts, forming connection to Gnd

The NOR gate



- NOR Gate
 - Same gate, now change A to 1
 - Output stays at 0
 - Two connections to Gnd , but that's fine

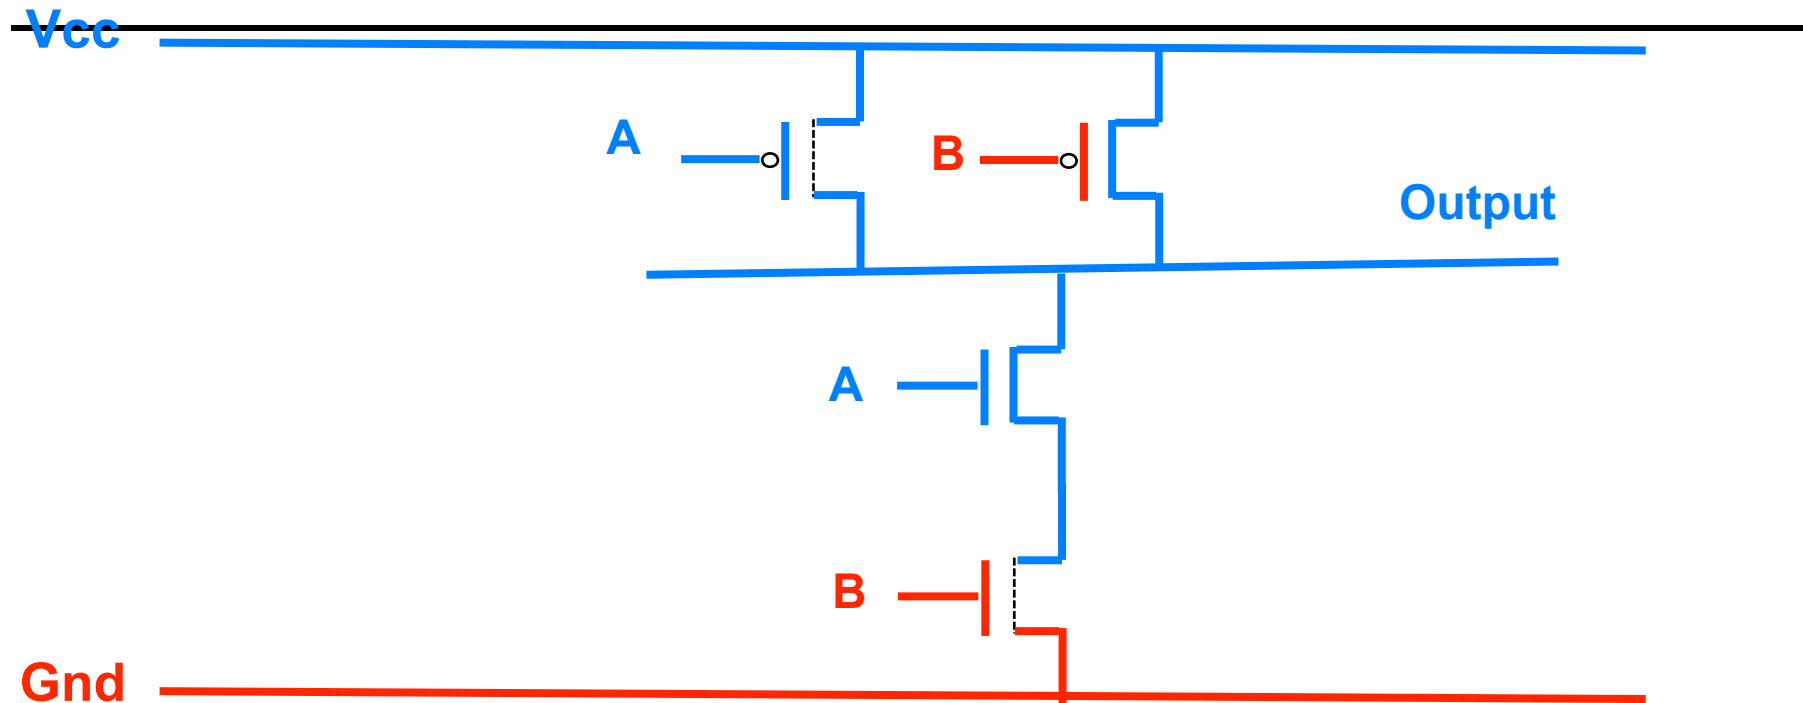


Let's build a more interesting gate


- I'll let you all try a 2-input NAND gate
 - Here is a **truth table** for NAND

A	B	Output
0	0	1
0	1	1
1	0	1
1	1	0

The NAND Gate










- The NAND Gate
 - PMOS in parallel (either at 0 results in a 1)
 - NMOS in series (both at 1 results in 0)



Boolean Gates


- Actually a bunch of standard logic gates:



How to keep them all straight?


Drew's Guide to Remembering your

- This one looks like it just points its input where to go
 - It just produces its input as its output
 - Called a buffer
 - (won't really worry much about these)




Drew's Guide to Remembering your

- This one looks like it just points its input where to go
 - It just produces its input as its output
 - Called a buffer




- A circle always means negate




Circle = NOT

Drew's Guide to Remembering Your

- And Gates have a straight edge, like an A (in AND)







OR Gates have a curved edge, like an O (in OR)



Drew's Guide to Remembering Your

- If we stick a circle on them...




Circle = NOT


- We get NAND (NOT-AND) and NOR (NOT-OR)
 - $\text{NAND}(a, b) = \text{NOT}(\text{AND}(a, b))$

Drew's Guide to Remembering Your

- XOR looks like OR (curved line)
 - But has two lines (like an X does)




- Can put a dot for XNOR
 - XNOR is 1-bit "equals" by the way




Multi-input gates

- So far gates have had 1 or 2 inputs
 - Can have more, though typically stop at 3 or 4
 - Symbols stay the same, just have more input lines



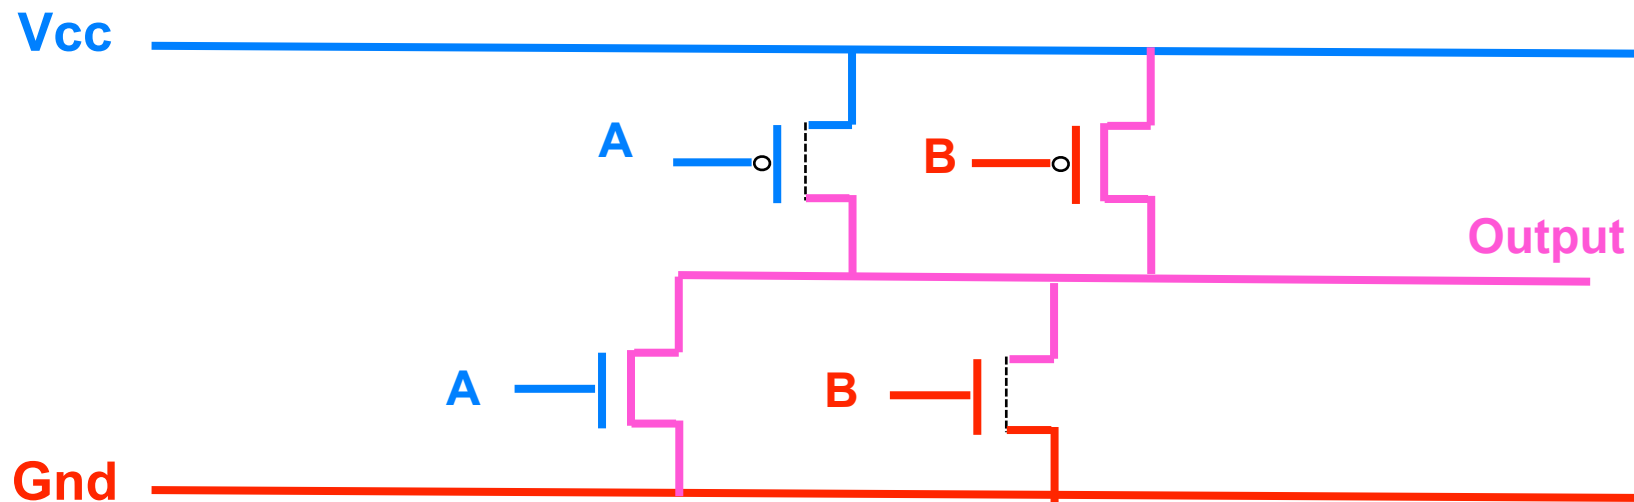
A	B	C	Ouput
0	0	0	1
0	0	1	0
0	1	0	0
0	1	1	0
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	0

Three input NOR Gate



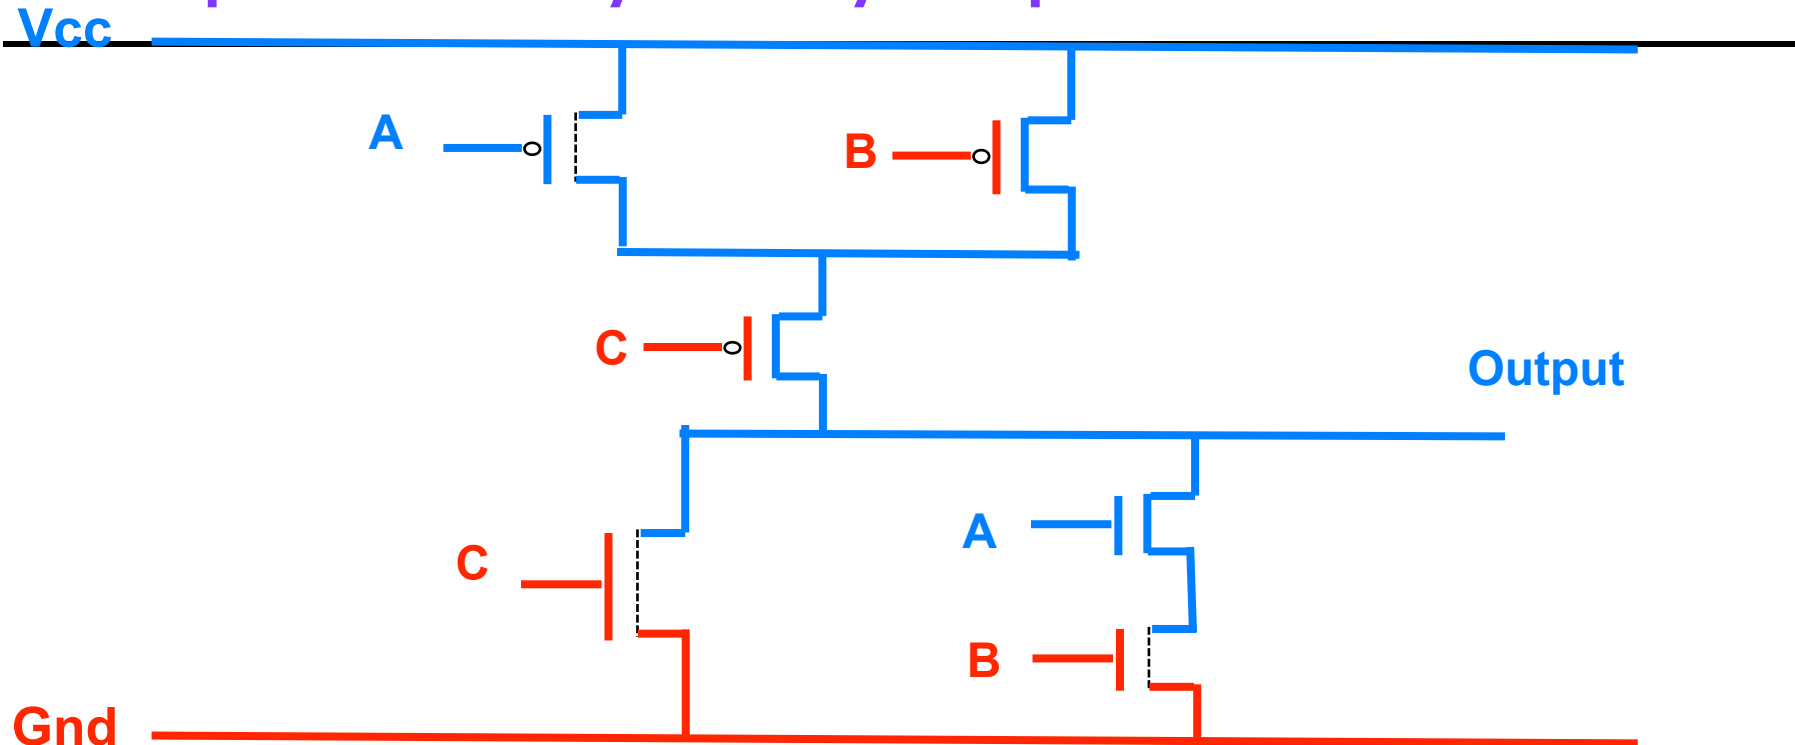
- Similar to two input, more transistors
 - Slightly slower

Complimentary: very important



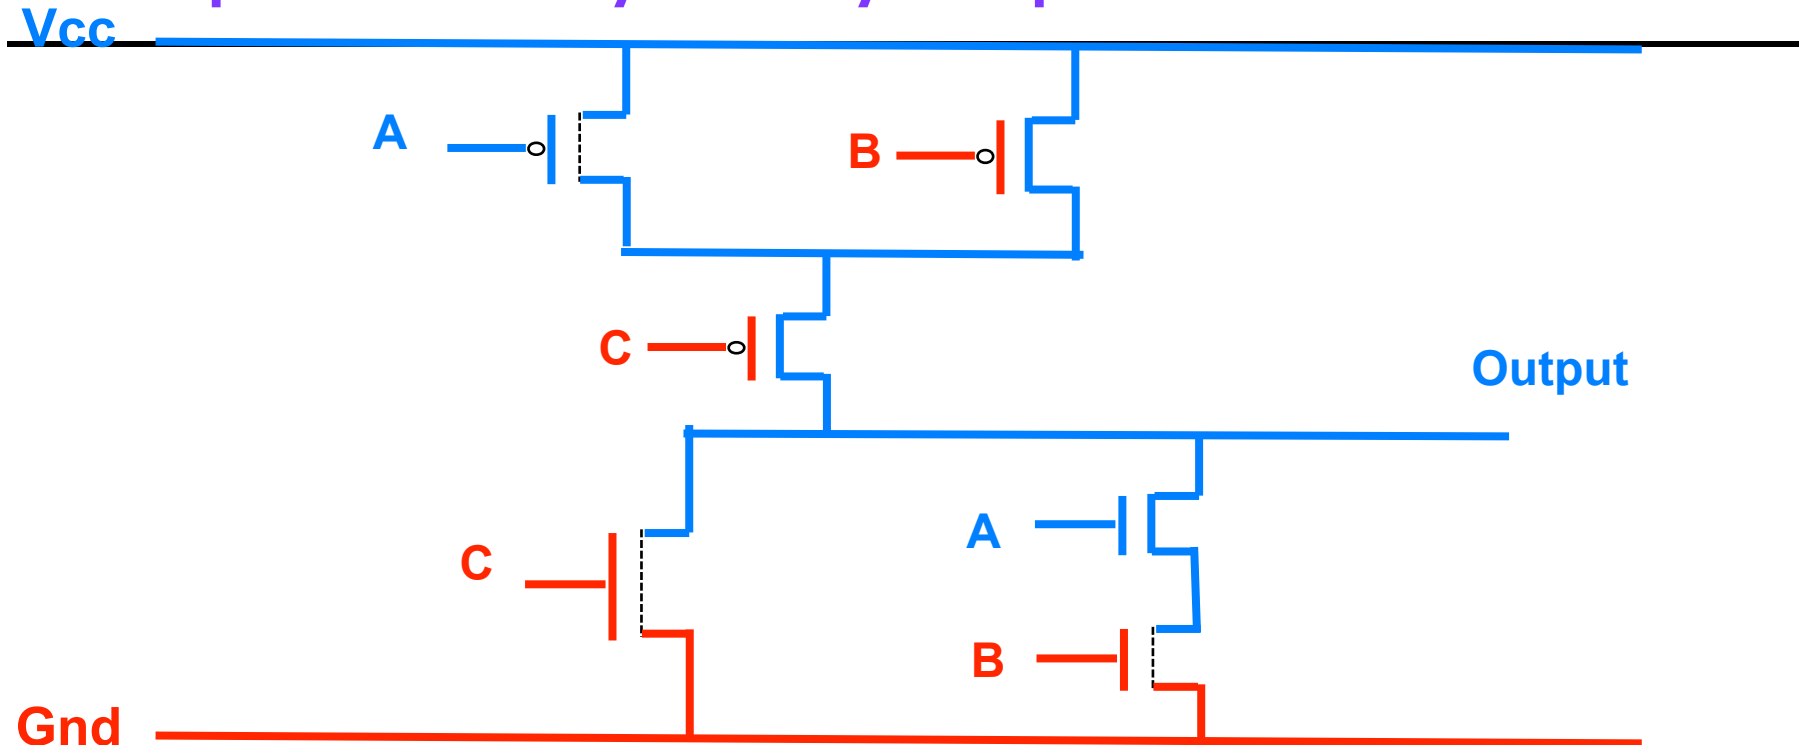
- Complementary nature: very important
 - Without it, we have a problem
 - Here: both PMOS and NMOS in parallel...
 - $A=1, B=0$ (or $A=0, B=1$) forms short-circuit
 - Chip catches fire ☹️

Complimentary: very important



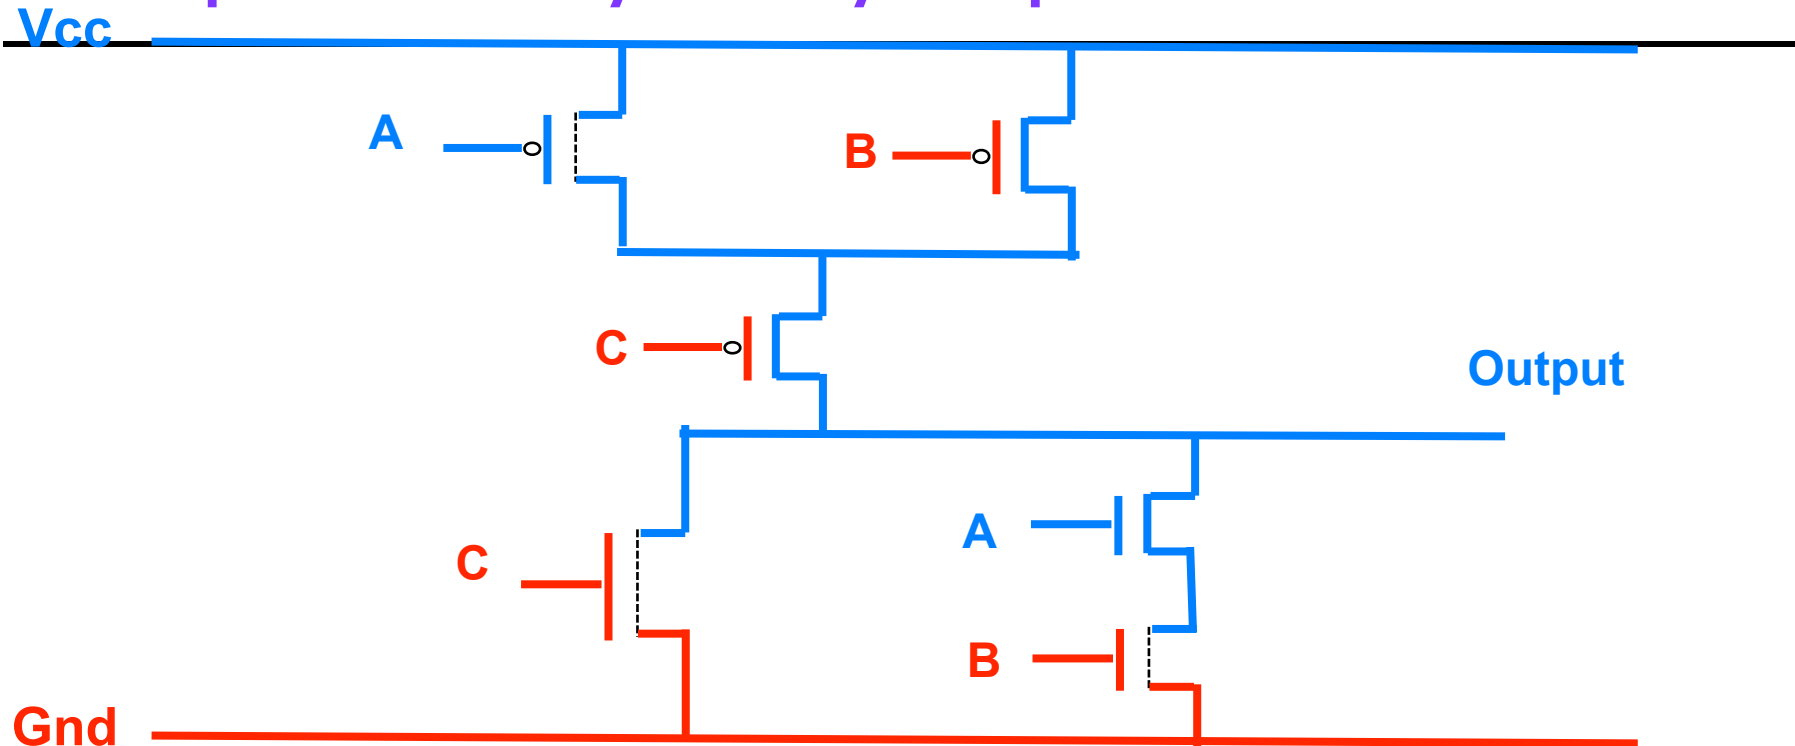
- With more than 2 inputs, can get very complicated
 - Are the PMOS and NMOS complimentary here?
 - We can go the other way: transistors \rightarrow formulas
 - Check if formulas logically equivalent

Complimentary: very important



- Everyone take a second to write down the formulas
 - PMOS: NOTs on inputs
 - NMOS: NOT around the outside

Complimentary: very important



- Everyone take a second to write down the formulas
 - PMOS: $((\text{Not } A) \text{ or } (\text{Not } B)) \text{ and } (\text{Not } C)$
 - NMOS: $\text{Not } (C \text{ or } (A \text{ and } B))$
 - $= (\text{Not } C) \text{ and } (\text{Not } (A \text{ and } B))$
 - $= (\text{Not } C) \text{ and } ((\text{Not } A) \text{ or } (\text{Not } B))$

What about... AND?


- Saw and did NAND, but what about AND?

- Truth table on the right...

A	B	Output
0	0	0
0	1	0
1	0	0
1	1	1

- Trying to do this causes problems
 - PMOS formula: NOTs on inputs
 - NMOS formula: NOT around outside
 - ... can't seem to find a formula which works (need more NOTs):
 - (**Not** (**Not** A)) and (Not (**Not** B))
 - **Not** ((**Not** A) or (**Not** B))
 - AND gate is really a couple gates squished together
 - Not (Nand (A,B))
 - Nor(Not A, Not B)

The AND Gate



- The AND gate
 - A NAND gate followed by a NOT gate
 - Also a good example of how gates connect together
 - Output of one gate goes to input of another

Speaking of transistors...

- Moore's law:
 - Transistor density doubles roughly every 18 months
 - Has been going on for many years (~1970)
 - Self-fulfilling prophecy?
 - Ending soon? We've heard that before, but....
 - Commonly stated as "computers get faster"... why?


Moore's law and speed

- How are size and speed related?
 - Historically: clock frequency ("How many MHz/GHz")
 - How fast do the transistors switch (more on this later)
 - Has leveled off: Diminishing returns, power in-efficient,...
 - Now: put more in same chip area
 - Larger caches
 - Better (bigger) predictors
 - ...
 - Future: ??????
 - Significant concern among micro-architects
 - Reason: power density

Power and Energy

- We won't focus too much on power and energy, but...
 - Very important concerns these days, so at least some mention
 - Energy costs money (power bills + cooling)
 - Water analogy: think total water pushed in/sucked out of system
 - Power: energy per time
 - Water analogy: How fast are we pumping in water
 - Power => Heat. Heat must be cooled, physical limits on cooling
- Previously:
 - Smaller transistors => lower voltages => less power/transistor
 - More transistors / area
 - Power density (power/area) held roughly constant

Going forwards: Logic Gates



- Going forwards, will mostly design from gates
 - Abstract away transistor level implementation
 - More on this next time...

Next Time...

- Next time, we'll delve into Combinatorial Logic
 - Putting gates together to do useful things
- Reading:
 - Continue to read Chapter 1 and Appendix C
- Homework:
 - Homework 1 will be out soon
 - Keep an eye on Piazza
 - I'll also post the code for the lab door on Piazza soon
- Recitation: Friday (here)
 - Bring laptops