

ECE 550: Fundamentals of Computer Systems and Engineering

Digital Arithmetic

Admin

- Homework
 - Homework 1
- Reading:
 - Chapter 3

Last Time in ECE 550....

- Who can remind us what we talked about last time?

Last Time in ECE 550....

- Who can remind us what we talked about last time?
 - Numbers
 - One hot
 - Binary
 - Hex
 - Digital Logic
 - Sum of products
 - Encoders
 - Decoders

Implementing Addition

- First, one bit addition.
 - Three inputs: Carry In (CI), A, B
 - Two outputs Carry Out (CO), Sum (S)
- Go around room for truth table:

CI	A	B	S	CO
0	0	0		
0	0	1		
0	1	0		
0	1	1		
1	0	0		
1	0	1		
1	1	0		
1	1	1		

Implementing Addition

- First, one bit addition.
 - Three inputs: Carry In (CI), A, B
 - Two outputs Carry Out (CO), Sum (S)
- Go around room for truth table:

CI	A	B	S	CO
0	0	0	0	0
0	0	1	1	0
0	1	0	1	0
0	1	1	0	1
1	0	0	1	0
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

Half Adder

- Ignore CI for a second (assume is 0)
 - Can simplify a lot and build "half adder"
 - Formula for S?
 - Formula for CO?


CI	A	B	S	CO
0	0	0	0	0
0	0	1	1	0
0	1	0	1	0
0	1	1	0	1
1	0	0	1	0
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

Half Adder

- Ignore CI for a second (assume is 0)
 - Can simplify a lot and build “half adder”
 - Formula for S? $A \text{ xor } B$
 - Formula for CO? $A \text{ and } B$

CI	A	B	S	CO
0	0	0	0	0
0	0	1	1	0
0	1	0	1	0
0	1	1	0	1
1	0	0	1	0
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

Half Adder



- Half adder:
- 1 XOR and 1 AND
- Can anyone guess why its called a **half** adder?

Implementing Addition

- Re-visit Truth table, but..
 - Use Half-Sum and Half-CO (results of Half-Adder)
- Go around room for truth table:

CI	Half-Sum	Half-CO	S	CO
0	0	0		
0	0	1		
0	1	0		
0	1	1		
1	0	0		
1	0	1		
1	1	0		
1	1	1		

Implementing Addition

- Re-visit Truth table, but..
 - Use Half-Sum and Half-CO (results of Half-Adder)
- Go around room for truth table:

CI	Half-Sum	Half-CO	S	CO
0	0	0	0	0
0	0	1	0	1
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	0	1
1	1	1	!!!	!!!

Implementing Addition

- Re-visit Truth table, but..
 - Use Half-Sum and Half-CO (results of Half-Adder)
- Go around room for truth table:

CI	Half-Sum	Half-CO	S	CO
0	0	0	0	0
0	0	1	0	1
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	0	1
1	1	1	0	1

Implementing Addition

- Formulas:
 - Sum?
 - CO?

CI	Half-Sum	Half-CO	S	CO
0	0	0	0	0
0	0	1	0	1
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	0	1
1	1	1	0	1

Implementing Addition

- Formulas:
 - Sum? $CI \text{ xor Half-Sum}$
 - CO? $(CI \text{ and Half-Sum}) \text{ OR Half-CO}$

CI	Half-Sum	Half-CO	S	CO
0	0	0	0	0
0	0	1	0	1
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	0	1
1	1	1	0	1

Implementing Addition



- Formulas:

- Sum? $CI \text{ xor Half-Sum}$
- CO? $(CI \text{ and Half-Sum}) \text{ OR Half-CO}$

What does this look like?


CI	Half-Sum	Half-CO	S	CO
0	0	0	0	0
0	0	1	0	1
0	1	0	1	0
0	1	1	1	1
1	0	0	1	0
1	0	1	1	1
1	1	0	0	1
1	1	1	0	1

Full Adder




- Full Adder
- 2 Half Adders + an OR Gate

Ripple Carry





- Full Adder = Add 1 Bit
 - Can chain together to add many bits
 - Upside: Simple
 - Downside?

Ripple Carry



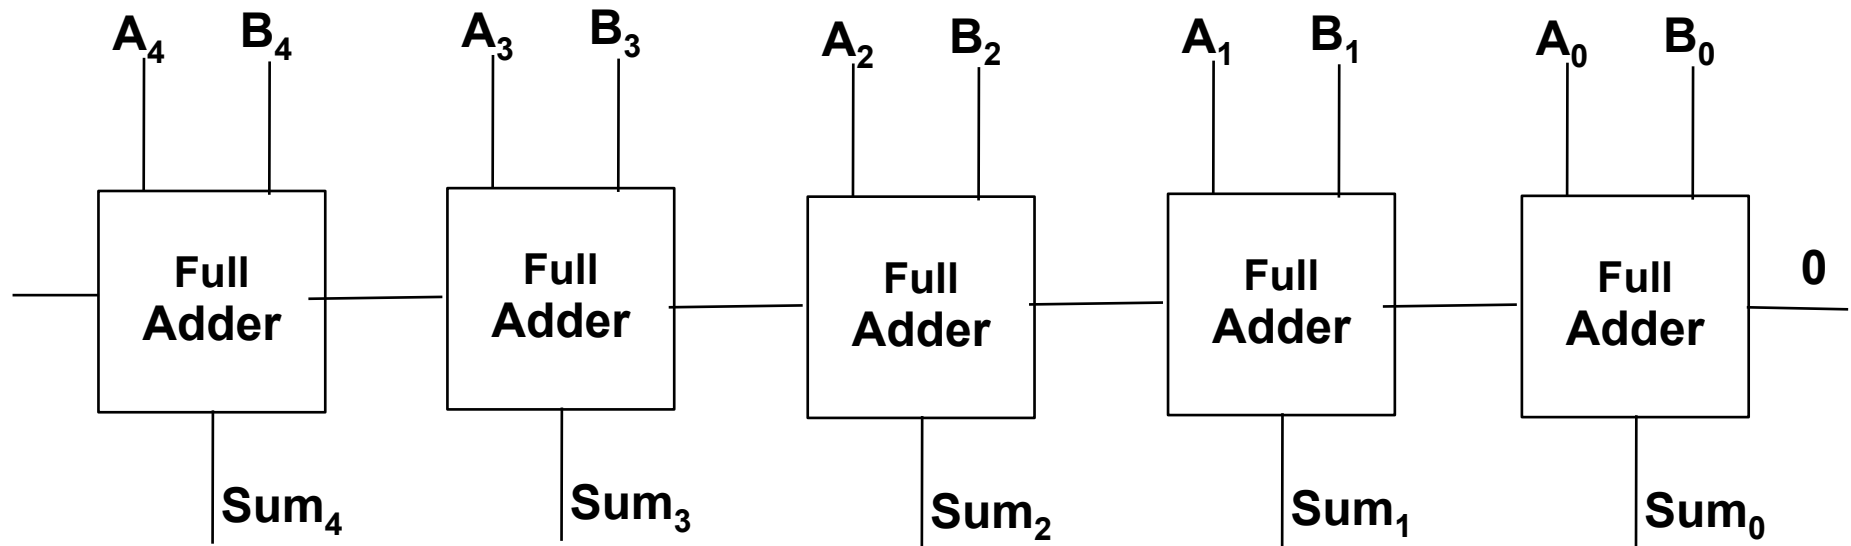
- Full Adder = Add 1 Bit
 - Can chain together to add many bits
 - Upside: Simple
 - Downside? Slow
 - Let's see why

Full Adder



- Cout depends on Cin
 - 2 "gate delays" through full adder for carry

Ripple Carry



- Carries form a chain
 - Need CO of bit N is CI of bit N+1
- For few bits (e.g., 4) no big deal
 - For realistic numbers of bits (e.g., 32, 64), slow

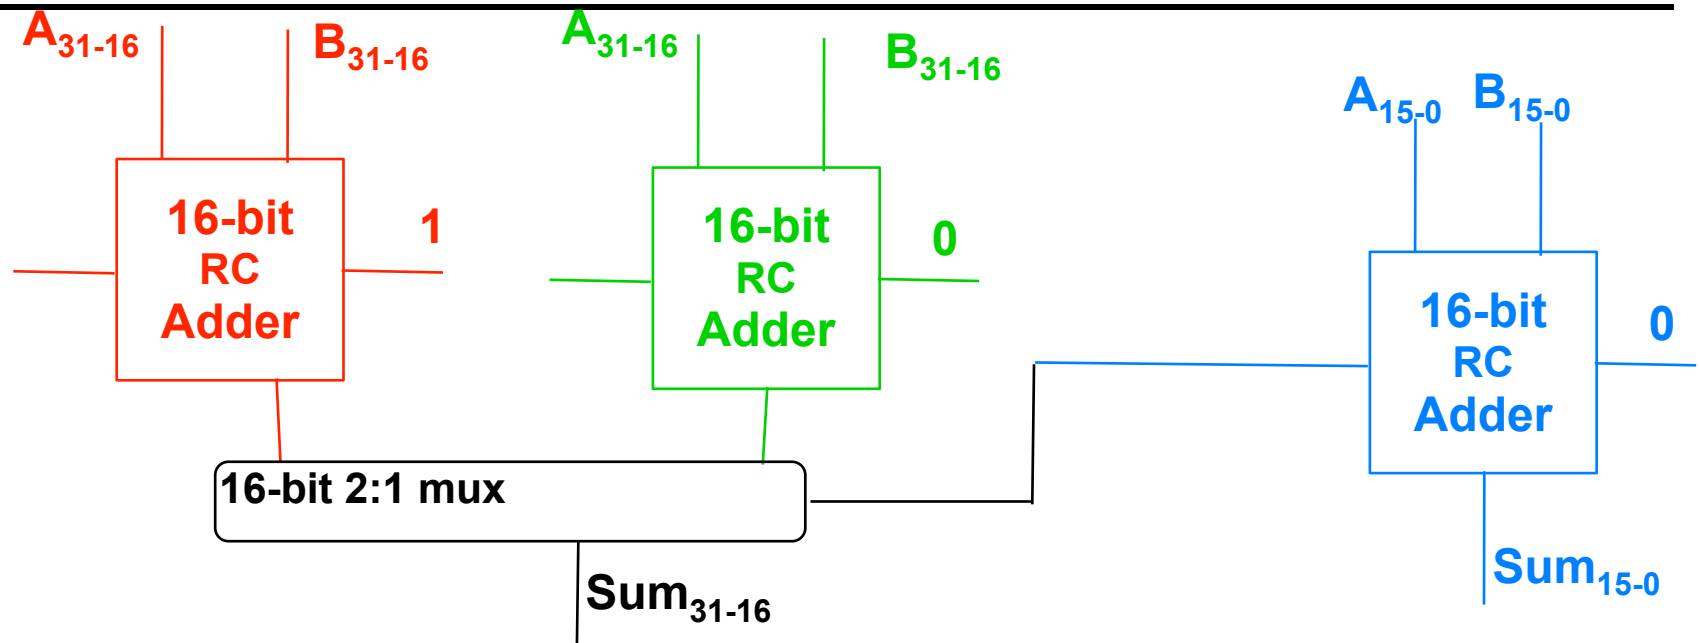
Adding

- Adding is important
 - Want to fit add in single clock cycle
 - (More on clocking soon)
 - Why? Add is ubiquitous
- Ripple Carry is slow
 - Maybe can do better?
 - But seems like C_{in} always depends on prev C_{out}
 - ...and C_{out} always depends on C_{in} ...

Hardware != Software

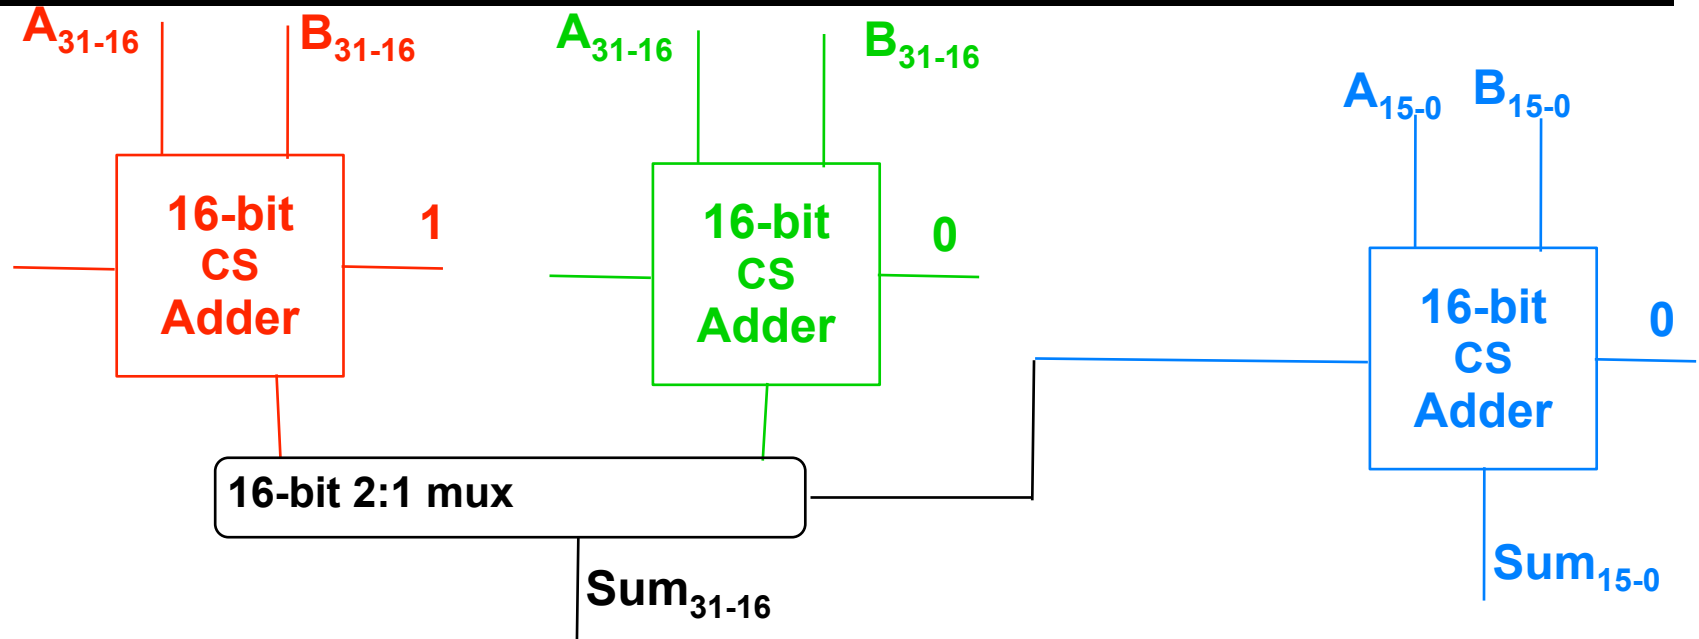
- If this were software, we'd be out of luck
 - But hardware is different
 - Parallelism: can do many things at once
 - Speculation: can guess

Carry Select



- Do three things at once (32 gates)
 - Add low 16 bits
 - Add high 16 bits assuming $CI = 0$
 - Add high 16 bits assuming $CI = 1$
- Then pick correct assumption for high bits (2—3 gates)

Carry Select



- Could apply same idea again
 - Replace 16-bit RC adders with 16-bit CS adders
 - Reduce delay for 16 bit add from 32 to 18
 - Total 32 bit adder delay = 20
- So... just go nuts with this right?

Tradeoffs


- Tradeoffs in doing this
 - Power and Area (\sim number of gates)
 - Roughly double every “level” of carry select we use
 - Less return on increase each time
 - Adding more mux delays
 - Wire delays increase with area
 - Not easy to count in slides
 - But will eat into real performance
- Fancier adders: recitation
 - Can do even better

Recall: Subtraction

- 2's complement makes subtraction easy:
 - Remember: $A - B = A + (-B)$
 - And: $-B = \sim B + 1$
 - ↑ that means flip bits ("not")
 - So we just flip the bits and start with CI = 1
 - Fortunate for us: **makes circuits easy**

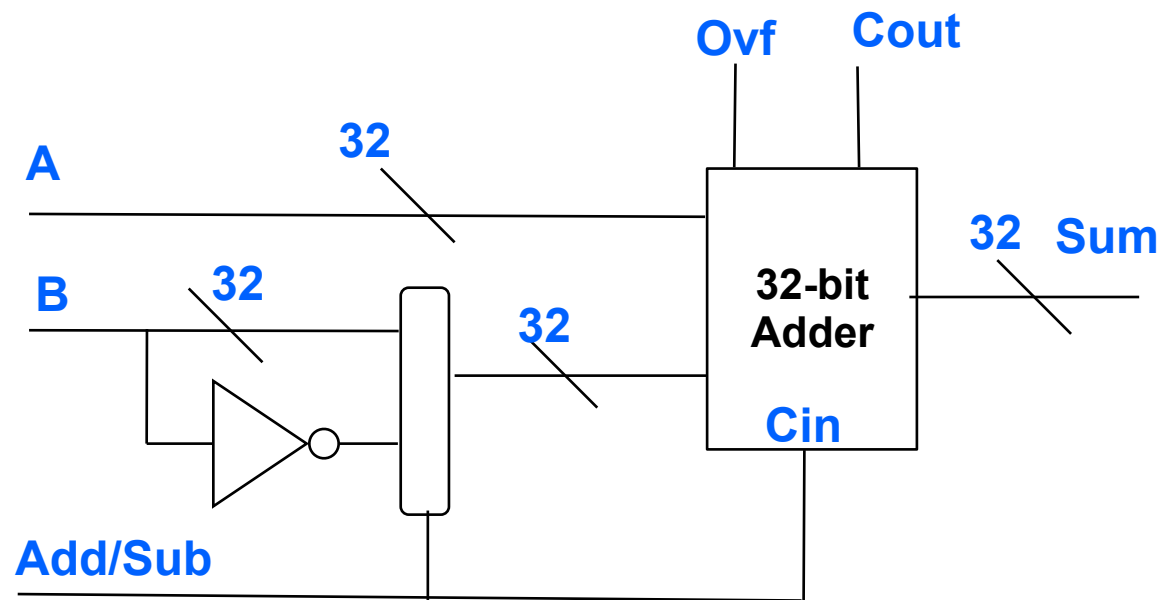
- $$\begin{array}{r} 0110101 \\ - 1010010 \\ \hline \end{array}$$
- $$\begin{array}{r} 0110101 \\ + 0101101 \\ \hline \end{array}$$

32-bit Adder/subtractor



- Inputs: A, B, Add/Sub (0=Add, 1 = Sub)
- Outputs: Sum, Cout, Ovf (Overflow)

32-bit Adder/subtractor



- By the way:
 - That thing has about 3,000 transistors
 - Aren't you glad we have abstraction?

Arithmetic Logic Unit (ALU)

- ALUs do a variety of math/logic
 - Add
 - Subtract
 - Bit-wise operations: And, Or, Xor, Not
 - Shift (left or right)
- Take two inputs (A,B) + operation (add,shift..)
 - Do a variety in parallel, then mux based on op

Bit-wise operations: SHIFT

- Left shift (\ll)
 - Moves left, bringing in 0s at right, excess bits “fall off”
 - $10010001 \ll 2 = 01000100$
 - $x \ll k$ corresponds to $x * 2^k$
- Logical (or unsigned) right shift (\gg)
 - Moves bits right, bringing in 0s at left, excess bits “fall off”
 - $10010001 \gg 3 = 00010010$
 - $x \gg k$ corresponds to $x / 2^k$ for unsigned x
- Arithmetic (or signed) right shift (\gg)
 - Moves bits right, bringing in (sign bit) at left
 - $10010001 \gg 3 = 11110010$
 - $x \gg k$ corresponds to $x / 2^k$ for signed x

Shift: Implementation...?

- Suppose an 8-bit number

$b_7b_6b_5b_4b_3b_2b_1b_0$

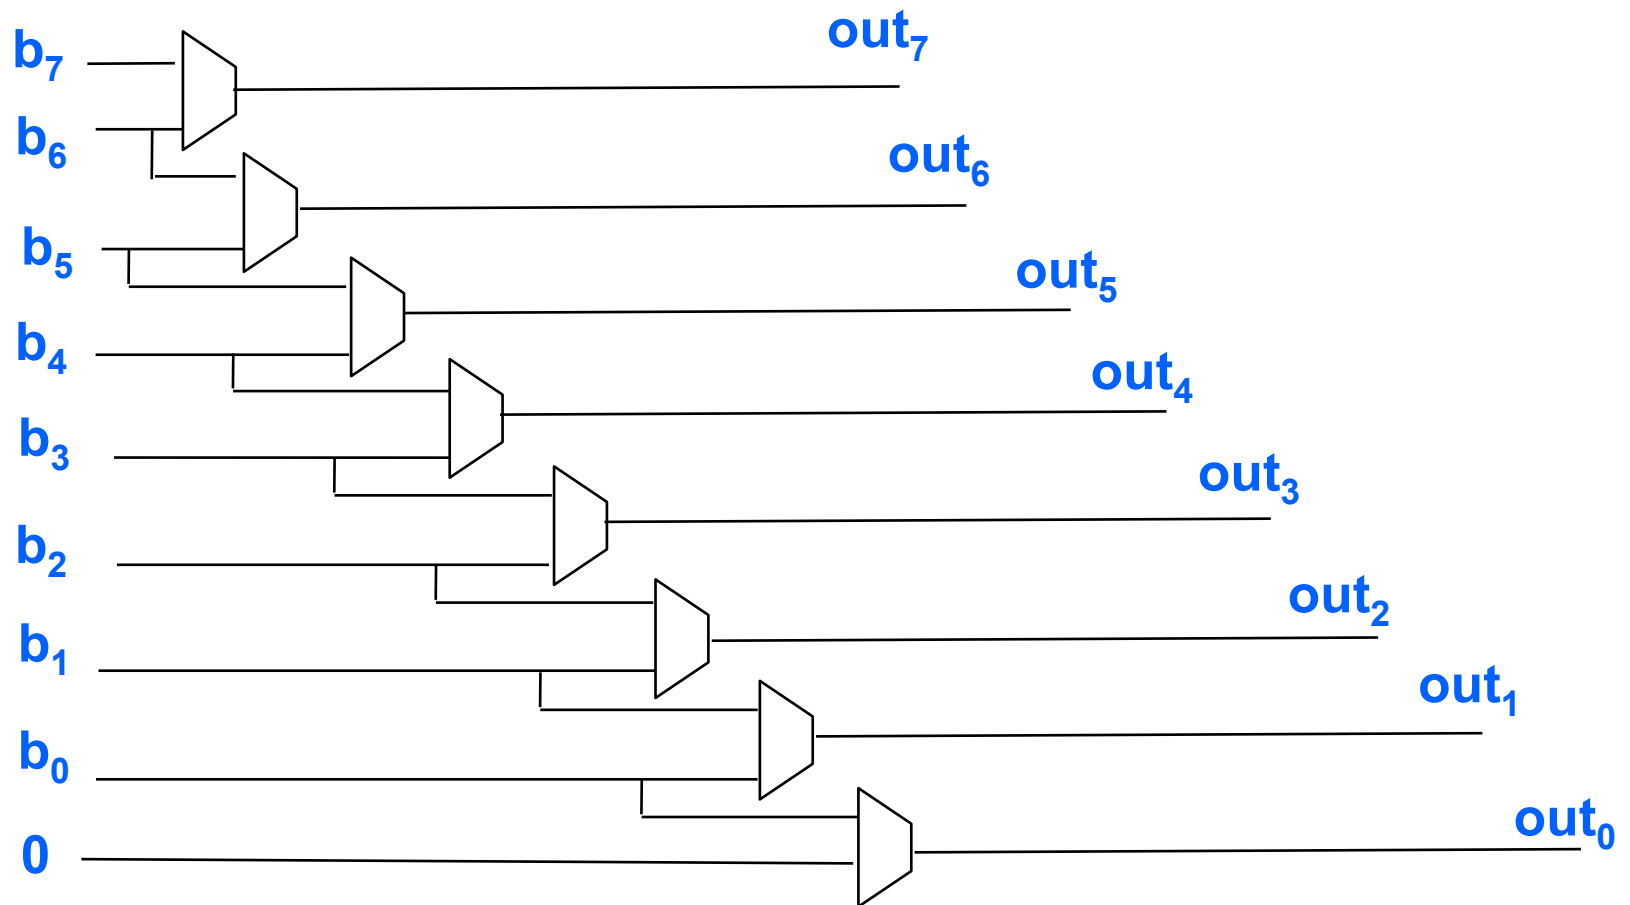
Shifted left by a 3 bit number

$s_2s_1s_0$

- Option 1: Truth Table?
 - 2048 rows? Not appealing

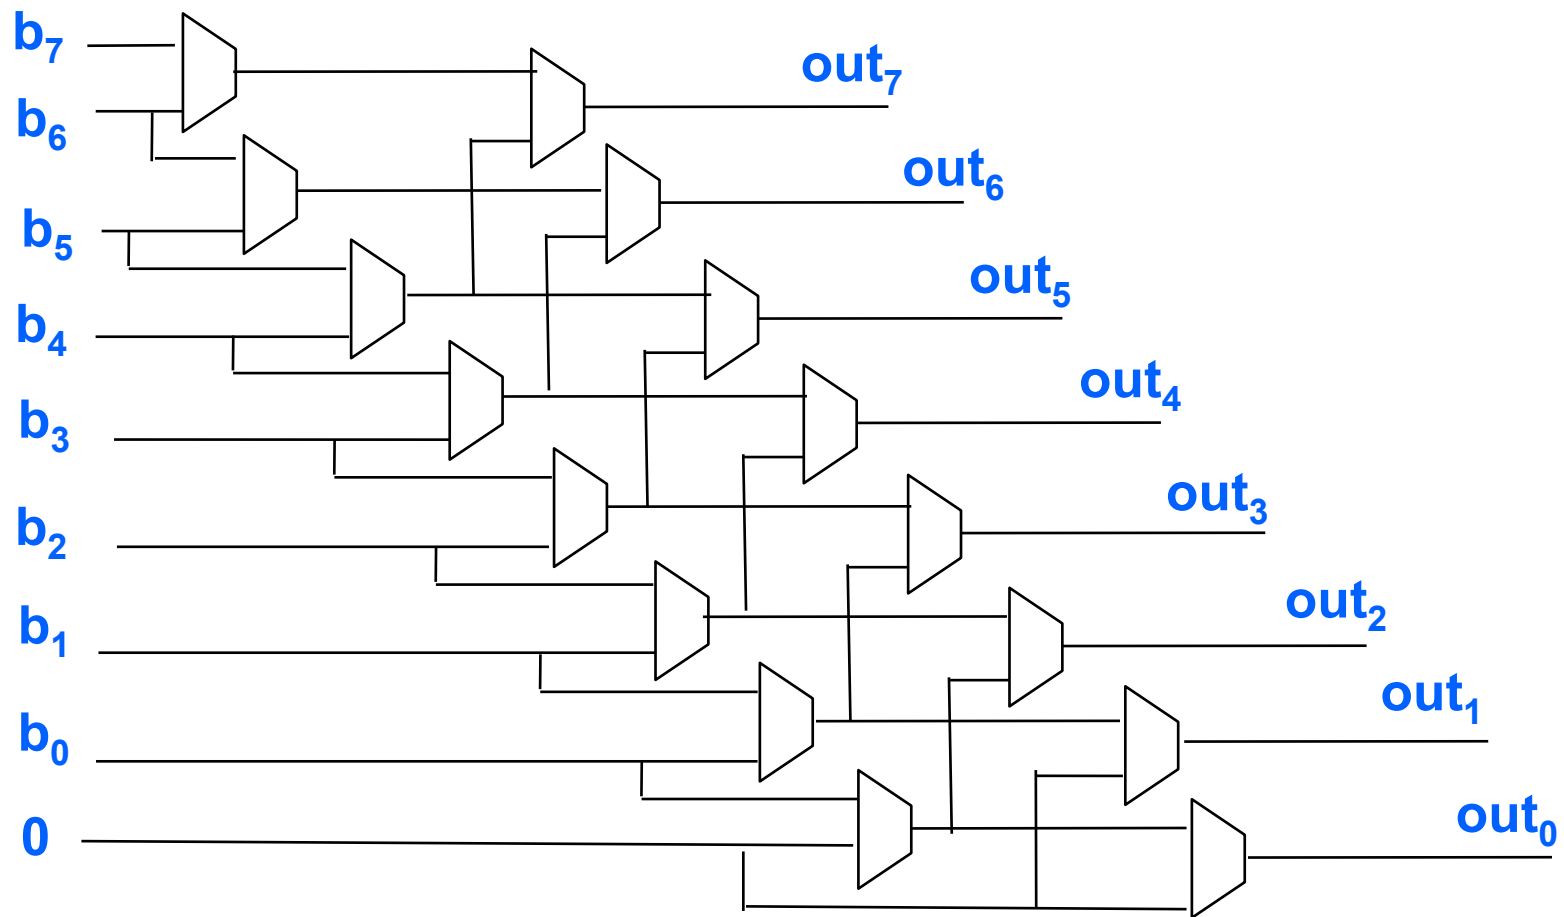
Lets simplify

- Simpler problem: 8-bit number shifted by 1 bit number (shift amount selects each mux)



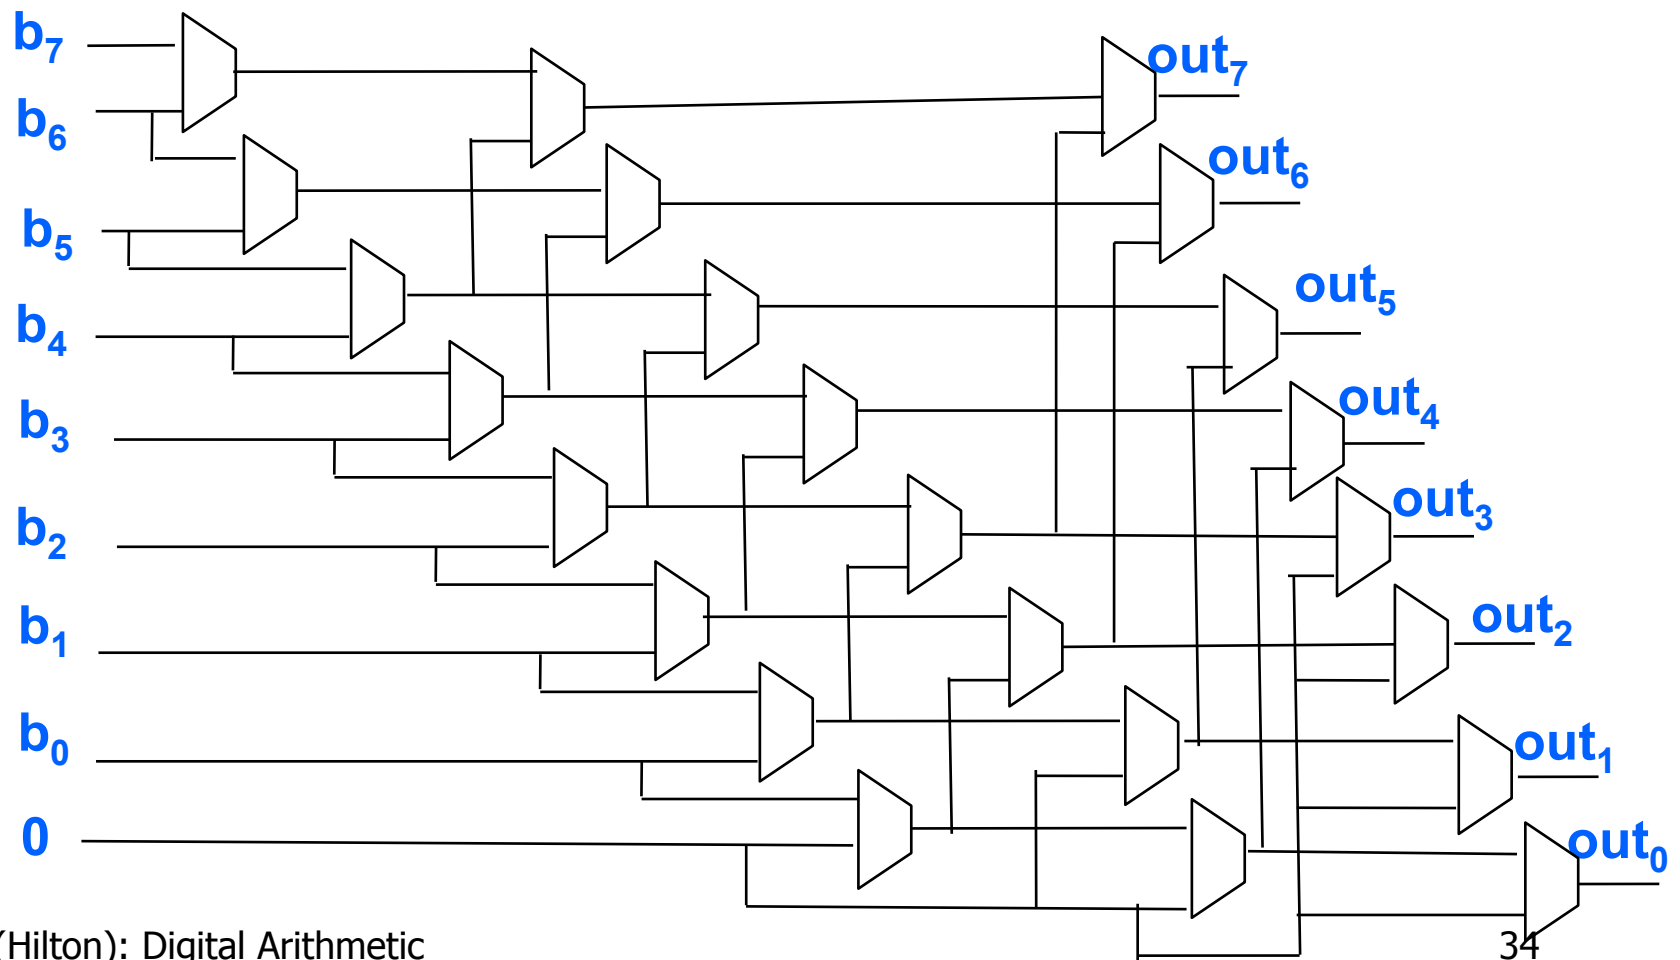
Lets simplify

- Simpler problem: 8-bit number shifted by 2 bit number (new muxes selected by 2nd bit)



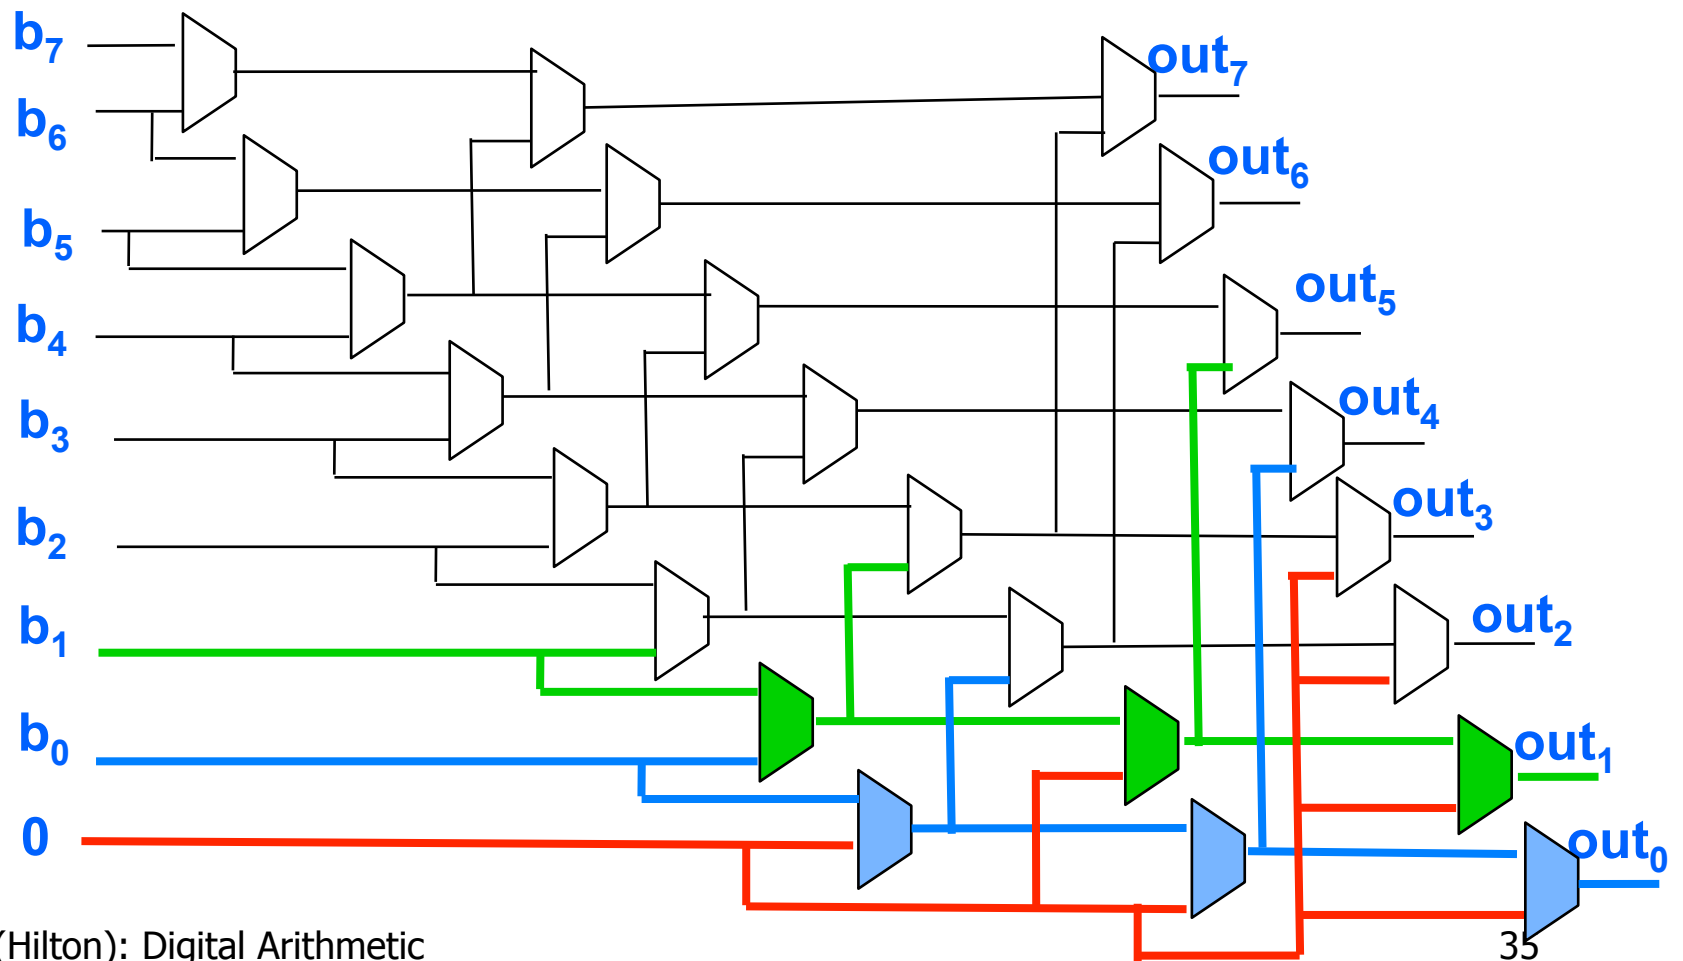
Now shifted by 3-bit number

- Full problem: 8-bit number shifted by 3 bit number (new muxes selected by 3rd bit)



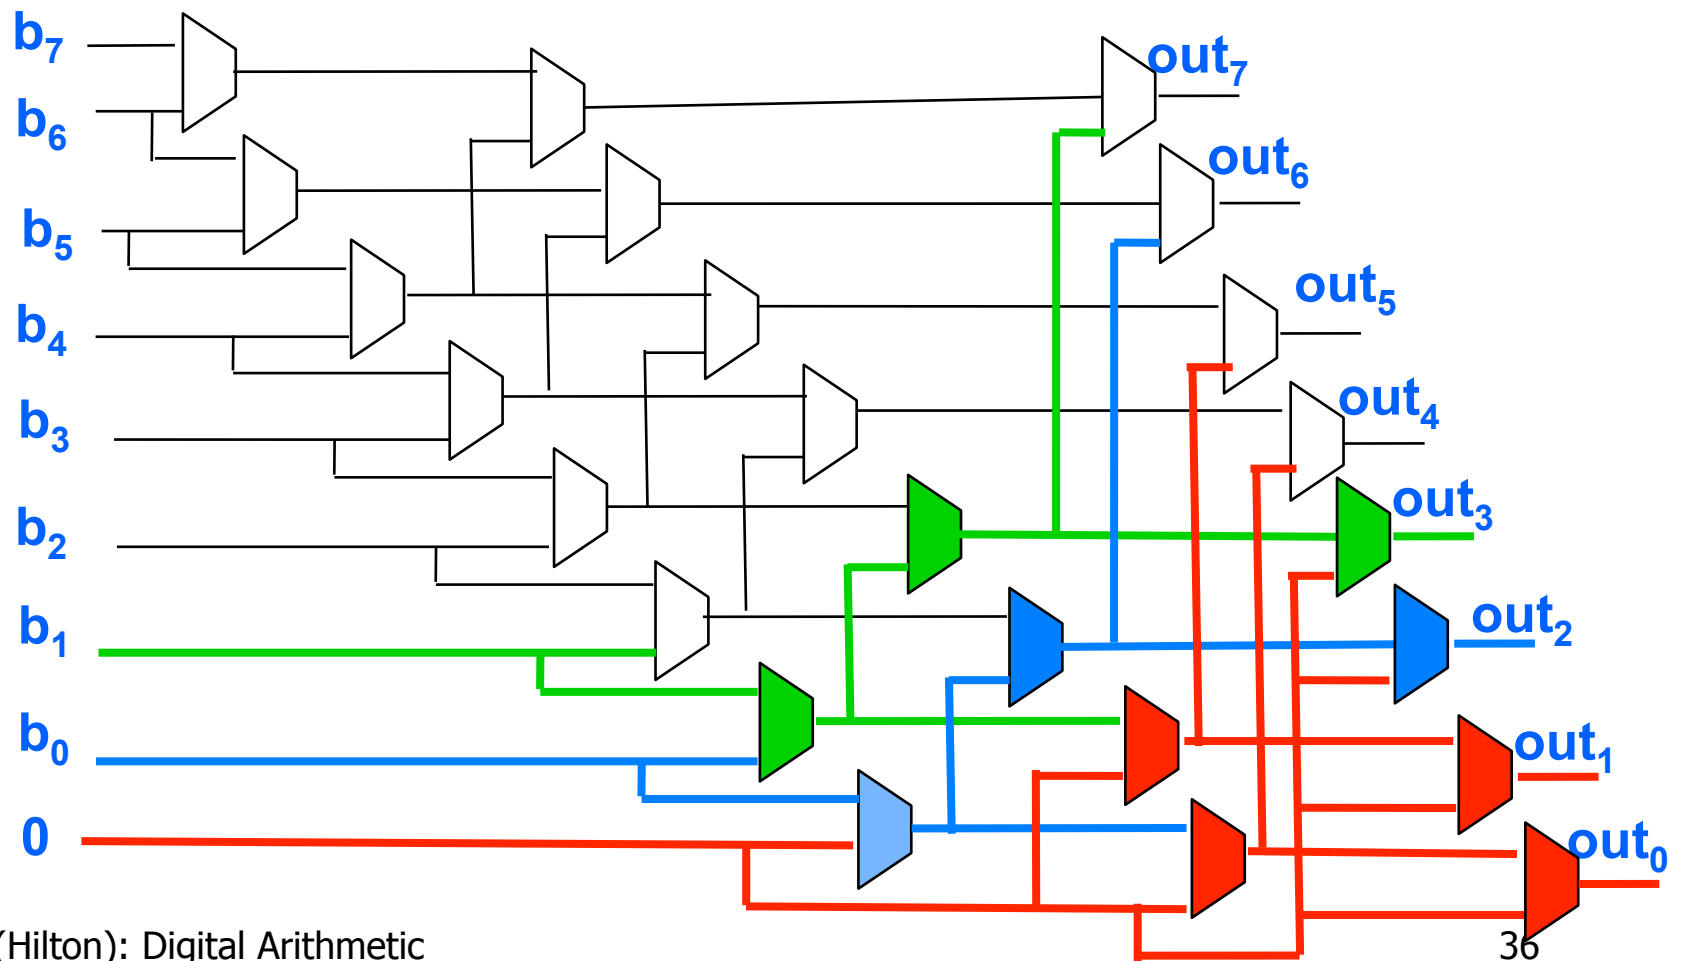
Now shifted by 3-bit number

- Shifter in action: shift by 000



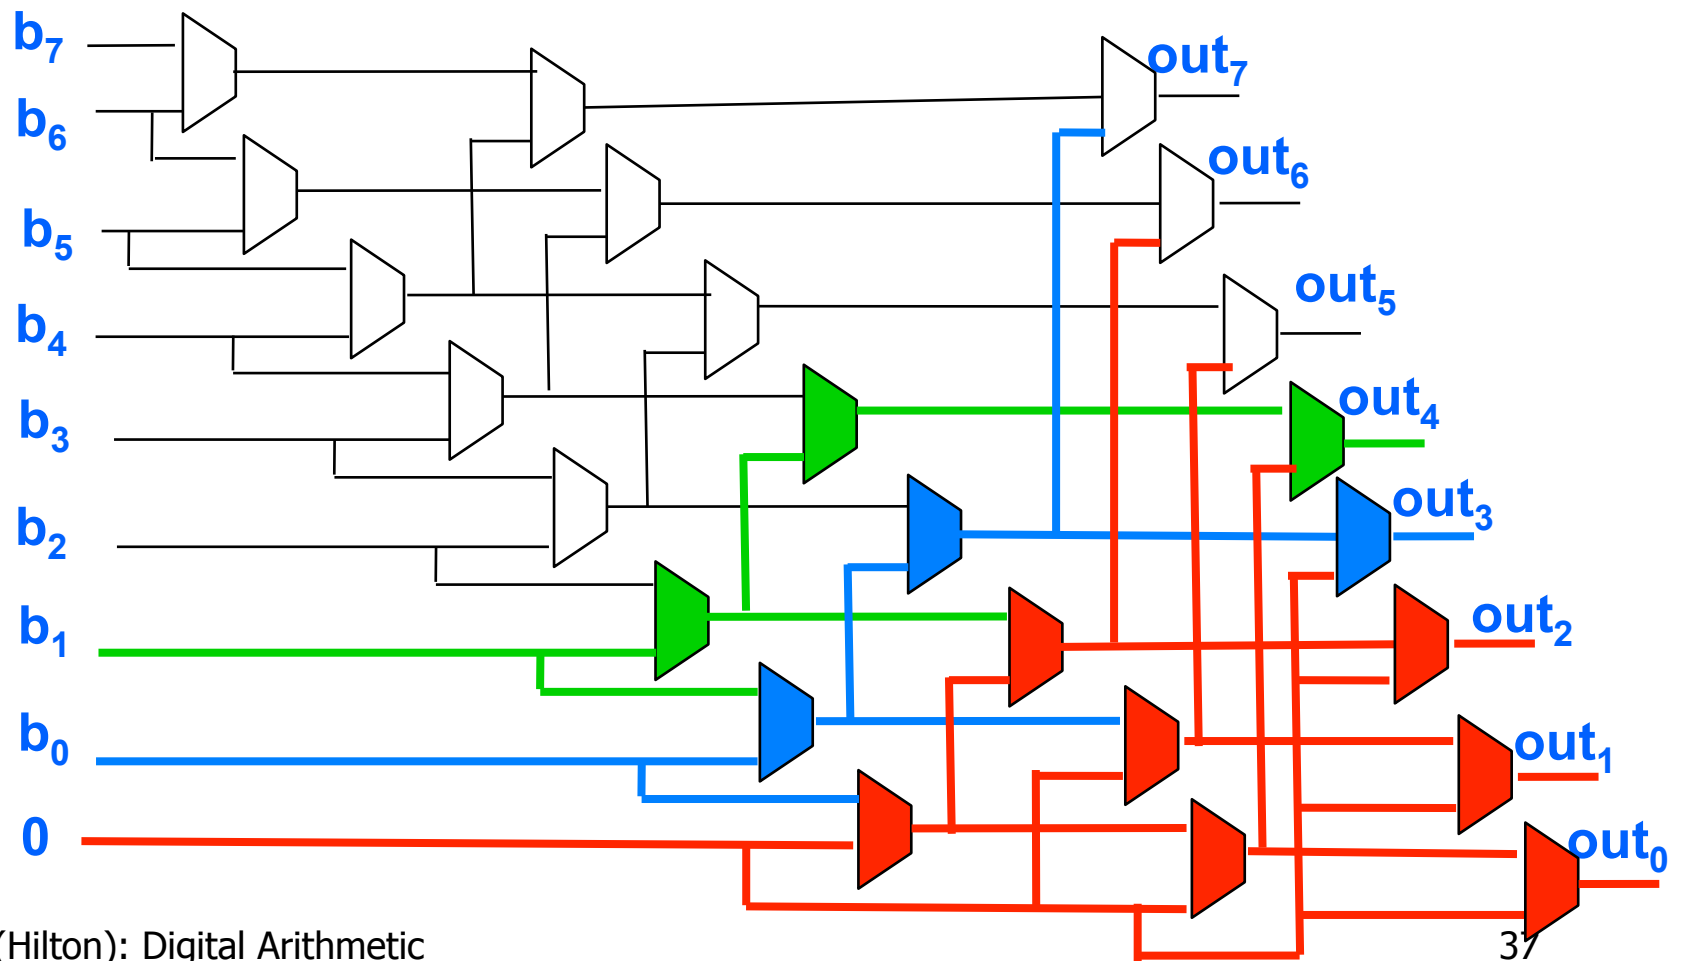
Now shifted by 3-bit number

- Shifter in action: shift by 010



Now shifted by 3-bit number

- Shifter in action: shift by 011



What About Non-integer Numbers?

- There are infinitely many real numbers between two integers
- Many important numbers are real
 - $\pi = 3.145\dots$
 - $\frac{1}{2} = 0.5$
- How could we represent these sorts of numbers?
 - Fixed Point
 - Rational
 - Floating Point (IEEE Single Precision)

Floating Point

- Think about scientific notation for a second:
- For example:
 $6.02 * 10^{23}$
- Real number, but comprised of ints:
 - 6 generally only 1 digit here
 - 2 any number here
 - 10 always 10 (base we work in)
 - 23 can be positive or negative
- Can we do something like this in binary?

Floating Point

- How about:
- $\pm X.YYYYYYY * 2^{\pm N}$
- Big numbers: large positive N
- Small numbers (<1): negative N
- Numbers near 0: small N
- This is “floating point” : most common way

IEEE single precision floating point

- Specific format called IEEE single precision:
- $\pm 1.YYYYYY * 2^{(N-127)}$
- “float” in Java, C, C++,...

- Assume X is always 1 (save a bit)
- 1 sign bit (+ = 0, 1 = -)
- 8 bit biased exponent (do N-127)
- Implicit 1 before binary point
- 23-bit mantissa (YYYYYY)

Binary fractions

- 1.YYYY has a binary point
 - Like a decimal point but in binary
 - After a decimal point, you have
 - tenths
 - hundredths
 - Thousandths
 -
- So after a binary point you have...

Binary fractions

- 1.YYYY has a binary point
 - Like a decimal point but in binary
 - After a decimal point, you have
 - Tenths
 - Hundredths
 - Thousandths
 -
- So after a binary point you have...
 - Halves
 - Quarters
 - Eights
 -

Floating point example

- Binary fraction example:
 - $101.101 = 4 + 1 + \frac{1}{2} + \frac{1}{8} = 5.625$
- For floating point, needs normalization:
 - $1.01101 * 2^2$
- Sign is +, which = 0
- Exponent = $127 + 2 = 129 = 1000\ 0001$
- Mantissa = $1.011\ 0100\ 0000\ 0000\ 0000\ 0000$

31 30 23 22 0

0 | 1000 0001 | 011 0100 0000 0000 0000 0000

Floating Point Representation

Example:

What floating-point number is:

0xC1580000?

Answer

What floating-point number is

0xC1580000?

1100 0001 0101 1000 0000 0000 0000 0000

Sign = 1 which is negative

Exponent = $(128+2)-127 = 3$

Mantissa = 1.1011

$-1.1011 \times 2^3 = -1101.1 = -13.5$

Trick question

- How do you represent 0.0?
 - Why is this a trick question?

Trick question

- How do you represent 0.0?
 - Why is this a trick question?
 - $0.0 = 000000000$
 - But need 1.XXXXXX representation?

Trick question

- How do you represent 0.0?
 - Why is this a trick question?
 - $0.0 = 000000000$
 - But need 1.XXXXX representation?
- Exponent of 0 is denormalized
 - Implicit 0. instead of 1. in mantissa
 - Allows 0000....0000 to be 0
 - Helps with very small numbers near 0
- Results in +/- 0 in FP (but they are "equal")

Other weird FP numbers

- Exponent = 1111 1111 also not standard
 - All 0 mantissa: $\pm \infty$
 $1/0 = +\infty$
 $-1/0 = -\infty$
 - Non zero mantissa: Not a Number (NaN)
 $\text{sqrt}(-42) = \text{NaN}$

Floating Point Representation

- Double Precision Floating point:

64-bit representation:


- 1-bit **sign**
 - 11-bit (biased) **exponent**
 - 52-bit **fraction** (with implicit 1).
- “double” in Java, C, C++, ...

S	Exp	Mantissa
1	11-bit	52 - bit

Danger: floats cannot hold all ints!

- Many programmers think:

- Floats can represent all ints
- NOT true



- First summer internship I had:

- Need some floats and some ints: just use floats!
- Bug in their code!
- Other developers shocked as I demonstrated problem...

- Doubles can represent all 32-bit ints

- (but not all 64-bit ints)



Wrap Up

- Implementation of Math
 - Addition/Subtraction
 - Shifting
- Floating Point Numbers
 - IEEE representation
 - Denormalized Numbers
- Next Time:
 - Storage
 - Clocking