


ECE 590.03 Fundamentals of Computer Systems and Engineering

IO Devices

ECE 590.03 IO

1

IO: Interacting with the outside world



Input and Output Devices

- Video
- Disk
- Keyboard
- Sound
- ...

ECE 590.03 IO

2


Communication with IO devices

- Processor needs to get info to/from IO device
 - Two ways:
 - In/out instructions
 - Read/write value to "io port"
 - Devices have specific port numbers
 - Memory mapped
 - Regions of physical addresses not actually in DRAM
 - But mapped to IO device
 - Stores to mapped addresses send info to device
 - Reads from mapped addresses get info from device

ECE 590.03 IO

3

A view of the world

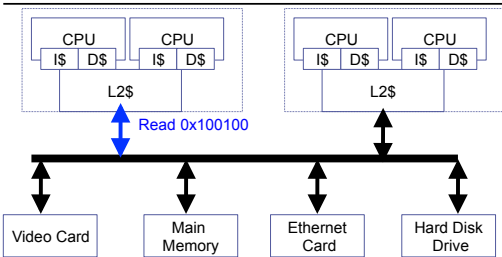


- 2 "socket" system (each with 2 cores)
- Real systems: more IO devices

ECE 590.03 IO

4

A view of the world

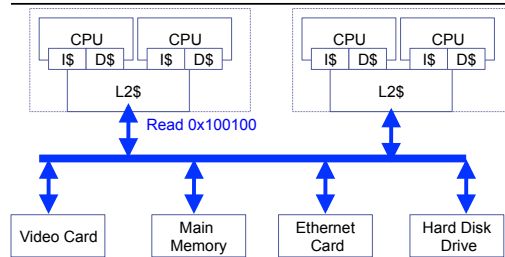


- Chip 0 requests read of 0x100100

ECE 590.03 IO

5

A view of the world

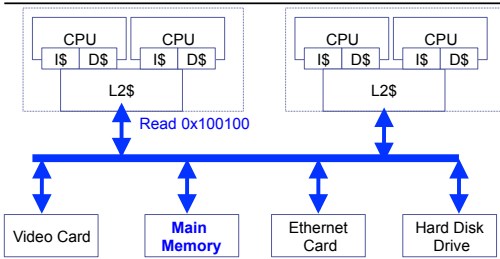


- Chip 0 requests read of 0x100100
- Request goes to all devices

ECE 590.03 IO

6

A view of the world




- Chip 0 requests read of 0x100100
- Request goes to all devices, which check address ranges

ECE 590.03 IO

7

A view of the world



- Other address ranges may be for a particular device

ECE 590.03 IO

8

Speaking of VGA video

- You all wrote a VGA controller early (hwk2)
 - Read a ROM with an image
 - Real ones: read a RAM
 - How to draw? CPU writes to physical memory mapped to video card RAM
 - Video card sees write and updates its internal RAM
 - The rest: FSM just like you did
- (Except 3D accelerators)

ECE 590.03 IO

9

Exploring Memory Mappings on Linux

- You can see what devices have what memory ranges on linux with `lspci -v` (at least those on the PCI bus)

```
00:02.0 VGA compatible controller: Intel Corporation Core Processor Integrated
Graphics Controller (rev 02)
Subsystem: Lenovo Device 215a
Flags: bus master, fast devsel, latency 0, IRQ 30
Memory at f2000000 (64-bit, non-prefetchable) [size=4M]
Memory at d0000000 (64-bit, prefetchable) [size=256M]
I/O ports at 1800 [size=8]
Capabilities: [90] Message Signalled Interrupts: Mask- 64bit- Queue=0/0
Enable+
Capabilities: [d0] Power Management version 2
Capabilities: [a4] PCIe advanced features <?>
Kernel driver in use: i915
Kernel modules: i915
```

ECE 590.03 IO

10

A simple "IO device" example

- Read (physical) address 0xFFFF1000 for "ready"
- If ready, read address 0xFFFF1004 for data value
 - IO device will go to next value automatically on read
- Write a value to 0xFFFF1008 to output it

```
read_dev:
la $t0, 0xFFFF1000
loop:
lw $t1, 0($t0)
beqz $t1, loop
lw $v0, 4($t0)
jr $ra
```

Who can remind us what this is called (last lecture)?

ECE 590.03 IO

11

A handful of questions...

- How do we use physical addresses?
 - Programs only know about virtual addresses right?
- What about caches?
 - Won't the first lw bring the current value of 0xFFFF1000 into the cache?
 - And then subsequent requests just hit the cache?

ECE 590.03 IO

12

A handful of questions...

- How do we use physical addresses?
 - Programs only know about virtual addresses right?
 - **Only OS accesses IO devices:**
 - OS knows about physical addresses, and can use them
- What about caches?
 - Won't the first lw bring the current value of 0xFFFF1000 into the cache?
 - And then subsequent requests just hit the cache?

ECE 590.03 IO

13


A handful of questions...

- How do we use physical addresses?
 - Programs only know about virtual addresses right?
 - Only OS accesses IO devices:
 - OS knows about physical addresses, and can use them
- What about caches?
 - Won't the first lw bring the current value of 0xFFFF1000 into the cache?
 - And then subsequent requests just hit the cache?
 - **Pages have attributes, including cacheability**
 - IO mapped pages marked non-cacheable
 - Also, prevent speculative loads (e.g., out-of-order)
 - Remember: speculative only fine as long as nobody knows

ECE 590.03 IO

14

Hard disks




- Viewed from above:
 - Disks are circular platters of spinning metal
 - Multiple tracks (concentric rings)
 - Each track divided into sectors
 - Modern disks: addressed by "logical block"
- (Real disks are actually circular...)

ECE 590.03 IO

15

Hard disks




- Read/written by "head"
 - Moves across tracks ("**seek**")
 - After seek completes, wait for proper sector to rotate under head.
 - Reads or writes magnetic medium by sensing/changing magnetic state (this takes time as the desired data 'spins under' the head)

ECE 590.03 IO

16

Hard disks




- Want to read data on blue curve (imagine circular arc)

ECE 590.03 IO

17

Hard disks




- Want to read data on blue curve (imagine circular arc)
 - First step: seek—move head over right track
 - Takes time (Tseek), disk keeps spinning

ECE 590.03 IO

18

Hard disks




- Want to read data on blue curve (imagine circular arc)
 - First step: seek—move head over right track
 - Takes time (Tseek), disk keeps spinning
 - Now head over right track... but data needs to move under head
 - Second step: wait (Trotate)

ECE 590.03 IO

19

Hard disks




- Want to read data on blue curve (imagine circular arc)
 - First step: seek—move head over right track
 - Takes time (Tseek), disk keeps spinning
 - Now head over right track... but data needs to move under head
 - Second step: wait (Trotate)
 - Third: as data comes under head, start reading

ECE 590.03 IO

20

Hard disks




- Want to read data on blue curve (imagine circular arc)
 - First step: seek—move head over right track
 - Takes time (Tseek), disk keeps spinning
 - Now head over right track... but data needs to move under head
 - Second step: wait (Trotate)
 - Third: as data comes under head, start reading
 - Takes time for data to pass under read head (Tread)

ECE 590.03 IO

21

Hard Disks: from the side



- Multiple platters, each with a head above and below
 - Two sided surface
 - Heads all stay together ("cylinder")
 - Heads not actually touching platters: just **very** close

ECE 590.03 IO

22

A few things about HDD performance

- Tseek:
 - Depends on how fast heads can move
 - And how far they have to go
 - OS may try to schedule IO requests to minimize Tseek
- Trotate:
 - Depends largely on how fast disk spins (RPM)
 - Also, how far around the data must spin, but usually assume avg
 - OS cannot keep track of position, nor schedule for better
- Tread:
 - Depends on RPM + how much data to read

ECE 590.03 IO

23

Disk Drive Performance

- Suppose on average
 - Tseek = 10 ms
 - Trotate = 3.0 ms
 - Tread = 5 usec/ 512-byte sector
- What is the average time to read one 512-byte sector?
 - $10 \text{ ms} + 3 \text{ ms} + 0.05 \text{ ms} = 13.05 \text{ ms}$
 - Reading 1 sector a a time: $512 \text{ byte} / 13.05 \text{ ms} \Rightarrow \sim 40\text{KB/sec}$

ECE 590.03 IO

24

Disk Drive Performance

- Suppose on average
 - Tseek = 10 ms
 - Trotate = 3.0 ms
 - Tread = 5 usec/ 512-byte sector
- What is the average time to read one 512-byte sector?
 - $10\text{ ms} + 3\text{ ms} + 0.005\text{ ms} = 13.005\text{ ms}$
 - Reading 1 sector a a time: 512 byte/ 13.005 ms => ~40KB/sec
- What is the avg time to read 1MB of (contiguous) data?
 - 1MB = 2048 sectors
 - $10 + 3 + 0.005 * 2048 = 23.24\text{ ms} \Rightarrow \sim 43\text{MB/sec}$

ECE 590.03 10

25

Disk Drive Performance

- Suppose on average
 - Tseek = 10 ms
 - Trotate = 3.0 ms
 - Tread = 5 usec/ 512-byte sector
- What is the average time to read one 512-byte sector?
 - $10\text{ ms} + 3\text{ ms} + 0.005\text{ ms} = 13.005\text{ ms}$
 - Reading 1 sector a a time: 512 byte/ 13.005 ms => ~40KB/sec
- What is the avg time to read 1MB of (contiguous) data?
 - 1MB = 2048 sectors
 - $10 + 3 + 0.005 * 2048 = 23.24\text{ ms} \Rightarrow \sim 43\text{MB/sec}$
- Larger contiguous reads: approach **100MB/sec**
 - **Amortize** Tseek + Trotate (key to good disk performance)

ECE 590.03 10

26

Disk Performance

- Hard disks have caches (spatial locality)
- OS will also buffer disk in memory
 - Ask to read 16 bytes from a file?
 - OS reads multiple KB, buffers in memory

ECE 590.03 10

27

Disk Performance

- Hard disks have caches (spatial locality)
- OS will also buffer disk in memory
 - Ask to read 16 bytes from a file?
 - OS reads multiple KB, buffers in memory
- “Defragmenting” (Windows):
 - Improve locality by putting blocks for same files near each other

ECE 590.03 10

28

Transferring the data to memory

- OS asks disk to read data
 - Disk read takes a long time (15 ms => millions of cycles)
 - Does OS poll disk for 15M cycles looking for data?

ECE 590.03 10

29

Transferring the data to memory

- OS asks disk to read data
 - Disk read takes a long time (15 ms => millions of cycles)
 - Does OS poll disk for 15M cycles looking for data?
 - No—disk interrupts OS when data is ready.

ECE 590.03 10

30

Transferring the data to memory

- OS asks disk to read data
 - Disk read takes a long time (15 ms => millions of cycles)
 - Does OS poll disk for 15M cycles looking for data?
 - No—disk interrupts OS when data is ready.
- Ready: version 1
 - Disk has data, needs it transferred to memory
 - OS does "memcpy" like routine:
 - Read hdd memory mapped IO
 - Write appropriate location in main memory
 - Repeat
 - For many KB to a few MB

DMA: Direct Memory Access

- Alternative: DMA
 - When OS requests disk read, sets up DMA
 - "Read this data from the disk, and put it in memory for me"
 - DMA controller handles "memcpy"
 - Ready (version 2.0): data is in memory
- Frees up CPU to do useful things

Hard disk: reliability

- Hard disks fail relatively easily
 - Spinning piece of metal
 - With head hovering <1mm from platter
- Hard drive failures: major pain..
 - Anyone ever have one?
 -

Reliability

- Solution to functionality problem?
 - Level of indirection
- Solution to performance problem?
 - Add a cache
- Solution to a reliability problem?
 - ...?

Reliability

- Solution to functionality problem?
 - Level of indirection
- Solution to performance problem?
 - Add a cache
- Solution to a reliability problem?
 - Add error checking and correction
 - For HDD's checking is easy: "wont read data"
 - Simplest correction: keep 2 copies

RAID: Reliability

- Redundant Array of In-expensive Disks (RAID)
 - Keep 2 hard-drives with identical copies of the data
 - One fails? Replace it, copy the other drive to it, resume
 - Can work from other drive while waiting for replacement
 - Performance?

RAID: Reliability

- Redundant Array of In-expensive Disks (RAID)
 - Keep 2 hard-drives with identical copies of the data
 - One fails? Replace it, copy the other drive to it, resume
 - Can work from other drive while waiting for replacement
 - Performance?
 - Writes to both drives in parallel (no cost)
 - Reads from **either** drive
 - Improve performance: twice the bandwidth

ECE 590.03 IO

37

RAID: Reliability

- Redundant Array of In-expensive Disks (RAID)
 - Keep 2 hard-drives with identical copies of the data
 - One fails? Replace it, copy the other drive to it, resume
 - Can work from other drive while waiting for replacement
 - Performance?
 - Writes to both drives in parallel (no cost)
 - Reads from **either** drive
 - Improve performance: twice the bandwidth
 - Downside?
 - Cost: need to buy 2x as many disks for 1x the space
 - Still: pretty popular (I have it on my home linux box)
 - Also very easy

ECE 590.03 IO

38

RAID: All sorts of things

- Mirroring data (prev slides): "RAID 1"
- Tons of other RAID configurations:
 - RAID 0: striping—performance, not reliability
 - Parity schemes: reduce overhead for num disks > 2
 - Still give reliability and good performance
- Many covered in detail in your book
 - Good to know they exist, may be good solution to a problem one day
 - Don't fret the obscure ones too much

ECE 590.03 IO

39

Other devices

- Wide variety of IO devices
 - Most basically work the same way from high-level
 - Read/write proper physical memory location(s)
- Reality: each device has its own protocol
 - Requires device driver: Software module that handles protocol details of specific device
 - Which memory locations to read/write etc
 - Example of?

ECE 590.03 IO

40

Other devices

- Wide variety of IO devices
 - Most basically work the same way from high-level
 - Read/write proper physical memory location(s)
- Reality: each device has its own protocol
 - Requires device driver: Software module that handles protocol details of specific device
 - Which memory locations to read/write etc
 - Example of
 - **Abstraction!**

ECE 590.03 IO

41

Next Up: OSes

- Working our way up the system
 - Just talked about how data is stored on disks...
 - Next time: how do we make a coherent filesystem?
- Followed up by various other bits of OS knowledge
 - How does the system boot?
 - How are programs scheduled?
 - For that matter where do they come from?

ECE 590.03 IO

42