

ECE 550

Fundamentals of Computer Systems and Engineering

Networking

Networking

- How do computers communicate?
 - Two computers connected by a direct wire?
 - Relatively straight forward: move bits across wire
 - Internet?
 - Many computers
 - All around the world
 - With other communications going on...
 - And un-reliable links
 - And tons of different systems, media, protocols...
- Pretty complicated, so how could we possibly manage it?

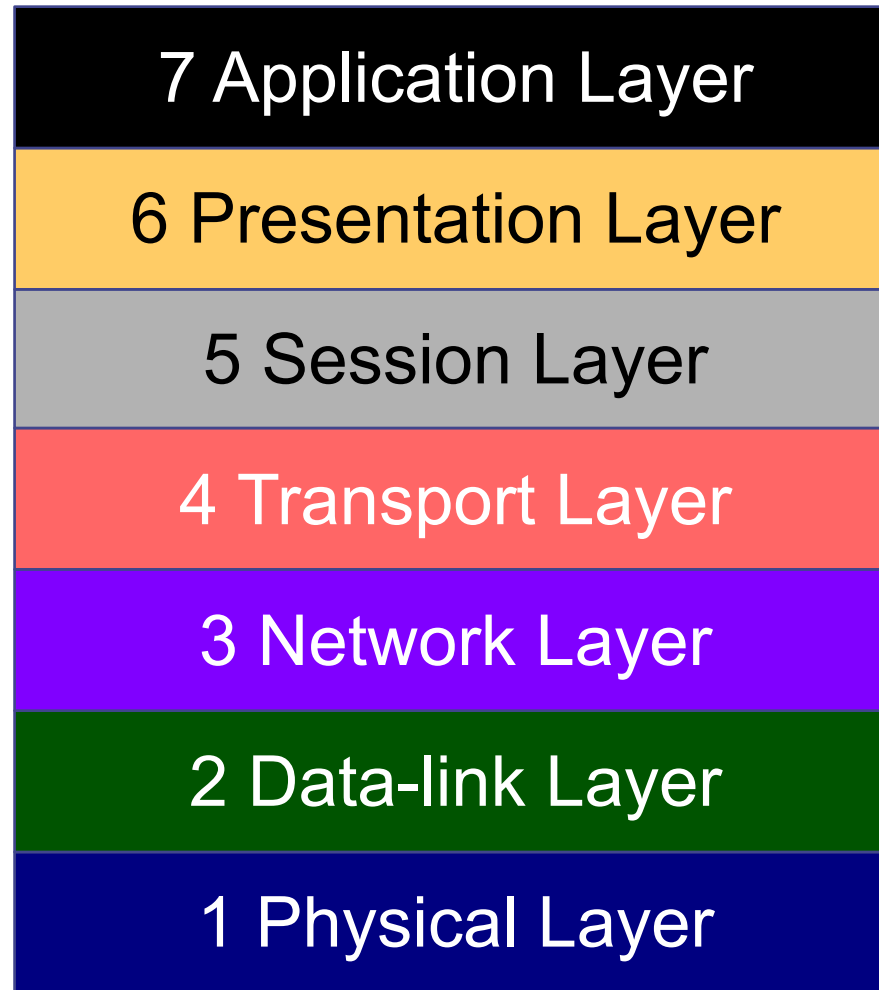
Networking

- How do computers communicate?
 - Two computers connected by a direct wire?
 - Relatively straight forward: move bits across wire
 - Internet?
 - Many computers
 - All around the world
 - With other communications going on...
 - And un-reliable links
 - And tons of different systems, media, protocols...
- Pretty complicated, so how could we possibly manage it?

Abstraction

(oh right, the answer to like...everything)

7-layer OSI model



- 7 layer networking stack
- (Theoretically) can change out any layer at a time

1 Physical Layer

- Defines physical how physical media work
 - Pin layout
 - Voltages
 - Timing Requirements
- Examples:
 - Cat 5 Cable (“Ethernet Cable”)
 - Wireless radio signal specifications
- Not much interesting to say here

2 Data-link Layer

- How to move bits across the wires in a meaningful way
 - Communication between two computers on same physical network
 - May include some error checking
- Example: Ethernet
 - Data transmitted in frames
 - Frame has:
 - **Pre-amble**: used to detect collisions
 - **Header**: source and destination MAC address
 - **Payload**: actual data
 - **CRC check**: detect corrupt data
 - Carrier Sense, Multiple Access, Collision Detect (**CSMA/CD**)
 - Carrier Sense: listen for if anyone else transmitting
 - Multiple Access: can wire up many computers to it
 - Collision Detect: two transmissions at once? Detect and retry



CSMACD

- Ethernet uses CSMACD for multiple systems on a network
 - Other options, but we won't go into them
 - Detection of collisions?
 - Pre-amble is fixed pattern
 - Network card senses medium while transmitting
 - Mismatch with expected? Collision
 - Collision happens?
 - Exponential backoff
 - Pick random number of time units
 - Retry
 - Fail again? Pick random number from 2x as big a range
- Analogy: crowded dinner party
 - Try to talk. Someone else talking? Wait. Try again. Fail again? Wait longer

Abstraction: Joys and Limitations

- Joys of abstraction:
 - Can build an Ethernet card without any info about higher layers
 - Will work with all of them
- Limitations:
 - 7-layer model's abstraction not perfect
 - Ethernet protocol imposes max limit on cable length
 - E.g., layer 2 constrains layer 1
 - This arises from the need to detect collisions before finishing sending

7-layer OSI model



- Reminder where we are so far

Our messages so far

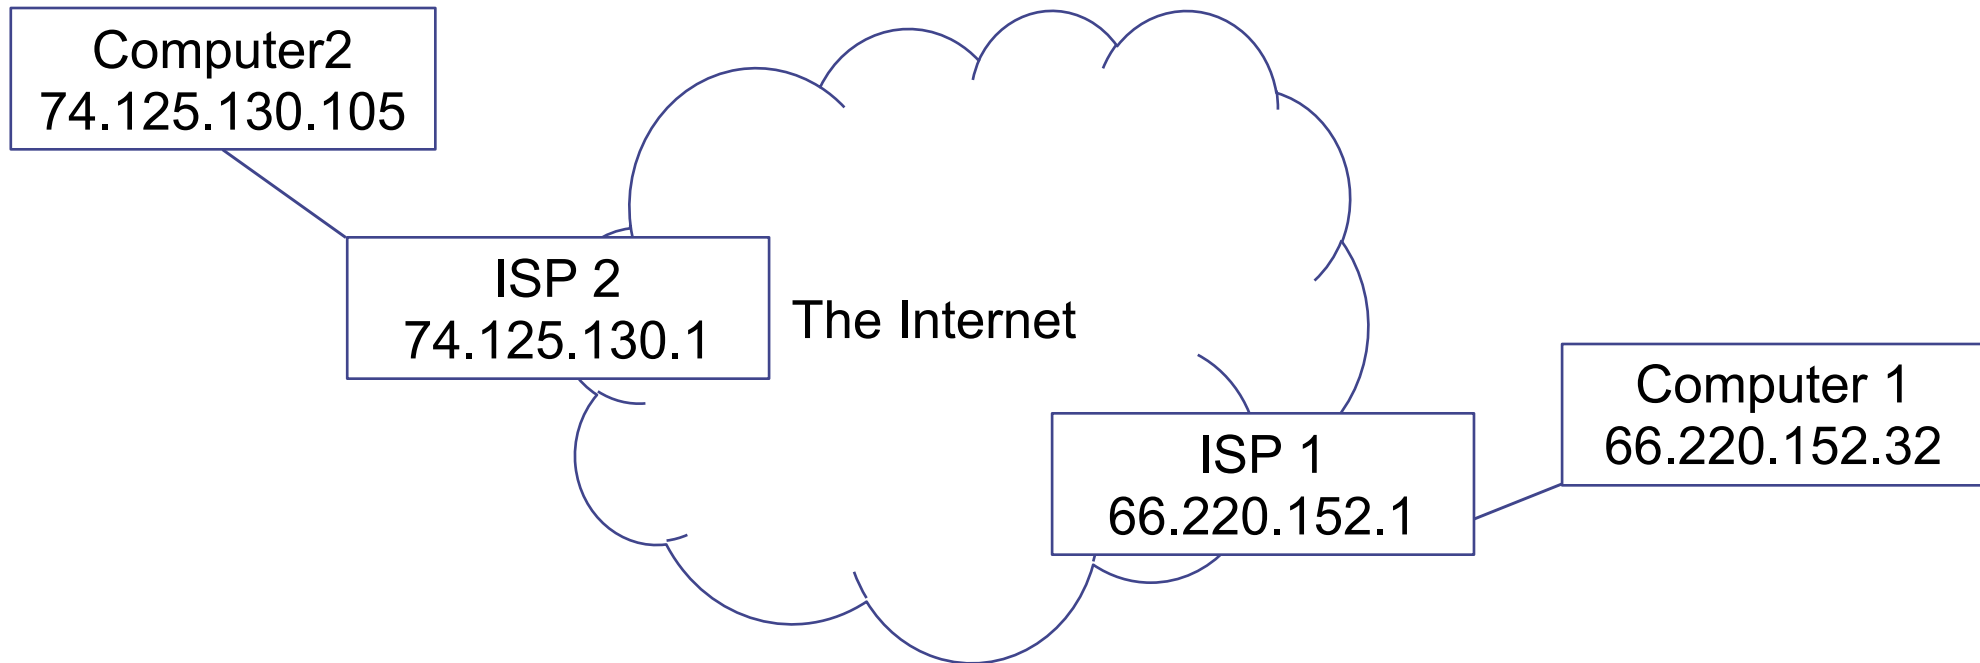


- Header says what network layer protocol the pay load is

3 Network Layer


- Layer 2 let's computers on **same** network talk
- Layer 3 let's computers talk across networks
 - Addressing
 - How do we specify what computer to talk to?
 - Routing
 - How do we get from here to there?
- Example: IP protocol
 - IPv4 and IPv6: pretty similar in most core regards
 - **Best effort delivery**
 - Addressing (IP addresses)
 - Routing
 - Analogy: Mailing a letter

IP



- Computer 1 wants to send data to Computer 2
 - For now, assume it knows IP address (we'll see DNS later)
 - Has direct connection to its ISP... but then what?
 - The internet is a big place after all..

Let's zoom in on ISP 1



- ISP has connections to a handful of other places
 - Generally very high bandwidth connections
 - Will send your data (packet) to one of these, but which one?

IP Routing

- IP addresses are hierarchical
 - May not know how to find 74.125.130.105...
 - But know which way to go to get to 74._____
 - Move one step closer
 - Within 74 network, know how to find 74.125
 - Then 74.125.130
 - Then find 74.125.130.105
- Analogy:
 - How do I get to 2200 Mission College Blvd, Santa Clara, CA?

IP Routing

- Analogy:
 - How do I get to 2200 Mission College Blvd, Santa Clara, CA?
 - I have no idea, but I can get you to I-40 West
 - Then you can ask someone else when you get to CA
 - Once in CA, you ask someone else:
 - “I only know its north of here, so take the 5 North and ask someone else”
 - Etc..
- This works because our physical addresses are hierarchical: Country, State, City, Street, Number

Routing Basics

- Routing is done with tables
 - CIDR notation: 40.1.0.0/16
 - Match first 16 bits of 40.1.0.0, ignore remaining 16 bits
 - Each number in IP addr is 8 bits, written in decimal
 - Find match, entry tells what link to send out on
 - Example
 - 40.0.0.0/8 => Link 0
 - 50.1.0.0/16 => Link 1
 - 50.2.0.0/16 => Link 2
 - 50.3.27.0/24 => Link 3
 - 50.3.42.0/24 => Link 1

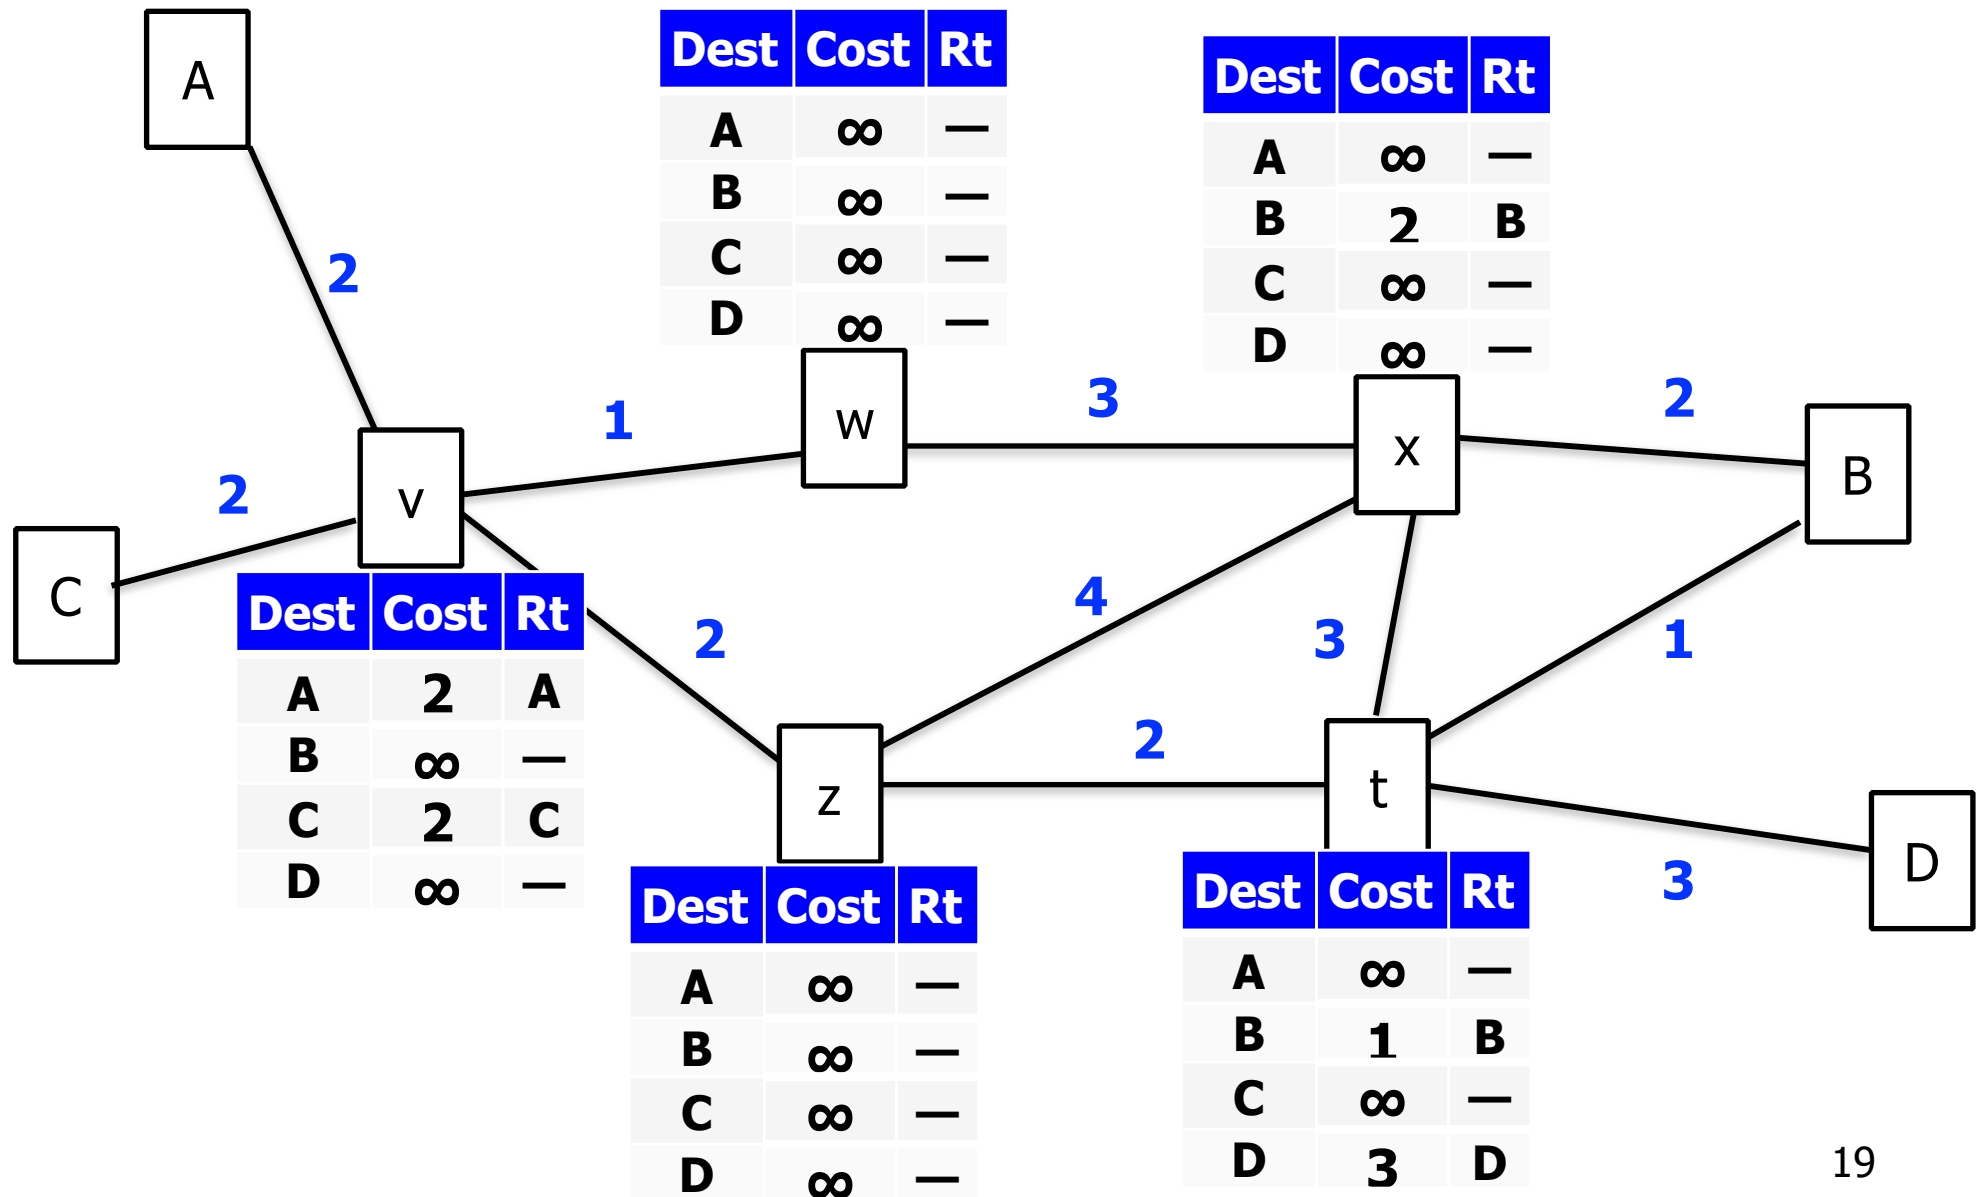
Routing: More Complex

- Approach one: Static Routing
 - Enter all routes
 - Let system run
 - Hope nothing goes down
 - Works fine for small networks
- Reality:
 - Network links/systems go down
 - Often multiple paths to same place
 - Changing traffic patterns = changing fastest route

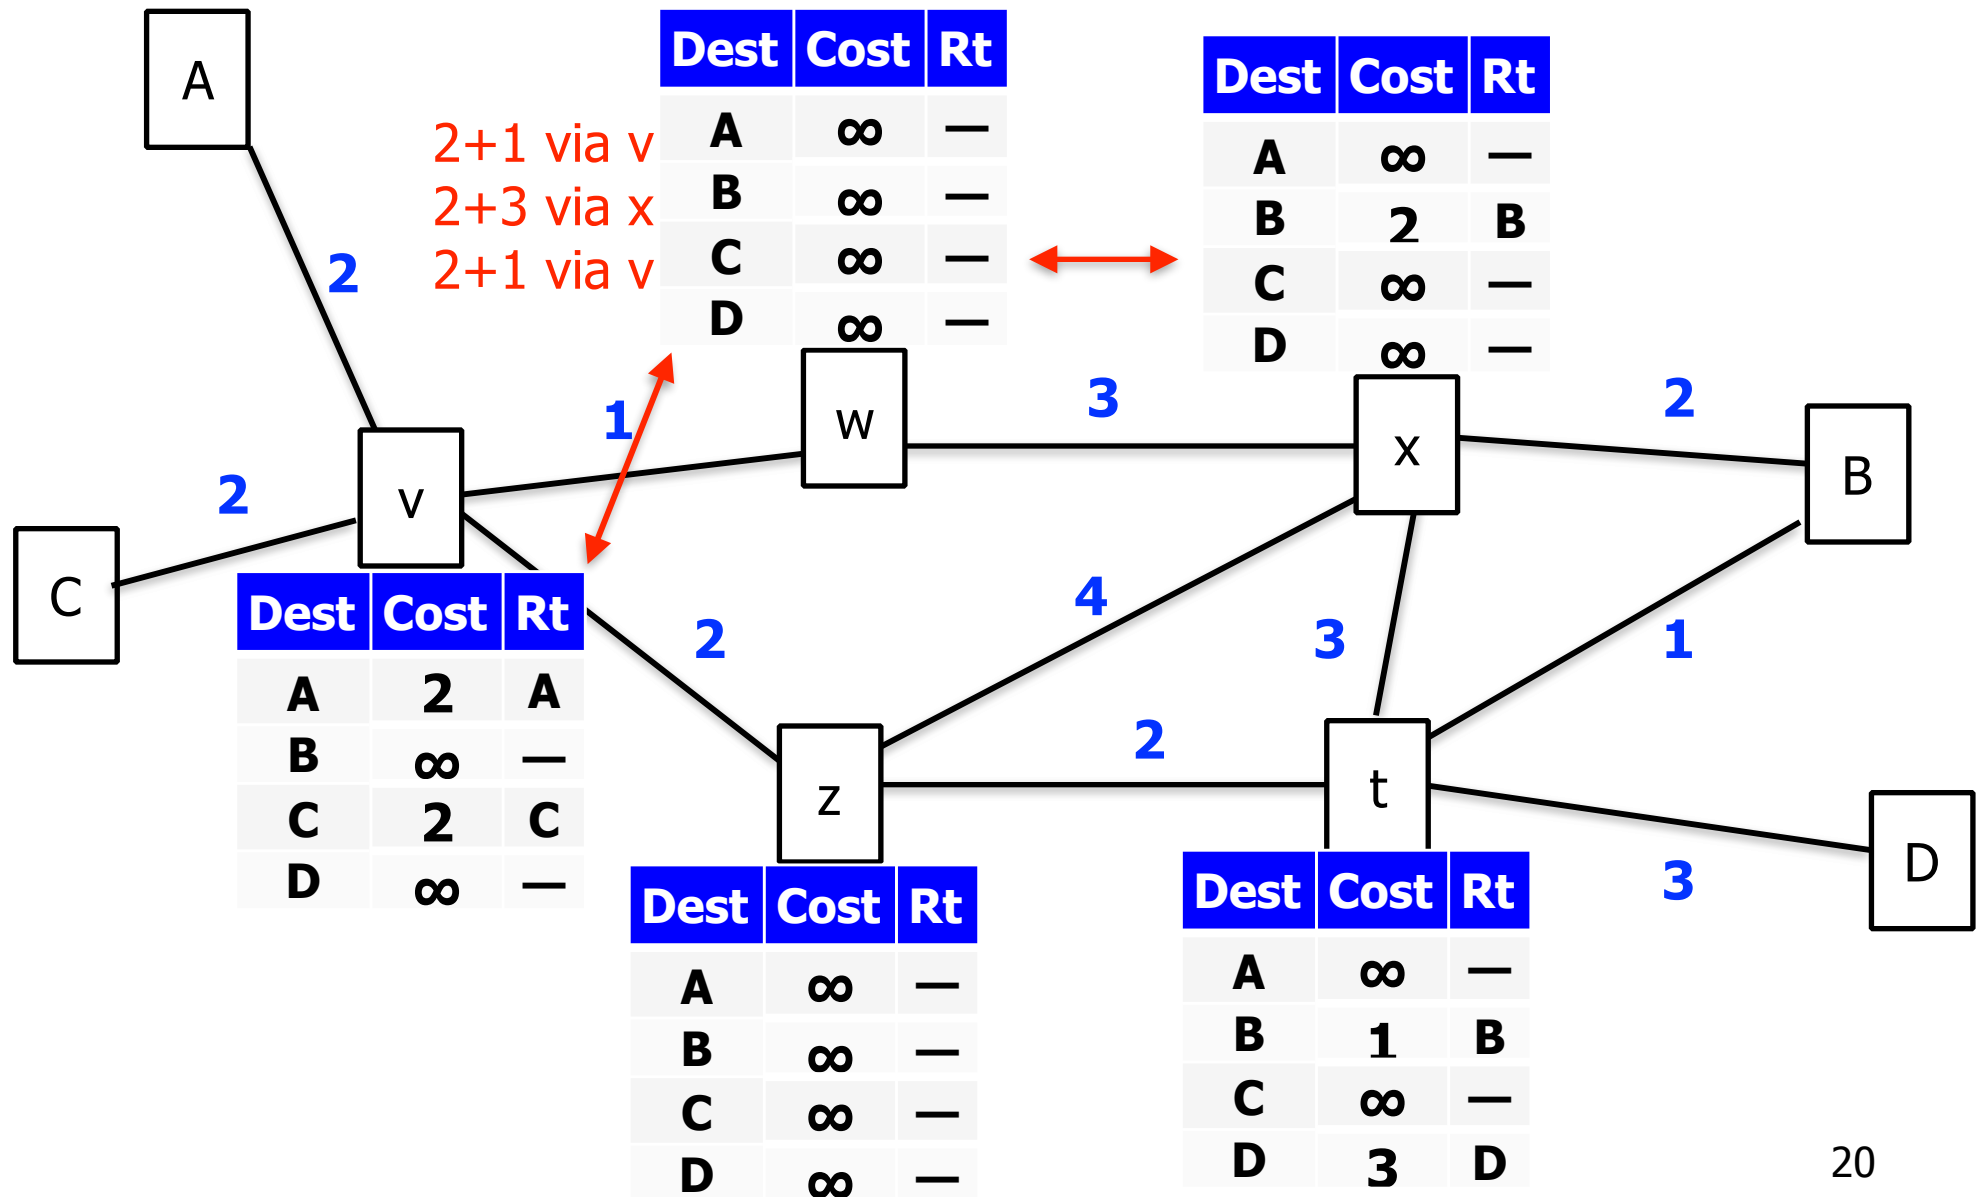
Distance Vector Protocols

- Routers
 - Know distances to immediate neighbors
 - Compute distance vector
 - How far to any destination from all known info
 - Transmit distance vectors to neighbors
 - Discover better (shorter) route? Update table
 - Now know more info, so repeat process

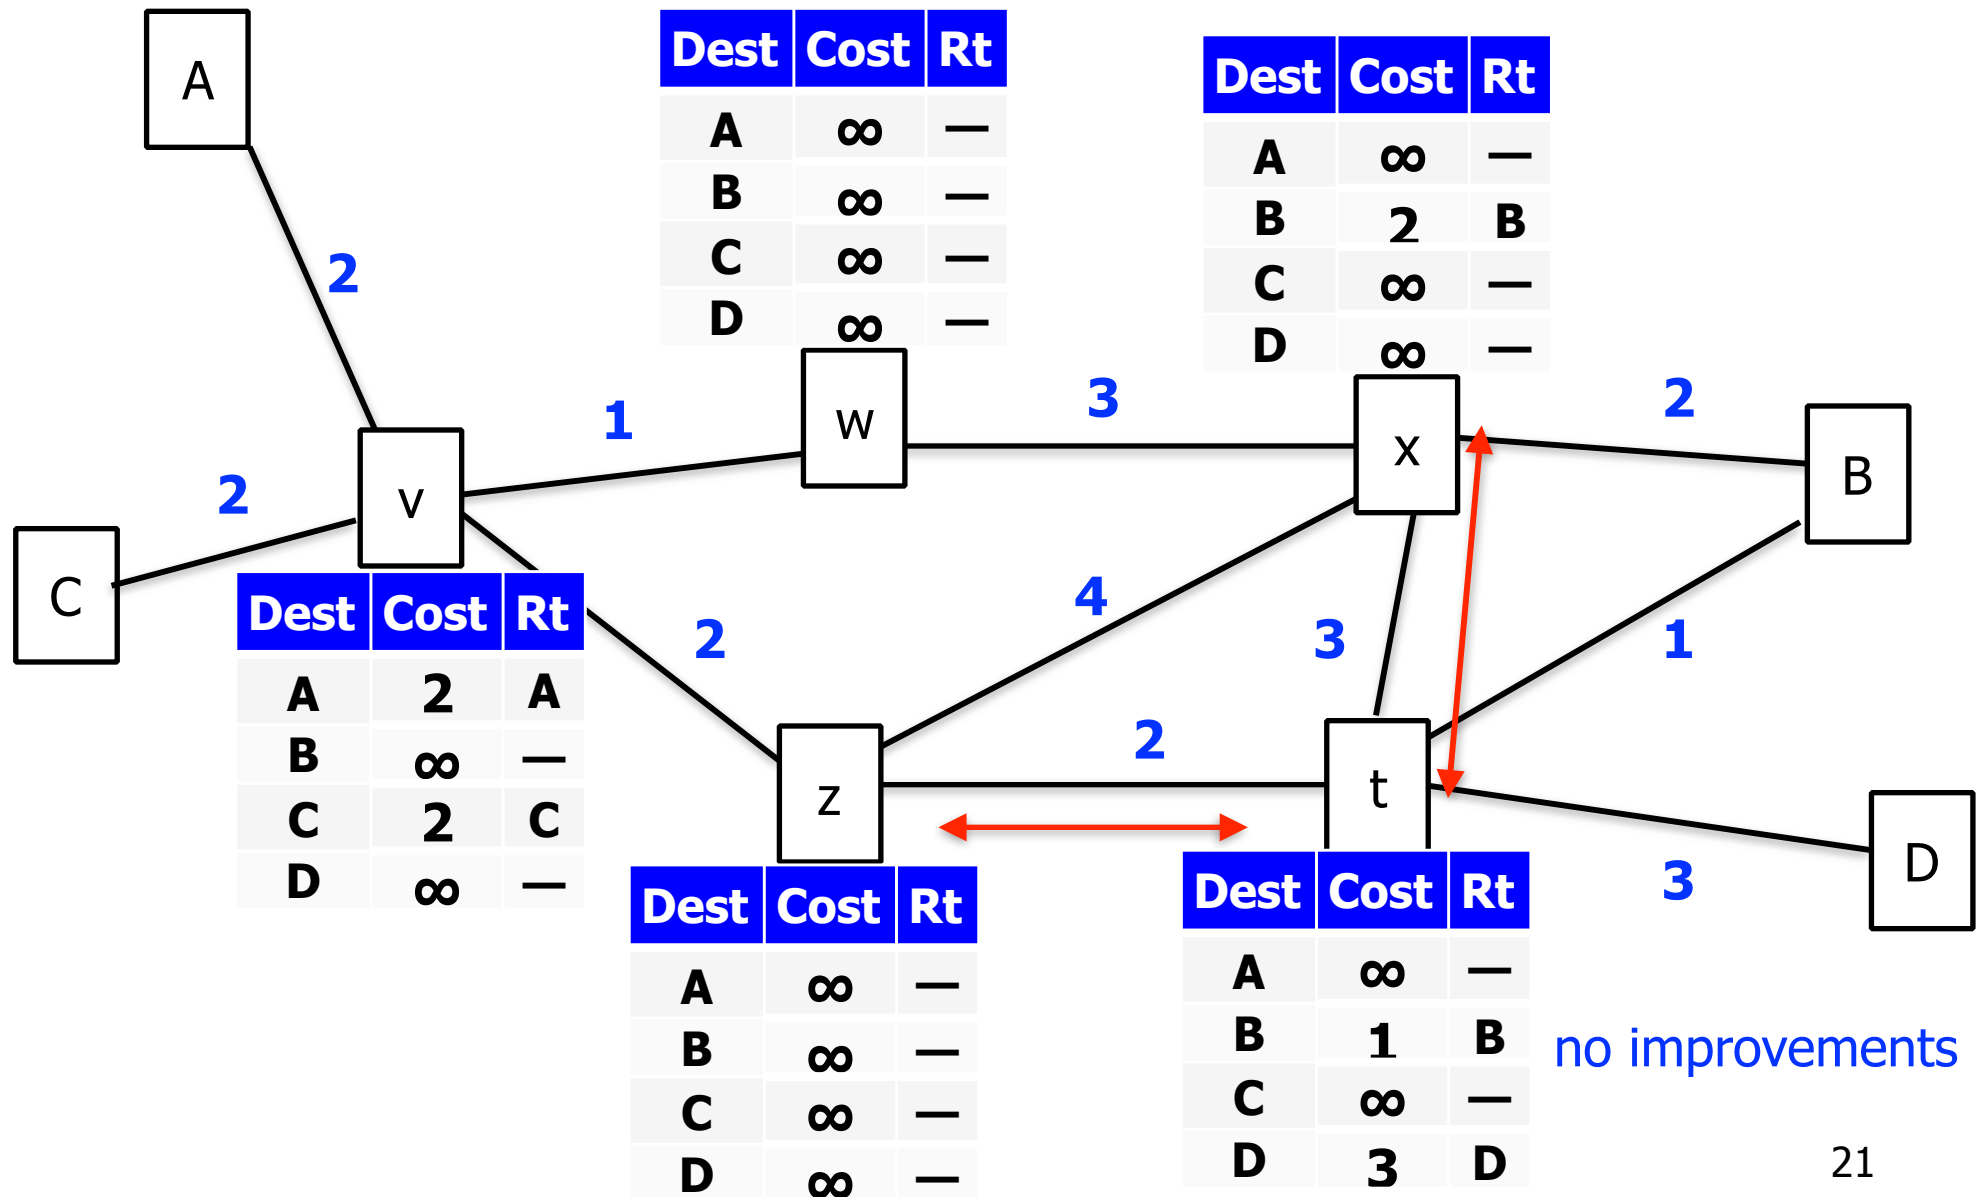
Distance Vector Routing



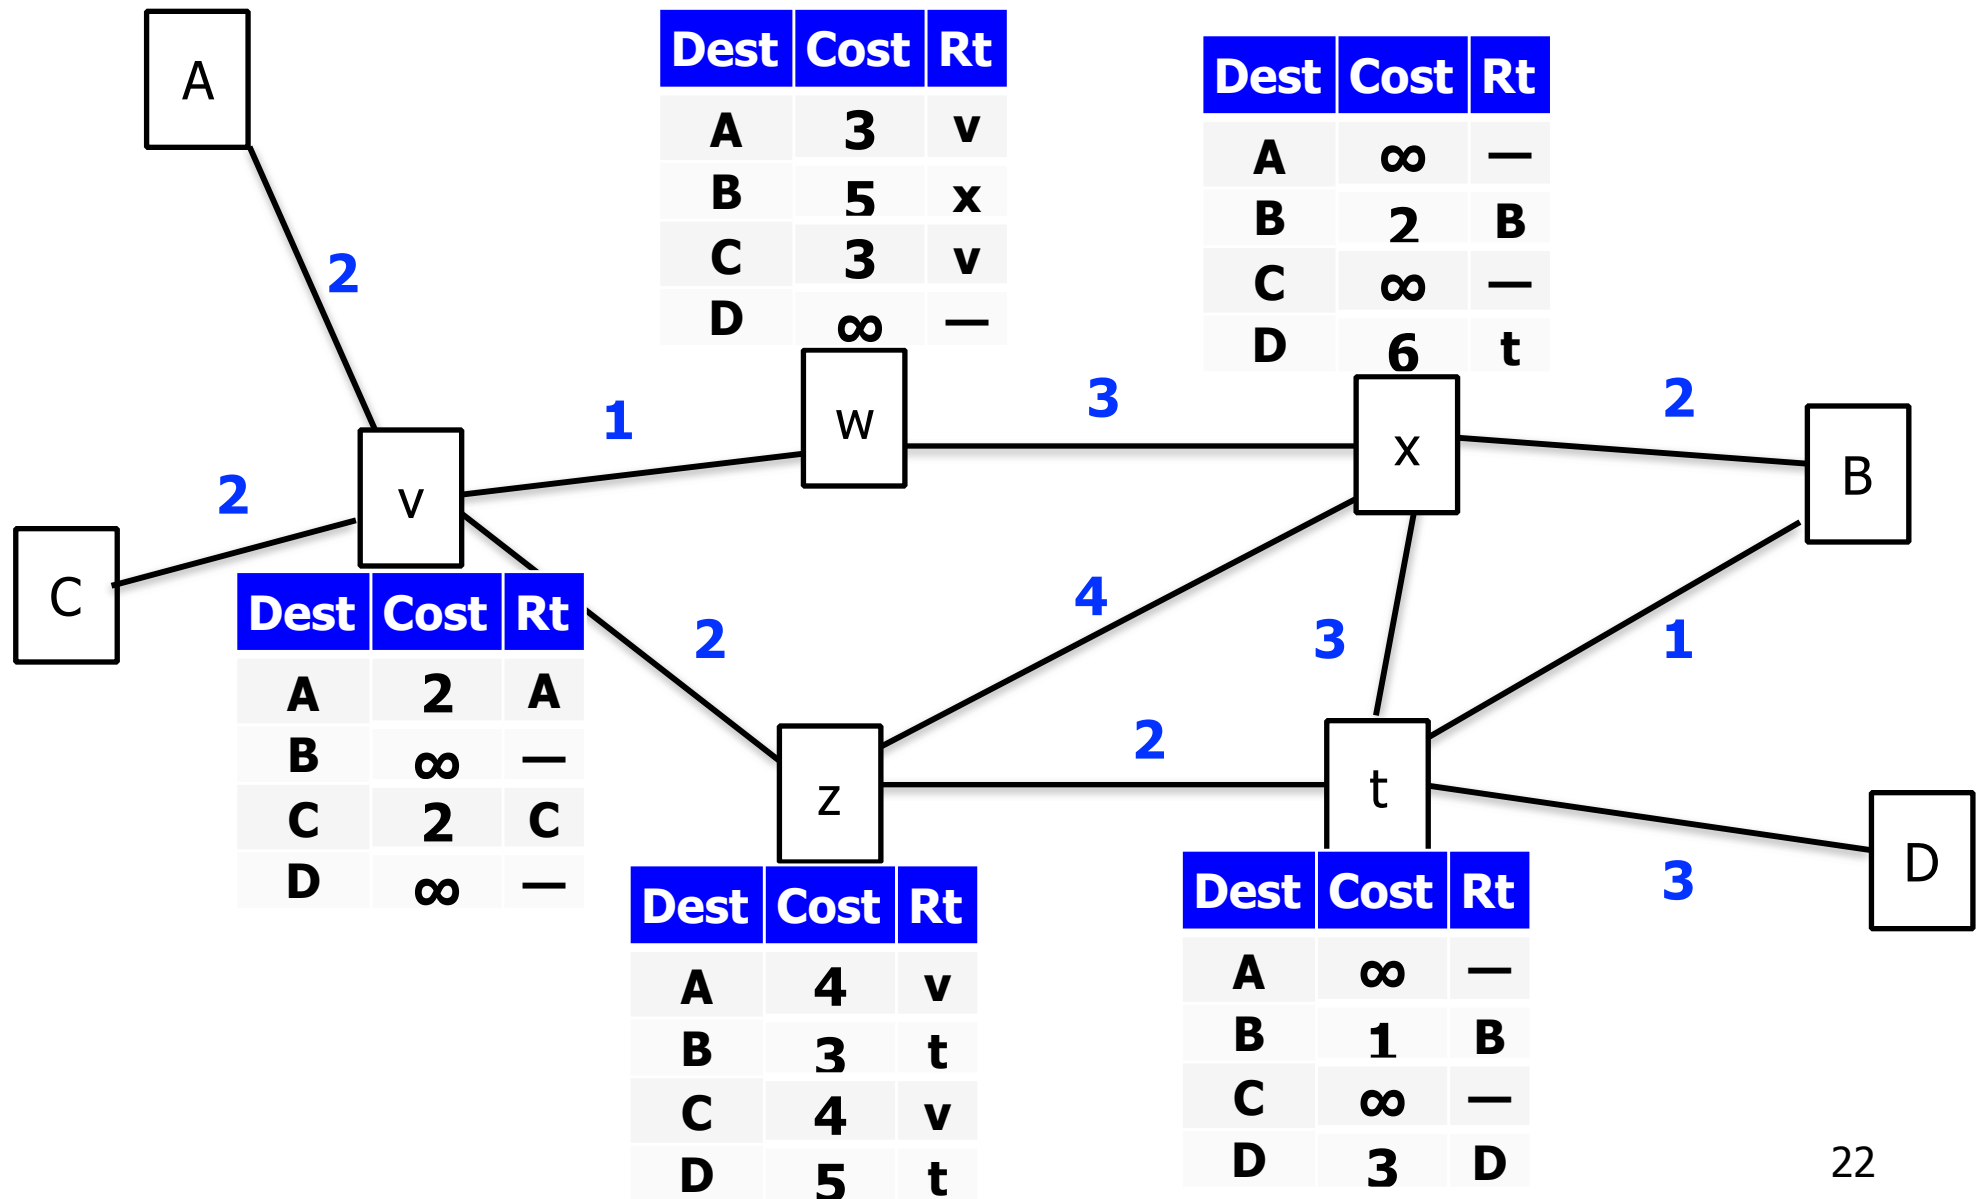
Distance Vector Routing



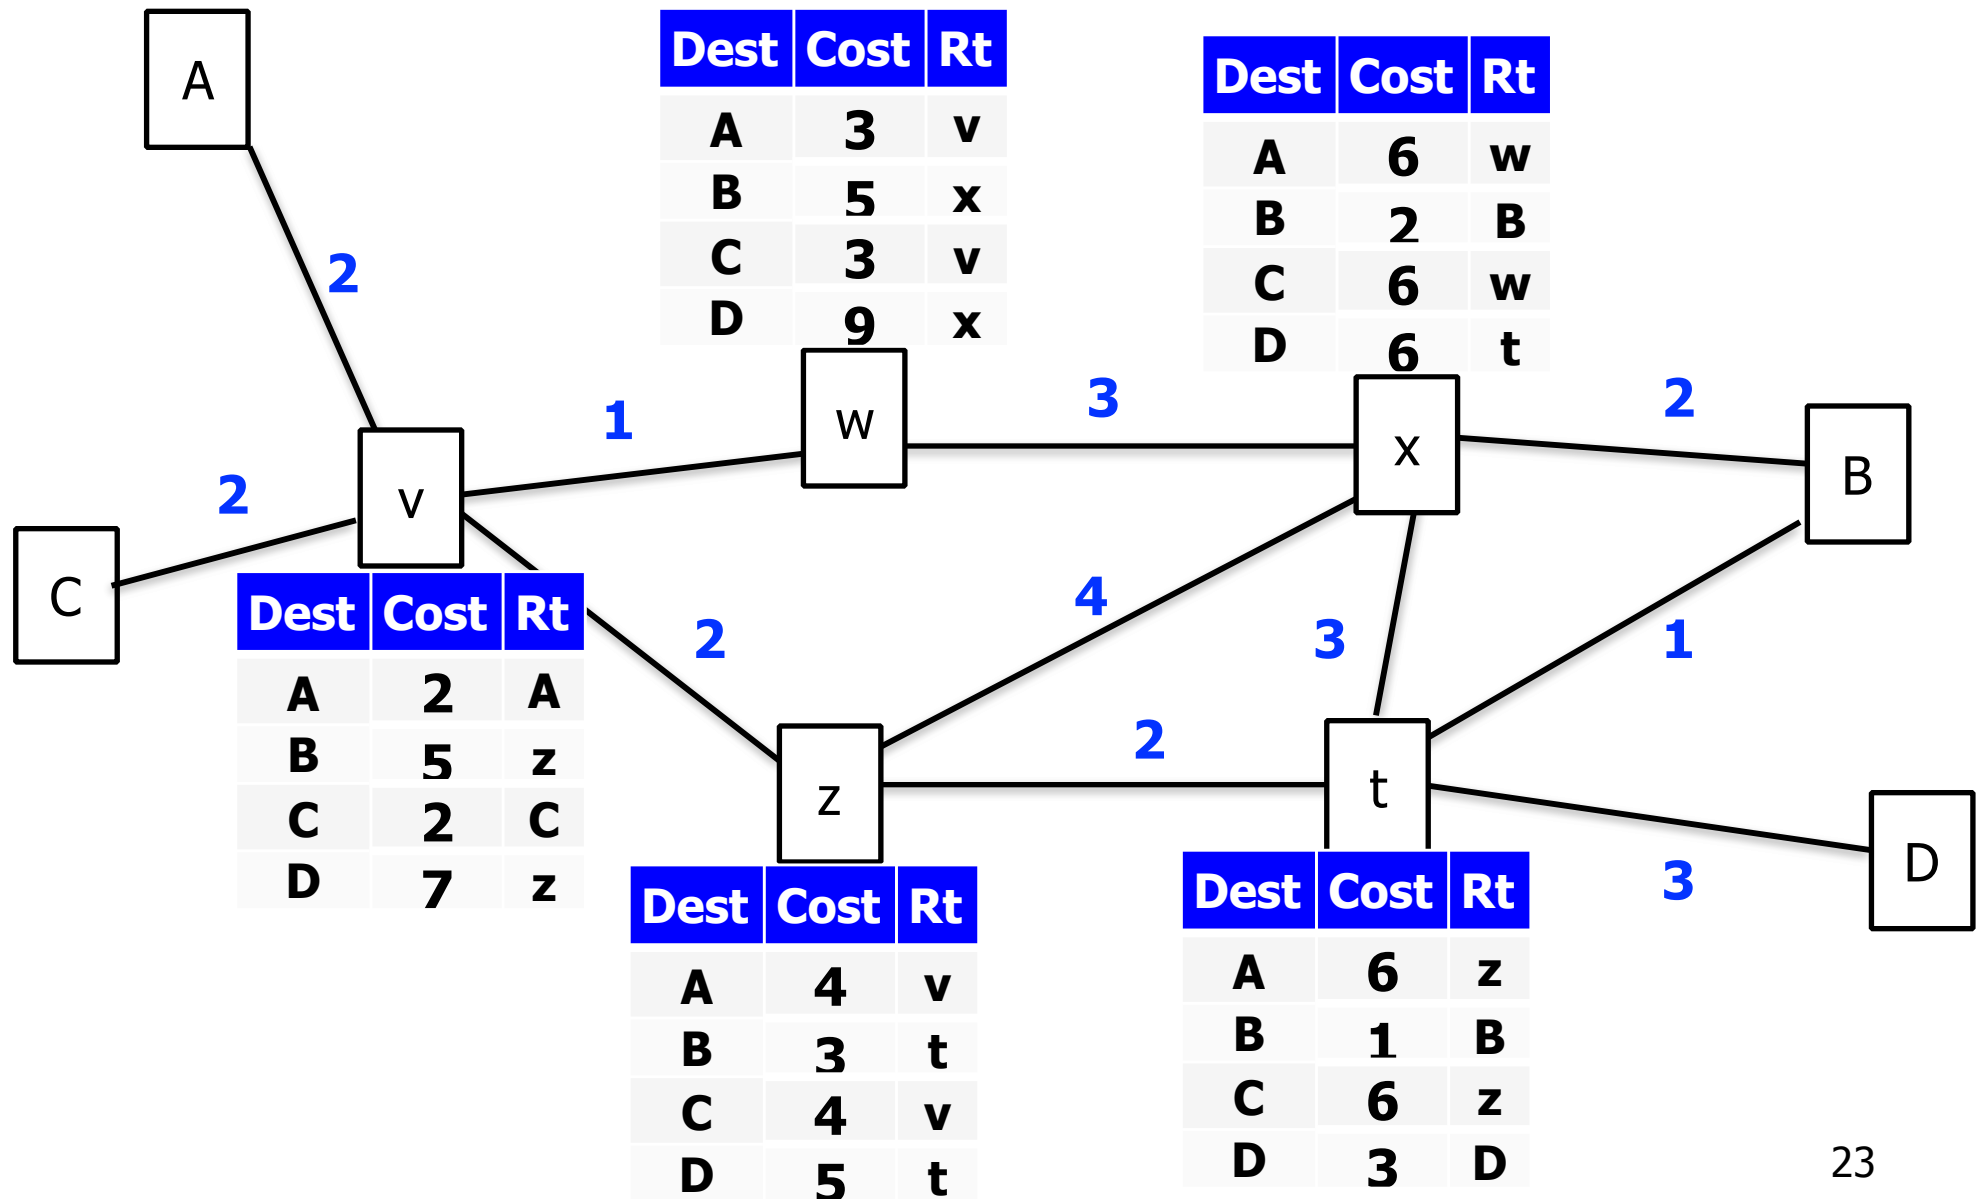
Distance Vector Routing



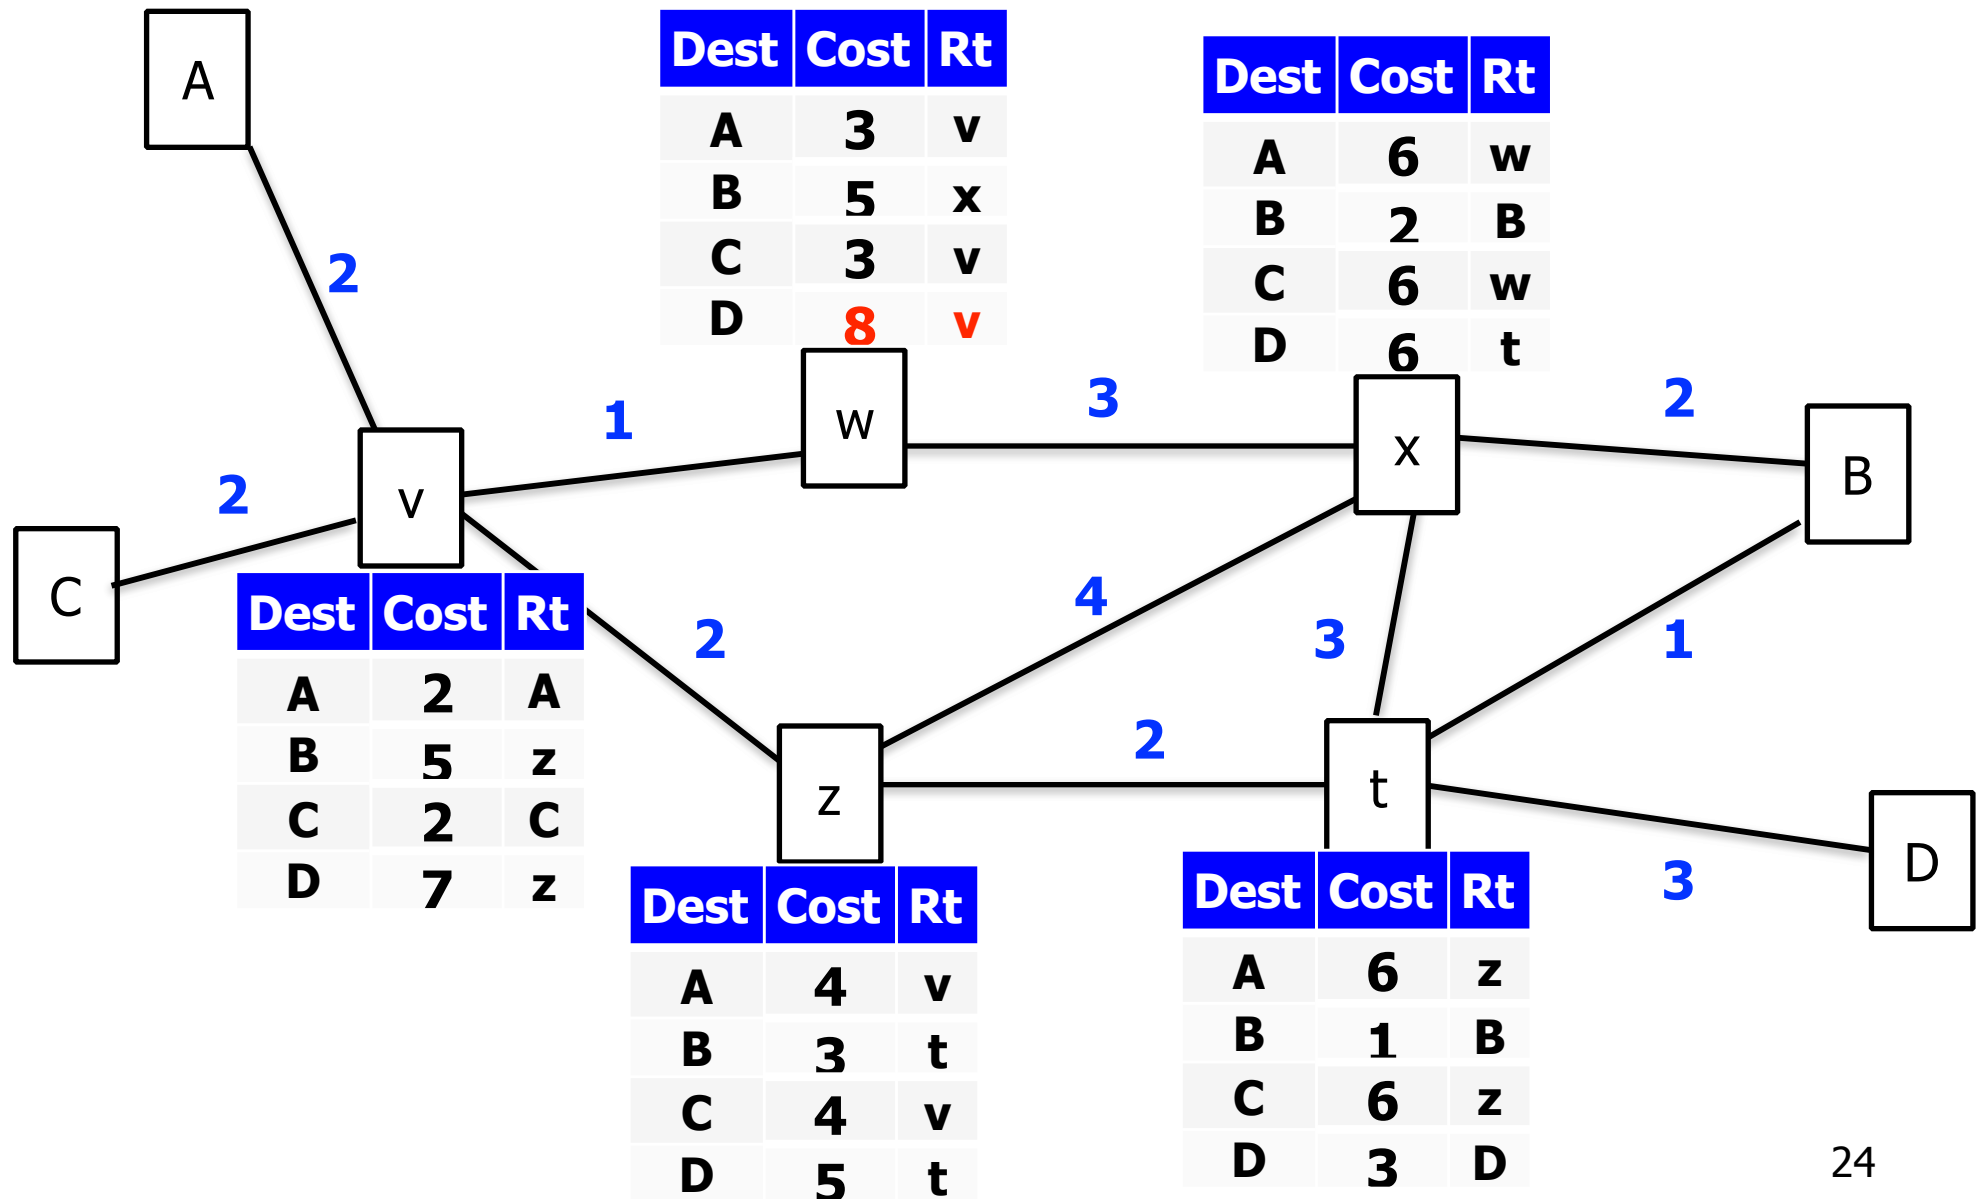
Distance Vector Routing



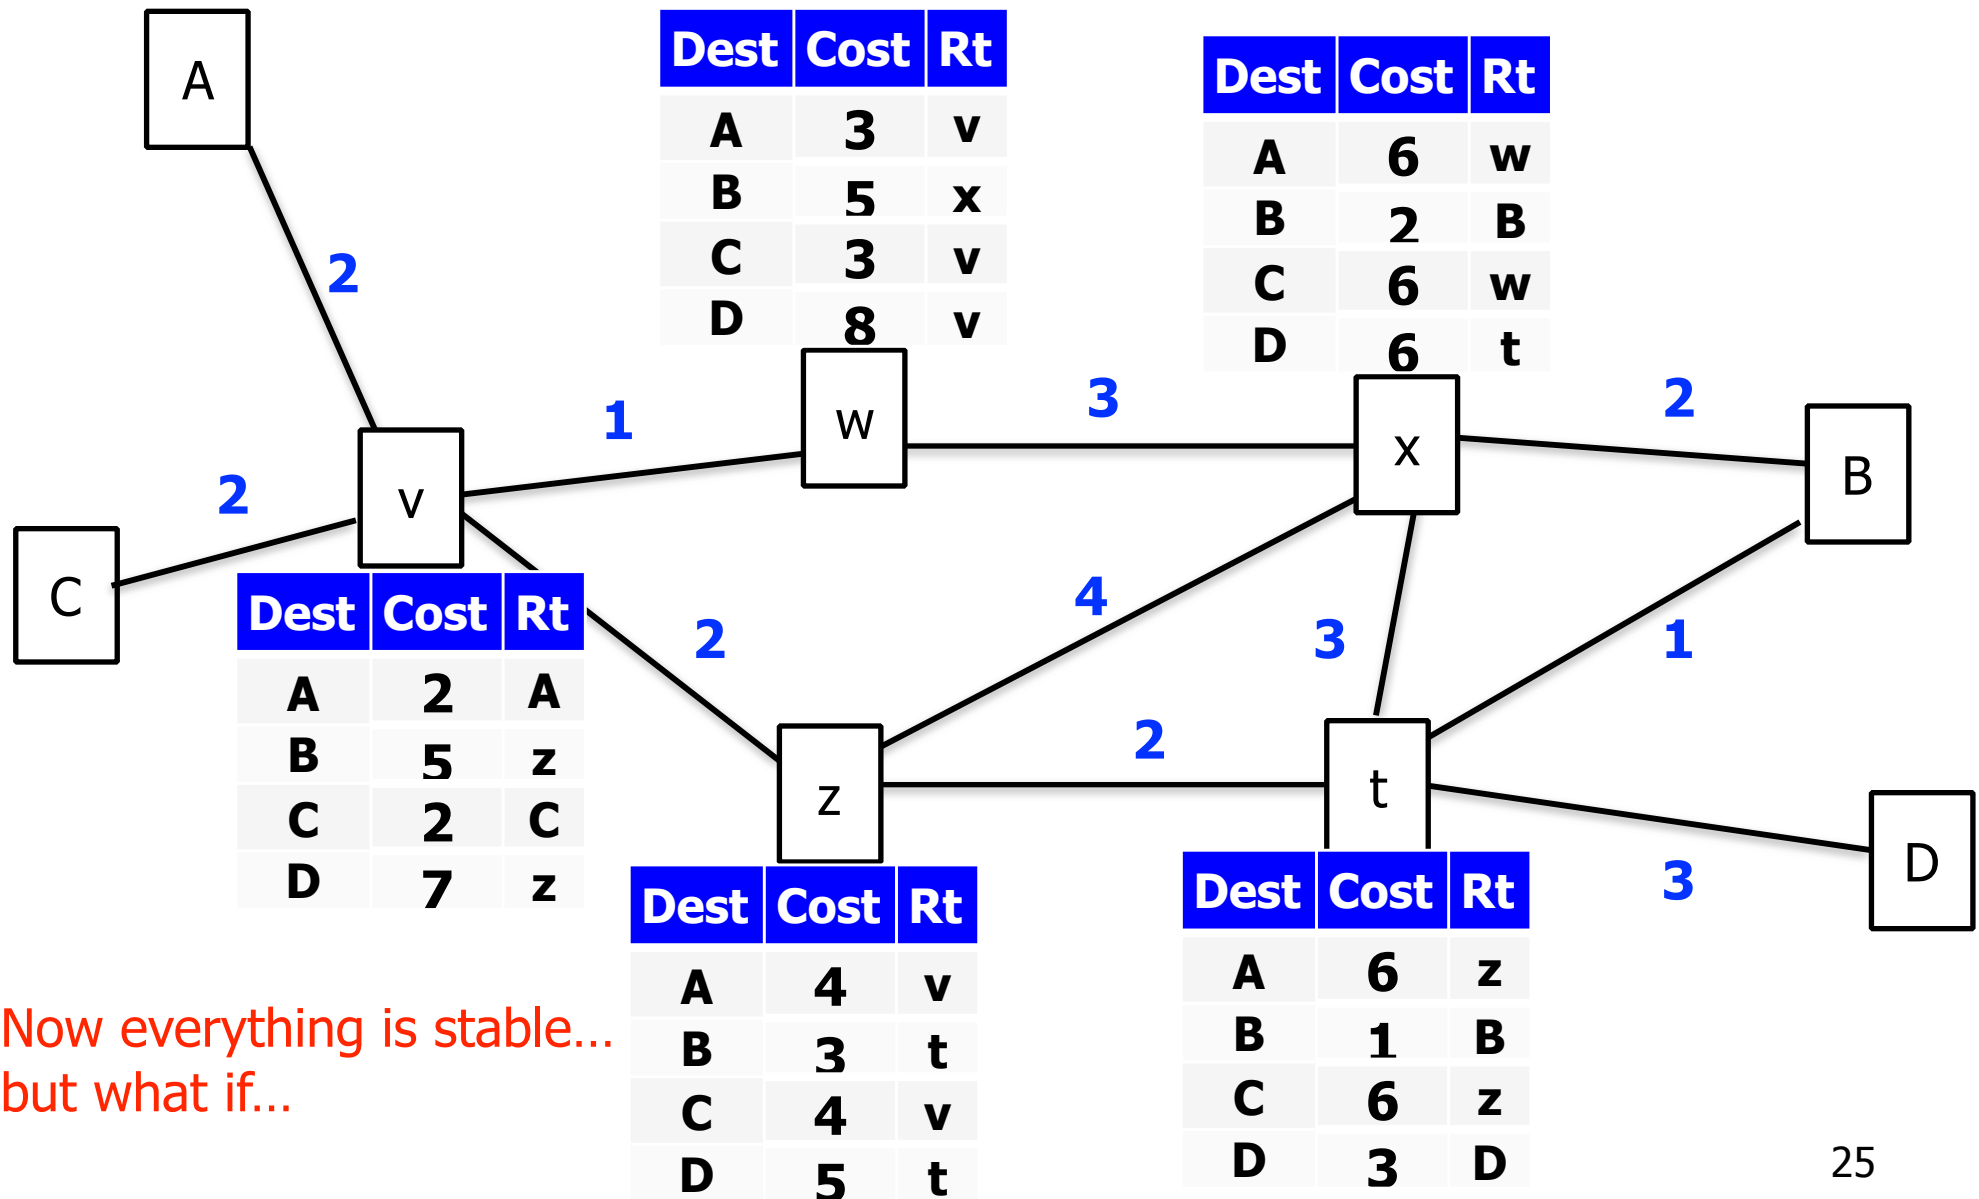
Distance Vector Routing



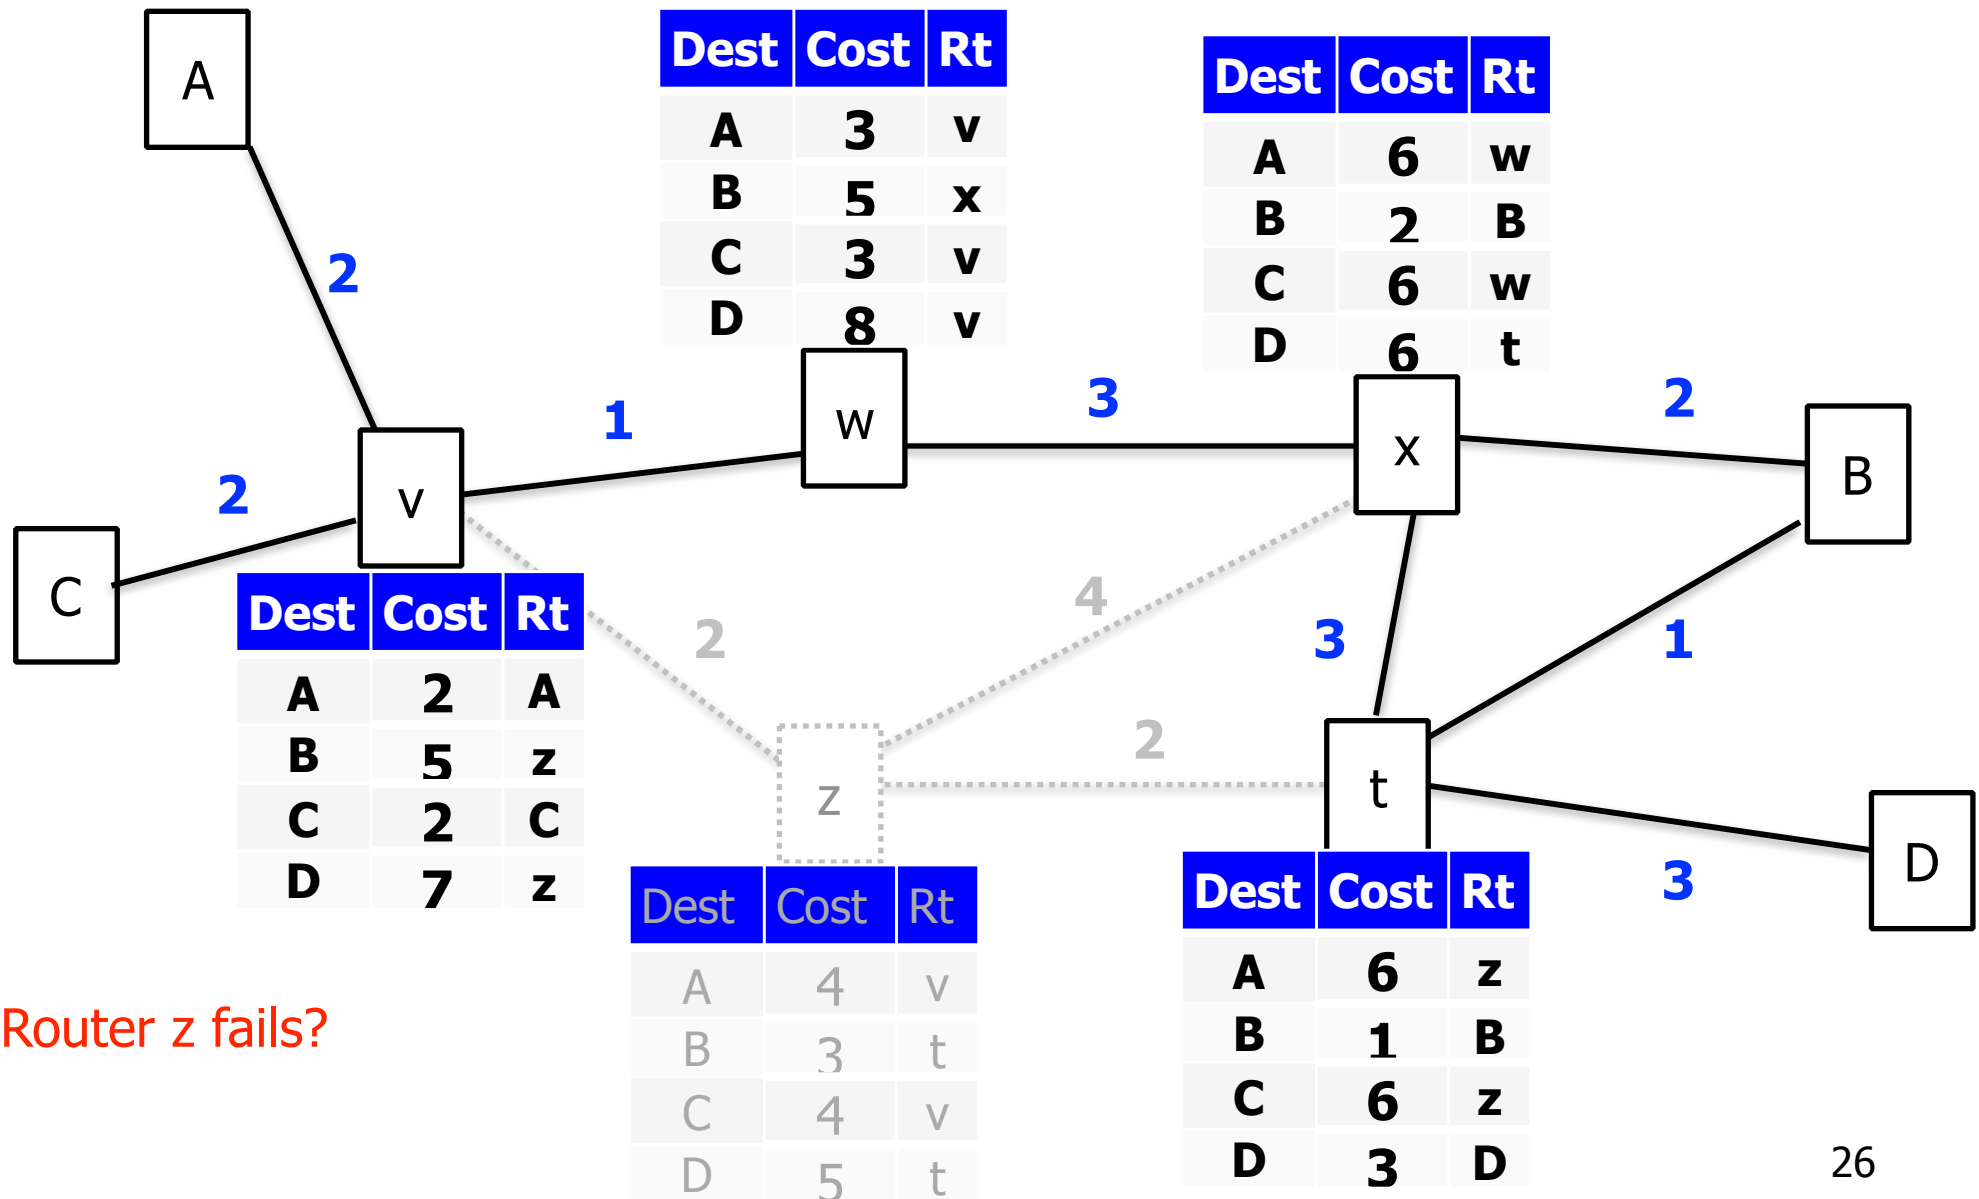
Distance Vector Routing



Distance Vector Routing

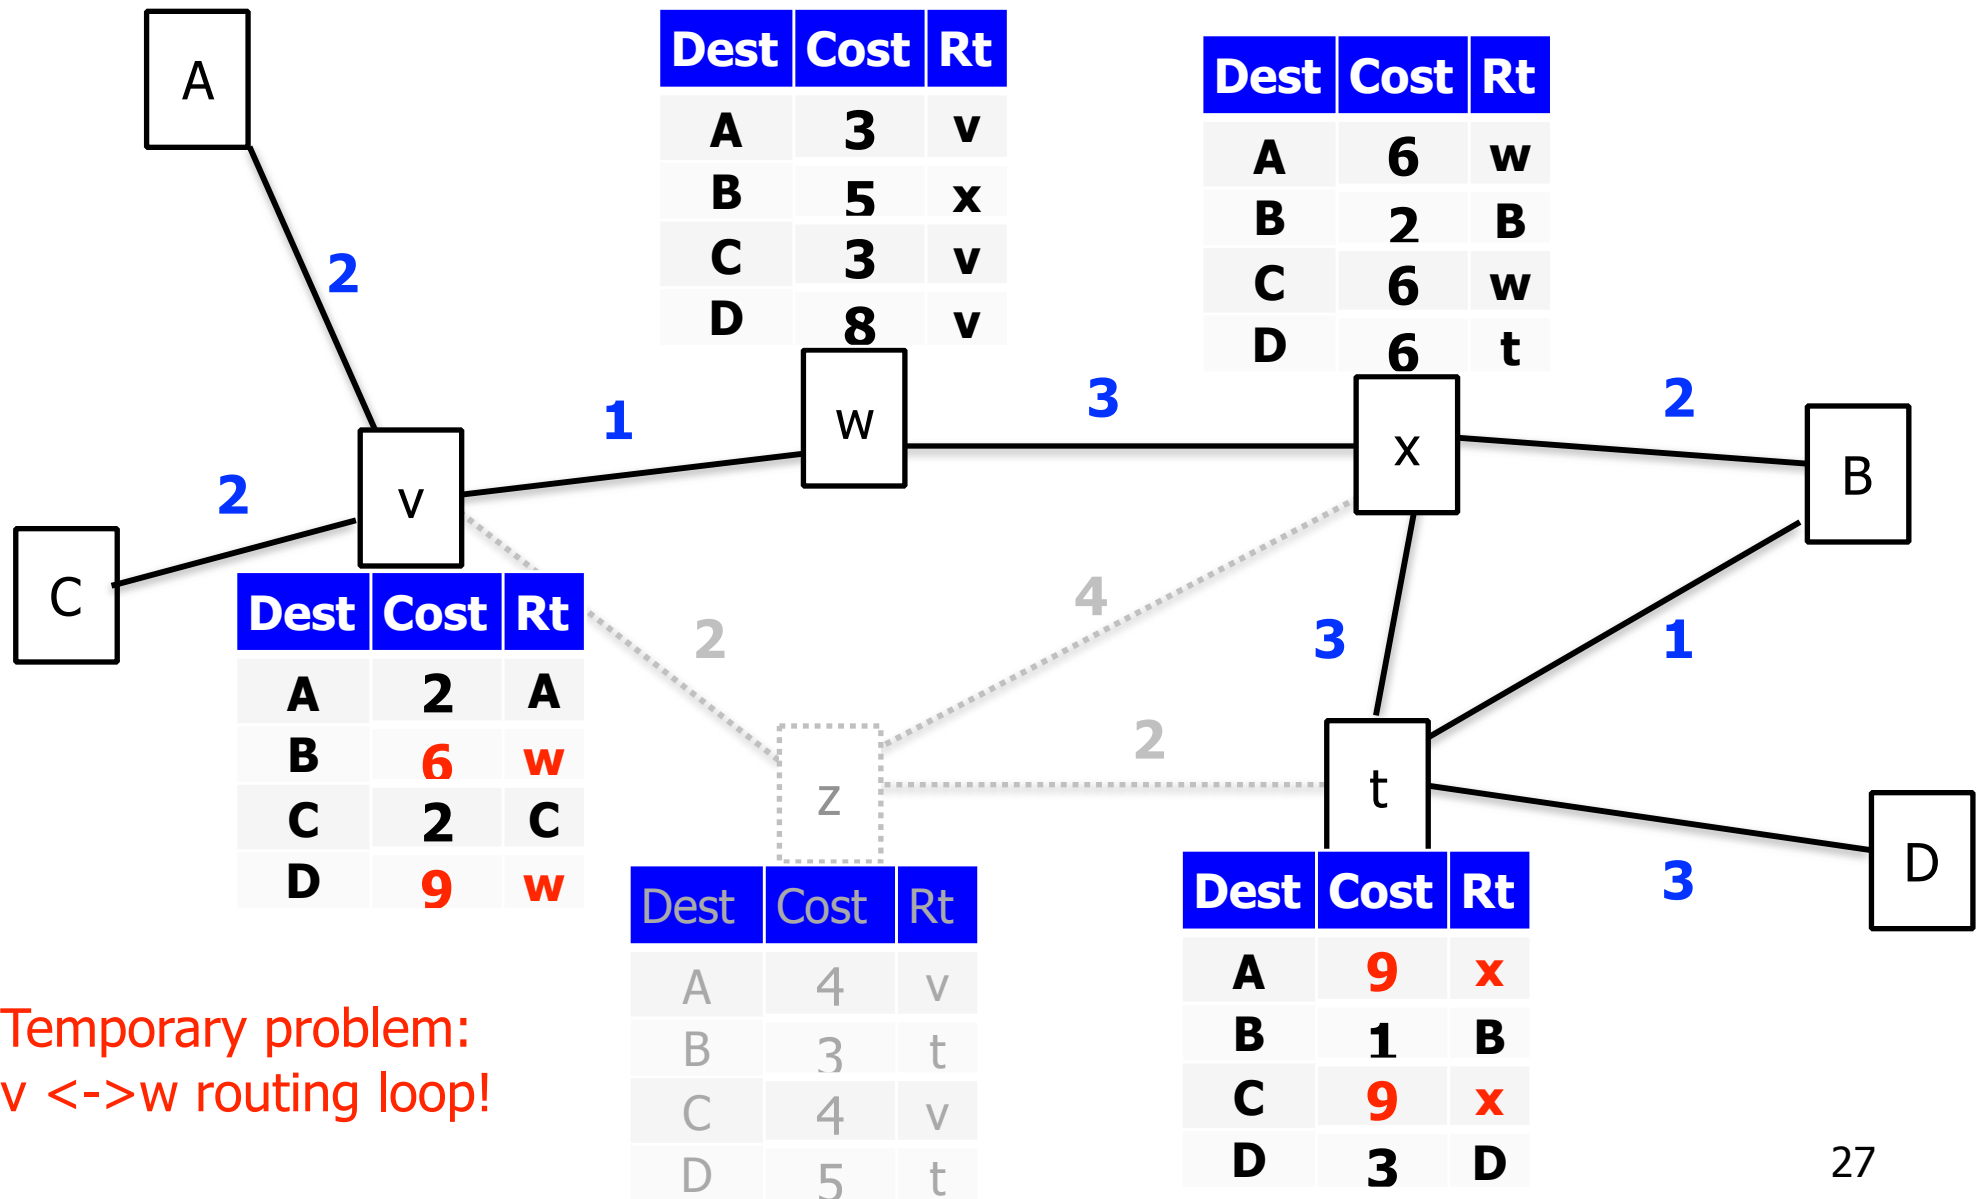


Distance Vector Routing

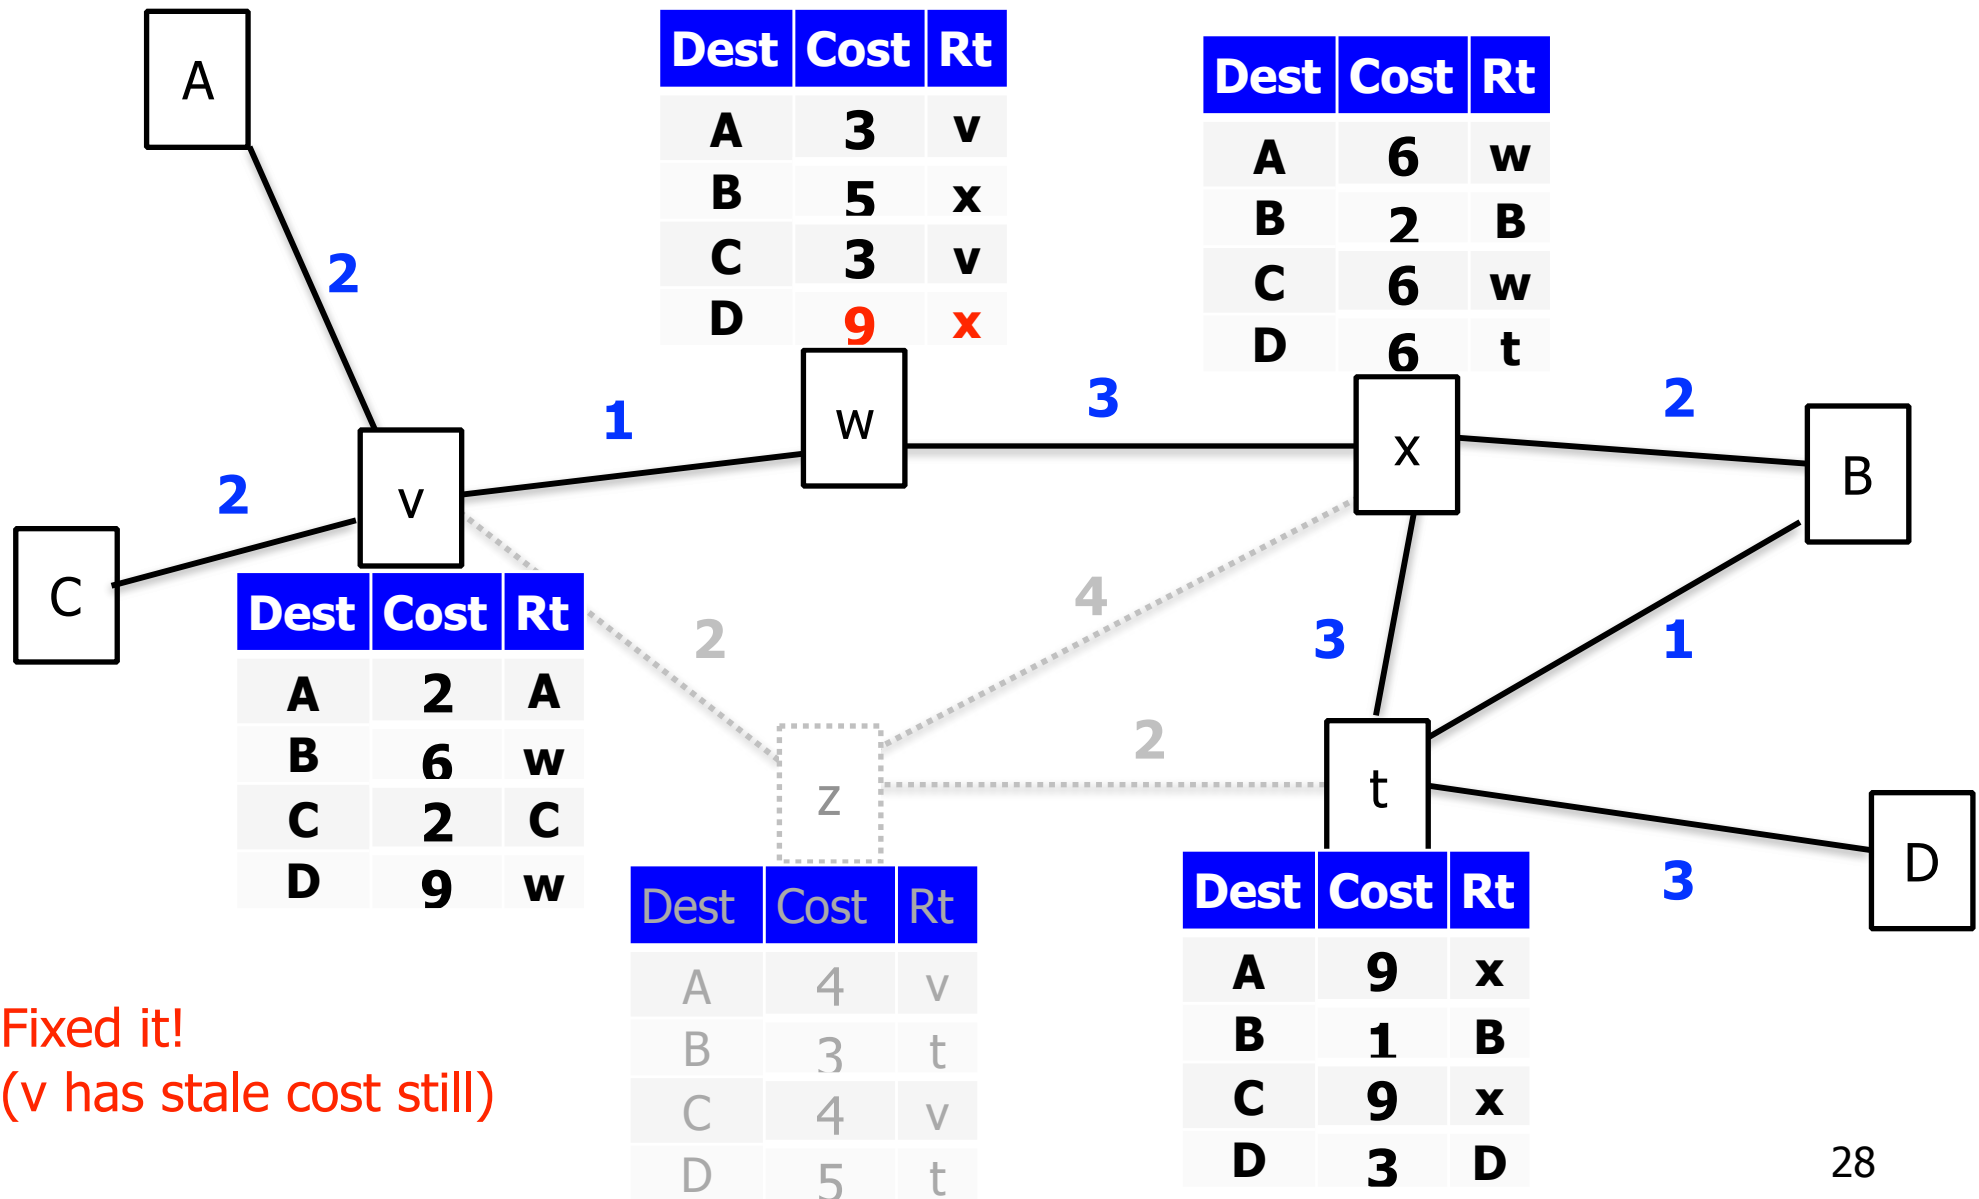


Router z fails?

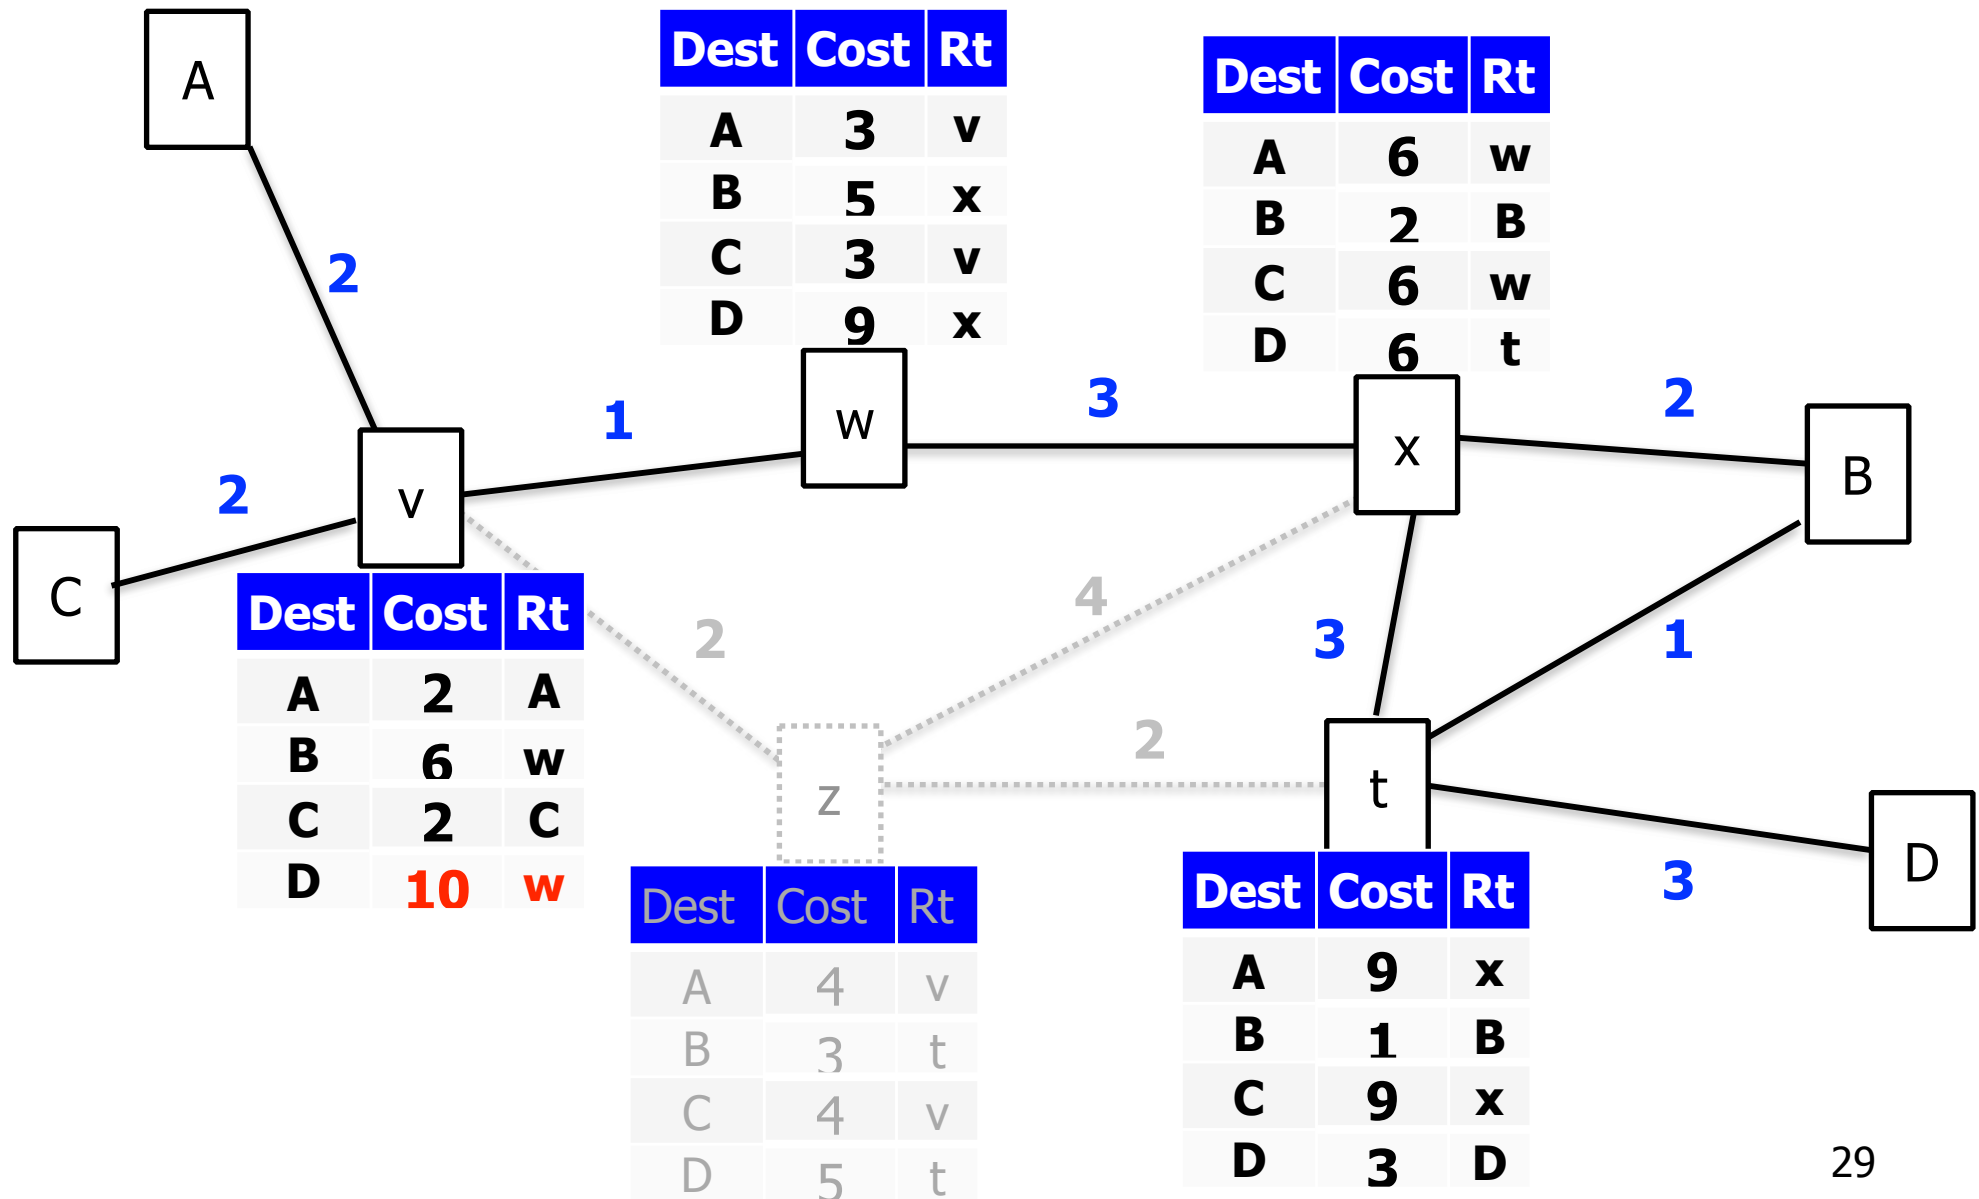
Distance Vector Routing



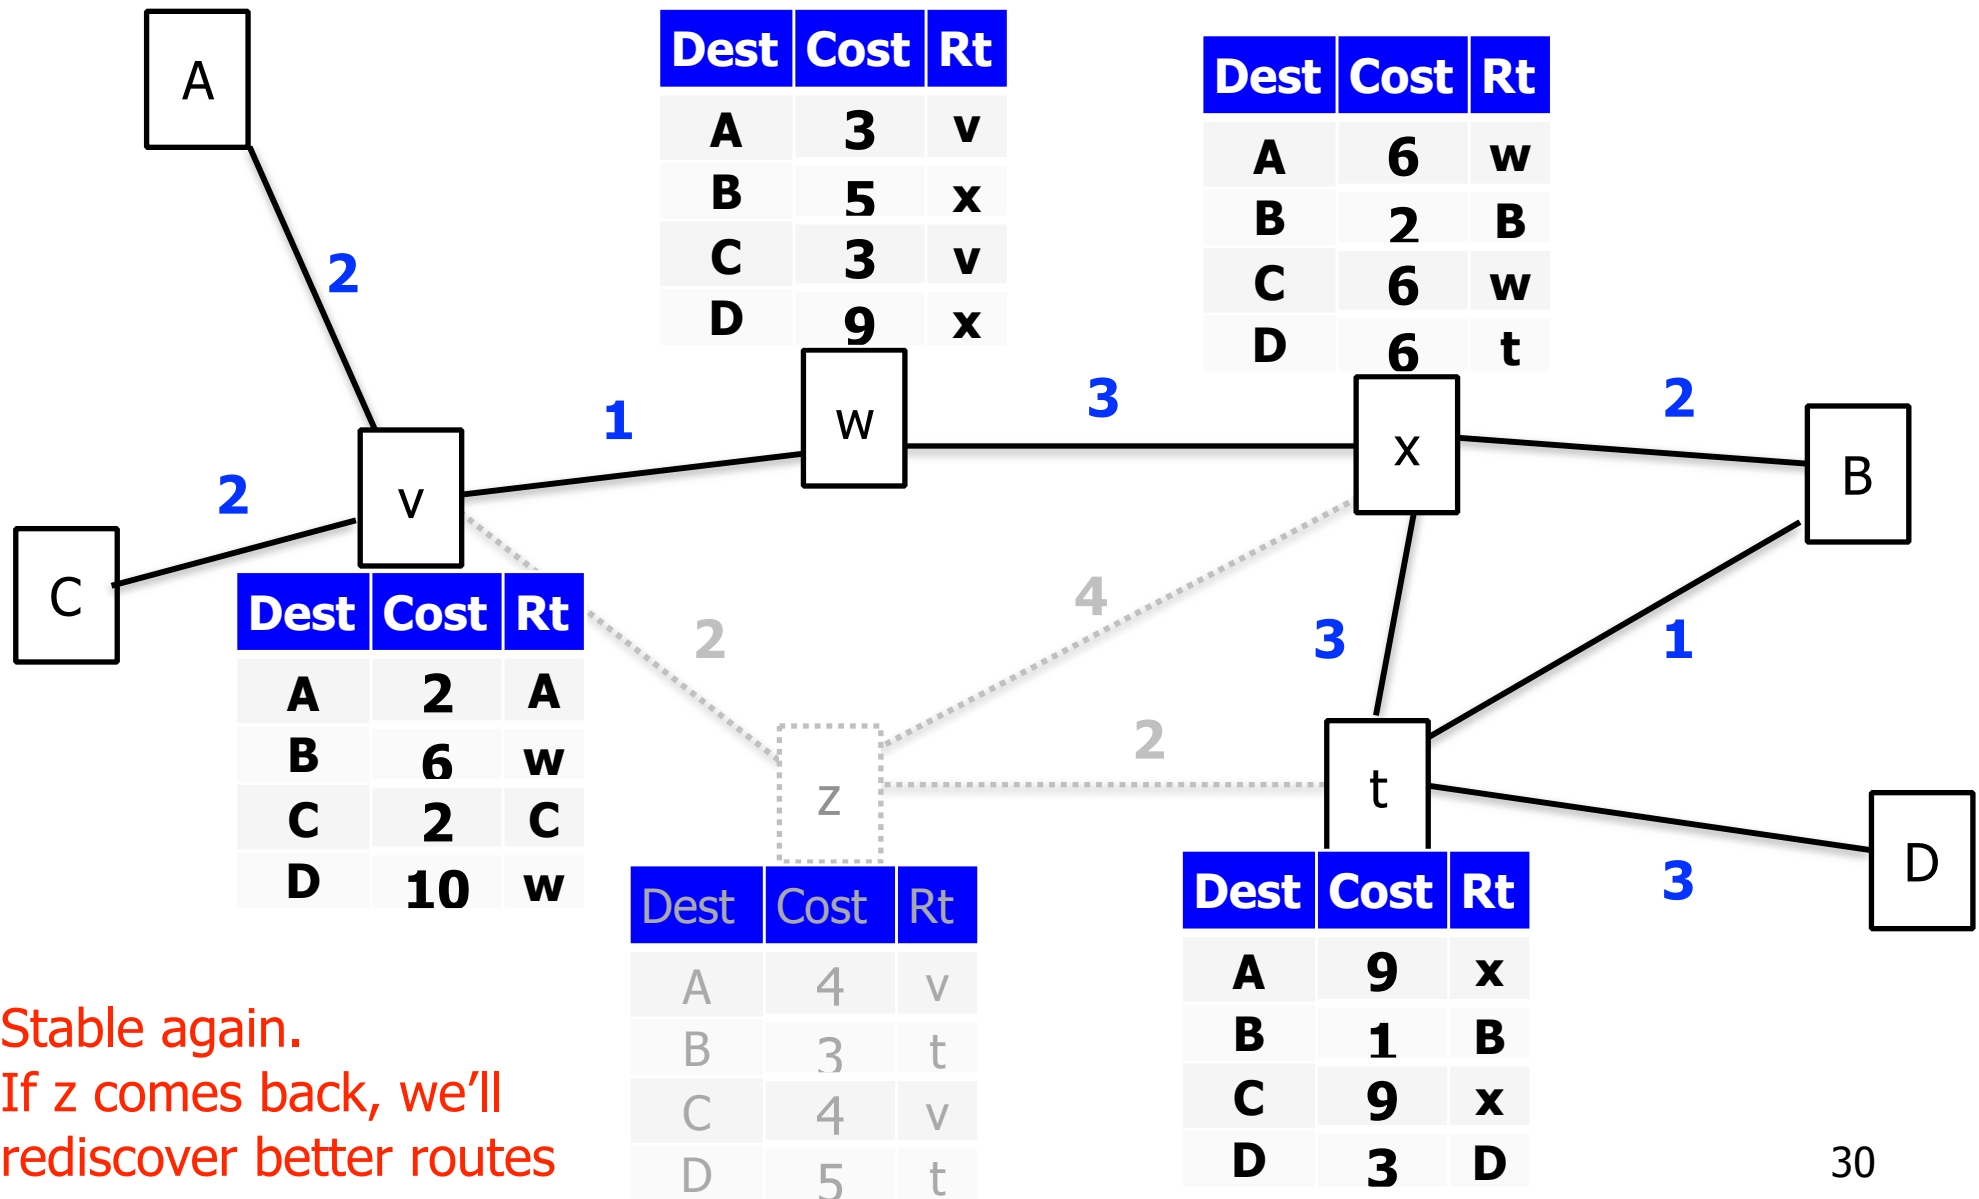
Distance Vector Routing



Distance Vector Routing

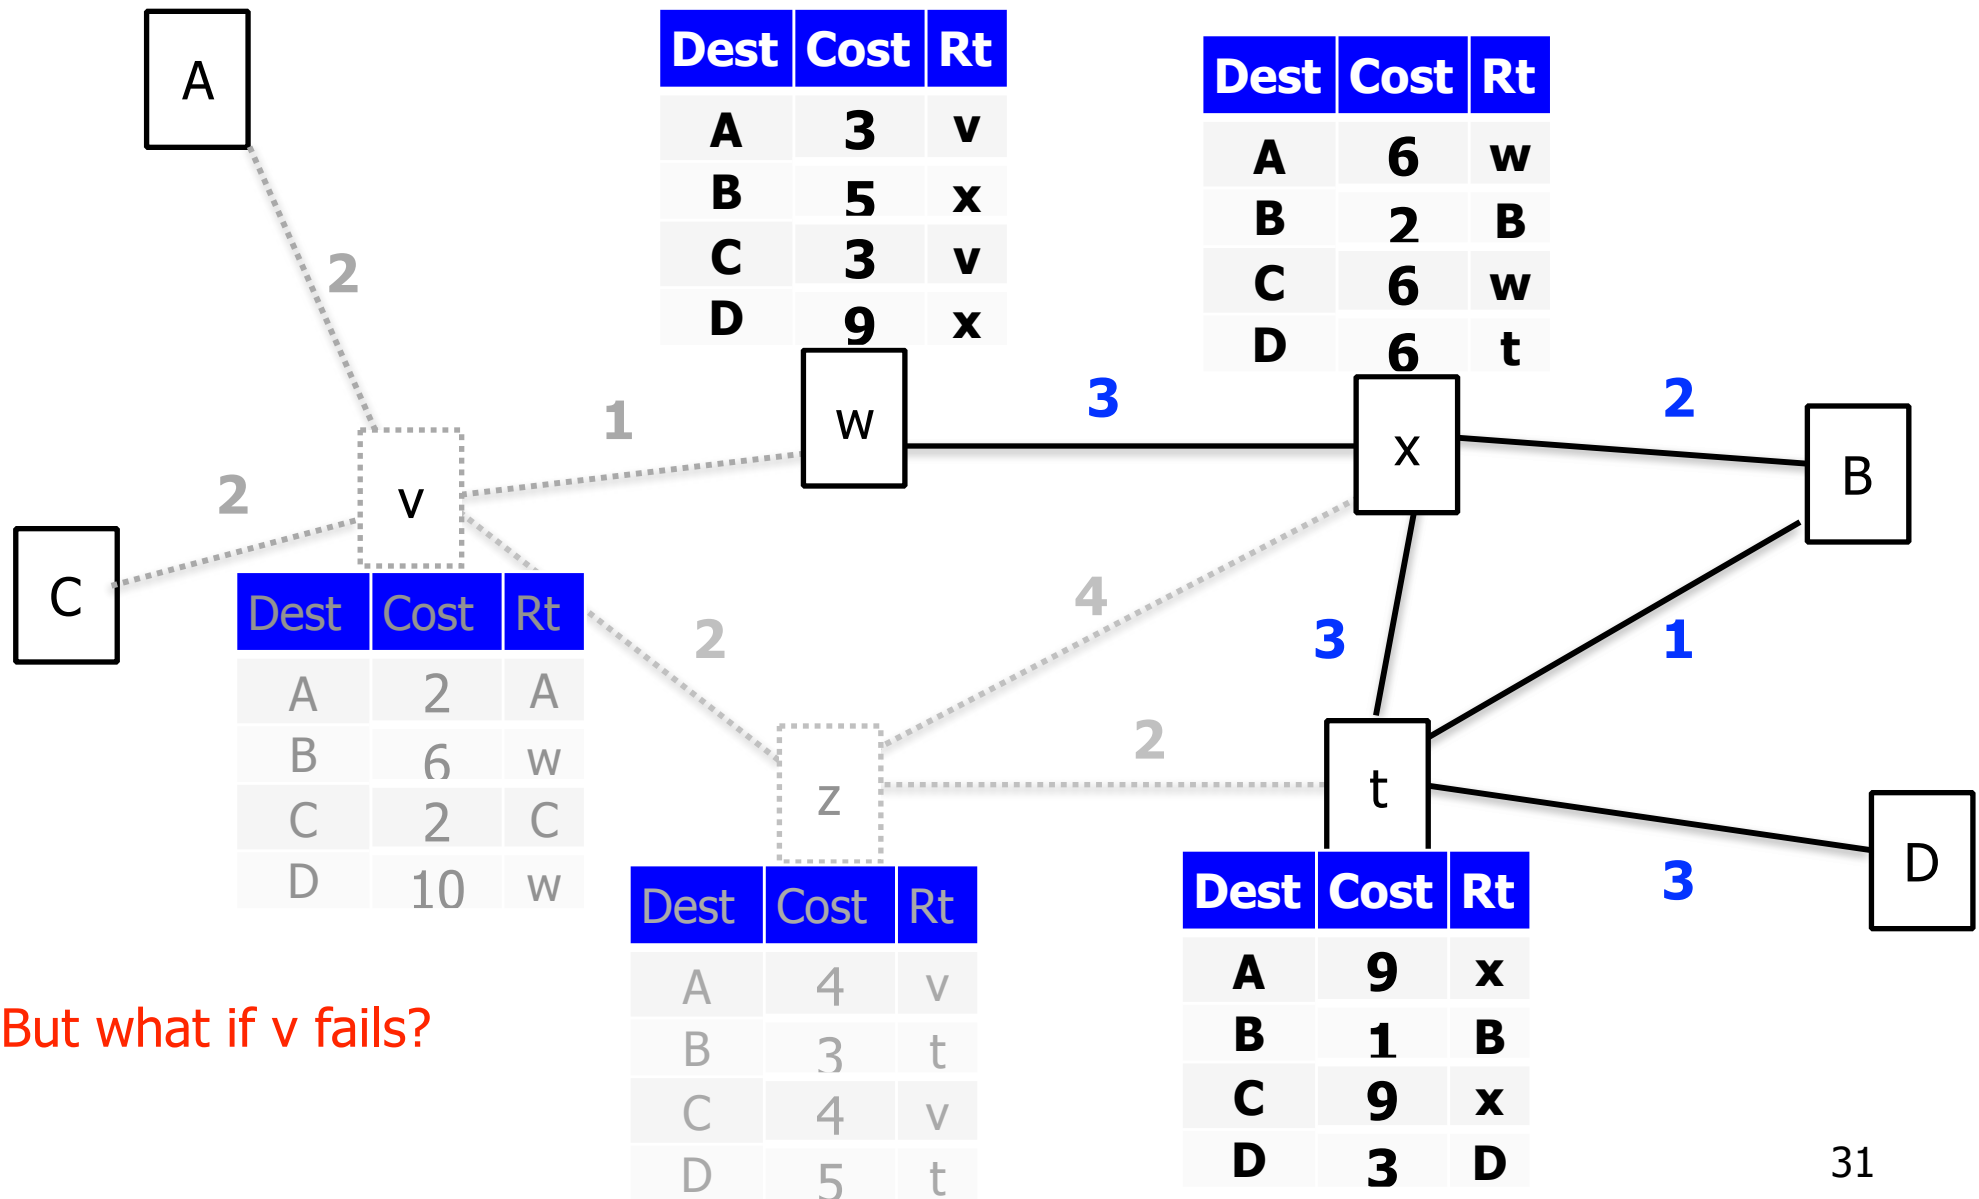


Distance Vector Routing

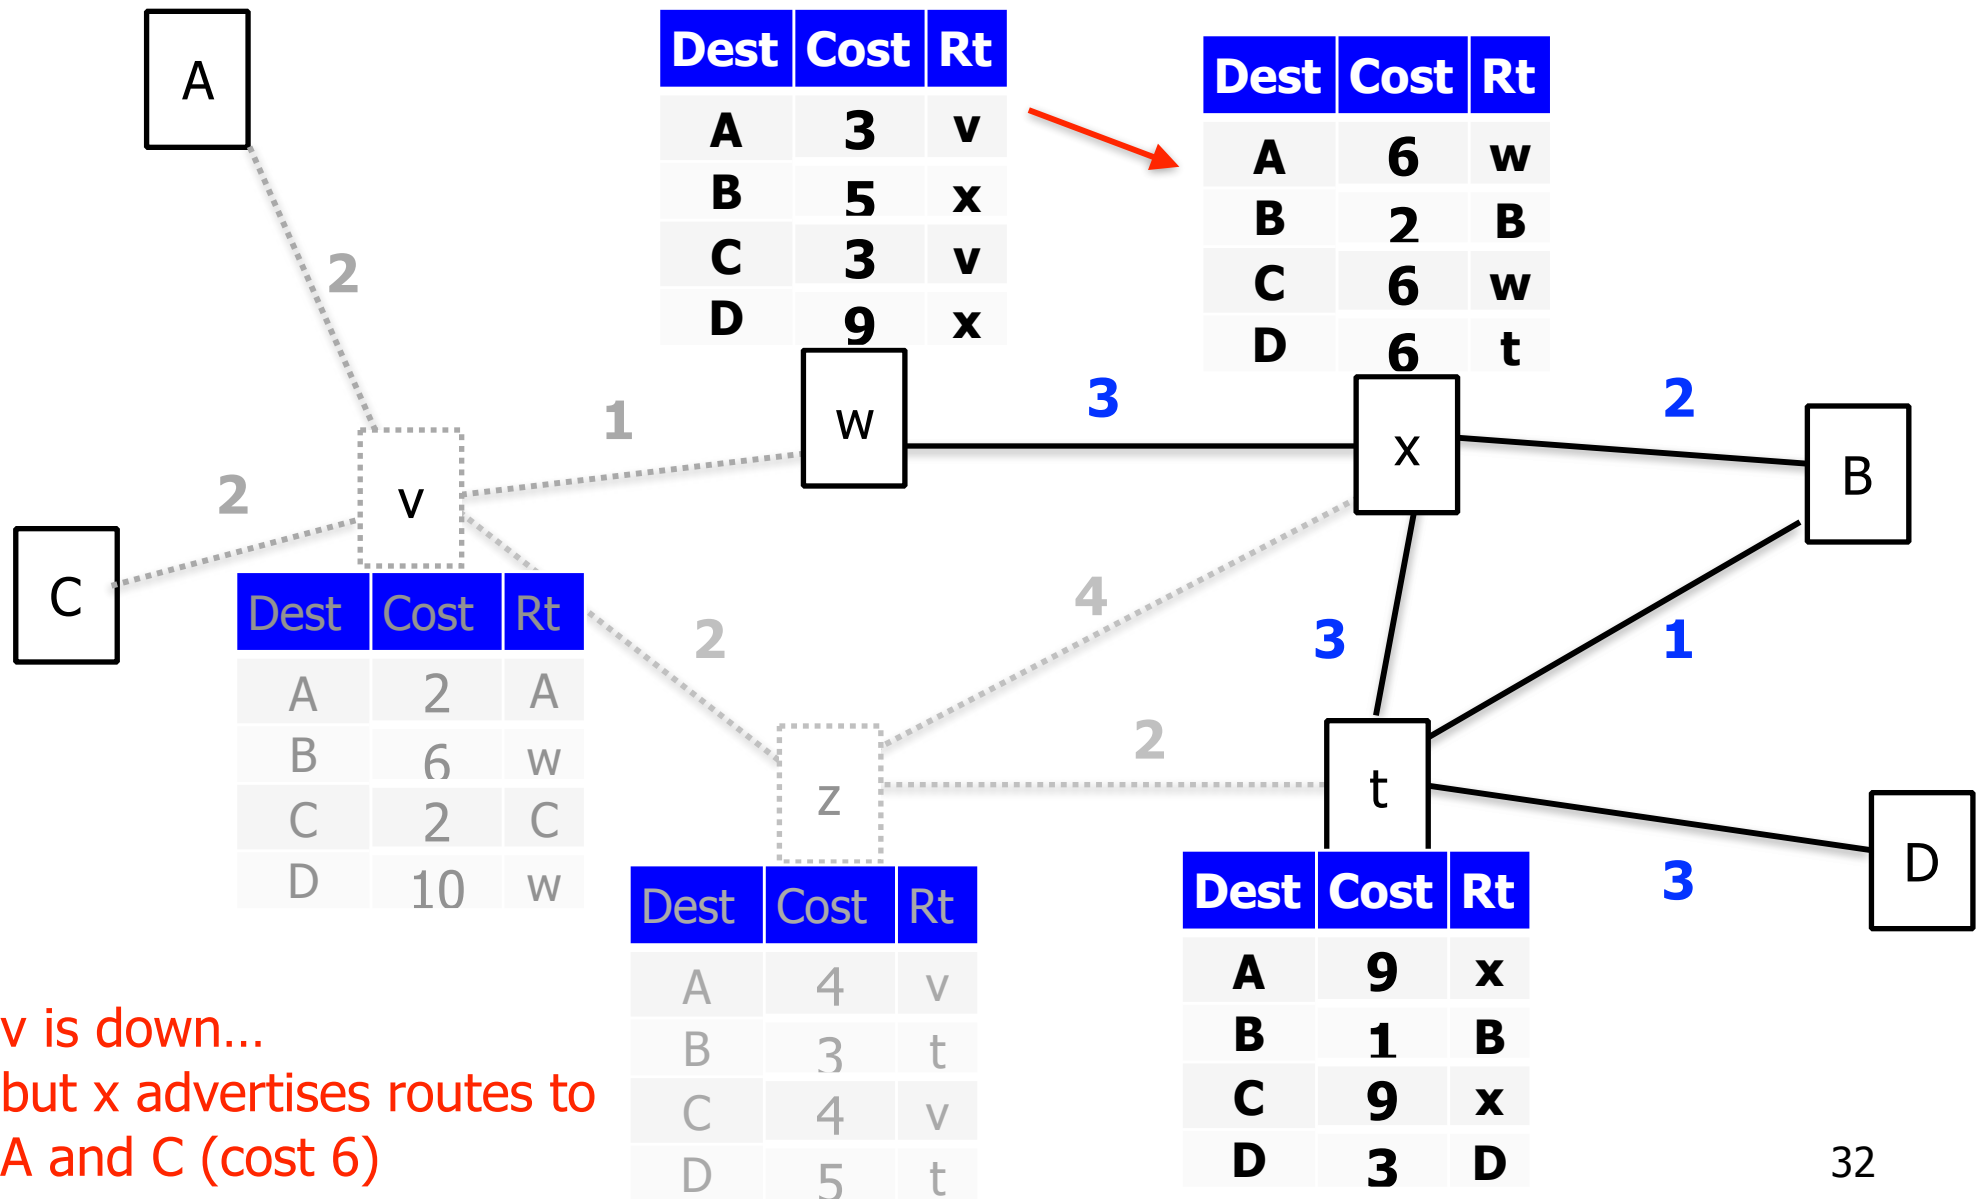


Stable again.
If z comes back, we'll
rediscover better routes

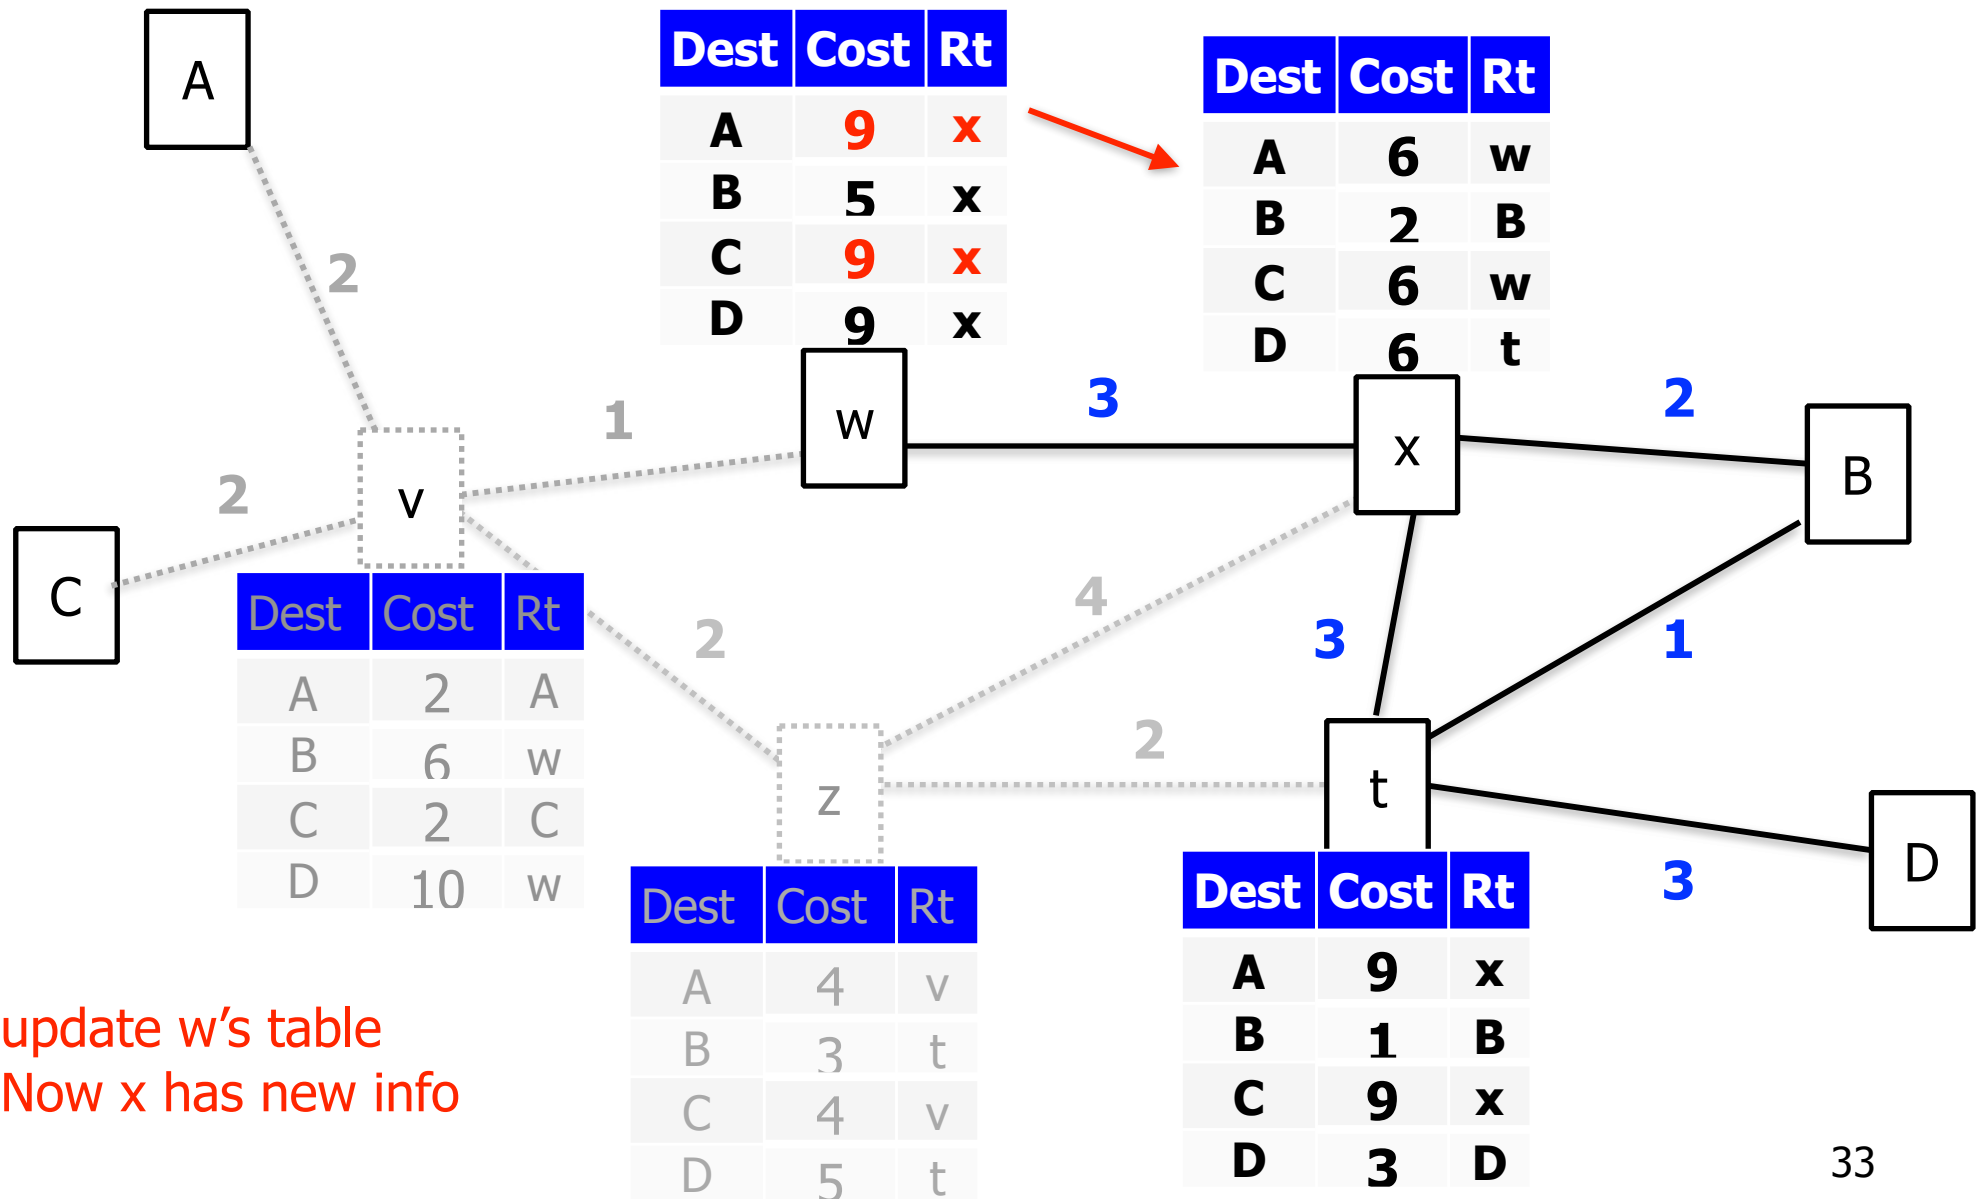
Distance Vector Routing



Distance Vector Routing

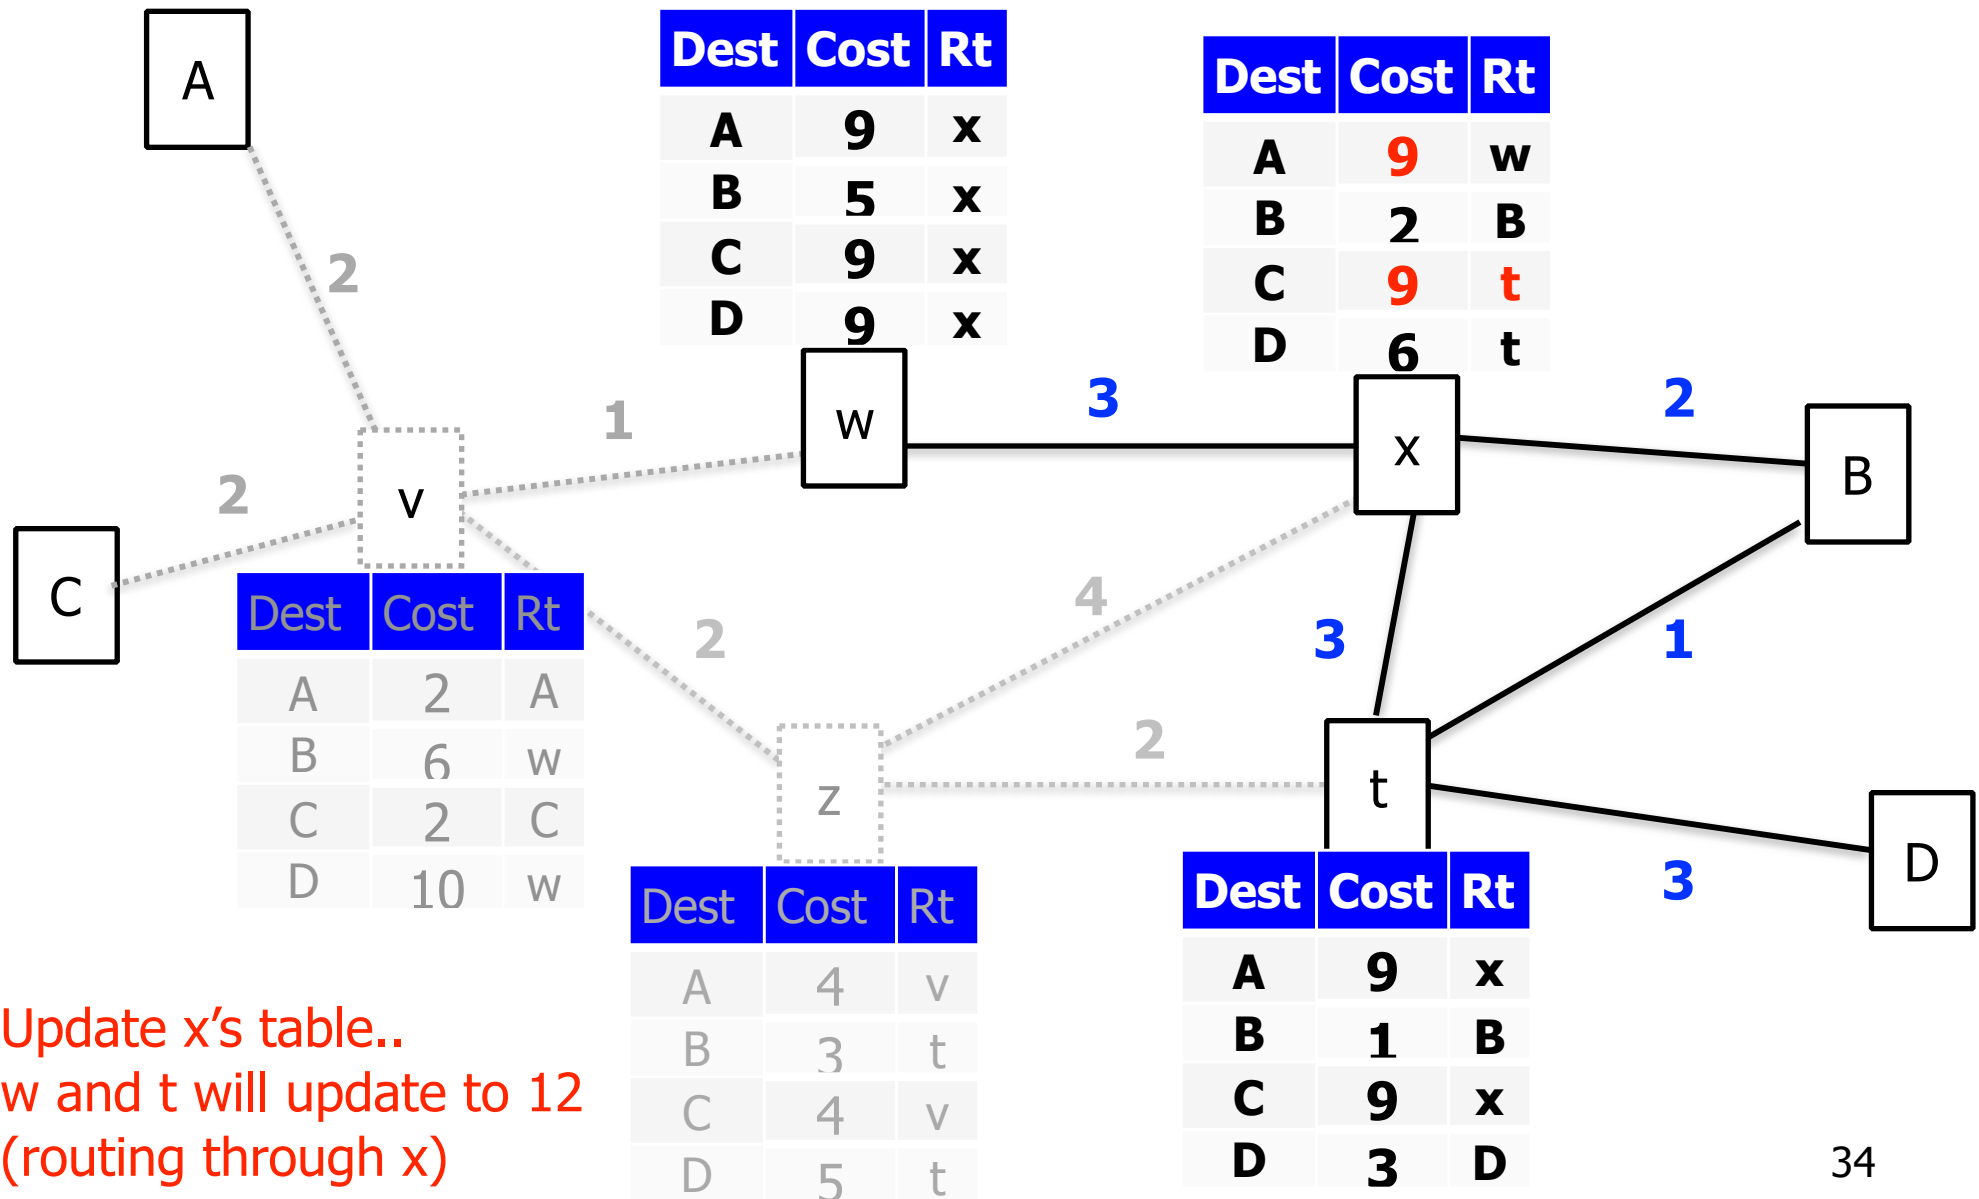


Distance Vector Routing



update w's table
Now x has new info

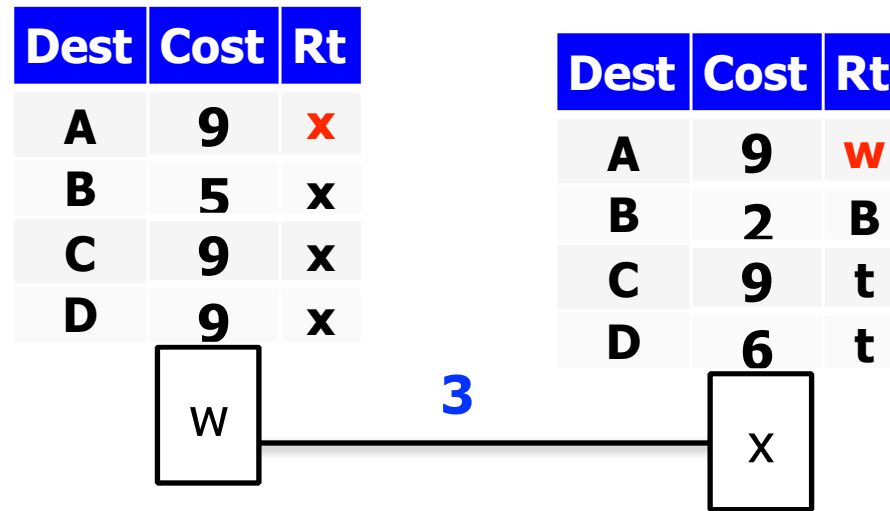
Distance Vector Routing



Count-to-infinity

- Algorithm will slowly “count to infinity”
 - Actually: count to max value it holds
 - Then throw away the route, concluding there is no path there
- Packets sent in meanwhile?
 - IP: Time To Live (TTL)
 - Starts at fixed value (e.g., 255)
 - Decrement every time the packet is forwarded
 - Packet dropped when $TTL == 0$
- Traceroute: uses TTL fields to probe to different distances
 - Uses ICMP protocol to get response on TTL
- Note: many fancier routing schemes, we aren't covering

Distance Vector Routing



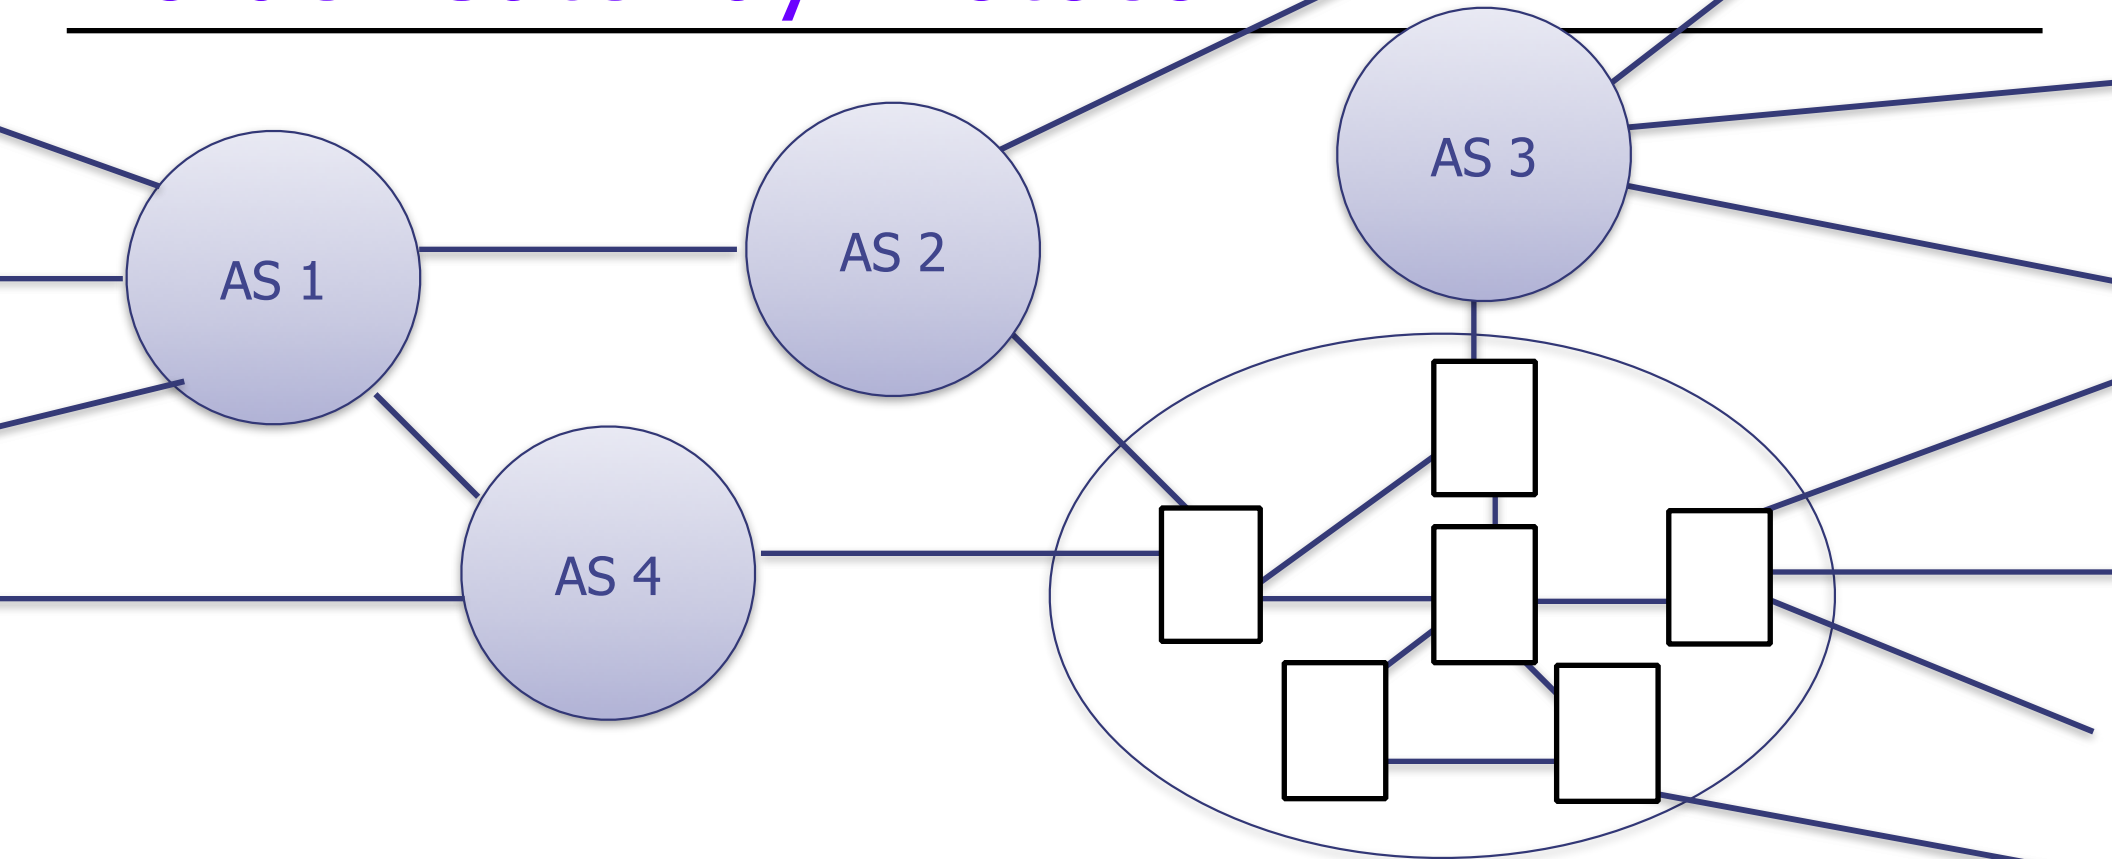
Obvious optimization:

- If x gets a route from w, it should not advertise that route back to w
- Called "split-horizon"
- Helps stabilize faster, but does not solve the problem (may still count to infinity)

Link State Protocols


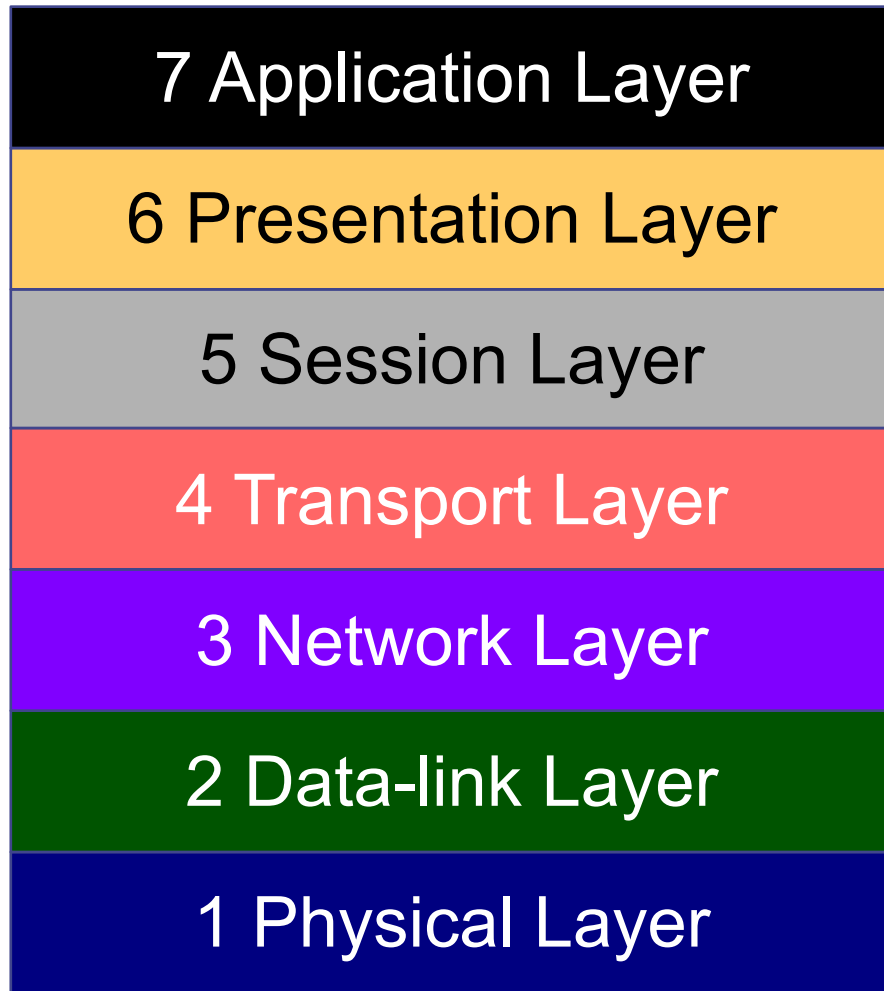
- Another option: link state protocols
 - Send info about direct connections to all routers
 - All routers build global pictures of network
 - Run graph algorithms to find shortest paths
 - E.g., Dijkstra's shortest path algorithm
- Global information is nice, but...
 - Complex for very large systems
 - How many routers on the internet?
 - Do they all exchange all their info and run Dijkstra's?
 - Of course not..
 - So... what do we do? **Use Abstraction... (and hierarchy)**

Border Gateway Protocol



- Divide internet up into Autonomous Systems (ASes)
 - Each AS can advertise routes to other ASes
 - Routing internal to AS is hidden from outside world
 - Can be Link State, Distance Vector, other...
 - We won't go into too many details

7-layer OSI model



- IP: hierarchical addresses + best effort delivery

Our messages so far



- IP Header has
 - Src IP
 - Dest IP
 - Payload type (what protocol)
 - Other info
- Side note: IP does fragmentation to fit within frame size
 - We aren't covering that


4: Transport Layer

- Reliability (if used)
 - Acknowledgements of data receipt
 - Retries of failed data
- Flow Control
 - Restrict rate of data sending
- Multiplexing/De-multiplexing data
 - E.g., Ports: identify which program some data is for
 - Keep data streams separate

5: Session Layer


- Concept of “a connection”
 - Establish/terminate
 - (OSI includes a variety of obscure features not often used)
- TCP: combines these two layers together
 - Sets up/terminates sessions
 - Has sequence numbering for packets
 - Acknowledges (ACKs) packets that are received
 - Establishes flow control (responds to congestion by throttling sending)

TCP




- We'll draw diagram with computers on each side
- Time goes down
- Three messages above (TCP's "3 way handshake")

TCP



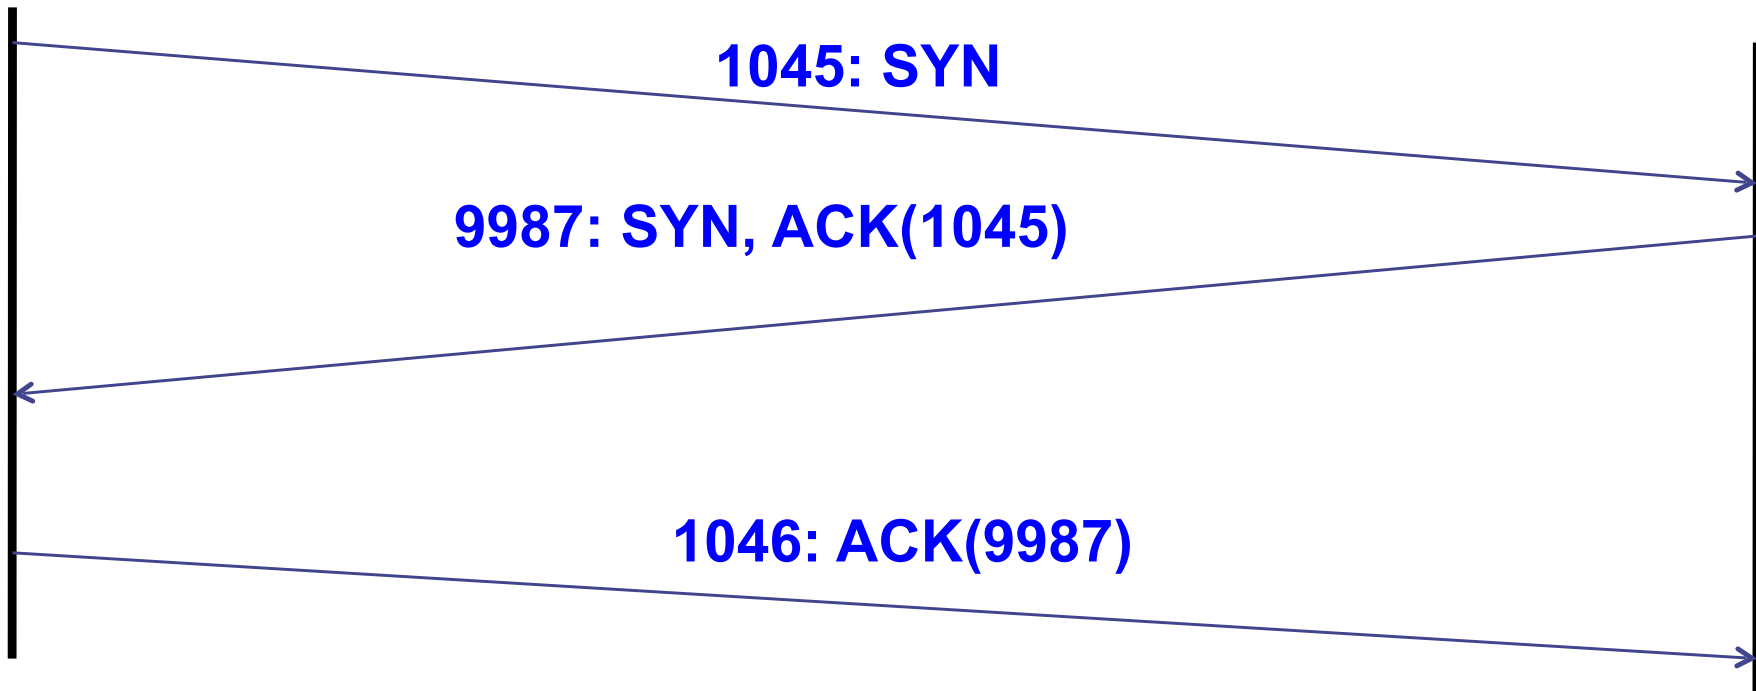
- To open a new connection:
 - 1 computer sends SYN ("Hi, lets talk")
 - All messages have sequence numbers including SYN
 - First sequence number of a new connection is random
 - TCP sequence numbers by byte

TCP




- Other computer
 - ACKs (Acknowledges) the message (says what sequence # it ACKs)
 - Also sends SYN "Hey sure, lets talk"

TCP




- First computer then ACKs this SYN
 - And probably sends data along with the ACK
- TCP control info (SYN, ACK, FIN): bits in TCP header
 - Packets can have multiple control bits on + carry data

TCP: Normal operation




- Data going right in blue
- ACKS coming left in green
 - note: ACK #ed by expected next data
- Sliding window (flow control)
 - Limit amount of un-ACKed data at a time

TCP: Re-ordered Data




- Data may get re-ordered in network
 - One packet takes one route, another takes another
- TCP: no problem
 - Sender re-orders data properly
 - Sends ACK for as much data as it has

TCP: Lost data




- Data may also get lost in the network
 - E.g., router is backlogged, can't handle it has to drop from queue
- TCP will re-send un-ACKed data after a timeout

TCP: Duplicate Data




- Receiver may get duplicate data
 - 7000-7499 gets lost
 - But 7500-7599 arrives
 - Then sender re-sends both: no ACK for either (why not?)
 - No problem: receiver drops duplicate (can tell: sequence #s)

TCP: Closing connection



- Connection closed with FIN message
 - Receiver ACKs
- Other side may close (with FIN) [typical]
 - Or remain open: can still send data
 - Side that closed cannot send, but should receive/ACK
 - FIN/ACK may be one message

TCP: Closing connection





- What if ACK for FIN gets lost?
 - FIN gets retried... but other side expects connection is closed?
- TCP has a state to handle this
 - Connection expected to be closed, but resources/state still held
 - Times out if no activity (assumes ACK got through if no retry)

Flow control: Sliding Window

- Problem:
 - Congestion -> dropped packets
 - Dropped packets -> Retries
 - Retries = duplicates of data -> More congestion
 - Vicious cycle...
- TCP implements **flow control** with a **sliding window**
 - Limitation of amount of un-ACKed data out at a time
 - Retry required? Shrink window
 - Assumes congestion, tries to avoid it
 - No retries in a while? Grow window back
 - Maybe it cleared up?

7-layer OSI model



- TCP: It's the coolest thing since memory got sliced into pages!

Our messages so far



- TCP header has
 - Source/Dest port
 - Sequence numbers
 - Control bits (SYN/ACK/FIN)
 - Check sum over data
 - Other stuff


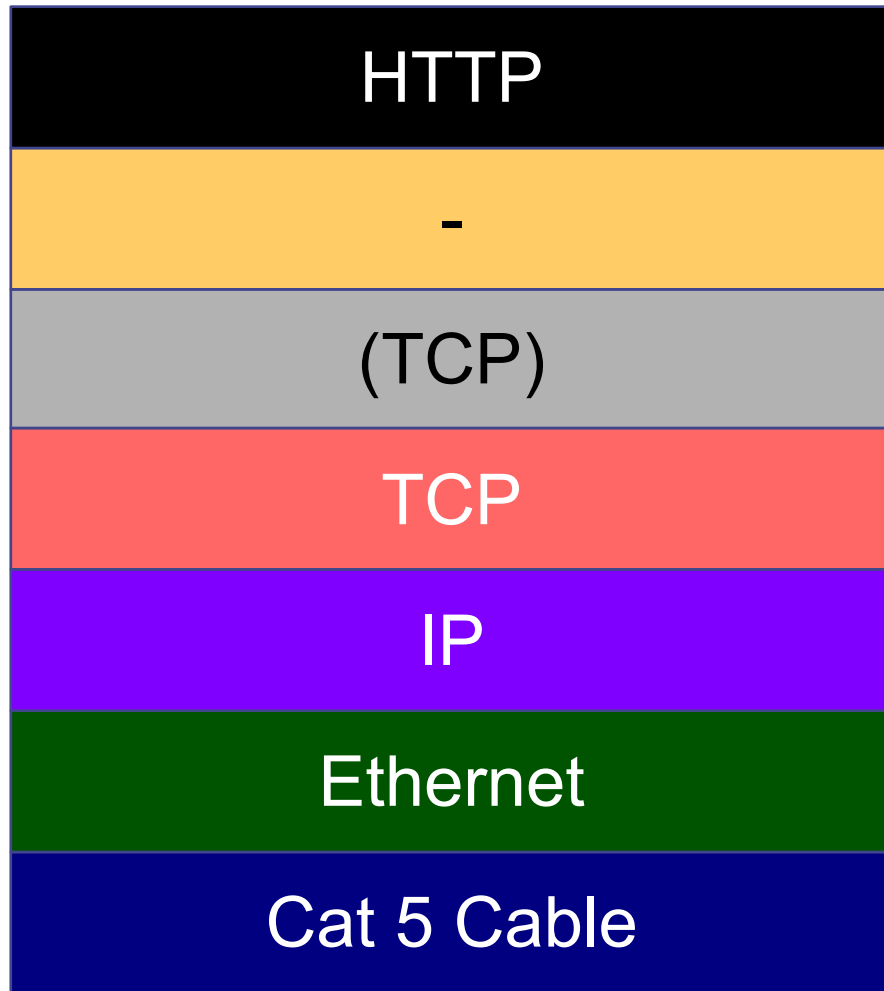
6: Presentation Layer

- Responsible for data formats
 - Examples
 - Character encoding schemes
 - Serialization of objects
- We're not really going to talk about it much

7: Application Layer



- Protocol specific to how applications want to communicate
 - Examples:
 - http
 - ftp
 - ssh
 - aim
 - SMTP
 - POP
 - ...
- Again, not going into this much...

7-layer OSI model



- Flexibility Example: Wired vs Wireless
 - Change out two layers, rest stay the same

7-layer OSI model



- Flexibility Example: A different application on top of both

Network programming

- Coding networking code in...
 - Java: Look in java.net, start with Socket
 - C:
 - socket()
 - connect()
 - accept()
 - bind()
 - listen()

When I say FIN, you say....?

FIN

When I say FIN, you say....?

ACK

Summary:

- Networking Overview
 - 7-layer model
 - Emphasis on IP (Layer 3) and TCP (Layers 4 and 5)
 - Not comprehensive, but...
 - You are now at least conversant enough to discuss the OSI stack at parties