

ECE 590.03 Computer Organization and Design

Pipelines

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

1

Admin

- Homework
 - Homework 6 out tonight
 - Due Friday (can use late days as needed)
 - Pencil and paper, no coding
- Reading:
 - Remainder of Chapter 4
- Recitation this week:
 - Review for final

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

2

Clock Period and CPI

- Single-cycle datapath
 - + Low CPI: 1
 - Long clock period: to accommodate slowest insn
- Multi-cycle datapath
 - + Short clock period
 - High CPI
- Can we have both low CPI and short clock period?
 - No good way to make a single insn go faster
 - + Insn latency doesn't matter anyway ... insn throughput matters
 - Key: **exploit inter-insn parallelism**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

3


Pipelining

- **Pipelining**: important performance technique
 - **Improves insn throughput rather than insn latency**
 - **Exploits parallelism at insn-stage level to do so**
 - Begin with multi-cycle design
- When insn advances from stage 1 to 2, next insn enters stage 1
- Individual insns take same number of stages
 - + **But insns enter and leave at a much faster rate**
 - Physically breaks "atomic" VN loop ... but must maintain illusion
- Automotive assembly line analogy

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

4


5 Stage Multi-Cycle Datapath



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

5

5 Stage Pipelined Datapath




- Temporary values (PC,IR,A,B,O,D) re-latched every stage
 - Why? 5 insns may be in pipeline at once, they share a single PC?
 - Notice, PC not latched after ALU stage (why not?)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

6

Pipeline Terminology



- Stages: **F**etch, **D**ecode, **X**ecute, **M**emory, **W**riteback
- Latches (pipeline registers): **PC**, **F/D**, **D/X**, **X/M**, **M/W**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

7

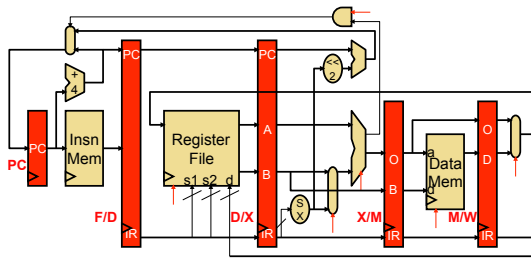
Some More Terminology

- **Scalar pipeline**: one insn per stage per cycle
 - Alternative: "superscalar" (take 552)
- **In-order pipeline**: insns enter execute stage in VN order
 - Alternative: "out-of-order" (take 552)
- **Pipeline depth**: number of pipeline stages
 - Nothing magical about five
 - Trend has been to deeper pipelines

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

8

Pipeline Example: Cycle 1



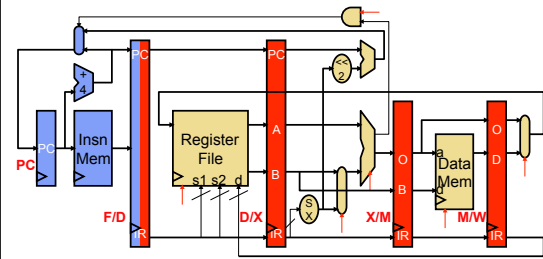
add \$3,\$2,\$1

- 3 instructions

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

9

Pipeline Example: Cycle 2




lw \$4,0(\$5)

add \$3,\$2,\$1

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

10

Pipeline Example: Cycle 3



sw \$6,4(\$7)

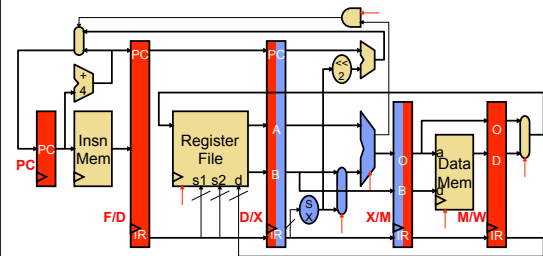
lw \$4,0(\$5)

add \$3,\$2,\$1

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

11

Pipeline Example: Cycle 4



sw \$6,4(\$7)

lw \$4,0(\$5)


add \$3,\$2,\$1

- 3 instructions

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

12

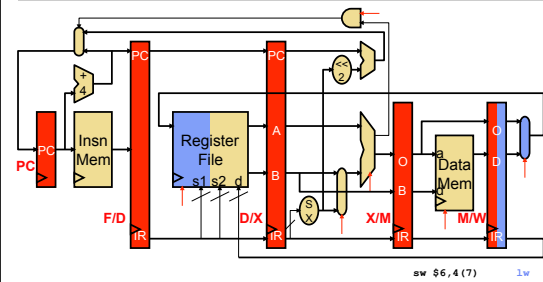
Pipeline Example: Cycle 5



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

13


Pipeline Example: Cycle 6



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

14

Pipeline Example: Cycle 7



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

15

Pipeline Diagram

- **Pipeline diagram:** shorthand for what we just saw
- Across: cycles
- Down: insns
- Convention: **X** means `lw $4, 0($5)` finishes execute stage and writes into X/M latch at end of cycle 4

	1	2	3	4	5	6	7	8	9
<code>add \$3, \$2, \$1</code>	F	D	X	M	W				
<code>lw \$4, 0(\$5)</code>		F	D	X	M	W			
<code>sw \$6, 4(\$7)</code>			F	D	X	M	W		

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

16


What About Pipelined Control?

- Should it be like single-cycle control?
 - But individual insn signals must be staged
- Should it be like multi-cycle control?
 - But all stages are simultaneously active
- How many different controllers are we going to need?
 - One for each insn in pipeline?
- Solution: use simple single-cycle control, but pipeline it
 - Single controller

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

17

Pipelined Control



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

18

Pipeline Performance Calculation

- Single-cycle
 - Clock period = 50ns, CPI = 1
 - Performance = 50ns/insn
- Multi-cycle
 - Branch: 20% (3 cycles), load: 20% (5 cycles), other: 60% (4 cycles)
 - Clock period = **12ns**, CPI = $(0.2*3 + 0.2*5 + 0.6*4) = 4$
 - Remember: latching overhead makes it 12, not 10
 - Performance = 48ns/insn
- Pipelined
 - Clock period = **12ns**
 - CPI = **1.5** (on average insn completes every 1.5 cycles)
 - Performance = **18ns/insn**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

19

Q1: Why Is Pipeline Clock Period ...

- ... > delay thru datapath / number of pipeline stages?
 - Latches (FFs) add delay
 - Pipeline stages have different delays, clock period is max delay
- Both factors have implications for ideal number pipeline stages

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

20

Q2: Why Is Pipeline CPI...

- ... > 1?
 - CPI for scalar in-order pipeline is **1 + stall penalties**
 - Stalls used to resolve hazards
 - **Hazard**: condition that jeopardizes VN illusion
 - **Stall**: artificial pipeline delay introduced to restore VN illusion
- Calculating pipeline CPI
 - **Frequency of stall * stall cycles**
 - Penalties add (stalls generally don't overlap in in-order pipelines)
 - $1 + \text{stall-freq}_1 * \text{stall-cyc}_1 + \text{stall-freq}_2 * \text{stall-cyc}_2 + \dots$
- Correctness/performance/MCCF
 - Long penalties OK if they happen rarely, e.g., $1 + 0.01 * 10 = 1.1$
 - Stalls also have implications for ideal number of pipeline stages

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

21


Dependences and Hazards

- **Dependence**: relationship between two insns
 - **Data**: two insns use same storage location
 - **Control**: one insn affects whether another executes at all
 - Not a bad thing, programs would be boring without them
 - Enforced by making older insn go before younger one
 - Happens naturally in single-/multi-cycle designs
 - But not in a pipeline
- **Hazard**: dependence & possibility of wrong insn order
 - Effects of wrong insn order cannot be externally visible
 - **Stall**: for order by keeping younger insn in same stage
 - Hazards are a bad thing: stalls reduce performance

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

22

Why Does Every Insn Take 5 Cycles?



- Could /should we allow **add** to skip M and go to W? No
 - It wouldn't help: peak fetch still only 1 insn per cycle
 - **Structural hazards**: imagine **add** follows **lw**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

23


Structural Hazards

- **Structural hazards**
 - Two insns trying to use same circuit at same time
 - E.g., structural hazard on regfile write port
- **To fix structural hazards**: proper ISA/pipeline design
 - Each insn uses every structure exactly once
 - For at most one cycle
 - Always at same stage relative to F

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

24

Data Hazards




- Let's forget about branches and the control for a while
- The three insn sequence we saw earlier executed fine...
 - But it wasn't a real program
 - Real programs have **data dependencies**
 - They pass values via registers and memory

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

25

Data Hazards




- Would this "program" execute correctly on this pipeline?
 - Which insns would execute with correct inputs?
 - add** is writing its result into **\$3** in current cycle
 - lw** read **\$3** 2 cycles ago → got wrong value
 - addi** read **\$3** 1 cycle ago → got wrong value
 - sw** is reading **\$3** this cycle → OK (regfile timing: write first half)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

26

Memory Data Hazards



- What about data hazards through memory? No
 - lw** following **sw** to same address in next cycle, gets right value
 - Why? DMem read/write take place in same stage
- Data hazards through registers? Yes (previous slide)
 - Occur because register write is 3 stages after register read
 - Can only read a register value 3 cycles after writing it

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

27

Fixing Register Data Hazards

- Can only read register value 3 cycles after writing it
- One way to enforce this: make sure programs don't do it
 - Compiler puts two independent insns between write/read insn pair
 - If they aren't there already
 - Independent means: "do not interfere with register in question"
 - Do not write it: otherwise meaning of program changes
 - Do not read it: otherwise create new data hazard
 - Code scheduling**: compiler moves around existing insns to do this
 - If none can be found, must use **nops**
- This is called **software interlocks**
 - MIPS**: Microprocessor w/out Interlocking Pipeline Stages

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

28

Software Interlock Example

```
add $3, $2, $1
lw $4, 0($3)
sw $7, 0($3)
add $6, $2, $8
addi $3, $5, 4
```

- Can any of last three insns be scheduled between first two
 - sw \$7, 0(\$3)**? No, creates hazard with **add \$3, \$2, \$1**
 - add \$6, \$2, \$8**? OK
 - addi \$3, \$5, 4**? No, **lw** would read **\$3** from it
 - Still need one more insn, use **nop**

```
add $3, $2, $1
add $6, $2, $8
nop
lw $4, 0($3)
sw $7, 0($3)
addi $3, $5, 4
```

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

29

Software Interlock Performance

- Same deal
 - Branch: 20%, load: 20%, store: 10%, other: 50%
- Software interlocks
 - 20% of insns require insertion of 1 **nop**
 - 5% of insns require insertion of 2 **nops**
- CPI is still 1 technically
 - But now there are more insns
 - #insns = $1 + 0.20 \cdot 1 + 0.05 \cdot 2 = 1.3$
 - **30% more insns (30% slowdown) due to data hazards**


ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

30

Hardware Interlocks


- Problem with software interlocks? Not compatible
 - Where does **3** in "read register 3 cycles after writing" come from?
 - From structure (depth) of pipeline
 - What if next MIPS version uses a 7 stage pipeline?
 - Programs compiled assuming 5 stage pipeline will break
- A better (more compatible) way: **hardware interlocks**
 - Processor detects data hazards and fixes them
 - Two aspects to this
 - Detecting hazards
 - Fixing hazards

Detecting Data Hazards




- Compare F/D insn input register names with output register names of older insns in pipeline
- Hazard =
- $$(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel (F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$$

Fixing Data Hazards




- Prevent F/D insn from reading (advancing) this cycle
 - Write **nop** into D/X.IR (effectively, insert **nop** in hardware)
 - Also reset (clear) the datapath control signals
 - Disable F/D latch and PC write enables (why?)
- Re-evaluate situation next cycle

Aside: Insert NOP/Reset Register



- Earlier: registers support separate clock, write enable
 - Useful for writes into register file
 - Also useful for implementing stalls
- Registers should also support **synchronous reset** (clear)
 - Useful for implementing stalls
 - Implement as additional hardwired 0 input to FF data mux
 - Resetting pipeline registers equivalent to inserting a NOP
 - If NOP is all zeros
 - If zero means "don't write" for all write-enable control signals
 - Design ISA/control signals to make sure this is the case

Hardware Interlock Example: cycle 1




lw \$4, 0(\$3) add \$3, \$2, \$1

$$(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel (F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$$

= 1

Hardware Interlock Example: cycle 2

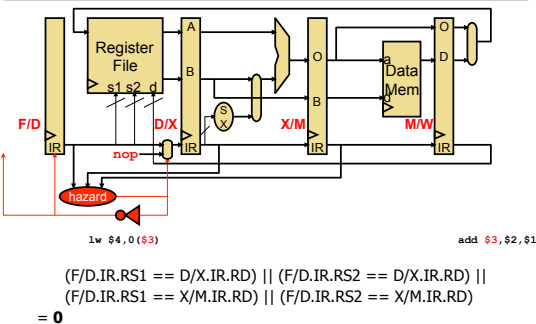


lw \$4, 0(\$3) add \$3, \$2, \$1

$$(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel (F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$$

= 1

Hardware Interlock Example: cycle 3



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

37

Pipeline Control Terminology

- Hardware interlock maneuver is called **stall** or **bubble**
- Mechanism is called **stall logic**
- Part of more general **pipeline control** mechanism
 - Controls advancement of insns through pipeline
- Distinguish from **pipelined datapath control**
 - Controls datapath at each stage
 - Pipeline control controls advancement of datapath control

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

38

Pipeline Diagram with Data Hazards

- Data hazard stall indicated with **d***
 - Stall propagates to younger insns

	1	2	3	4	5	6	7	8	9
add \$3, \$2, \$1	F	D	X	M	W				
lw \$4, 0(\$3)		F	d*	d*	D	X	M	W	
sw \$6, 4(\$7)					F	D	X	M	W

- This is not good (why?)

	1	2	3	4	5	6	7	8	9
add \$3, \$2, \$1	F	D	X	M	W				
lw \$4, 0(\$3)		F	d*	d*	D	X	M	W	
sw \$6, 4(\$7)			F	D	X	M	W		



ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

39


Hardware Interlock Performance

- Same deal
 - Branch: 20%, load: 20%, store: 10%, other: 50%
- Hardware interlocks: same as software interlocks
 - 20% of insns require 1 cycle stall (I.e., insertion of 1 `nop`)
 - 5% of insns require 2 cycle stall (I.e., insertion of 2 `nops`)
- $$CPI = 1 * 0.20 * 1 + 0.05 * 2 = 1.3$$
 - So, either CPI stays at 1 and #insns increases 30% (software)
 - Or, #insns stays at 1 (relative) and CPI increases 30% (hardware)
 - Same difference
- Anyway, we can do better

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

40

Observe




- Technically, this situation is broken
 - `lw $4, 0($3)` has already read `$3` from regfile
 - `add $3, $2, $1` hasn't yet written `$3` to regfile
- But fundamentally, everything is OK
 - `lw $4, 0($3)` hasn't actually used `$3` yet
 - `add $3, $2, $1` has already computed `$3`

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

41

Bypassing




- Bypassing**
 - Reading a value from an intermediate (μarchitectural) source
 - Not waiting until it is available from primary source
 - Here, we are bypassing the register file
 - Also called **forwarding**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

42

WX Bypassing




- What about this combination?
 - Add another bypass path and MUX input
 - First one was an **MX** bypass
 - This one is a **WX** bypass

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

43

ALUinB Bypassing




- Can also bypass to ALU input B

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

44

WM Bypassing?




- Does WM bypassing make sense?
 - Not to the address input (why not?)
 - But to the store data input, yes

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

45

Bypass Logic



- Each MUX has its own, here it is for MUX ALUinA
 - (D/X.IR.RS1 == X/M.IR.RD) ==> 0
 - (D/X.IR.RS1 == M/W.IR.RD) ==> 1
 - Else ==> 2

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

46

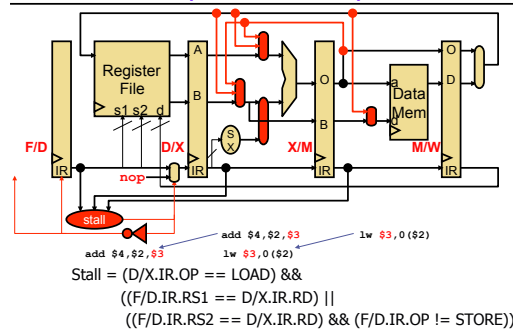
Bypass and Stall Logic

- Two separate things
 - Stall logic controls pipeline registers
 - Bypass logic controls MUXs
- But complementary
 - For a given data hazard: if can't bypass, must stall
- Slide #43 shows **full bypassing**: all bypasses possible
 - Is stall logic still necessary?

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

47

Yes, Load Output to ALU Input



Stall = (D/X.IR.OP == LOAD) &&
 ((F/D.IR.RS1 == D/X.IR.RD) ||
 ((F/D.IR.RS2 == D/X.IR.RD) && (F/D.IR.OP != STORE)))

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

48

Pipeline Diagram With Bypassing

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	D	X	M	W			
addi \$6,\$4,1			F	d*	D	X	M	W	


- Use compiler scheduling to reduce load-use stall frequency
 - Like software interlocks, but for performance not correctness

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	D	X	M	W			
sub \$8,\$3,\$1			F	D	X	M	W		
addi \$6,\$4,1				F	D	X	M	W	

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

49

Control Hazards




- Control hazards**
 - Must fetch post branch insns before branch outcome is known
 - Default: assume "not-taken" (at fetch, can't tell it's a branch)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

50

Branch Recovery



- Branch recovery:** what to do when branch is actually taken
 - Insns that will be written into F/D and D/X are wrong
 - Flush them**, i.e., replace them with **nops**
 - + They haven't had written permanent state yet (regfile, DMem)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

51

Branch Recovery Pipeline Diagram

	1	2	3	4	5	6	7	8	9
addi \$3,\$0,1	F	D	X	M	W				
bnez \$3,targ		F	D	X	M	W			
sw \$6,4(\$7)			F	D					
targ: addi \$6,\$7,1				F					
targ: addi \$8,\$7,1					F	D	X	M	W

- Convention: don't fill in flushed insns
- Taken branch penalty is 2 cycles

	1	2	3	4	5	6	7	8	9
addi \$3,\$0,1	F	D	X	M	W				
bnez \$3,targ		F	D	X	M	W			
targ: addi \$6,\$7,1					F	D	X	M	W

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

52


Branch Performance

- Back of the envelope calculation
 - Branch: 20%**, load: 20%, store: 10%, other: 50%
 - 75% of branches are taken**
- $CPI = 1 + 0.20 * 0.75 * 2 = 1.3$
 - Branches cause 30% slowdown**
 - How do we reduce this penalty?

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

53

Fast Branch



- Fast branch:** can decide at D, not X
 - Test must be comparison to zero or equality, no time for ALU
 - New taken branch penalty is 1
 - Additional insns (s1t) for more complex tests, must bypass to D too
 - 25% of branches have complex tests that require extra insn
 - $CPI = 1 + 0.20 * 0.75 * 1(\text{branch}) + 0.20 * 0.25 * 1(\text{extra insn}) = 1.2$

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

54

Speculative Execution

- Speculation: "risky transactions on chance of profit"
- **Speculative execution**
 - Execute before all parameters known with certainty
 - **Correct speculation**
 - + Avoid stall, improve performance
 - **Incorrect speculation (mis-speculation)**
 - Must abort/flush/squash incorrect insns
 - Must undo incorrect changes (recover pre-speculation state)
 - The "game": $[\%_{\text{correct}} * \text{gain}] - [(1 - \%_{\text{correct}}) * \text{penalty}]$
- **Control speculation:** speculation aimed at control hazards
 - Unknown parameter: are these the correct insns to execute next?

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

55

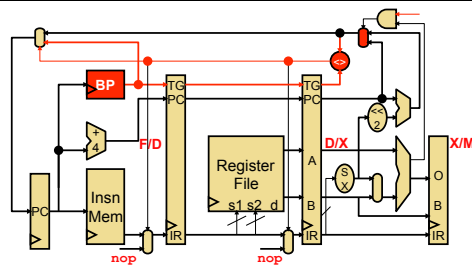
Control Speculation Mechanics

- Guess branch target, start fetching at guessed position
 - Doing nothing is implicitly guessing target is PC+4
 - Can actively guess other targets: **dynamic branch prediction**
- Execute branch to verify (check) guess
 - Correct speculation? keep going
 - Mis-speculation? Flush mis-speculated insns
 - Hopefully haven't modified permanent state (Regfile, DMem)
 - + Happens naturally in in-order 5-stage pipeline
- "Game" for in-order 5 stage pipeline
 - $\%_{\text{correct}} = ?$
 - Gain = 2 cycles
 - + Penalty = 0 cycles → **mis-speculation no worse than stalling**

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

56

Dynamic Branch Prediction



- **Dynamic branch prediction:** guess outcome
 - Start fetching from guessed address
 - Flush on **mis-prediction** (notice new recovery circuit)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

57

Branch Prediction: Short Summary

- Key principle of micro-architecture:
 - Programs do the same thing over and over (why?)
- Exploit for performance:
 - Learn what a program did before
 - Guess that it will do the same thing again
- Quick intro to details

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

58

Branch Prediction 10K feet

- Two (separate) tasks:
 - Predict taken/not taken
 - Predict taken target

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

59

Branch Prediction 10K feet

- Two (separate) tasks:
 - Predict taken/not taken
 - Predict taken target
- High level solution (both tasks):
 - SRAM "array" to remember most recent behaviors
 - Kind of like a cache, indexed by PC bits, but different
 - Typically no next level (but can have 2 levels)
 - Can skip tag, or use partial tag
 - Predictor: OK to be wrong (as long as we fix it)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

60

Branch Target Buffer (BTB)

- Branch Target Buffer
 - SRAM array, holds recent taken targets
 - Example: 4K entries, direct mapped
 - Can be set-associative
 - Each entry holds partial PC (low order bits)
 - Assume high bits unchanged (why?)
 - Example: 16 bits

0	01F3
1	4242
2	1234
.....
4097	4242

Branch Target Buffer (BTB)

- Branch Target Buffer
 - SRAM array, holds recent taken targets
 - Example: 4K entries, direct mapped
 - Can be set-associative
 - Each entry holds partial PC (low order bits)
 - Assume high bits unchanged (why?)
 - Example: 16 bits
- Prediction of taken target:
 - Use PC bits 2—13 to index BTB (why these bits?)
 - Replace PC bits 2—17 with value in BTB

0	01F3
1	4242
2	1234
.....
4097	4242

Branch Target Buffer (BTB)

- Branch Target Buffer
 - SRAM array, holds recent taken targets
 - Example: 4K entries, direct mapped
 - Can be set-associative
 - Each entry holds partial PC (low order bits)
 - Assume high bits unchanged (why?)
 - Example: 16 bits
- Prediction of taken target:
 - Use PC bits 2—13 to index BTB (why these bits?)
 - Replace PC bits 2—17 with value in BTB
- Update (how do values get into predictor?)
 - At execute, if branch is taken write target into BTB
 - Use PC bits 2—13 to index for write also (same entry)

0	01F3
1	4242
2	1234
.....
4097	4242


Target Prediction: BTB collisions

- PCs may collide in BTB
 - Example: 0x10000000 and 0x20000000 (both index 0)
 - Could use tags (or partial tags)
 - Better to just guess "not taken" than "taken to bogus target"
 - Why?

Target Prediction: BTB collisions

- PCs may collide in BTB
 - Example: 0x10000000 and 0x20000000 (both index 0)
 - Could use tags (or partial tags)
 - Better to just guess "not taken" than "taken to bogus target"
 - Why?
 - What if 0x10000000 is a branch, and 0x20000000 is not?
 - Pipeline may predict bogus next PC for non-branch
 - Fine as long as detected/fixed (extra checking)
 - Usually checked in decode if possible
 - Alternative: pre-decode bits
 - Add bits in I\$ to say "is this a branch"
 - Know if not a branch while predicting
 - Bits set on I\$ fill path (examine bits coming from L2)

Our branch predictor (so far)



- Missing piece (???): Direction predictor
 - Should we use the taken target (from BTB) or not?

Direction Prediction

- Need to predict “taken” (T) or “not taken” (N)
 - This is typically the hard part, by the way
- Simplest approach: just guess “same as last time”
 - Actually, kind of not bad:
 - Loops: almost always right (taken)
 - Error checks: almost always right (no error)
 - ...etc..
- Implementation:
 - SRAM, indexed by PC bits
 - 1 bit per entry: 1 = taken, 0 = not taken
 - No tags.
 - Collisions? Meh—they happen

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

67

Direction Prediction: Example

- Consider:

```
for (int i = 0; i < 10000000; i++) {  
  for (int j = 0; j < 6; j++) {  
    //stuff  
  }  
}
```

Branches outcomes:
TTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNT...

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

68

Direction Prediction: Example

- Consider:

```
for (int i = 0; i < 10000000; i++) {  
  for (int j = 0; j < 6; j++) {  
    //stuff  
  }  
}
```

Branches outcomes:
TTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNT...

Predictions:
NTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTT...

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

69

Direction Prediction: Can we do better?

Branches outcomes:
TTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNT...

Predictions:
NTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTT...

- Problem:
 - A little too quick to react
 - One-off difference causes **two** mis-predictions
- Solution:
 - Slow down changes in prediction: 2-bit counters
 - T (11), t (10), n (00), N (01)
 - “Strongly” (T/N) and “weakly” (t/n) taken/not taken
 - Updates: taken-> increment, not taken -> decrement

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

70

Direction Prediction: Can we do better?

Branches outcomes:
TTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNT...

Predictions:
NTTTTTNTTTTTNTTTTTNTTTTTNTTTTTNTTTTTT...

tTTTTTtTTTTTtTTTTTtTTTTTtTTTTTtTTTTTt...

- Problem:
 - A little too quick to react
 - One-off difference causes **two** mis-predictions
- Solution:
 - Slow down changes in prediction: 2-bit counters
 - T (11), t (10), n (00), N (01)
 - “Strongly” (T/N) and “weakly” (t/n) taken/not taken
 - Updates: taken-> increment, not taken -> decrement

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

71

Can we do even better still?

- Our branches have a very regular pattern
 - 6Ts, then 1 N
 - We really should be able to get them all right... right?
- Real predictors use **history**
 - Take recent branch outcomes (NTTTTT = 0111111)
 - XOR with PC to form table index
 - Same PC, different history -> different index -> different counter
 - Would predict previous example perfectly
- Also useful for correlation of branches
 - Nearby branches with related outcomes (why is this common?)

ECE 590 (Hilton): Pipelines [Slides adapted from A. Roth's]

72

Direction Prediction: Continued..

- Real direction predictors more complex even still
 - Multiple tables with choosers (hybrid history schemes)
- Research ideas too
 - Late 90s/early 2000s: think up bpred idea, publish, repeat
- Big impediment to performance/hard to get well
- Also research ideas for how to get around it
 - Control Independence: predicting reconvergence point easier

Predicting returns

- Previous things don't work well on "return" instructions
 - jr \$ra
 - Why not?

Predicting returns

- Previous things don't work well on "return" instructions
 - jr \$ra
 - Why not?
 - Functions called from many places
 - Previous place to return to, not always current place to return to...
 - But should be predictable: why?

Predicting returns

- Previous things don't work well on "return" instructions
 - jr \$ra
 - Why not?
 - Functions called from many places
 - Previous place to return to, not always current place to return to...
 - But should be predictable: why?
 - Matches up with jal's PC +4
 - In stack-like fashion
 - So....

Predicting returns

- Previous things don't work well on "return" instructions
 - jr \$ra
 - Why not?
 - Functions called from many places
 - Previous place to return to, not always current place to return to...
 - But should be predictable: why?
 - Matches up with jal's PC +4
 - In stack-like fashion
 - So....
- "Return Address Stack" (aka "Link Stack")
 - Predictor tracks a stack of recent jals
 - Encounter a jr \$ra? Pop stack for predicted target

Branch Prediction Performance

- Dynamic branch prediction
 - Simple predictor: branches predicted with 75% accuracy
 - $CPI = 1 + 0.20 * 0.25 * 2 = 1.1$
 - More advanced predictor: 95% accuracy
 - $CPI = 1 + 0.20 * 0.05 * 2 = 1.02$
- Branch mis-predictions still a big problem though
 - Pipelines are long: typical mis-prediction penalty is 10+ cycles
 - Pipelines have full bypassing: compiler schedules the rest
 - Pipelines are superscalar (later)

Pipelining And Exceptions

- Pipelining makes exceptions more complex
 - 5 insns in pipeline at once
 - Exception happens, how do you know which insn caused it?
 - Exceptions propagate along pipeline in latches
 - Two exceptions happen, how do you know which one to take first?
 - One belonging to oldest insn
 - When handling exception, have to flush younger insns
 - Piggy-back on branch mis-prediction machinery to do this
 - What about multi-cycle operations?
- Just FYI

Pipeline Depth

- No magic about 5 stages, trend had been to deeper pipelines
 - 486: 5 stages (50+ gate delays / clock)
 - Pentium: 7 stages
 - Pentium II/III: 12 stages
 - Pentium 4: 22 stages (~10 gate delays / clock) **"super-pipelining"**
 - Core1/2: 14 stages
- Increasing **pipeline depth**
 - + Increases clock frequency (reduces period)
 - But decreases IPC (increases CPI)
 - Branch mis-prediction penalty becomes longer
 - Non-bypassed data hazard stalls become longer
 - At some point, CPI losses offset clock gains, question is when?
 - 1GHz Pentium 4 was slower than 800 MHz PentiumIII
 - What was the point? People by frequency, not frequency * IPC

Real pipelines...

- Real pipelines fancier than what we have seen
 - Superscalar: multiple instructions in a stage at once
 - Out-of-order: re-order instructions to reduce stalls
 - SMT: execute multiple threads at once on processor
 - Side by side, sharing pipeline resources
 - Multi-core: multiple pipelines on chip
 - Cache coherence: No stale data
- Learn more?
 - Take ECE 552 next Fall!