

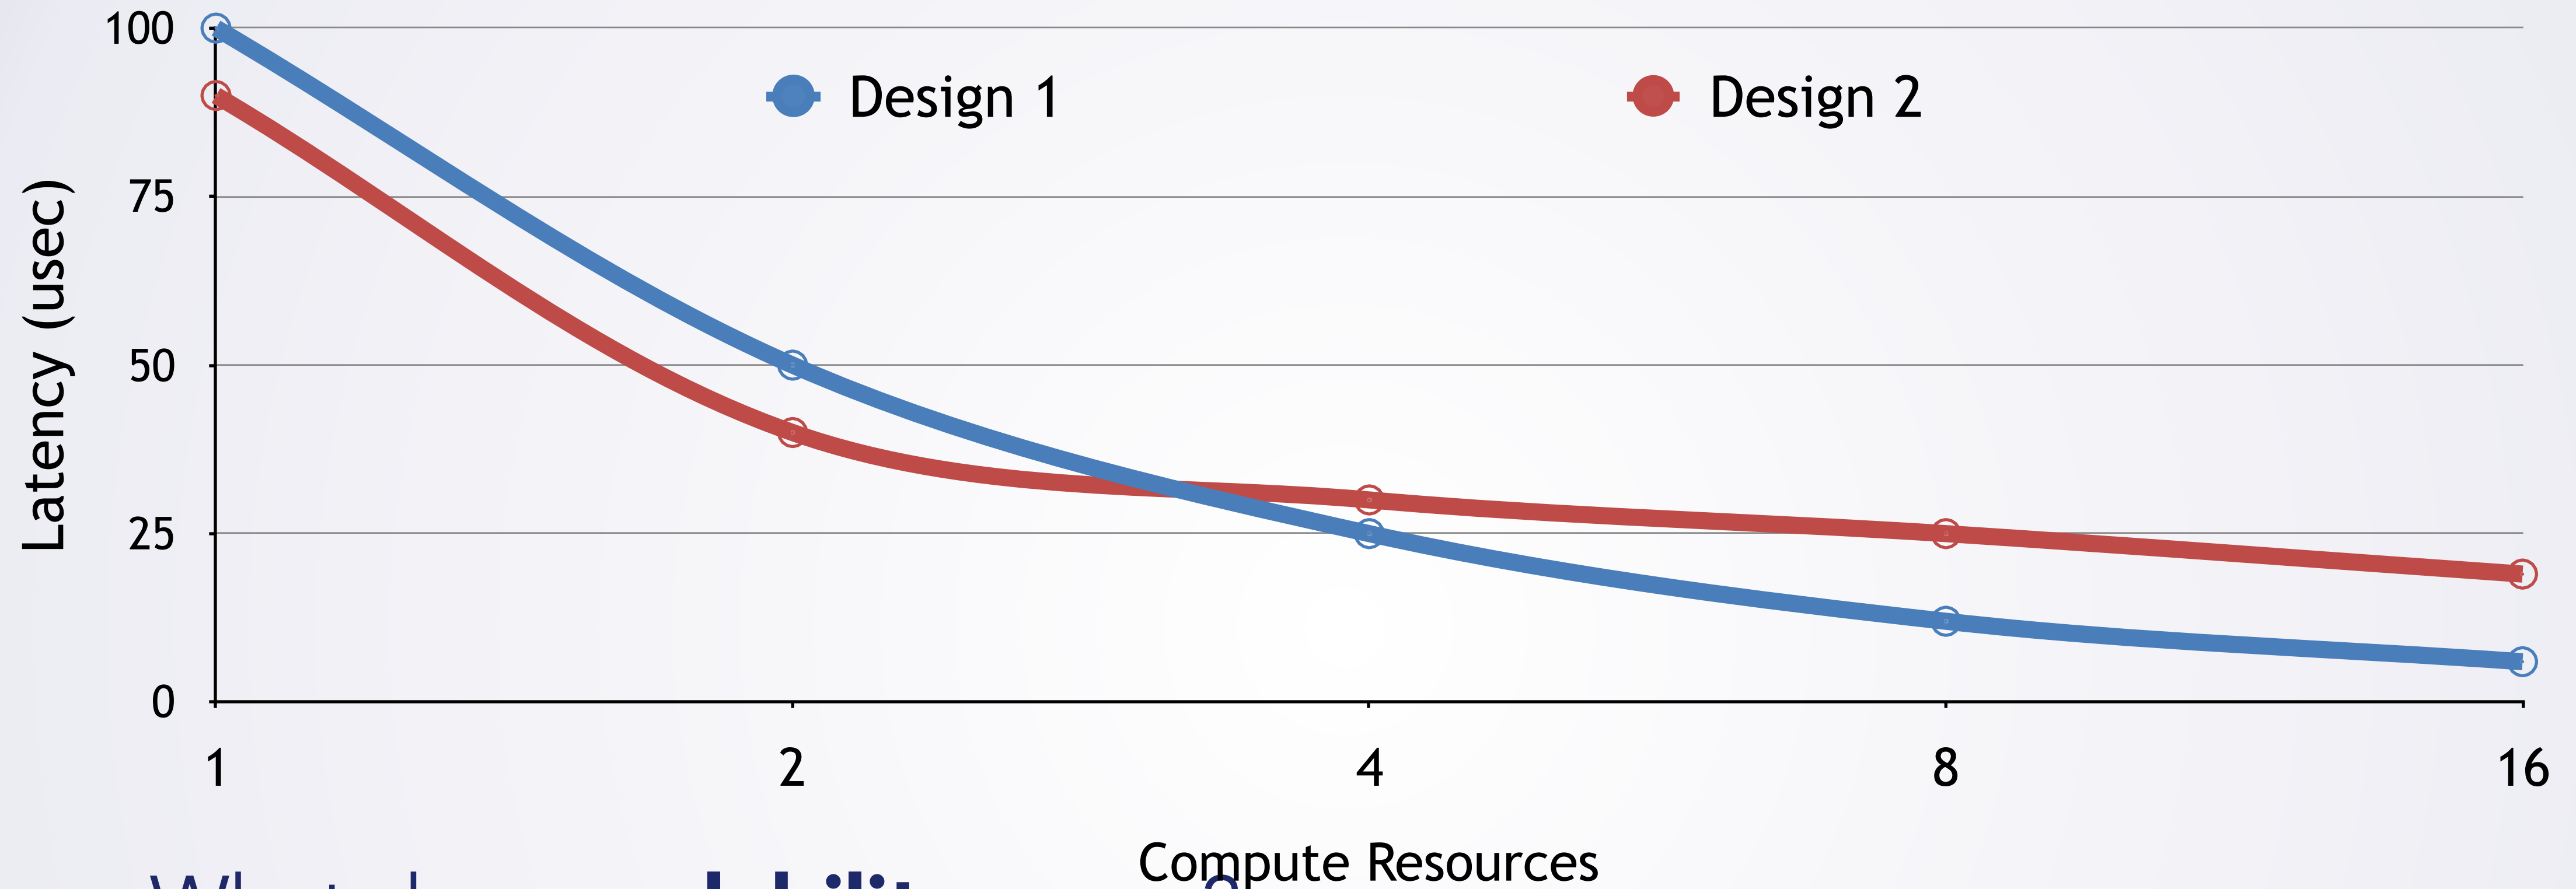
# Engineering Robust Server Software

## Scalability

# Intro To Scalability


- What does **scalability** mean?

# Intro To Scalability



- What does **scalability** mean?
  - How does performance change with resources?

# Intro To Scalability




- What does **scalability** mean?
  - How does performance change with resources?
  - How does performance change with load?

# Scalability Terms

- **Scale Out:** Add more nodes
  - More computers
- **Scale Up:** Add more stuff in each node
  - More processors in one node
- **Strong Scaling:** How does time change for fixed problem size?
  - Do 100M requests, add more cores -> speedup?
- **Weak Scaling:** How does time change for fixed (problem size/core)?
  - Do  $(100 \cdot N)M$  requests, with  $N$  cores -> speedup?

# Amdahl's Law

$$\text{Speedup (N)} = \frac{S + P}{S + \frac{P}{N}}$$






# Amdahl's Law

$$\text{Speedup (N)} = \frac{S + P}{S + \frac{P}{N}} = \frac{6 + 4}{6 + \frac{4}{N}}$$

Serial Portion:  $S = 6$




Parallel Portion:  $P = 4$

# Amdahl's Law

$$\text{Speedup (N)} = \frac{S + P}{S + \frac{P}{N}} = \frac{6 + 4}{6 + \frac{4}{2}} = \frac{10}{8}$$

Serial Portion:  $S = 6$



Parallel Portion:  $P = 4$


- $10/8 = 1.25x$  speedup = 25% increase in **throughput**.
- $8/10 = 0.8x$  = 20% reduction in **latency**



# Amdahl's Law

$$\text{Speedup (N)} = \frac{S + P}{S + \frac{P}{N}} = \frac{6 + 4}{6 + \frac{4}{4}} = \frac{10}{7}$$

Serial Portion:  $S = 6$




Parallel Portion:  $P = 4$

- $10/7 = 1.42x$  speedup = 42% increase in **throughput**.
- $7/10 = 0.7x$  = 30% reduction in **latency**

# Amdahl's Law

$$\text{Speedup (N)} = \frac{S + P}{S + \frac{P}{N}} = \frac{6 + 4}{6 + \frac{4}{\infty}} = \frac{10}{6}$$

Serial Portion:  $S = 6$




Speedup  $\propto x$

Parallel Portion:  $P = 4$


- $10/6 = 1.67x$  speedup = 67% increase in **throughput**.
- $6/10 = 0.6x$  = 40% reduction in **latency**

# Amdahl's Law



- Anne Bracy: "Don't try to speed up brushing your teeth"
  - What does she mean?

# Why Not Perfect Scalability?



- Why don't we get (Nx) speedup with N cores?
  - What prevents ideal speedups?

# Impediments to Scalability

- Shared Hardware
  - Functional Units
  - Caches
  - Memory Bandwidth
  - IO Bandwidth
  - ...
- Data Movement
  - From one core to another
- Blocking
  - Locks (and other synchronization)
  - Blocking IO




# Impediments to Scalability

- Shared Hardware
  - Functional Units
  - Caches
  - Memory Bandwidth
  - IO Bandwidth
  - ...
- Data Movement
  - From one core to another
- Blocking
  - Blocking IO
  - Locks (and other synchronization)

Let's talk about these for now

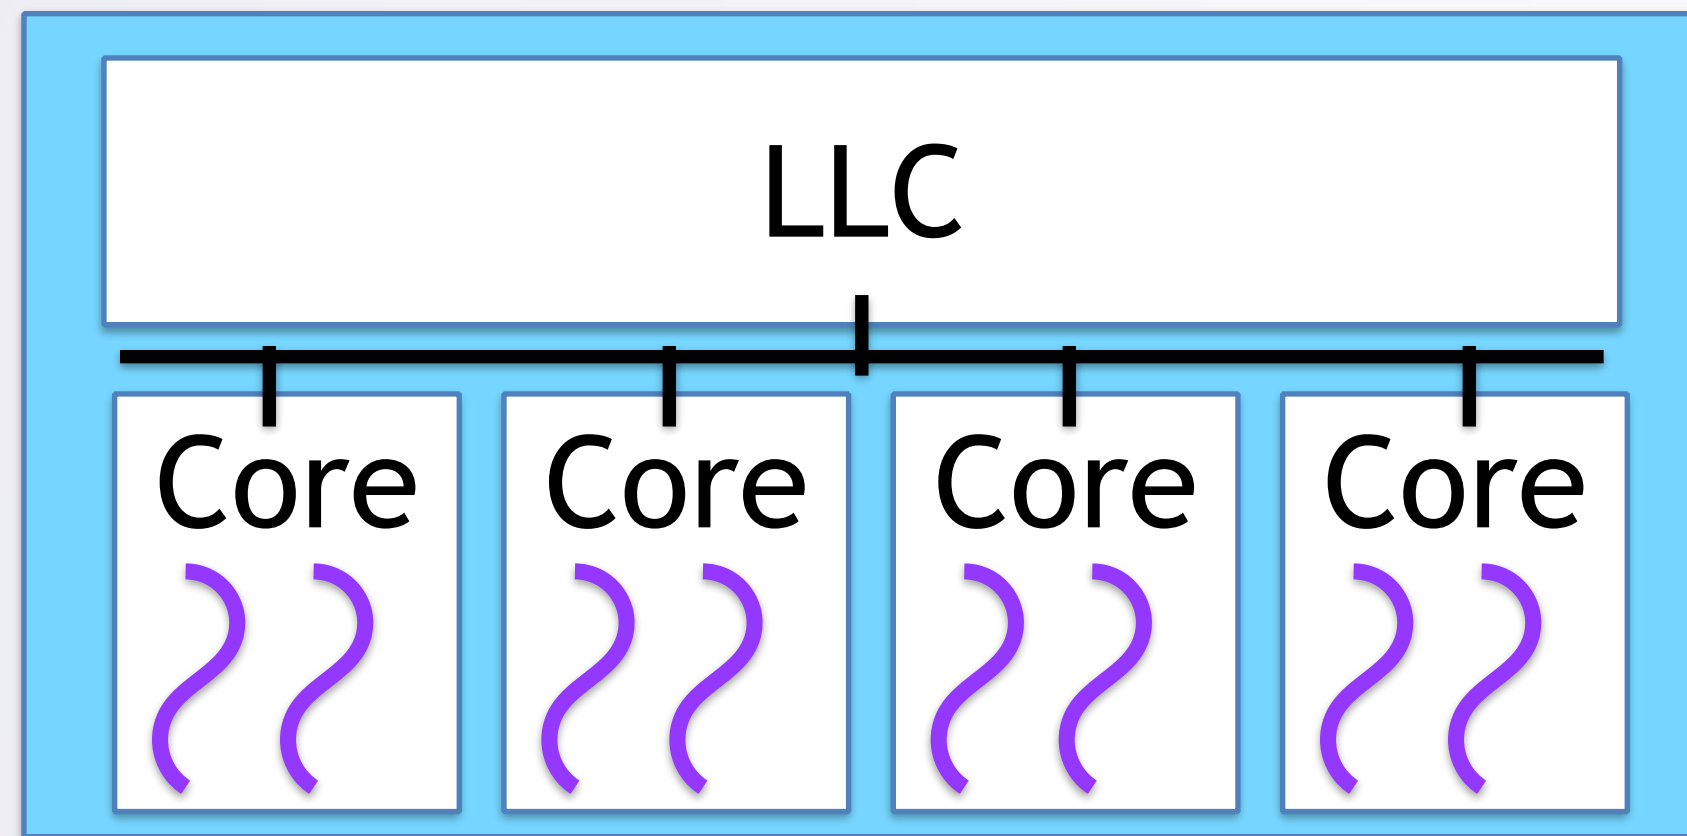


# Hypothetical System



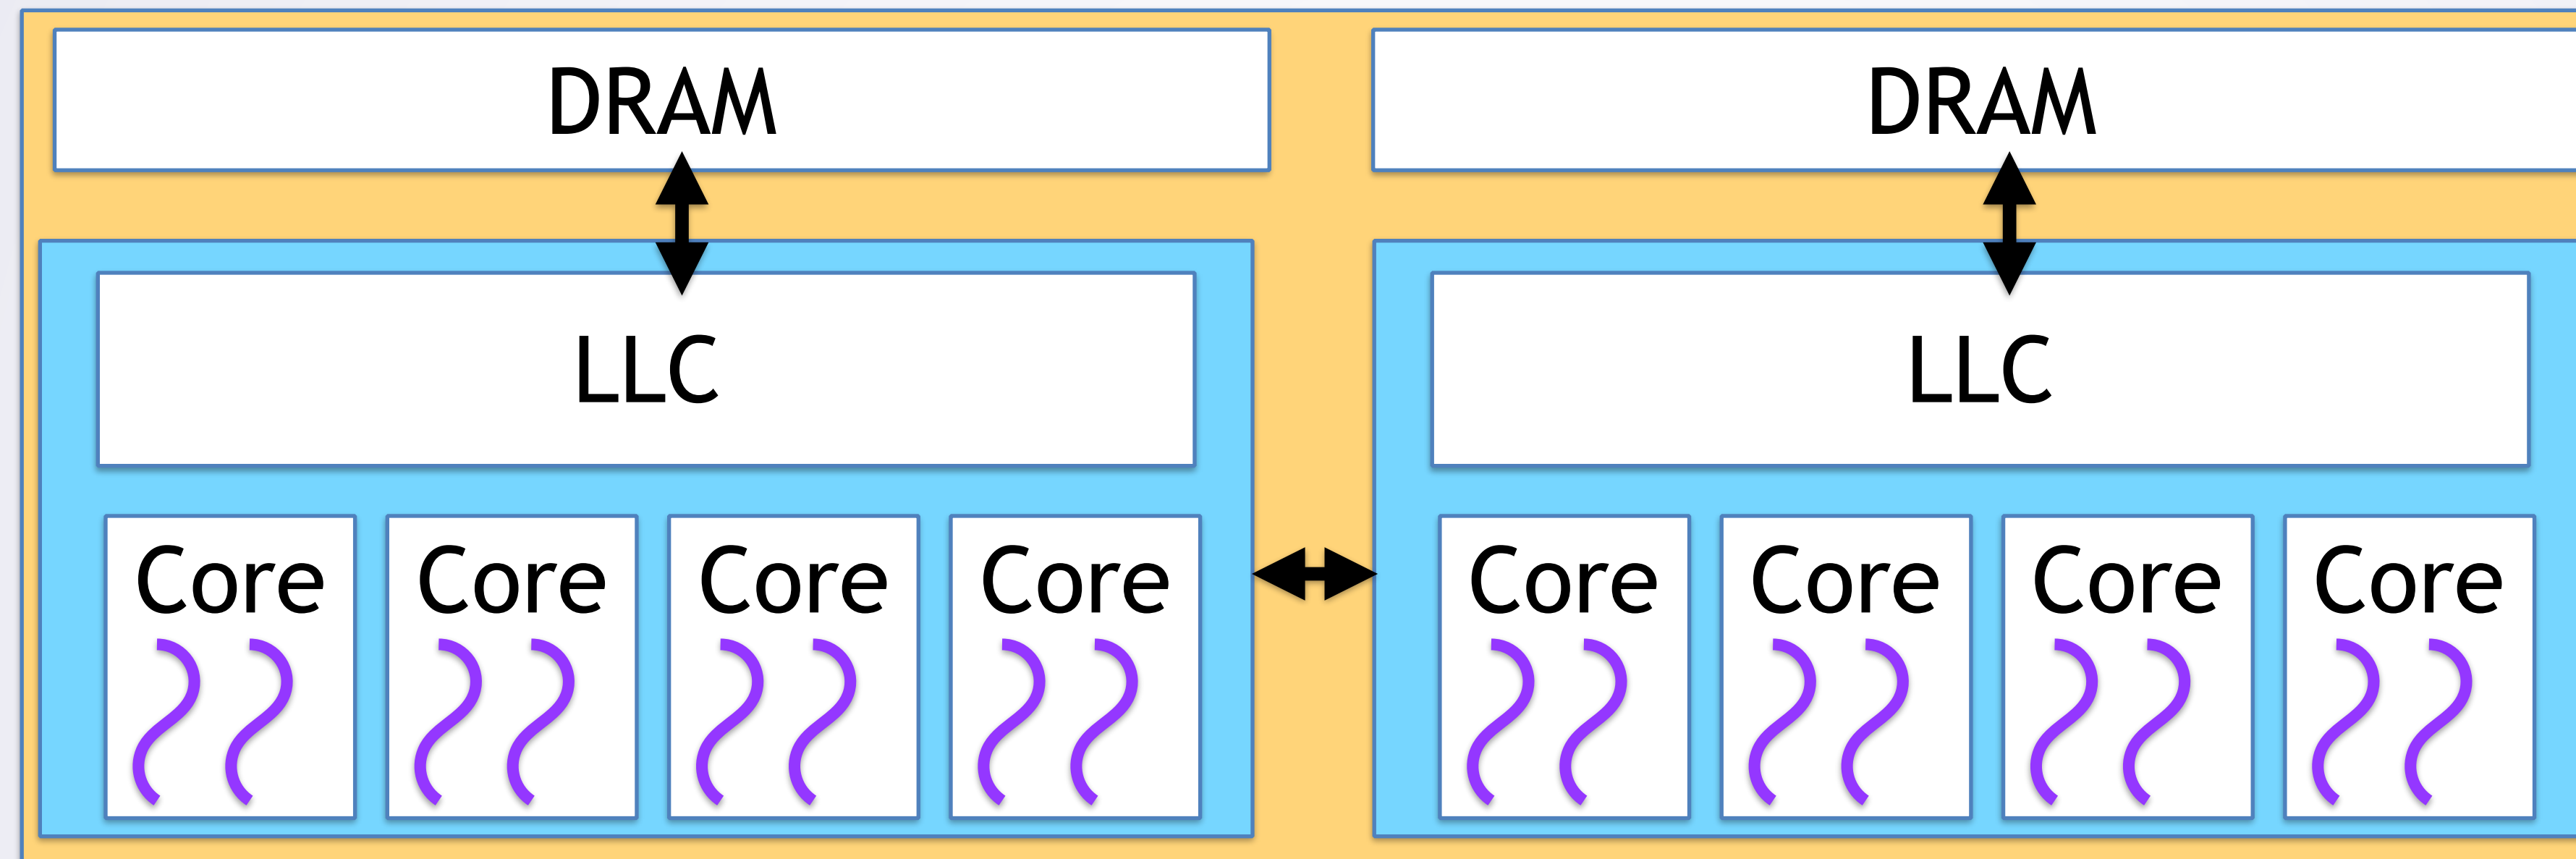
A core has 2 threads (2-way SMT)  
- Also private L1 + L2 caches (not shown)

# Hypothetical System



4 cores share an LLC  
- Connected by on chip interconnect


# Hypothetical System



We have a 2 socket node

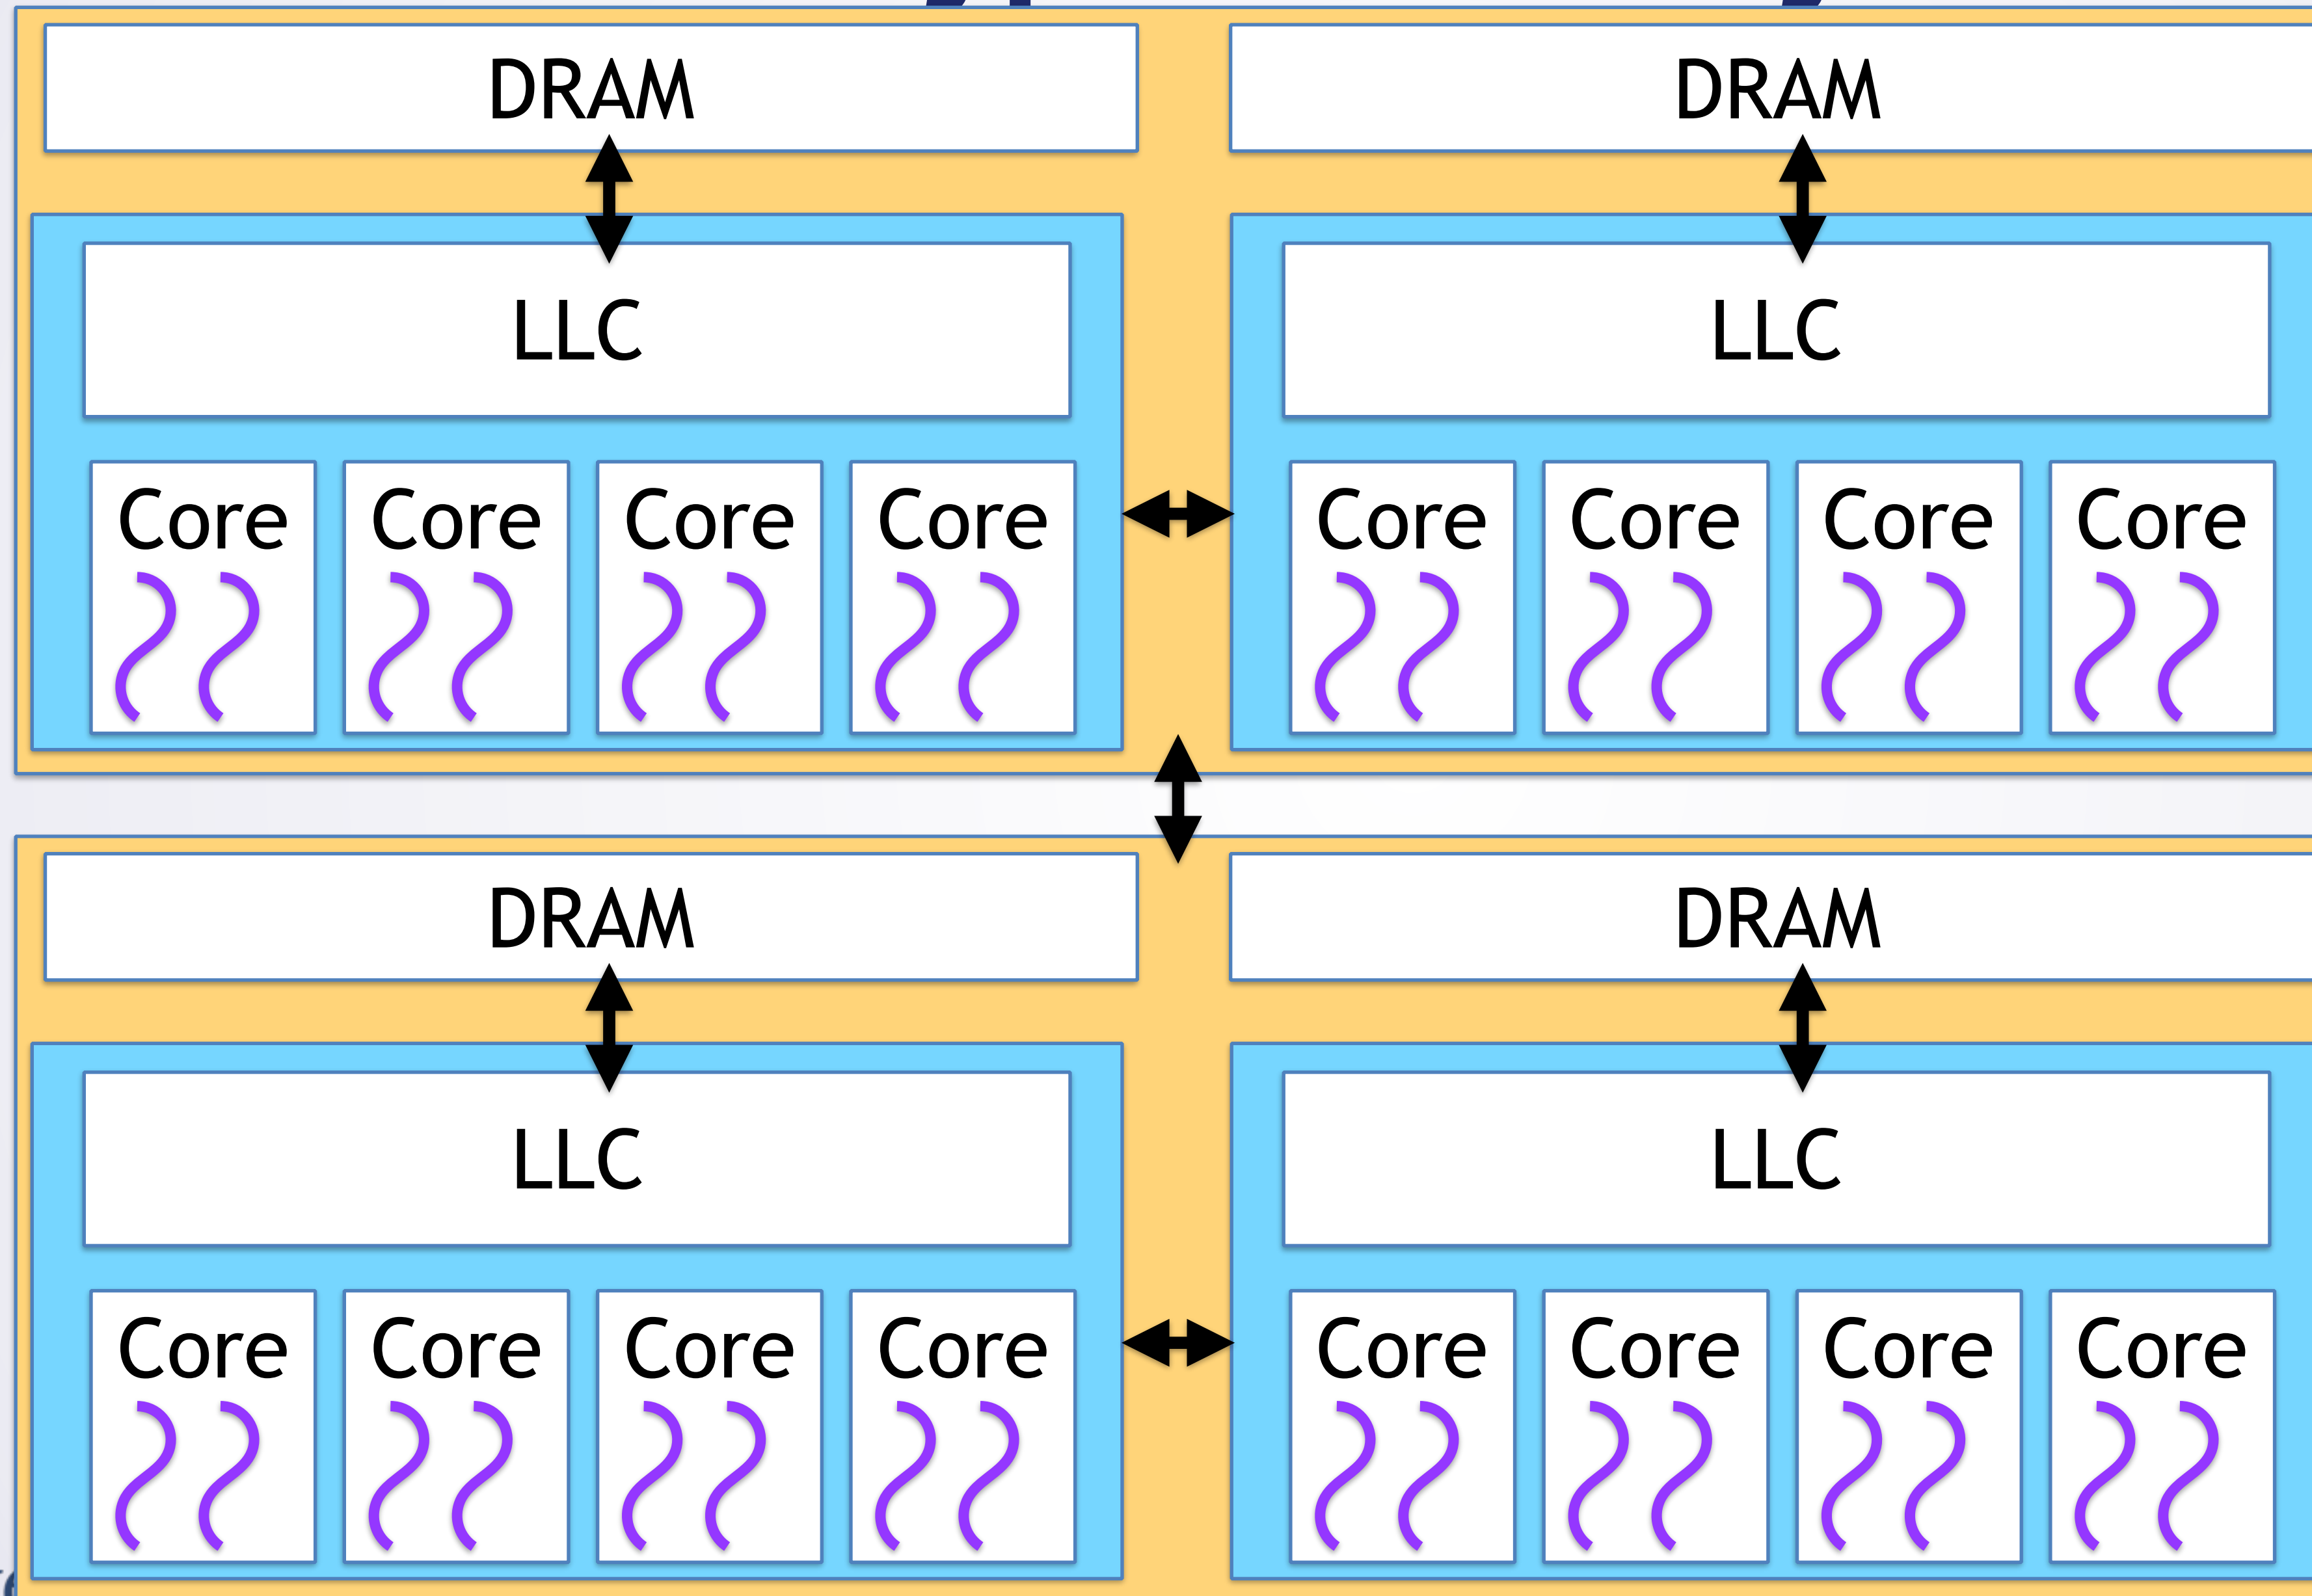
- Has 2 chips
- DRAM
- Also some IO devices (not shown)

# Hypothetical System



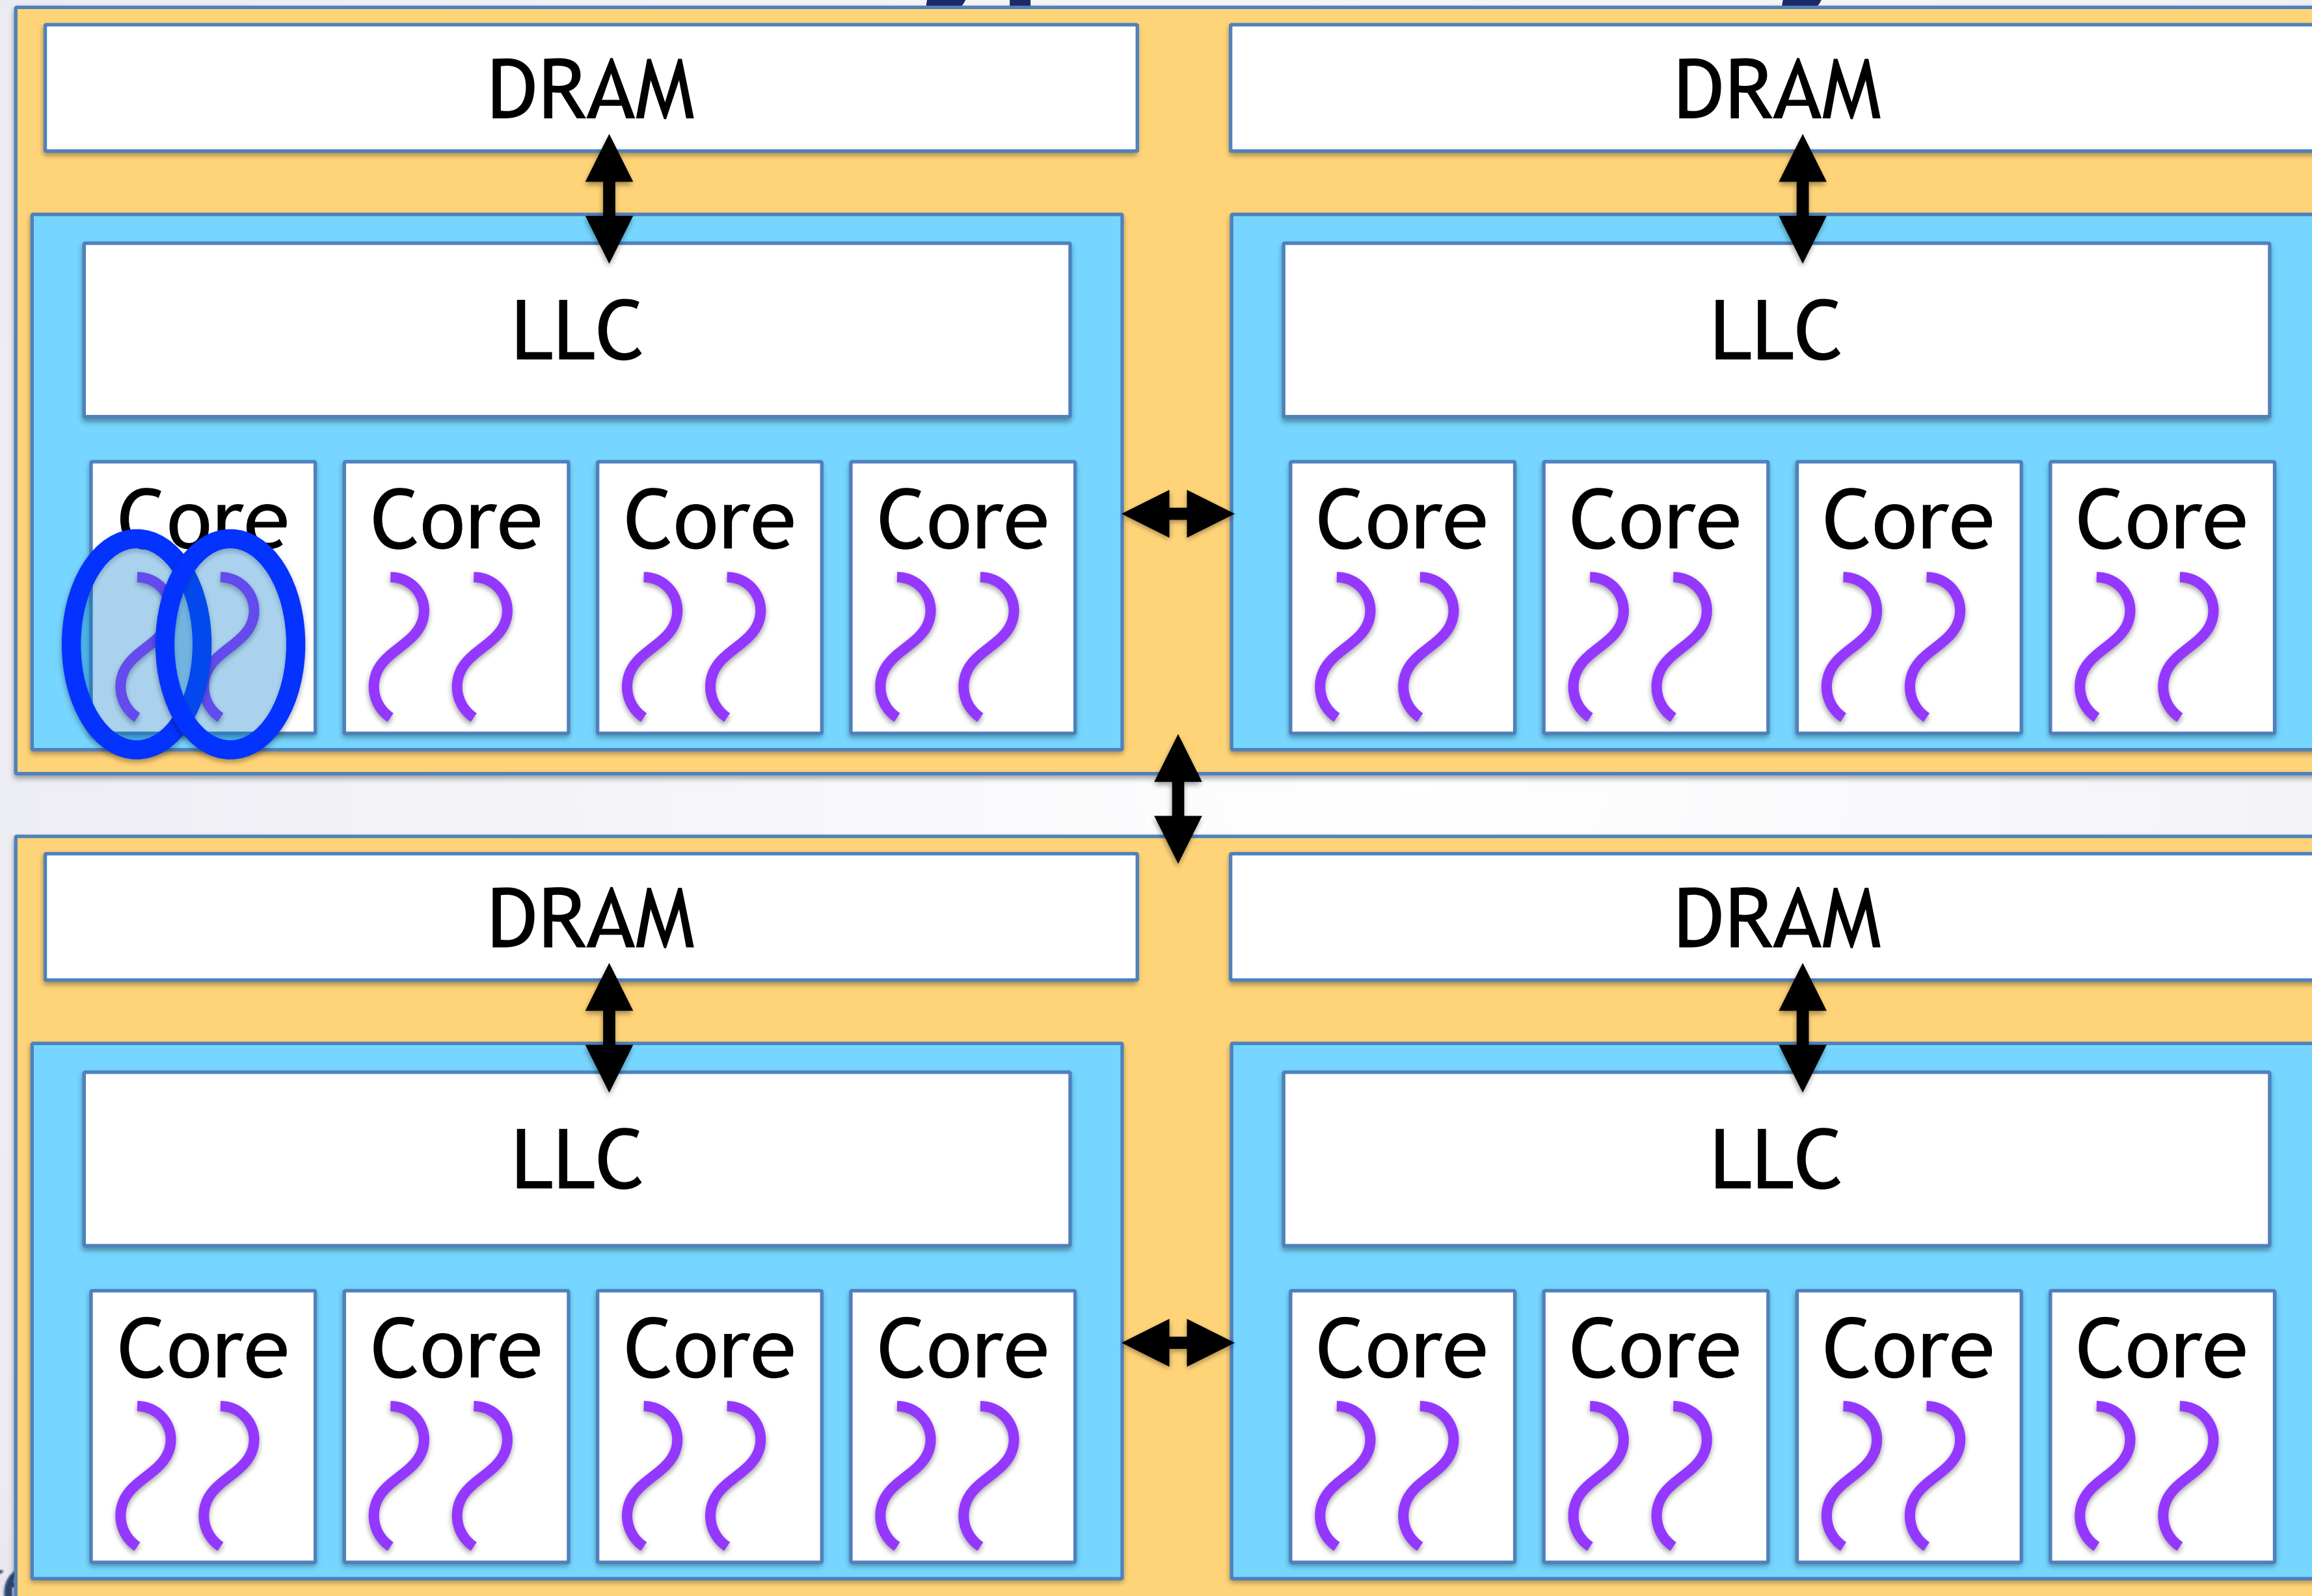
We have 2 nodes

# Hypothetical System



Suppose we have 2 requests: where best to run them?

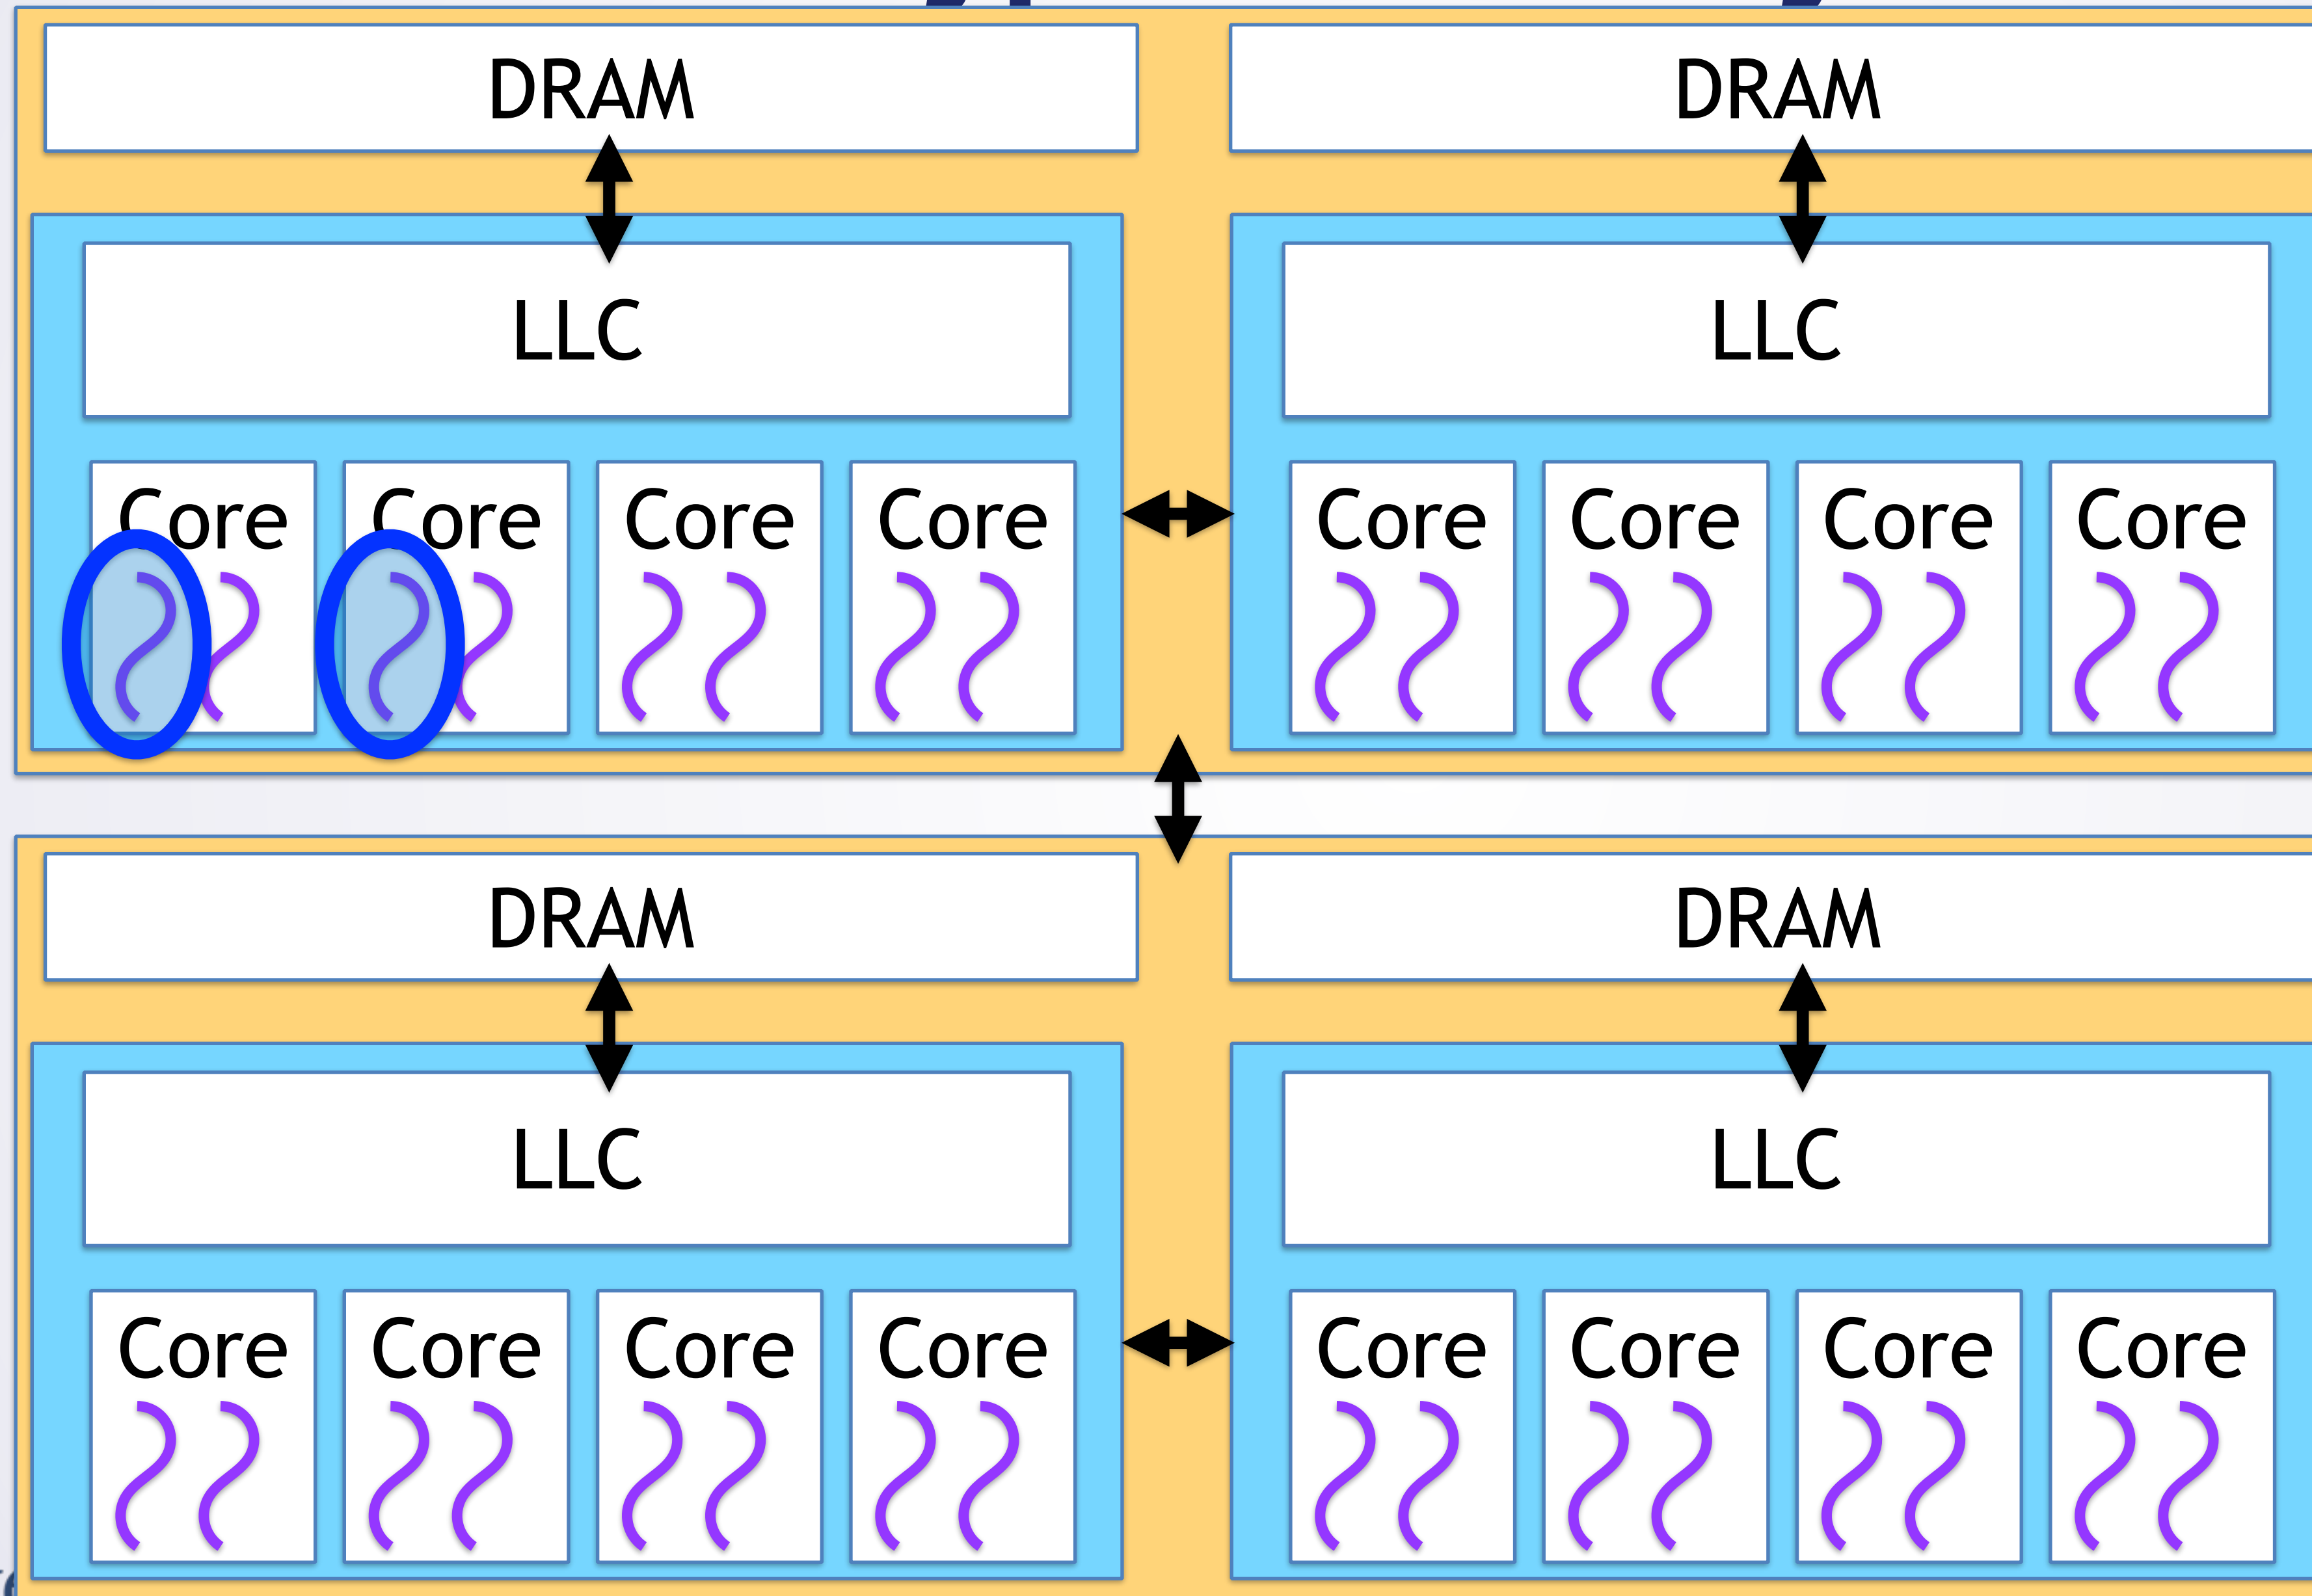
# Hypothetical System



Different threads  
on same core?

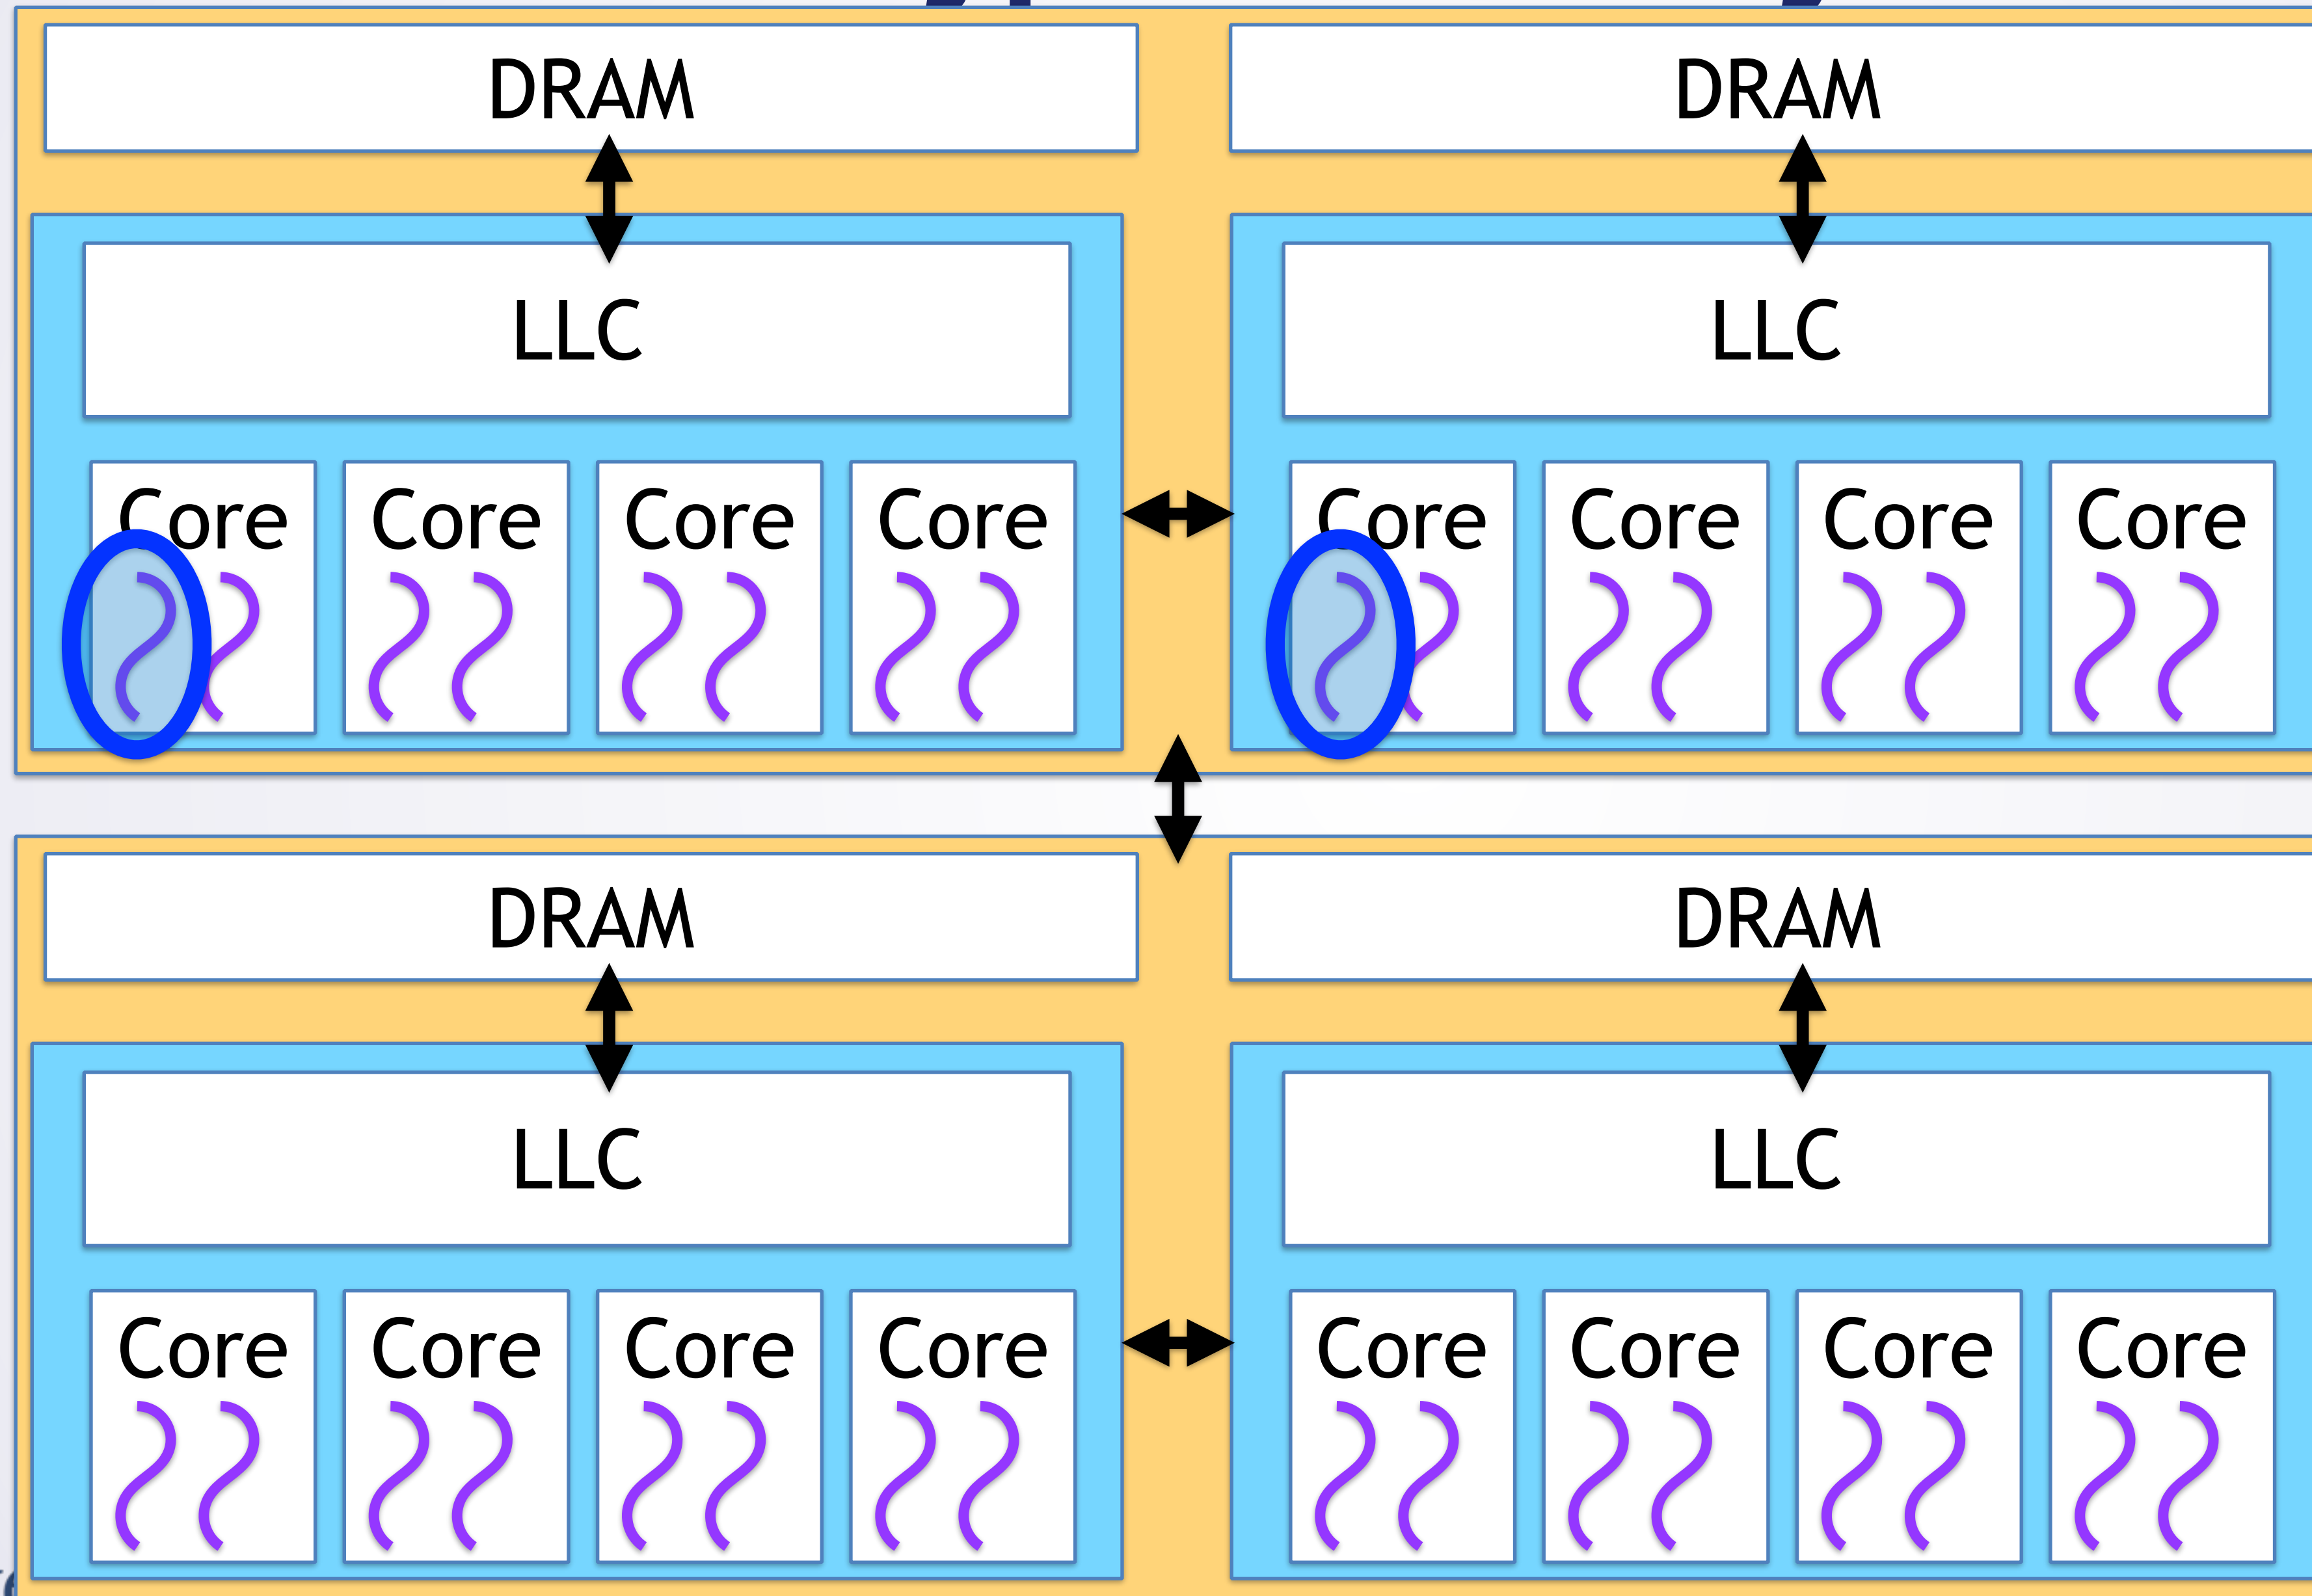


# Hypothetical System



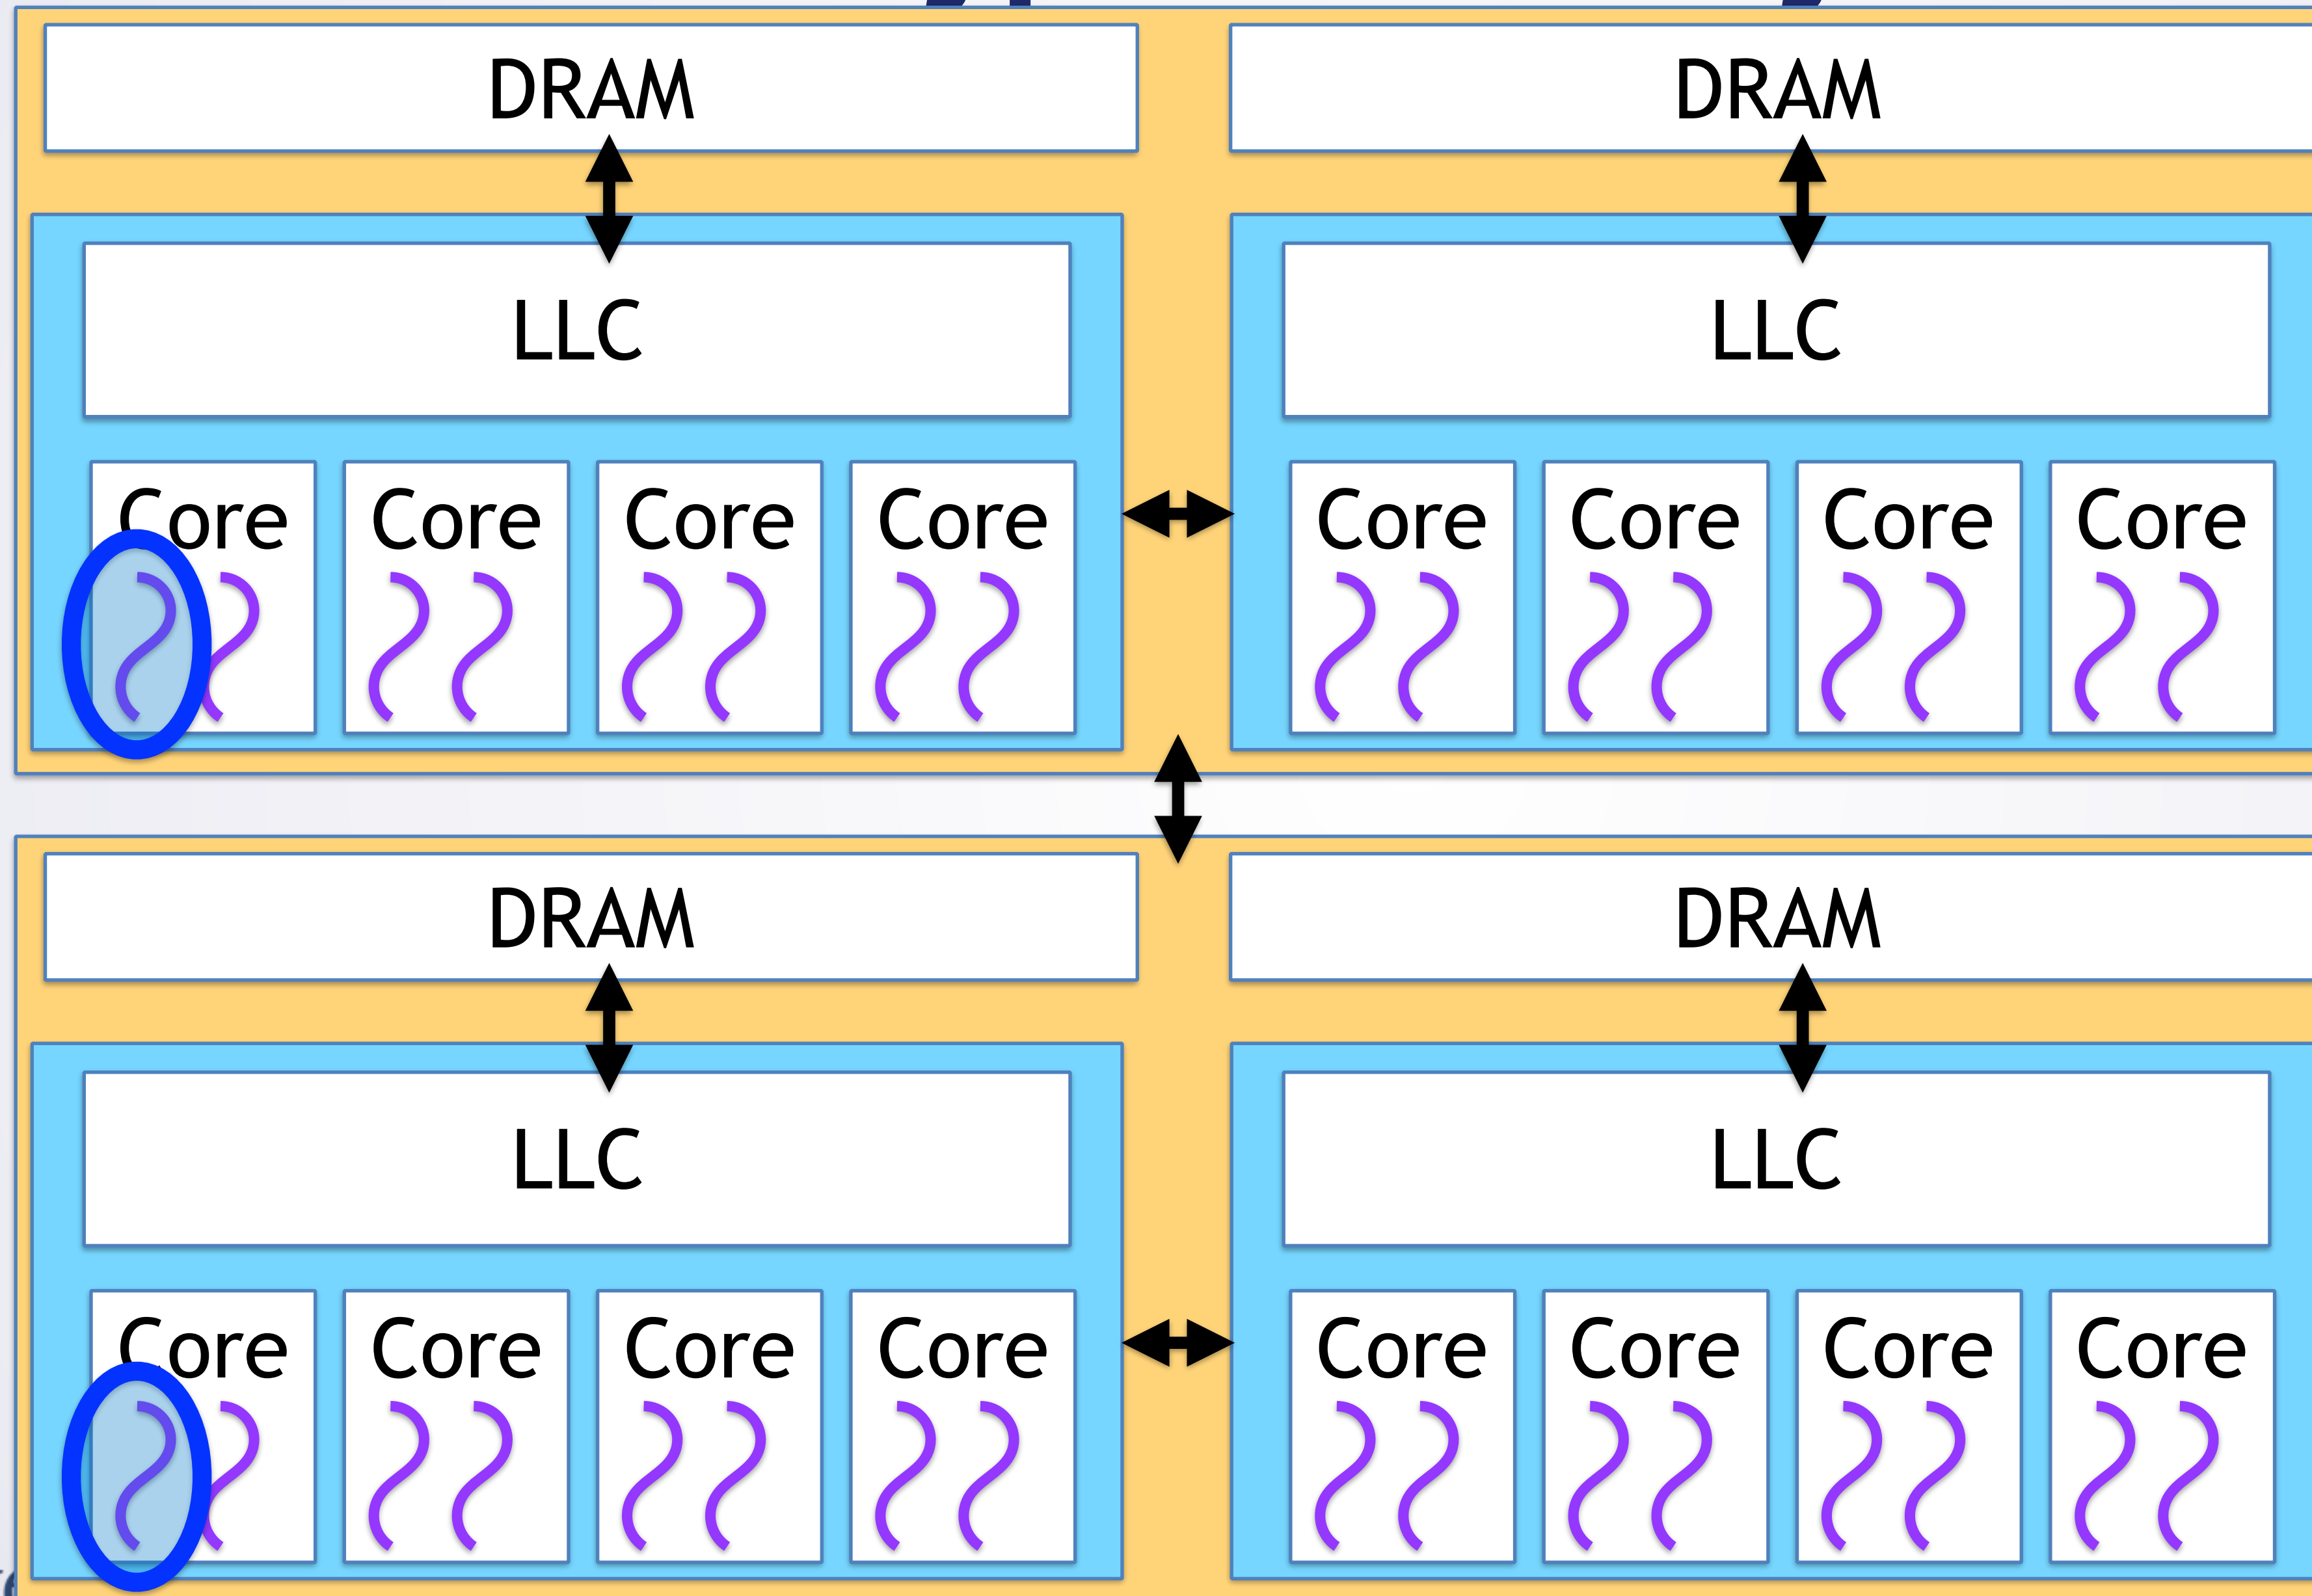
Different cores  
on same chip?

# Hypothetical System



Different chips on same node?

# Hypothetical System




Different nodes?

# How To Control Placement?

- Within a node: `sched_setaffinity`
  - Set mask of CPUs that a thread can run on
  - SMT contexts have different CPU identifiers
  - In pthreads, library wrapper: `pthread_setaffinity_np`
- Across nodes: depends..
  - Daemons running on each node? Direct requests to them
  - Startup/end new services? Software management
  - Load balancing becomes important here

# Tradeoff: Contention vs Locality



- Trade off:
  - Contend for shared resources?
  - Longer/slower communication?



# Tradeoff: Contention vs Locality

Increasing Communication Latency



Same Core

Same Chip

Same Node

Different Node

Loads + Stores  
Same Cache  
1s of cycles

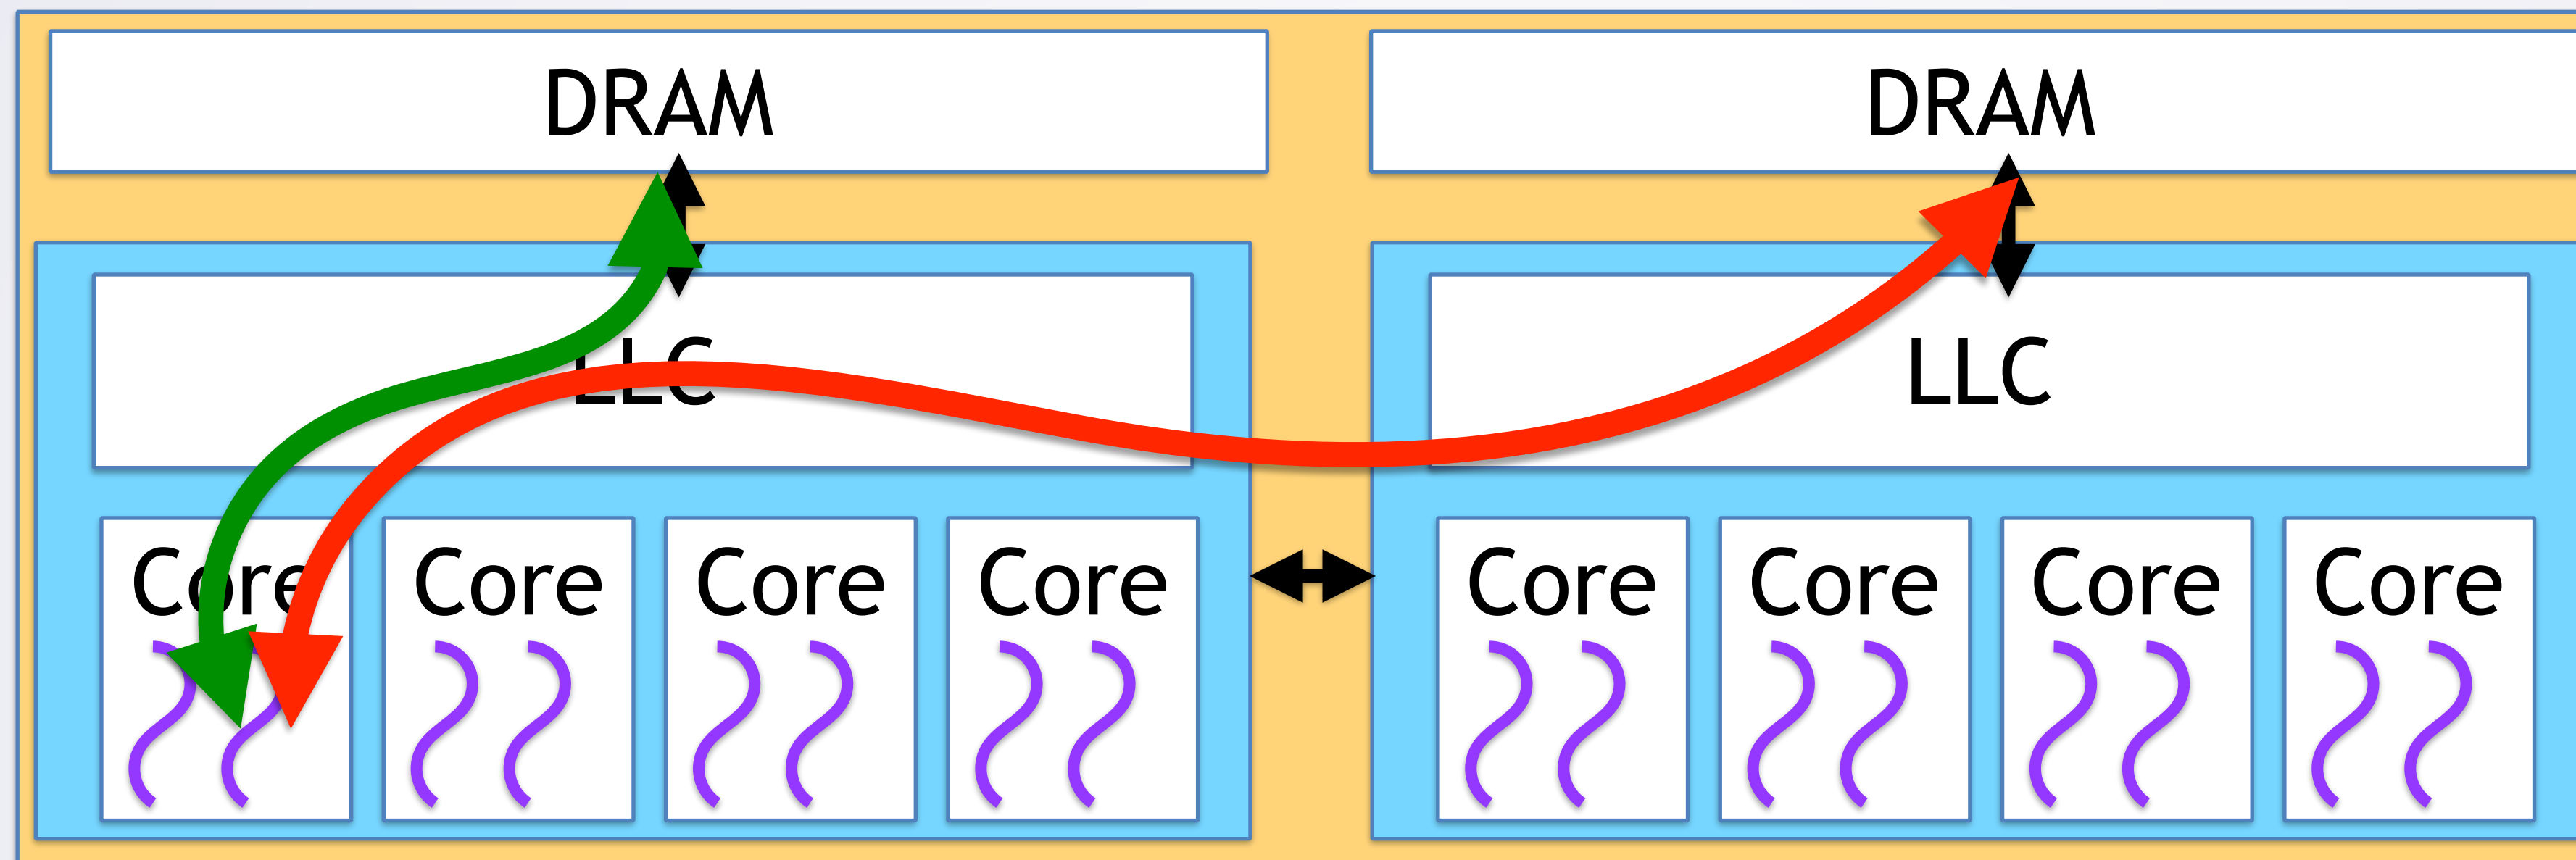
Loads + Stores  
On Chip Coherence  
10s of cycles

Loads + Stores  
Off Chip Coherence  
100s cycles

IO Operations  
Network  
Ks-Ms of cycles

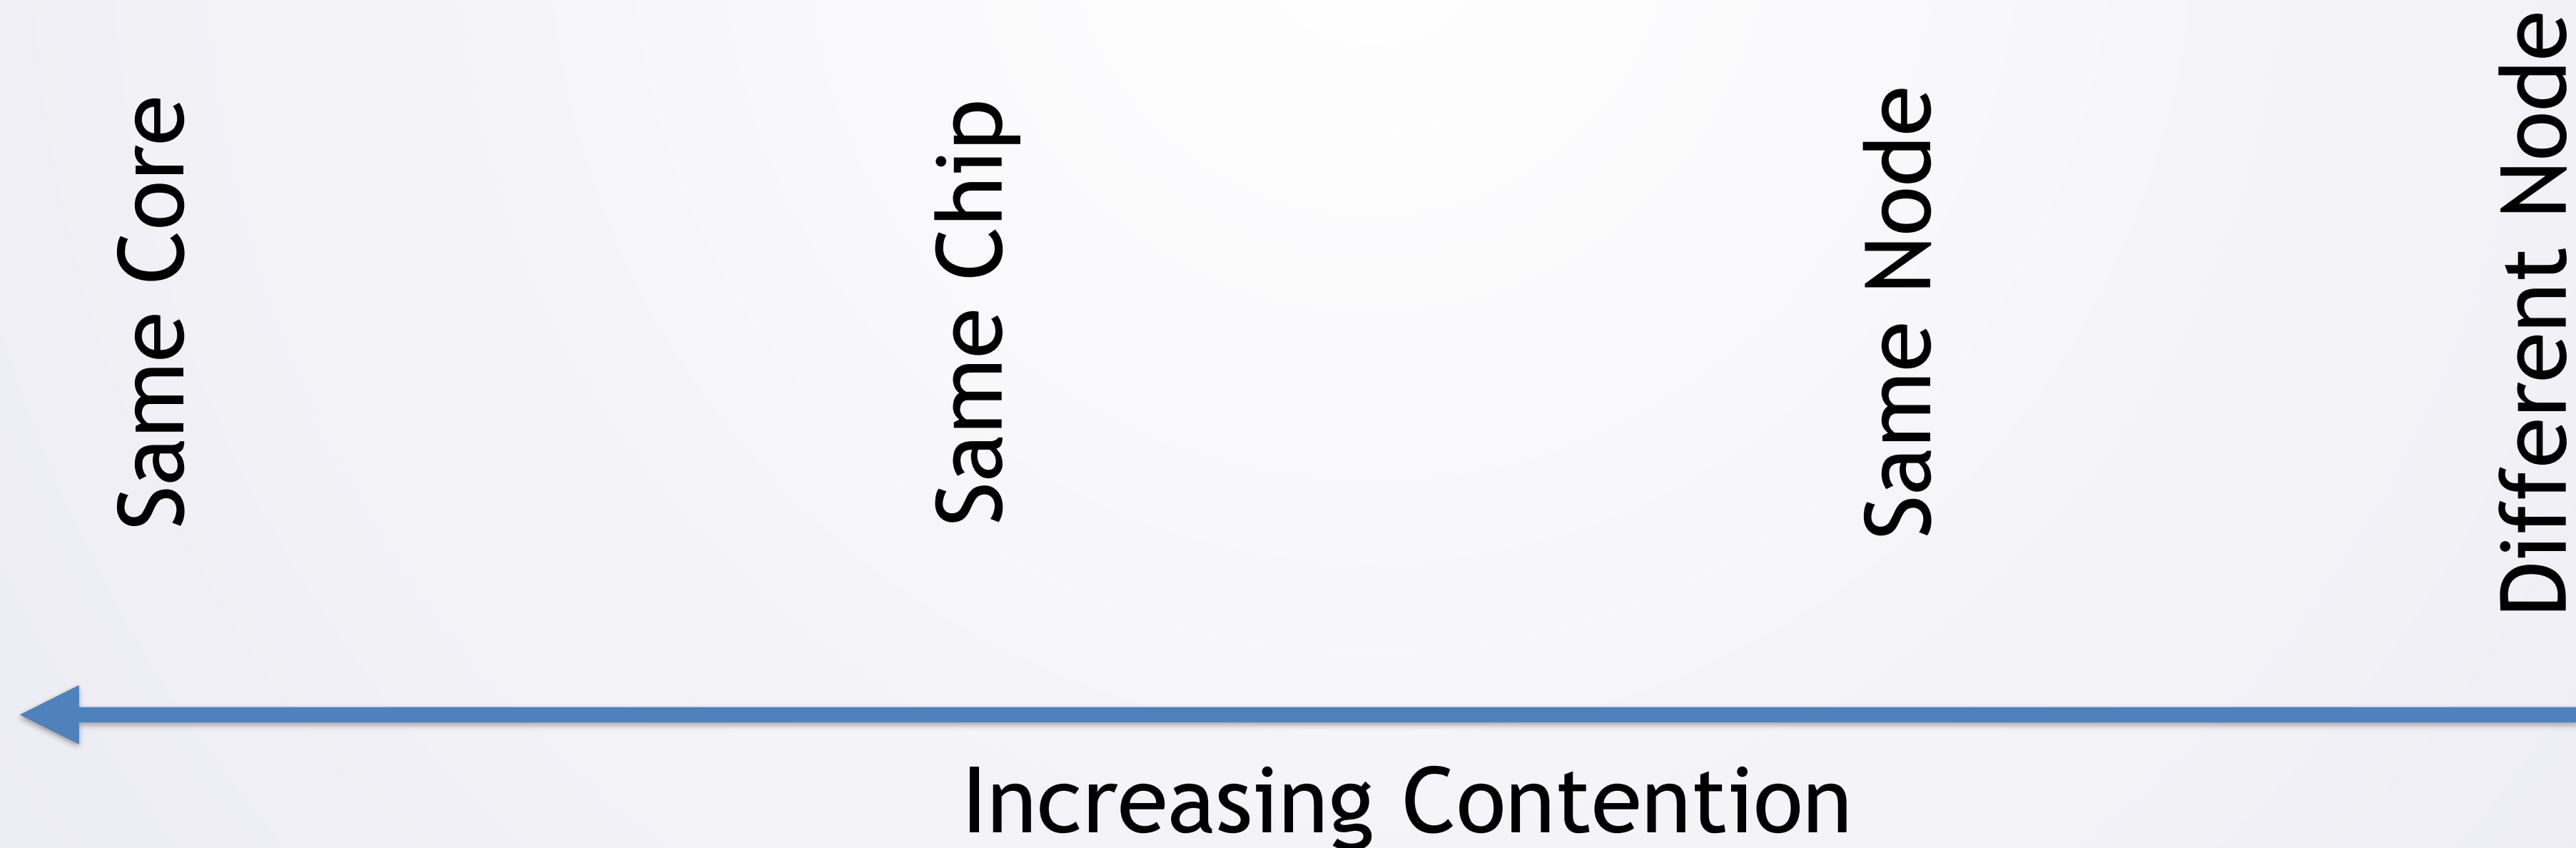


# NUMA

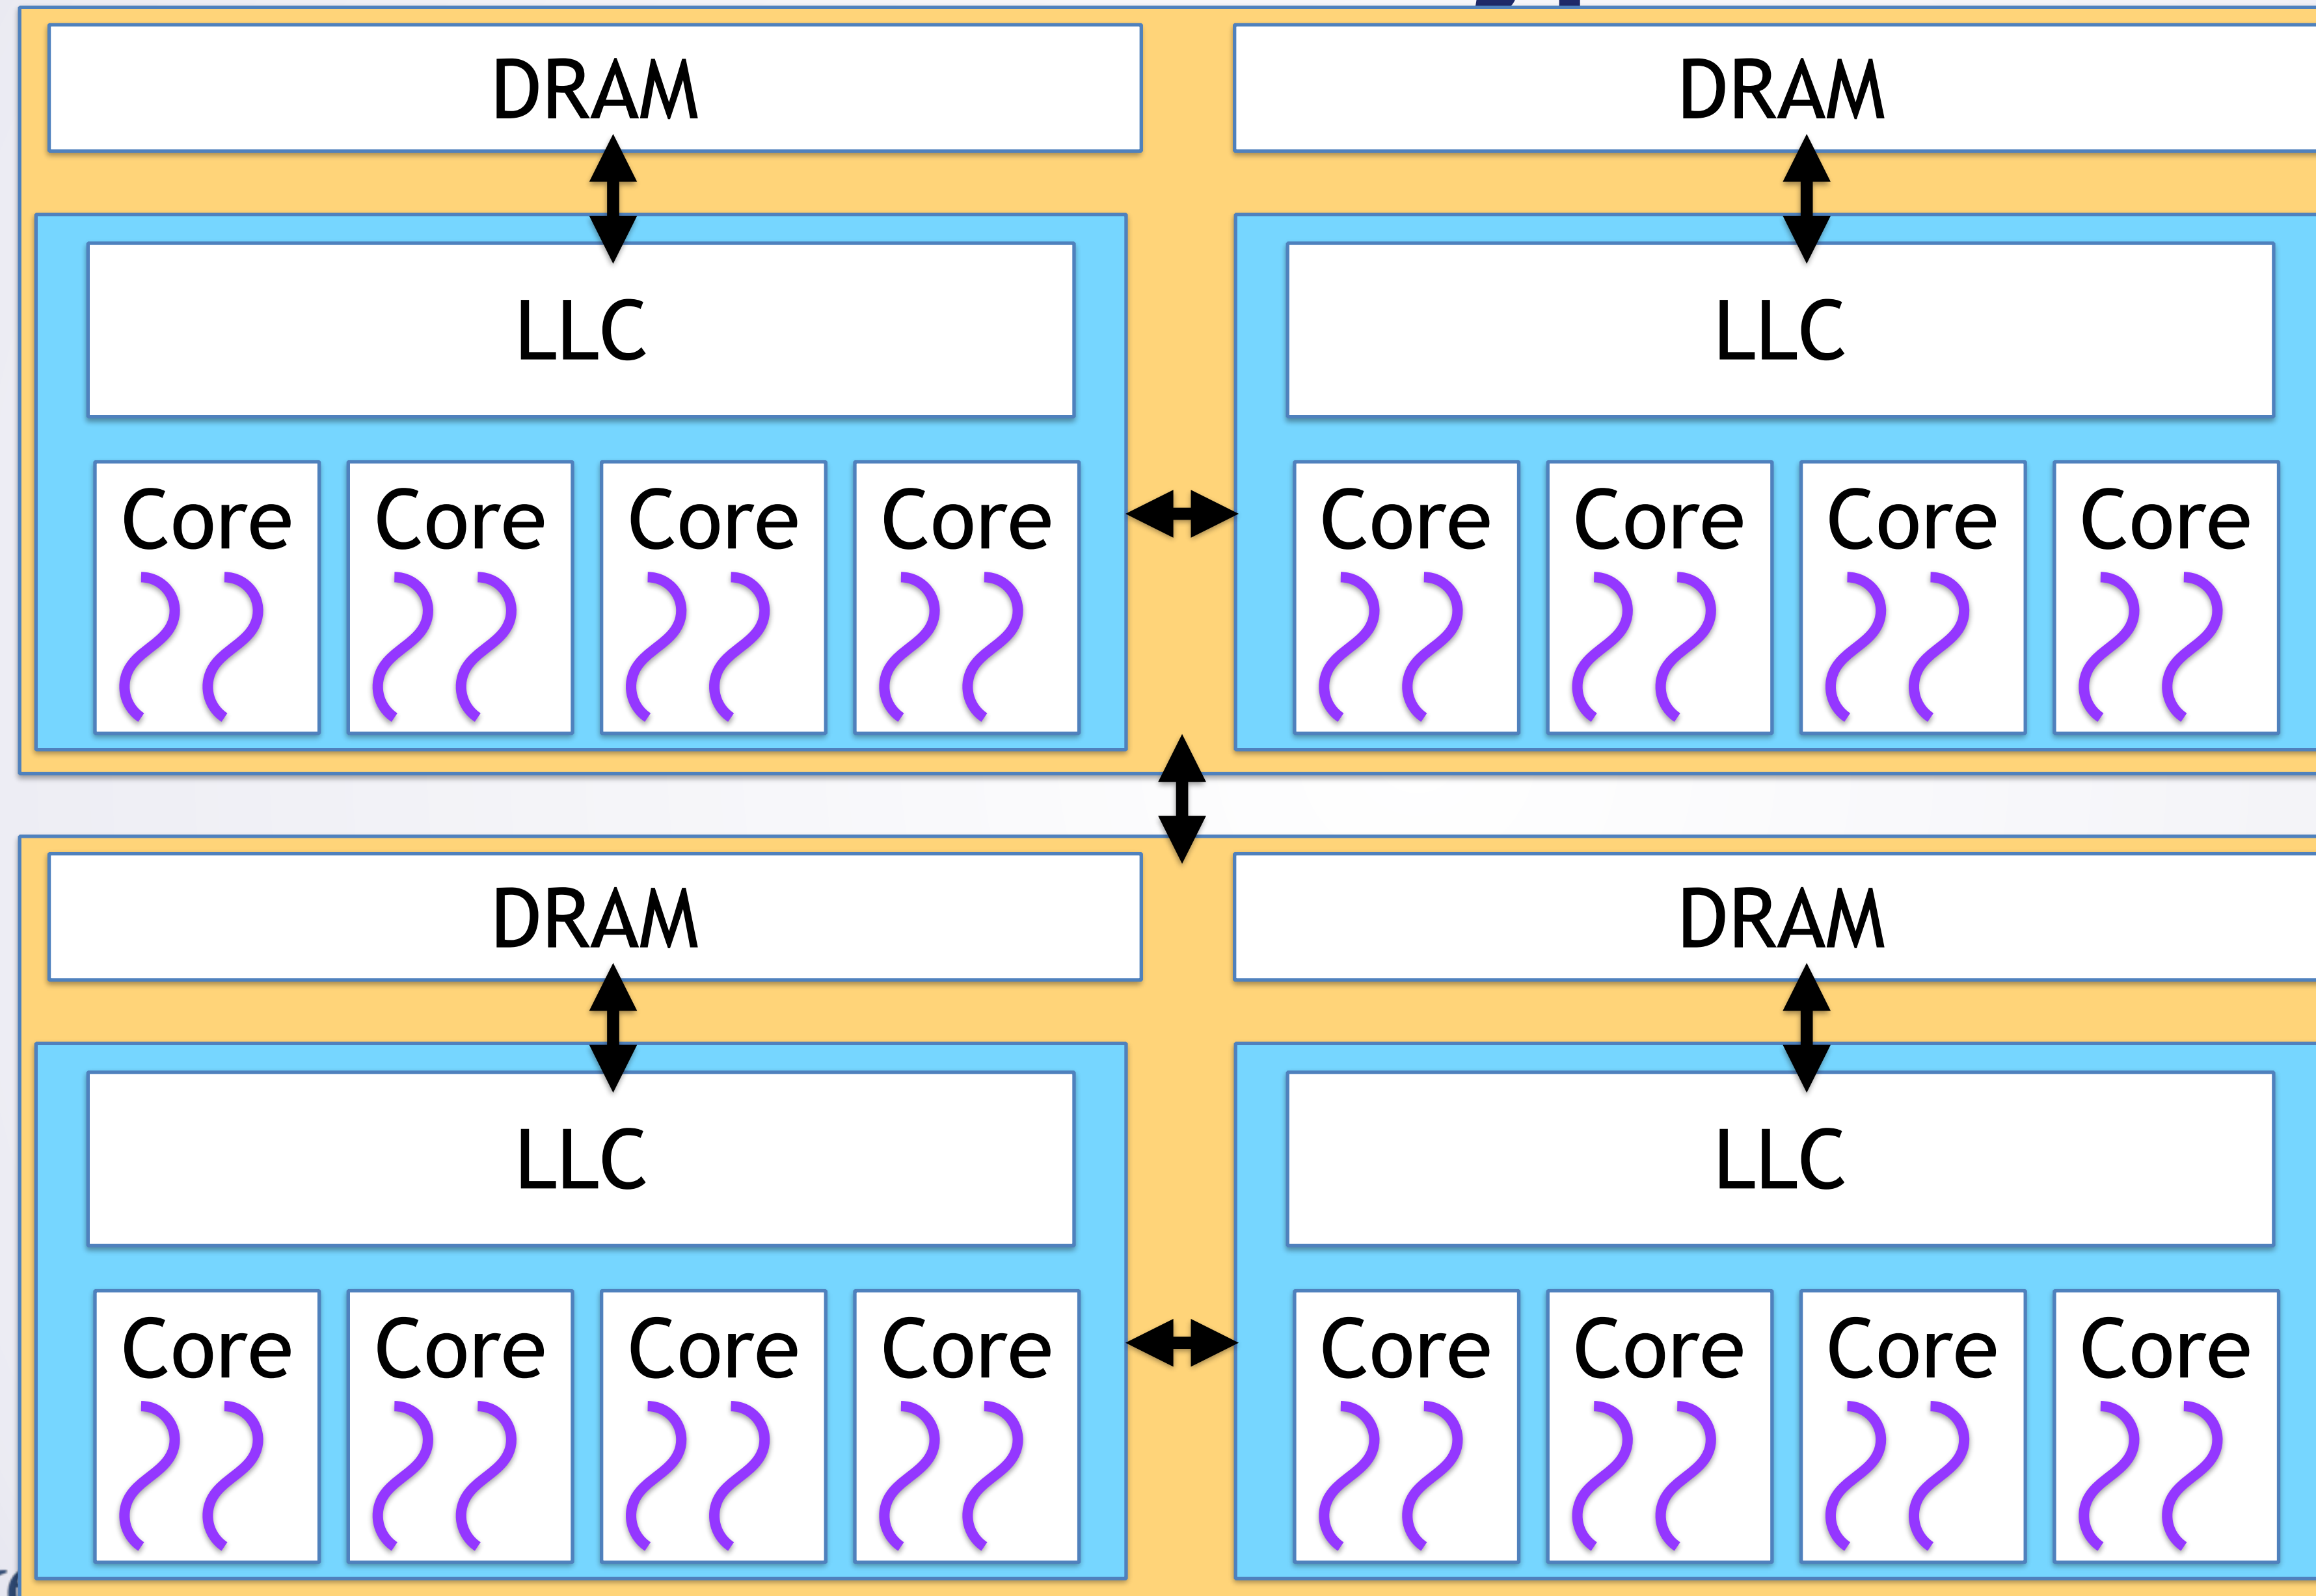


- Non Uniform Memory Access (NUMA—technically, ccNUMA)
  - Memory latency differs depending on physical address
  - `migrate_pages`, `mbind`: control physical memory placement

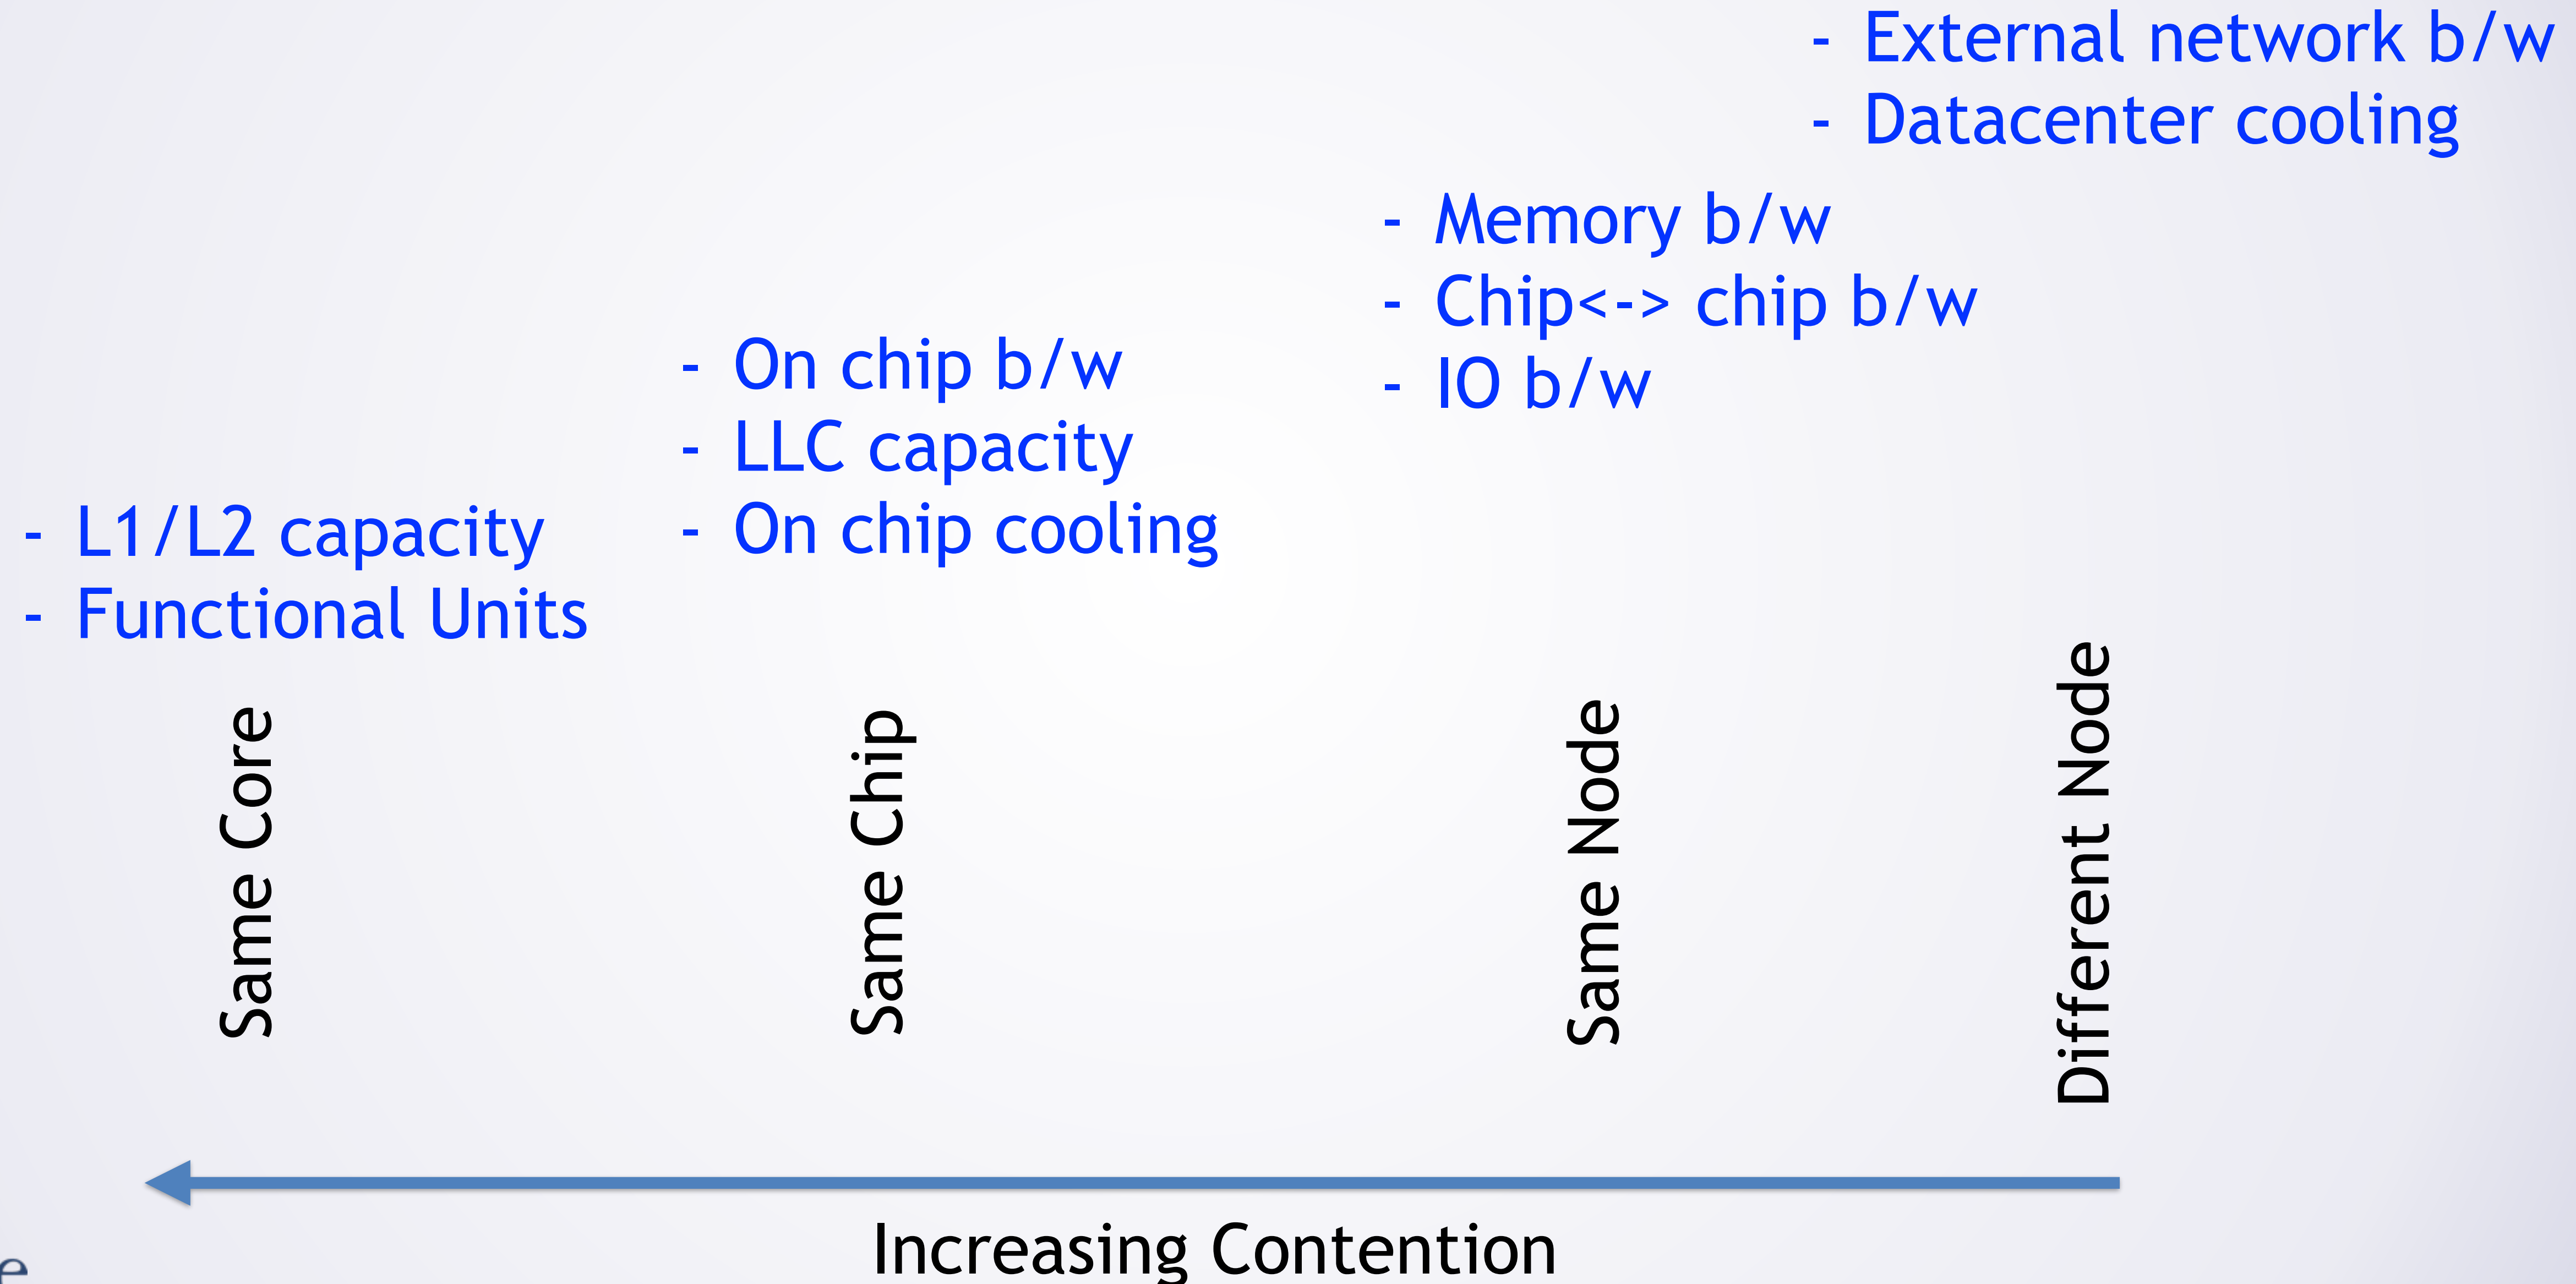
# Tradeoff: Contention vs Locality



# Re-examine Our Hypothetical System



# Tradeoff: Contention vs Locality



# Interactions Between Resource Contention

- Suppose two threads need + are sensitive to:
  - LLC Capacity
  - Memory bandwidth
- What happens when we run them together?

# Interactions Between Resource Contention

- Suppose two threads need + are sensitive to:
  - LLC Capacity
  - Memory bandwidth
- What happens when we run them together?
  - Contention for LLC -> more cache misses
    - Slows down program, but also...



# Interactions Between Resource Contention

- Suppose two threads need + are sensitive to:
  - LLC Capacity
  - Memory bandwidth
- What happens when we run them together?
  - Contention for LLC -> more cache misses
    - Slows down program, but also...
  - Increases memory bandwidth demands
    - Which we already need and are contending for :(

# Interactions Between Resource Contention

- Suppose two threads need + are sensitive to:
  - LLC Capacity
  - Memory bandwidth
- What happens when we run them together?
  - Contention for LLC -> more cache misses
    - Slows down program, but also...
  - Increases memory bandwidth demands
    - Which we already need and are contending for :(
- Interactions can make contention even worse!
  - Is there a flip side?

# Improved Utilization

- Can improve utilization of resources
  - One thread executes while another stalls
  - One thread uses FUs that the other does not need
  - Pair large cache footprint with small cache footprint
  - Shared code/data: one copy in cache

# Performance/Scalability 1

- So what can we do?
  - Profile code and understand its behavior/resource usage
  - Optimize code to improve its performance
  - Transform code to improve resource usage (e.g. cache space)
  - Pair threads with complementary resource usage

# Performance/Scalability 1

- So what can we do?
  - Profile code and understand its behavior/resource usage
  - Optimize code to improve its performance
  - Transform code to improve resource usage (e.g. cache space)
  - Pair threads with complementary resource usage
- Sounds complicated?
  - Learn more about hardware (e.g., ECE 552)
  - Take Performance/Optimization/Parallelism



# Impediments to Scalability

- Shared Hardware
  - Functional Units
  - Caches
  - Memory Bandwidth
  - IO Bandwidth
  - ...
- Data Movement
  - From one core to another
- Blocking
  - Blocking IO
  - Locks (and other synchronization)

**Let's talk about this next**



# Never Block

- Critical principle: never block
  - Why not?

# Never Block

- Critical principle: never block
  - Why not?
- Can't we just throw more threads at it?
  - One thread per request (or even a few per request)
  - Just block whenever you want

# Never Block

- Critical principle: never block
  - Why not?
- Can't we just throw more threads at it?
  - One thread per request (or even a few per request)
  - Just block whenever you want
- Nice in theory, but has overheads
  - Context switching takes time
  - Switching threads reduces temporal locality
    - Threads not blocked? May thrash if too many
  - Threads use resources

# Non-Blocking IO

- IO operations often block (we never want to block)
  - Can use non-blocking IO

# Non-Blocking IO

- IO operations often block (we never want to block)
  - Can use non-blocking IO

- Set FD to non-blocking using fcntl:

```
int x = fcntl(fd, F_GETFL, 0);  
x |= O_NONBLOCK;  
fcntl(fd, F_SETFL, x);
```

- Now reads/writes/etc won't block
  - Just return immediately if can't perform IO immediately
  - Note: not magic
    - **ONLY** means that IO operation returns without waiting

# Non-Blocking IO: Continued

```
int x = read (fd, buffer, size);  
if (x < 0) {  
    if (errno == EAGAIN) {  
        //no data available  
    }  
    else {  
        //error  
    }  
}
```



# Non-Blocking IO: Continued

```
while (size > 0) {  
    int x = read (fd, buffer, size);  
    if (x < 0) {  
        if (errno == EAGAIN) {  
            //no data available  
        }  
        else {  
            //error  
        }  
    }  
    else {  
        buffer += x;  
        size -= x;  
    }  
}
```

What if we just wrap this up in a while loop?

# Non-Blocking IO: Continued

```
while (size > 0) {  
    int x = read (fd, buffer, size);  
    if (x < 0) {  
        if (errno == EAGAIN) {  
            //no data available  
        }  
        else {  
            //error  
        }  
    }  
    else {  
        buffer += x;  
        size -= x;  
    }  
}
```

What if we just wrap this up in a while loop?

Now we just made this blocking!

We are just doing the blocking ourselves...

# Busy Wait

- This approach is **worse** than blocking IO
  - Why?

# Busy Wait

- This approach is **worse** than blocking IO
  - Why?
- Busy waiting
  - Code is "actively" doing nothing
  - Keeping CPU busy, consuming power, contending with other threads
- Blocking IO:
  - At least OS will put thread to sleep while it waits

# So What Do We Do?

- Need to do **something else** while we wait
  - Like what?

# So What Do We Do?

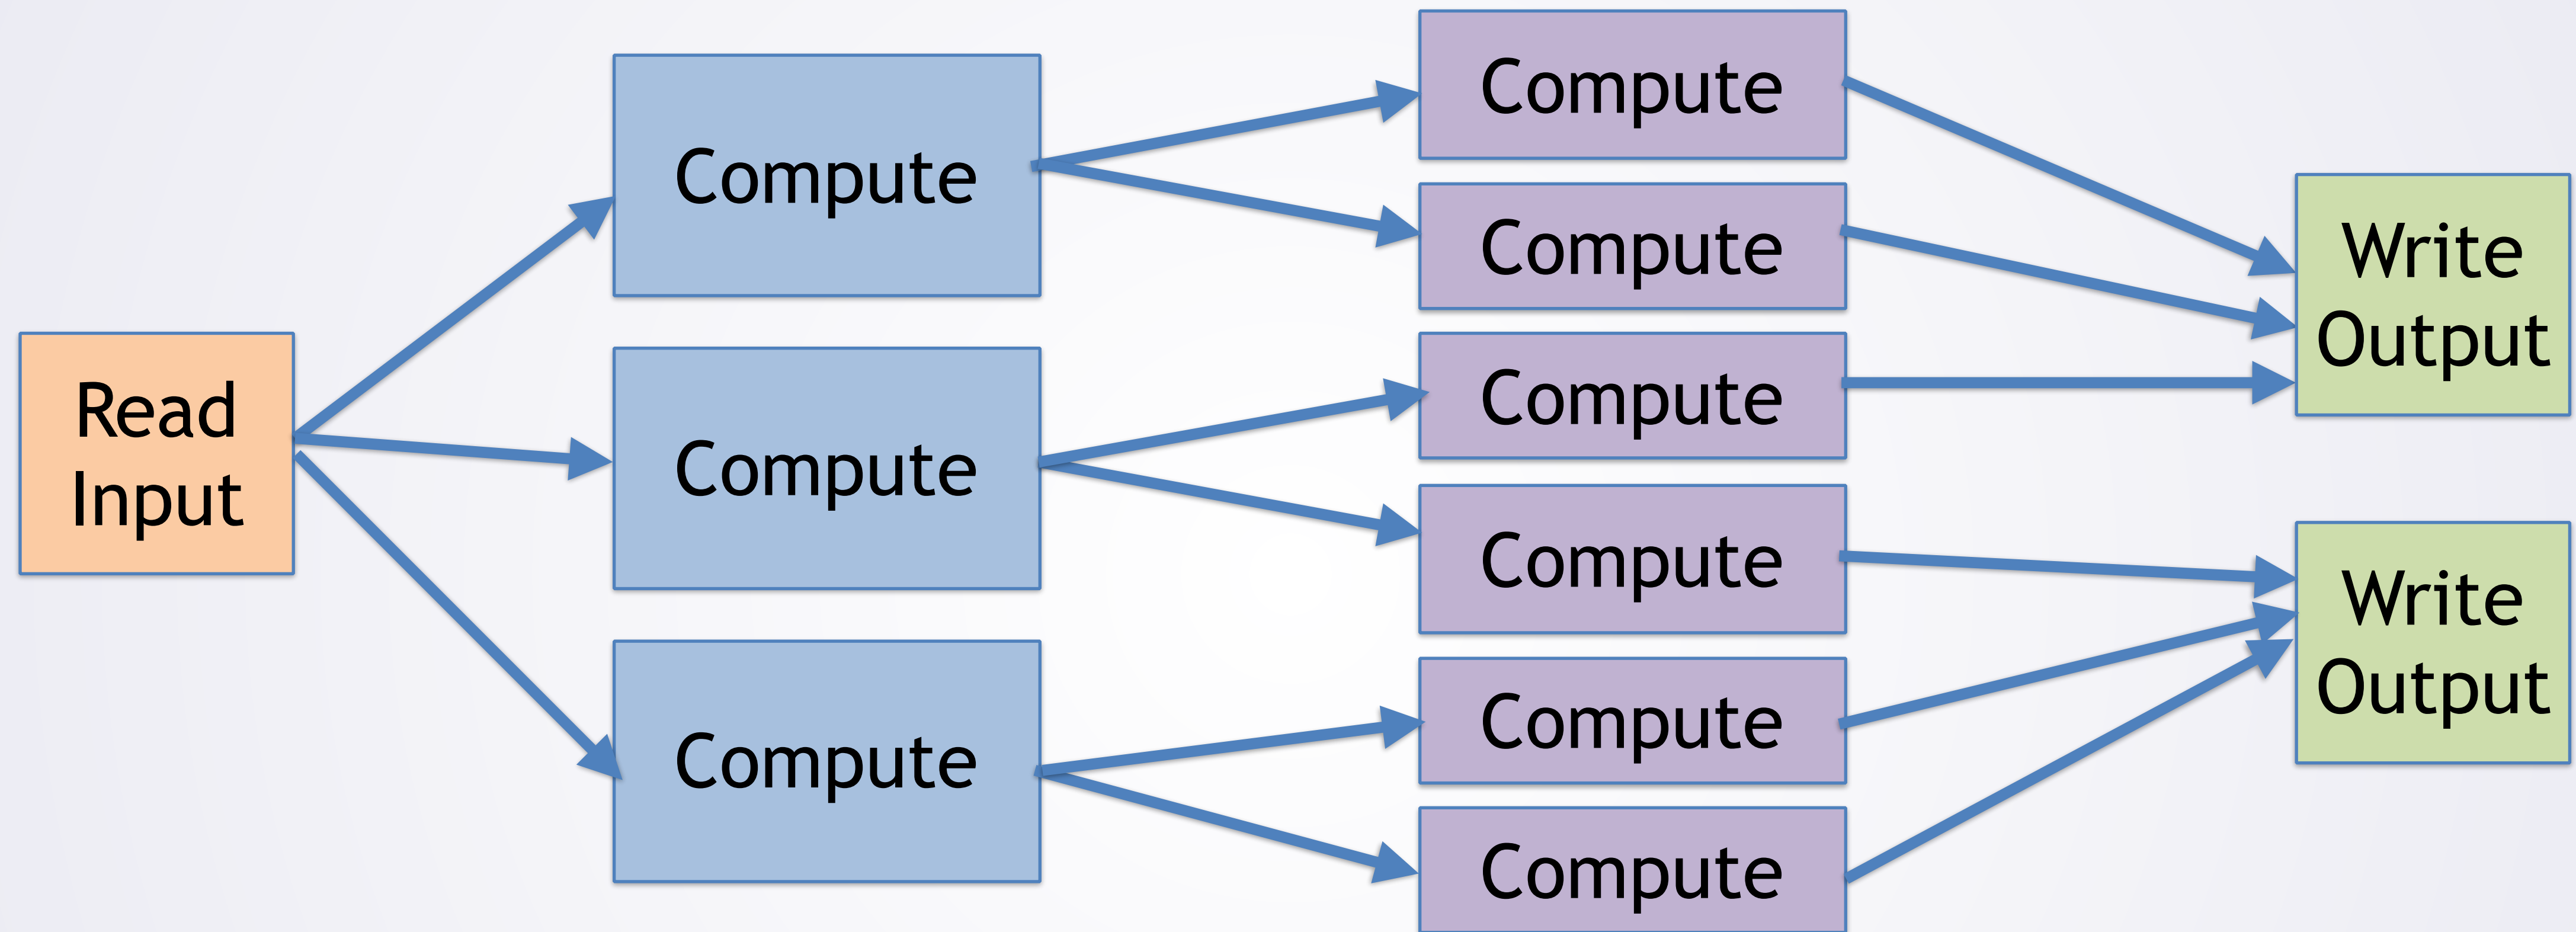
- Need to do **something else** while we wait
  - Like what?
- It depends....
  - On what?



# So What Do We Do?

- Need to do **something else** while we wait
  - Like what?
- It depends....
  - On what?
- On what our server does
- On what the demands on it are
- On the model of parallelism we are using
  - Who can name some models of parallelism? [AoP Ch 28 review]

# Pipeline Parallelism



- When would this be appropriate?
- What do our IO threads do for "something else"?

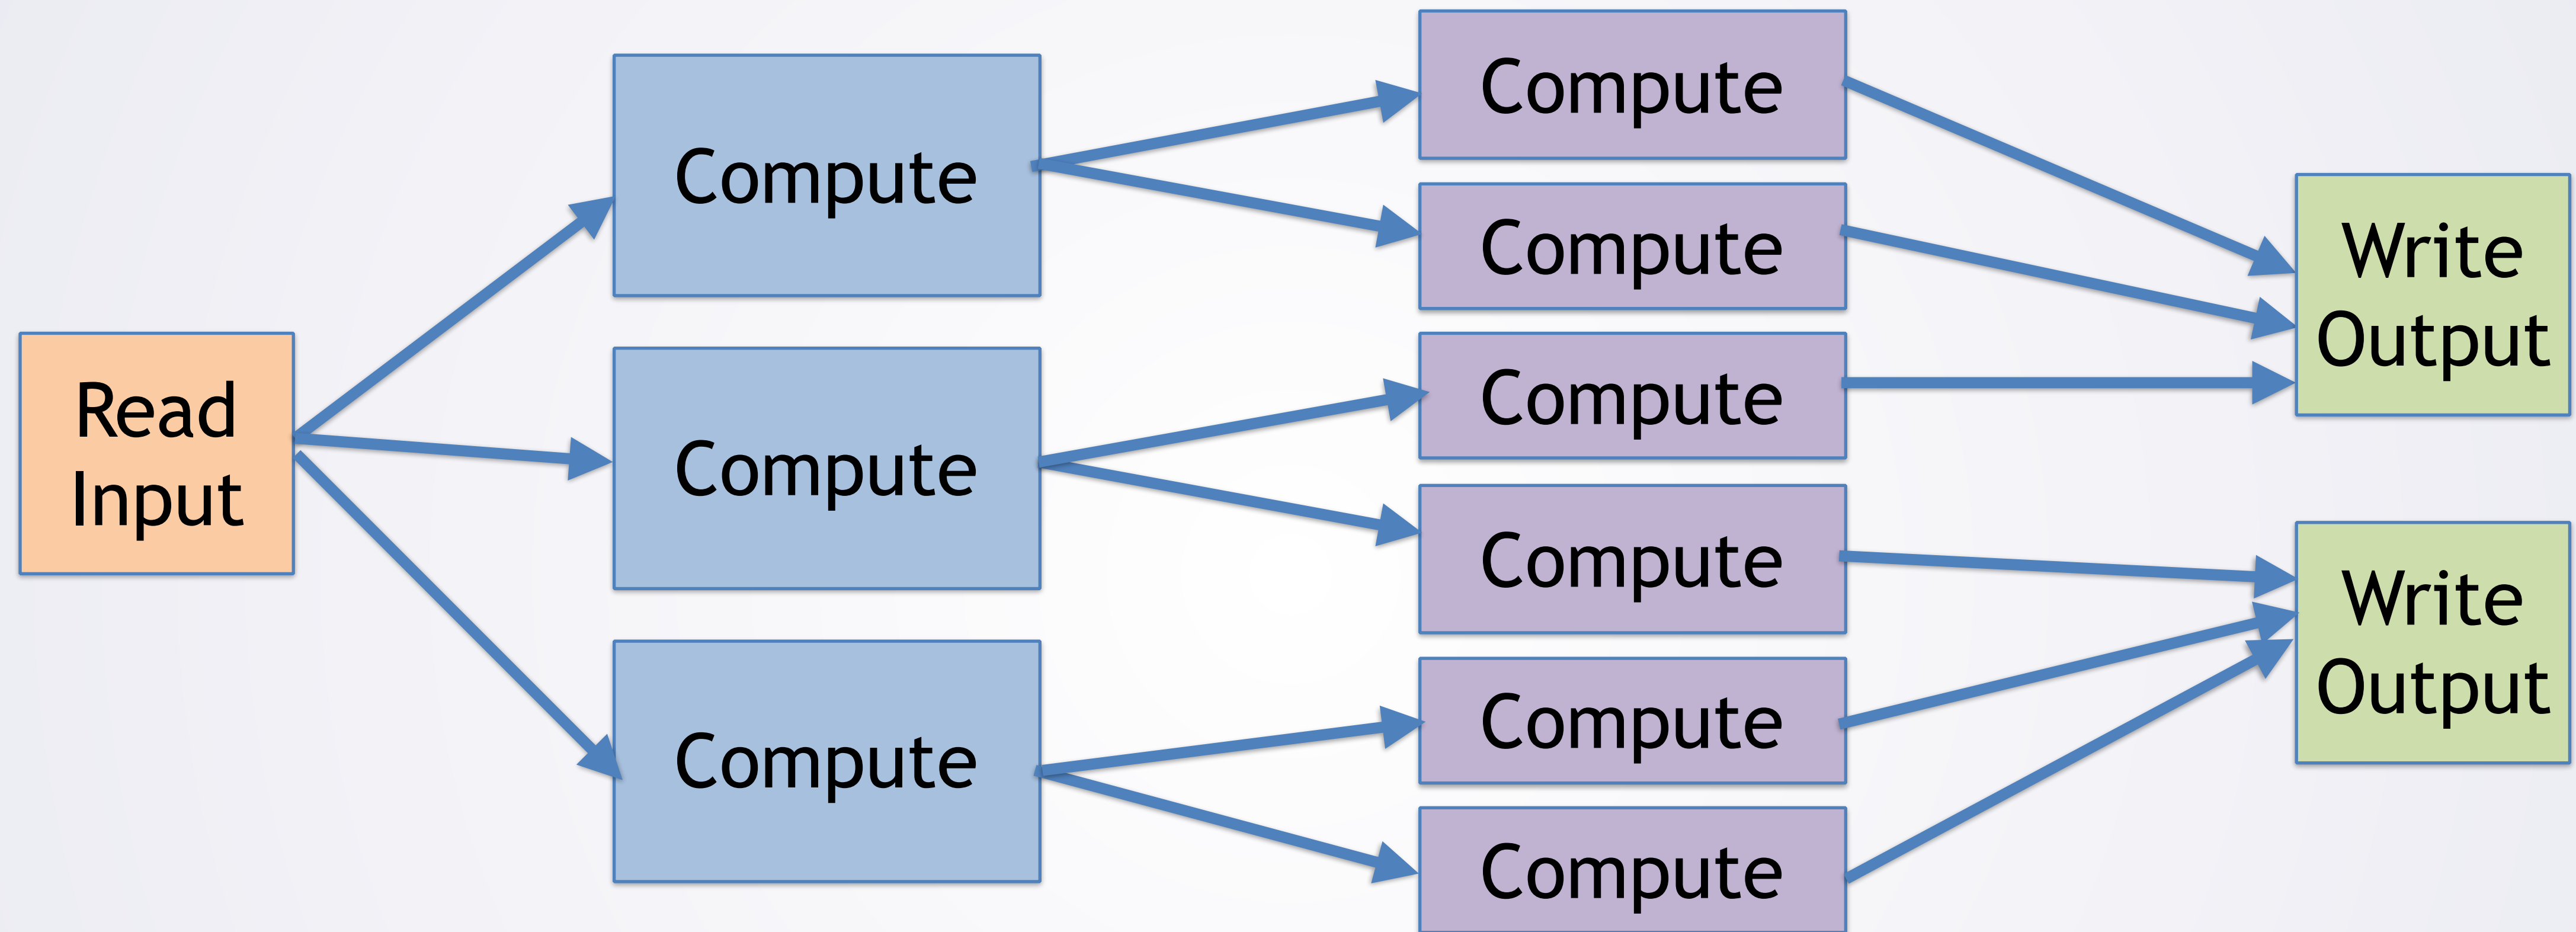
# Pipeline Parallelism

- When appropriate: Can keep IO thread(s) busy
  - Heavy IO to perform
  - Might have one thread do reads and writes
- What is "something else"?
  - Other IO requests
  - Do whichever one is ready to be done

# Pipeline Parallelism

- When appropriate: Can keep IO thread(s) busy
  - Heavy IO to perform
  - Might have one thread do reads and writes
- What is "something else"?
  - Other IO requests
  - Do whichever one is ready to be done
- Making hundreds of read/write calls to see which succeeds = inefficient
  - Use `poll` or `select`

# Pipeline Parallelism



- What can you say about data movement in this model?
- What can you say about load balance?



# Another Option

- Could have one thread work on many requests

```
while(1) {
```

```
    Accept new requests
```

```
    Do any available reads/writes
```

```
    Do any available compute
```

```
}
```

- What can you say about data movement in this model?
- What can you say about load balance?



# A Slight Variant

- Slightly different inner loop:

```
while(1) {
```

Accept new requests

For each request with anything to do

Do any available IO for that request

Do any compute for that request

```
}
```

- What can you say about data movement in this model?
- What can you say about load balance?