


ECE590-02
Engineering Robust Server Software
Spring 2017

IO Performance and Scalability

Tyler Bletsch
Duke University


Basics of large-scale storage systems and their performance

Average person's view of storage




storage = computer does it !!

Average engineer's view of storage



Performance and Scalability

- What is **performance**?
 - Speed of IO operations
 - Measured in **IO Operations Per Second (IOPS)** or **throughput (MB/s)**
- Key insight: We can increase performance (and capacity) by adding drives!
 - How? Recall from ECE550 that RAID allows the use of multiple drives in a way that maintains reliability with aggregating performance
- So we can just keep adding drives forever, right?
 - Well, eventually the case is full or you're out of ports...



Performance and Scalability

- What is **scalability**?
 - A measure of how big can you grow a system
- Hardware example:
a traditional desktop PC has a fixed number of drive bays and drive connections
- Software example:
a database may use a data structure whose performance becomes unacceptable at a certain number of records (e.g. if it relies on a linear search)





Image from <http://s57.photobucket.com/user/junfeng/media/mediaserver2.jpg.html>

Storage in enterprise-scale environments

- On a laptop/desktop, you have Directly Attached Storage (DAS)
 - Poor scalability!
- Storage at large scale is separated from the compute servers



High disk density

- Disks installed in enclosures separate from storage servers, connected with high speed bus (SAS or Fibre Channel)

A disk shelf with 24 3.5-inch drives



From <http://imn.de/product/show/1377/ds4246-ds4246-shelf>

Two disk shelves and a storage server in a rack




From <http://www.talking-it.com/2014/storage/nimble/installing-a-nimble-es1-expansion-shelf-hard-bit/>

- Result: Vastly improved **scalability** of physical storage

Storage in enterprise-scale environments

- Moving up the stack...




The storage server

- Storage servers aggregate the disks and apply RAID
 - *They also do a LOT of other things to enhance storage!*
See my Enterprise Storage Architecture course to learn more...
- How to connect the aggregated storage to compute servers?
Two methodologies:
 - **SAN: Storage Area Network**
 - Storage server divides storage into “virtual block devices”
 - Clients make “read block”/“write block” requests just like to a hard drive, but they go to the storage server
 - Examples: Fibre Channel and iSCSI
 - **NAS: Network-Attached Storage**
 - Storage server runs a file system to create abstraction of files/directories
 - Clients make open/close/read/write requests just like to the OS’s local file system
 - Examples: NIFS and CIFS (also known as ‘Windows shares’)


SAN: Storage Area Network

- Split the aggregated storage into virtual drives called Logical Units (LUNs)
- Clients make read/write requests for blocks of “their” drive(s)
- Storage server translates request for block 50 of client 2 to actual block 4000



NAS: Network-Attached Storage

- Put a file system on the storage server so it has the concept of files and directories
- Clients make open/close/read/write requests for files on the remote file system



The story so far

- **Performance** is how many IOPS or MB/s
- Can be increased by adding drives
- Every system has limits to its ability to grow; the *ability* to *scale* up a system is referred to as **scalability**
- Scalability of physical storage is the amount we can add drives to a system
 - Direct-attach systems have poor physical storage scalability; NAS and SAN schemes are much better in this regard
- Two questions:
 - How else could we influence performance?
 - How much performance do we need?

Influencing IO performance

- Quantity of disks (as stated previously)
- Access pattern
 - If on HDDs, performance is driven by seek time:
better to read few big chunks of data than many small ones
- Caching effects
 - Can we add caches (RAM or SSD cache) or to server or client?
- Type of media
 - Higher RPM HDDs
 - Solid State Drives (SSDs) – flash solid-state storage


How to size storage systems

The problem

- **Workload characterization:** Determining the IO pattern of an application (or suite of applications)
 - We do so by measuring it, known as **workload profiling**
- **Storage sizing:** Determining how much hardware you need to serve a given application (or suite of applications)
- The challenge of characterization and sizing
 - Storage is a complex system!
 - Danger: high penalty for underestimating needs...


Two kinds of metrics

- Inherent **access pattern metrics**:
 - Based on the code
- Resulting **performance metrics**:
 - The performance observed when those access patterns hit the storage system
- Sometimes difficult to separate:
 - Common one that's hard to tell: *IOPS*
 - Did we see 50 IOPS because the workload only made that many requests, or because the storage system could only respond that fast?
 - Was storage system mostly idle? Then IOPS was limited by workload.




Access pattern metrics

- Random vs. sequential IO
 - Often expressed as random%
 - Alternatives: average distance, seek distance histogram, etc.
- IO size
- IOPS
 - If controller/disk utilization was low, then IOPS represent storage *demand* (the rate the app asked for)
 - Alternative metric: inter-arrival time (average, histogram, etc.)
- Reads vs. writes
 - Often expressed as read%
 - May also split all of the above by read vs. write (read access pattern often different from write pattern)
- Breaking down application: can we identify separate threads?
 - Is it 50% random, or is there one 100% random thread and one 100% sequential thread?











Performance metrics

- IOPS (if storage system was bottleneck)
 - Alternative metric: IO latency (average, histogram, etc.)
 - Alternative metric: throughput (for sequential workloads)
- Queue length: number of IO operations outstanding at a time
 - A measure of IO parallelism




Example of metrics



- Metrics for “DVDStore”, a web store benchmark.
 - Random workload (seek distance \neq 0)
 - IO size = 8k
 - Short read queue, long write queue
 - Reasonable latency (within usual seek time)
 - Seek distance for writes is biased positive (likely due to asynchronous write flushing doing writes in positive order to minimize write seek distance)

How to get these metrics?

- **Profiling:** Run the workload and measure
- Two problems:
 1. How to “run”?
 - Most workloads interact with users
 - Need user behavior to get realistic access pattern!
 - Where to get users?
 - App already in production? Use actual users
 - If not, fake it: **synthetic load generation** (extra program pretends to be users)
 - What about so-called **benchmarks**?
 2. How to “measure”? We’ll see in a bit...




Benchmarks: How to “run”

- **Benchmark:** a program used to generate load in order to measure resulting performance. Various types:
 - **The application itself:** You literally run the real app with a synthetic load generator.
 - Example: Microsoft Exchange plus LoadGen
 - **Application-equivalent:** Implements a realistic task from scratch, often with synthetic load generation built in.
 - Example: DVDStore, an Oracle benchmark that literally implements a web-based DVD store.
 - **Task simulator:** Generate an access pattern commonly associated with a certain *type* of workload
 - Example: Swingbench DSS, which generates database requests consistent with computing long-running reports
 - **Synthetic benchmark:** Generate a mix of load with a specific pattern
 - Example: IOZone, which runs a block device at a given random%, read%, IO size, etc.

Methods of profiling: How to “measure”

- App instrumentation
 - Requires code changes
- Kernel instrumentation
 - Can hook at system call level (e.g. strace) or block IO level (e.g. blktrace).
 - Can also do arbitrary kernel instrumentation, hook anything (e.g., systemtap)
- Hypervisor instrumentation
 - Hypervisor sees all I/O by definition
 - Example: vscsiStats in VMware ESX
- Storage controller instrumentation
 - Use built-in performance counters
 - Basically this is kernel instrumentation on the storage controller kernel
- User-level metrics (e.g. latency to load an email)
 - These don't directly help understand storage performance, but they *are* the metrics that users actually care about



Sizing



- Now we know how workload acts; need to decide how much storage gear we need to buy
- Will present basic rules, but there are complicating factors:
 - Effects of storage efficiency features?
 - Effects of various caches?
 - CPU needs of the storage controller?
 - Result when multiple workloads are combined on one system?
- Real-world sizing of enterprise workloads:
 - For commercial apps, ask the vendor – companies with big, expensive, scalable apps have sizing teams that write sizing guides, tools, etc.
 - On the storage system side, ask the system vendor – companies with big, expensive, scalable storage systems have sizing teams too.

Disk array sizing

- Recall: In a RAID array, performance is proportional to number of disks; this includes IOPS
- Each disk “provides” some IOPS: $IOPS_{disk}$
- Our workload profile tells us: $IOPS_{workload}$
- Compute $\frac{IOPS_{workload}}{IOPS_{disk}}$: get number of disks needed
- Add overhead: RAID parity disks, hot spares, etc.
- Add safety margin: 20% minimum, >50% if active/active
- Note: this works for SSDs too, $IOPS_{disk}$ is just way bigger

Characterizing disks

- Use synthetic benchmark to find performance in the extremes (100% read, 100% write, 100% seq, 100% random, etc.)
- Results for Samsung 850 Evo 2TB SSD:



Interpolation-based sizing

- For large/complex storage deployments with mixed workloads, simple IOPS math may break down
- Alternative: **measurement with interpolation**
 - Beforehand:
 - foreach (synthetic benchmark configuration with access pattern **a**)
 - foreach (storage system configuration **s**)
 - set up storage **s**, generate IO pattern **a**, record metrics as **M[a,s]**
 - Later, given real workload with access pattern **a_{given}** and performance requirements **M_{required}**
 - Find points **a,s** in table where **a** is near **a_{given}** and performance **M[a,s] > M_{required}**
 - Deploy a storage system based on the constellation of corresponding **s** values.
 - Can interpolate storage configurations **s** (with risk)
 - Pessimistic model: Can just pick from systems where **a** was clearly “tougher” and performance **M[a,s]** was still sufficient
- Why do all this? Because **s** can include ALL storage config parameters (storage efficiency, cache, config choices, etc.)

Combining workloads

- Rare to have one storage system handle just ONE workload; shared storage on the rise
- Can we simply add workload demands together?
 - Sometimes...it's complicated.
 - Example that **works**: two random workloads run on separate 3-disk RAID5s will get similar performance running together on one 6-disk RAID5
 - Example that **doesn't**: a random workload plus a sequential workload wrecks performance of the sequential workload
 - Random IOs will "interrupt" big sequential reads that would otherwise be combined by OS/controller.

Workload combining

RAID5 config

Workload	LUN configuration	IOPS	IO latency	Application Metric
Random → DVDStore	2+1	130	100	6132 TPM
Random → OLTP	2+1	141	30	5723 TPM
DVDStore (Shared)	5+1	144 😞	30	7630 TPM
OLTP (Shared)	5+1	135 😞	30	5718 TPM

Table 1. Comparison of DVDStore and OLTP when run in isolation and shared mode

Workload	LUN configuration	Throughput	95% tile latency	Application Metric
Random → DVDStore	2+1	130 IOPS	100	6132 TPM
Sequential → DSS	2+1	44 MB/s	30	6 completed transactions
DVDStore (Shared)	5+1	164 IOPS 😊	15	7630 TPM
DSS (Shared)	5+1	31 MB/s 😞	1	3 completed transactions

Table 2. Comparison of DVDStore and DSS when run in isolation and shared mode

- “OLTP” = “Online Transaction Processing” (normal user-activity-driven database)
- “DSS” = “Decision Support System” (long-running report on a database)
- Table 2: DVDStore benefits a little from twice as many disks to help with latency, but DSS’s sequential IO gets wrecked by the random interruptions to its stream

Conclusion

- To **characterize** a workload, we must **profile** it
 - Run it (generating user input if needed)
 - Measure IO metrics in app/kernel/hypervisor/controller
- Can use workload profile for **sizing**:
to identify storage gear needed
 - Basic rule: provision enough disks for the IOPS you need
 - Past that, look for published guidance from software/hardware vendor
 - Failing that, use successive experiments with differing gear to identify performance trends