

Concurrency and Synchronization

ECE 650
Systems Programming & Engineering
Duke University, Spring 2016

Concurrency


- Multiprogramming
 - Supported by most all current operating systems
 - More than one “unit of execution” at a time
- Uniprogramming
 - A characteristic of early operating systems, e.g. MS/DOS
 - Easier to design; no concurrency
- What do we mean by a “unit of execution”?

ECE 650 – Fall 2016

2

Process vs. Thread

- Process vs. Thread




- A process is –
 - Execution context
 - Program counter (PC)
 - Stack pointer (SP)
 - Registers
 - Code
 - Data
 - Stack

ECE 650 – Fall 2016

3

Process vs. Thread

- Process vs. Thread



- A thread is –
 - Execution context
 - Program counter (PC)
 - Stack pointer (SP)
 - Registers

ECE 650 – Fall 2016

4

Process vs. Thread

- Process: unit of allocation
 - resources, privileges, etc.
- Thread: unit of execution
 - PC, SP, registers
- Thread is a unit of control within a process
- Every process has one or more threads
- Every thread belongs to one process

ECE 650 – Fall 2016

5

Process Execution

- When we execute a program
 - OS creates a process
 - Contains code, data
 - OS manages process until it terminates
 - We will talk more about these aspects of process management
 - E.g. scheduling, system calls, etc.
- Every process contains certain information
 - ID
 - Process state (for scheduling purposes)
 - Program counter, stack pointer, CPU registers
 - Memory management info, files, I/O

ECE 650 – Fall 2016

6

Process Execution (2)

- A process is created by the OS via system calls
 - fork(): make exact copy of this process and run
 - Forms parent/child relationship between old/new process
 - Return value of fork indicates the difference
 - Child returns 0; parent returns child's PID
 - exec(): can follow fork() to run a different program
 - Exec takes filename for program binary from disk
 - Loads new program into the current process's memory
- A process may also create & start execution of threads
 - Many ways to do this
 - System call: clone(); Library call: pthread_create()

ECE 650 – Fall 2016

7

Back to Concurrency...

- We have multiple units of execution, but single resources
 - CPU, physical memory, IO devices
 - Developers write programs as if they have exclusive access
- OS provides illusion of isolated machine access
 - Coordinates access and activity on the resources

ECE 650 – Fall 2016

8

How Does the OS Manage?

- Illusion of multiple processors
 - Multiplex threads in time on the CPU
 - Each virtual "CPU" needs a structure to hold:
 - Program Counter (PC), Stack Pointer (SP)
 - Registers (Integer, Floating point, others...?)
 - How switch from one CPU to the next?
 - Save PC, SP, and registers in current state block
 - Load PC, SP, and registers from new state block
 - What triggers switch?
 - Timer, voluntary yield, I/O, other things
- We will talk about other management later in the course
 - Memory protection, IO, process scheduling

ECE 650 – Fall 2016

9

Concurrent Program

- Two or more threads execute concurrently
 - Many ways this may occur...
 - Multiple threads time-slice on 1 CPU with 1 hardware thread
 - Multiple threads at same time on 1 CPU with N HW threads
 - Simultaneous multi-threading
 - Multiple threads at same time on M CPUs with N HW threads
 - Chip multi-processor (CMP) or Symmetric multi-processor (SMP)
- Cooperate to perform a task
- How do threads communicate?
 - Recall they share a process context
 - Code, static data, "heap"
 - Can read and write the same memory
 - variables, arrays, structures, etc.

ECE 650 – Fall 2016

10

Motivation for a Problem

- What if two threads want to add 1 to shared variable?
 - x is initialized to 0

`x = x + 1;` May get compiled into:
(x is at mem location 0x8000)

- A possible interleaving:


P1	P2
<code>ldw r1, 0(0x8000)</code>	<code>ldw r1, 0(0x8000)</code>
<code>addi r1, r1, 1</code>	
<code>stw r1, 0(0x8000)</code>	<code>addi r1, r1, 1</code>
	<code>stw r1, 0(0x8000)</code>

- At the end, x will have a value of 1 in memory!!

ECE 650 – Fall 2016

11

Another Example – Linked List



- Two concurrent threads (A & B) want to add a new element to list
 1. A executes first three instructions & stalls for some reason (eg. cache miss)
 2. B executes all 4 instructions
 3. A eventually continues and executes 4th instruction
 - Item added by thread B is lost!

ECE 650 – Fall 2016

12

Race Conditions

- The example problems occur due to race conditions
- Race Condition
 - Result of computation by concurrent threads depends on the precise timing of the execution of an instruction sequence by one thread relative to another
- Sometimes result may be correct, sometimes incorrect
 - Depends on execution timing
 - Non-deterministic result
- Need to avoid race conditions
 - Programmer must control possible execution interleaving of threads

Mutual Exclusion

- Previous examples show problem of multiple processes or threads performing read/write ops on shared data
 - Shared data = variables, array locations, objects
- Need mutual exclusion!
 - Enforce that only one thread at a time in a code section
 - This section is also called a *critical section*
 - Critical section is set of operations we want to execute *atomically*
- Provided by lock operations


```
lock(x_lock);
x = x + 1;
unlock(x_lock);
```
- Also note => this isn't only an issue on parallel machines
 - Think about multiple threads time-sharing a single processor
 - What if a thread is interrupted after load/add but before store?

Mutual Exclusion

- Interleaving with proper use of locks (mutex)

<p>P1</p> <pre>lock(x_lock) ldw r1, 0(8000) addi r1, r1, 1 stw r1, 0(8000) unlock(x_lock)</pre>	<p>P2</p> <pre>lock(x_lock) ldw r1, 0(8000) addi r1, r1, 1 stw r1, 0(8000) unlock(x_lock)</pre>
--	--

- At the end, x will have a value of 2 in memory

Global Event Synchronization

- BARRIER (name, nprocs)
 - Thread will wait at barrier call until nprocs threads arrive
 - Built using lower level primitives
 - Separate phases of computation
- Example use:
 - N threads are adding elements of an array into a sum
 - Main thread is to print sum
 - Barrier prevents main thread from printing sum too early
- Use barrier synchronization only as needed
 - Heavyweight operation from performance perspective
 - Exposes load imbalance in threads leading up to a barrier

Point-to-point Event Synchronization

- A thread notifies another thread so it can proceed
 - E.g. when some event has happened
 - Typical in producer-consumer behavior
 - Concurrent programming on uniprocessors: semaphores
 - Shared memory parallel programs: semaphores or monitors or variable flags

flag

P0:
S1: datum = 5;
S2: datumReady = 1;

P1:
S3: while (!datumReady) {};
S4: print datum

monitor

P0:
S1: datum = 5;
S2: signalReady;

P1:
S3: waitReady;
S4: print datum

Lower Level Understanding

- How are these synchronization operations implemented?
 - Mutexes, monitors, barriers
- An attempt at mutex (lock) implementation

```
void lock (int *lockvar) {
    while (*lockvar == 1) {} ; // wait until released
    *lockvar = 1; // acquire lock
}


void unlock (int *lockvar) {
    *lockvar = 0;
}
```

In machine language, it looks like this:

```
lock: ld R1, @lockvar // R1 = lockvar
      bnz R1, lock // jump to lock if R1 != 0
      st @lockvar, #1 // lockvar = 1
      ret // return to caller
unlock: st @lockvar, #0 // lockvar = 0
      ret // return to caller
```

Problem

- Unfortunately, this attempted solution is incorrect
- The sequence of ld, bnz, and sti are not atomic
 - Several threads may be executing it at the same time
- It allows several threads to enter the critical section simultaneously



Software-only Solutions

- Peterson's Algorithm (mutual exclusion for 2 threads)

```
int turn;
int interested[n]; // initialized to 0

void lock (int process, int lvar) { // process is 0 or 1
    int other = 1 - process;
    interested[process] = TRUE;
    turn = process;
    while (turn == process && interested[other] == TRUE) {}
}
// Post: turn != process or interested[other] == FALSE

void unlock (int process, int lvar) {
    interested[process] = FALSE;
}
```

- Exit from lock() happens only if:
 - interested[other] == FALSE: either the other process has not competed for the lock, or it has just called unlock()
 - turn != process: the other process is competing, has set the turn to itself, and will be blocked in the while() loop

Help From Hardware

- Software-only solutions have drawbacks
 - Tricky to implement – think about more than 2 threads
 - Need to consider different solutions for different memory consistency models
- Most processors provide atomic operations
 - E.g. Test-and-set, compare-and-swap, fetch-and-increment
 - Provide atomic processing for a set of steps, such as
 - Read a location, capture its value, write a new value
 - Test-and-set
 - Instruction supported by HW
 - Write to a memory location and return its old value as a single atomic operation

Programming with Pthreads

- POSIX pthreads
 - Found on most all modern POSIX-compliant OS
 - Also Windows implementations
 - Allows a process to create, spawn, and manage threads
- How to use it:
 - Add #include <pthread.h> to your C/C++ source code
 - When compiling with gcc, add -lpthread to your list of libraries
 - gcc -o p_test p_test.c -lpthread
 - Instrument the code with pthread function calls to:
 - Create threads
 - Wait for threads to complete
 - Destroy threads
 - Synchronize across threads
 - Protect critical sections

Pthread Thread Creation

- Create a pthread:
 - int pthread_create(pthread_t *thread, pthread_attr_t *attr, void *(*start_routine)(void *), void *arg);
- Arguments:
 - pthread_t *thread_name – thread object (contains thread ID)
 - pthread_attr_t *attr – attributes to apply to this thread
 - void *(*start_routine)(void *) – pointer to function to execute
 - void *arg – arguments to pass to above function

- Example:

```
pthread_t *thrd;
pthread_create(&thrd, NULL, &do_work_for_n, NULL);
```

Pthread Destruction

- pthread_join(pthread_t thread, void **value_ptr)
 - Suspends the calling thread
 - Waits for successful termination of the specified thread
 - value_ptr is optional data passed from terminating thread's exit
- pthread_exit(void *value_ptr)
 - Terminates a thread
 - Provides value_ptr to any pending pthread_join() call

Pthread Mutex

- pthread_mutex_t lock;
- Initialize a mutex; 2 ways:
 - int pthread_mutex_init(pthread_mutex_t* mutex, const pthread_mutexattr_t* mutex_attr);
 - pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
 - Initialized with default pthread mutex attributes
 - This is typically good enough
- Operate on the lock:
 - int pthread_mutex_lock(pthread_mutex_t* mutex);
 - int pthread_mutex_trylock(pthread_mutex_t* mutex);
 - int pthread_mutex_unlock(pthread_mutex_t* mutex);

Pthread Barrier

- pthread_barrier_t barrier;
- Initialize a barrier; 2 ways:
 - int pthread_barrier_init(pthread_barrier_t* barrier, const pthread_barrierattr_t* barrier_attr, unsigned int count);
 - pthread_barrier_t barrier = PTHREAD_BARRIER_INITIALIZER(count);
 - Initialized with default pthread barrier attributes
 - This is typically good enough
- Operation on a barrier:
 - int pthread_barrier_wait(pthread_barrier_t* barrier);

Pthread Example (Matrix Mul)

```
double **a, **b, **c;
int numthreads, matrixSize;

int main(int argc, char *argv[]) {
    int i, j;
    int *p;
    pthread_t *threads;

    // Initialize numthreads, matrixSize; allocate and init a/b/c matrices
    // ...

    // Allocate thread handles
    threads = (pthread_t *) malloc(numthreads * sizeof(pthread_t));

    // Create threads
    for (i = 0; i < numthreads; i++) {
        p = (int *) malloc(sizeof(int));
        *p = i;
        pthread_create(&threads[i], NULL, worker, (void *)p);
    }

    for (i = 0; i < numthreads; i++) {
        pthread_join(threads[i], NULL);
    }

    printMatrix(c);
}
```

Pthread Example (Matrix Mul) cont.

```
void mm(int myid) {
    int i, j, k;
    double sum;
    // compute bounds for this thread
    int startrow = myid * matrixSize / numthreads;
    int endrow = (myid+1) * (matrixSize / numthreads) - 1;

    // matrix mult over the strip of rows for this thread
    for (i = startrow; i <= endrow; i++) {
        for (j = 0; j < matrixSize; j++) {
            sum = 0.0;
            for (k = 0; k < matrixSize; k++) {
                sum = sum + a[i][k] * b[k][j];
            }
            c[i][j] = sum;
        }
    }
}

void *worker(void *arg){
    int id = *((int *) arg);
    mm(id);
    return NULL;
}
```