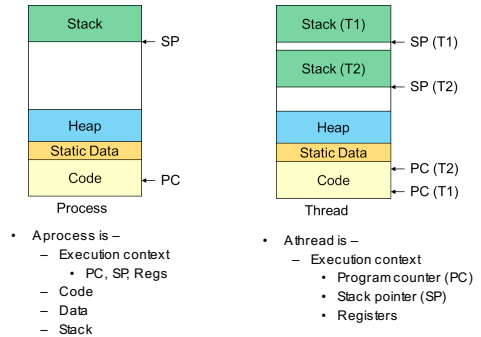


Inter-process Communication (IPC)

ECE 650
Systems Programming & Engineering
Duke University, Spring 2016

Recall Process vs. Thread



ECE 650 – Fall 2016

2

Cooperation

- Two or more threads or processes execute concurrently
- Sometimes cooperate to perform a task
 - Sometimes independent; not relevant for IPC discussion
- How do they communicate?
 - Threads of the same process: Shared Memory
 - Recall they share a process context
 - Code, static data, "heap"
 - Can read and write the same memory
 - variables, arrays, structures, etc.
 - What about threads of different processes?
 - They do not have access to each other's memory

ECE 650 – Fall 2016

3

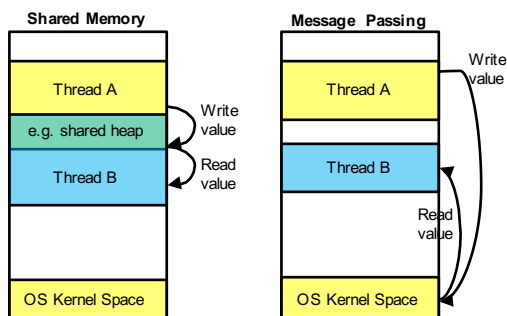
Models for IPC

- Shared Memory
 - E.g. what we've discussed for threads of same process
 - Also possible across processes
 - E.g. memory mapped files (mmap)
- Message Passing
 - Use the OS as an intermediary
 - E.g. Files, Pipes, FIFOs, Messages, Signals

ECE 650 – Fall 2016

4

Models for IPC



ECE 650 – Fall 2016

5

Shared Memory vs. Message Passing

- Shared Memory
 - Advantages
 - Fast
 - Easy to share data (nothing special to set up)
 - Disadvantages
 - Need synchronization! Can be tricky to eliminate race conditions
- Message Passing
 - Advantages
 - Trust not required between sender / receiver (receiver can verify)
 - Set of shared data is explicit
 - Is synchronization needed?
 - Disadvantages
 - Explicit programming support needed to share data
 - Performance overhead (e.g. to copy messages through OS space)

ECE 650 – Fall 2016

6

Shared Memory Across Processes

- Different OSes have different APIs for this
- UNIX
 - System V shared memory (shmget)
 - Allows sharing between arbitrary processes
 - <http://www.fdp.org/LDP/lpg/node21.html>
 - Shared mappings (mmap on a file)
 - Different forms for only related processors or unrelated processes (via filesystem interaction)
 - POSIX shared memory (shm_open + mmap)
 - Sharing between arbitrary processes; no overhead of filesystem IO
- Still requires synchronization!

mmap

- ```
#include <sys/mman.h>
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset)
```
- Creates new mapping in virtual address space of caller
    - addr: starting address for mapping (or NULL to let kernel decide)
    - length: # bytes to map starting at "offset" of the file
    - prot: desired memory protection of the mapping
      - PROT\_EXEC, PROT\_READ, PROT\_WRITE, PROT\_NONE
    - flags: are updates to mapping are visible to other processes?
      - MAP\_SHARED, MAP\_PRIVATE
      - Other flags can be added, e.g. MAP\_ANON (more later)
    - fd: file descriptor for open file
      - Can close the "fd" file after calling mmap()
    - Return value is the address where the mapping was made

## mmap operation

- Kernel takes an open file (given by FD)
- Maps that into process address space
  - In unallocated space between stack & heap regions
  - Thus also maps file into physical memory
  - Creates one-to-one correspondence between a memory address and a word in the file
- Useful even apart from the context of IPC
  - Allows programmer to read/write file contents without read(), write() system calls
- Multiple (even non-related) processes can share mem
  - They open & mmap the same file

## munmap

- ```
#include <sys/mman.h>
int munmap(void *addr, size_t length);
```
- Removes mapping from process address space
 - addr: address of the mapping
 - length: # bytes in mapped region
- ```
#include <sys/mman.h>
int msync(void *addr, size_t length, int flags);
```
- Removes mapping from process address space
    - addr: address of the mapping
    - length: # bytes in mapped region

## Synchronization

- Semaphores

```
#include <fcntl.h> /* For O_* constants */
#include <sys/stat.h> /* For mode constants */
#include <semaphore.h>

sem_t *sem_open(const char *name, int oflag, mode_t mode, int value);
or
int sem_init(sem_t *sem, int pshared, unsigned int value);

sem_t *mutex;
mutex = sem_open("my_sem_name", O_CREAT | O_EXCL,
MAP_SHARED, 1);

sem_wait(mutex);
//Critical section
sem_post(mutex);
```

## Example

- Show code & run in class

## Taking it Further

- This required some work
  - Create file in file system
  - Open the file & initialize it (e.g. with 0's)
- There is a better way if just sharing mem across a fork()
  - Anonymous memory mapping
- Use mmap flags of MAP\_SHARED | MAP\_ANON
  - File descriptor will be ignored (also offset)
  - Memory initialized to 0
  - Alternative approach: open /dev/zero & mmap it
- Can anonymous approach work across non-related processes?

## Message Passing

- Messages between processes, facilitated by OS
- Several approaches:
  - Files
    - Can open the same file between processes
    - Communicate by reading and writing info from the file
    - Can be difficult to coordinate
  - Pipes
  - FIFOs
  - Messages (message passing)
  - Signals

## Pipes

```
#include <unistd.h>
int pipe(int pipefd[2]);
```

- Creates a unidirectional channel (pipe)
- Can be used for IPC between processes / threads
- Returns 2 file descriptors
  - pipefd[0] is the read end
  - pipefd[1] is the write end
- Kernel support
  - Data written to write end is buffered by kernel until read
  - Data is read in same order as it was written
  - No synchronization needed (kernel provides this)
  - Must be related processes (e.g. children of same parent)

## Example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

#define N 1024

int main(int argc, char *argv[]) {
 int pipefd[2];
 char data_buffer[N];

 pipe(pipefd);
 int id = fork();
 if (id == 0) { //child
 write(pipefd[1], "hello", 6);
 } else {
 read(pipefd[0], data_buffer, 6);
 printf("Received data: %s\n", data_buffer);
 }

 return 0;
}
```

## More Complex Uses of Pipes

- Can use pipes to coordinate processes
- For example, chain output of one process to input of next
  - E.g. command pipes in UNIX shell!
- Requires 1 additional (very useful) piece

```
#include <unistd.h>
int dup2(int oldfd, int newfd);
```
- Creates a copy of an open file descriptor into a new one
  - After closing the new file descriptor if it was open

## UNIX Pipes Example

- Show code & run in class

## UNIX FIFOs

- Similar to a pipe
  - Also called a "named pipe"
- Persist beyond lifetime of the processes that create them
- Exist as a file in the file system

```
#include <sys/types.h>
#include <sys/stat.h>
int mkfifo(const char *pathname, mode_t mode);
```
- pathname points to the file
- Mode specifies the FIFO's permissions (similar to a file)

## UNIX FIFOs (2)

- After FIFO is created, processes must open it
  - By default, first open blocks until a second process also opens
- Since FIFOs persist, they can be re-used
- No synchronization needed (like pipes, OS handles it)

## Multiple Producers

- Multiple producers problem:
  - What if >1 producers and 1 consumer
  - Producers are performing write(...)
  - Consumer is performing (blocking) read(...)
  - What if consumer is blocked, but other IPC channels have data?
- Would like to be notified if one channel is ready

## Select

```
#include <sys/select.h>
int select(int nfd, fd_set *readfds, fd_set *writefds,
 fd_set *exceptfds, struct timeval *timeout);
```

- nfd = number of file descriptors to monitor
- readfds, writefds, exceptfds are bit vectors of file descriptors to check
- timeout is a maximum time to wait
- Macros are available to work with bit sets:
  - FD\_ZERO(&fds), FD\_SET(n, &fds), FD\_CLEAR(n, &fds)
  - int FD\_ISSET(n, &fds); //useful after select() returns

## Poll

```
#include <poll.h>
int poll(struct pollfd *fds, nfds_t nfds, int timeout);
```

- nfds = number of file descriptors to monitor
- fds is an array of descriptor structures
  - File descriptors, desired events, returned events
- timeout is a maximum time to wait
- Returns number of descriptors with events