

Protection & Security

ECE 650
Systems Programming & Engineering
Duke University, Spring 2016

Protection

- OS manages resources for users & user processes
 - Files, memory regions, IO channels, CPU
- Protection is a critical part of this management
 - Ensure that resources can only be used with proper authorization from the OS
- Reasons for protection
 - Prevent users from malicious access to resources
 - Ensure processes use system resources only as consistent with allowed policies
- Protection is a "mechanism"
 - Mechanism to enforce policies that define how resources should be used
 - As opposed to a "policy" (definition of how resources should be used)
 - Policies may adapt and change over time (or between different applications)
 - Thus mechanisms should be general to allow flexibility

ECE 650 – Fall 2016

2

Basics of Protection

- Most protection mechanisms based on key principle
 - Principle of least privilege
 - Users, processes, etc. have the minimum level of access to resources and privileges needed to accomplish intended task
 - "Need to know basis"
- Minimizes damage from failed or compromised pieces
 - They can only affect a minimal set of components in the system
- OS designs provide support for this
 - System calls and services for apps to specify fine-grained permissions & controls
 - Apps enable and disable permissions as needed
 - Also applies to users (separate accounts, permissions)

ECE 650 – Fall 2016

3

Protection Domains

- Think of all resources as an object
 - E.g. in UNIX – "everything is a file"
 - Hardware objects: CPU, memory spaces, disk, keyboard, display
 - Software objects: files, directories, programs
- Different objects have different possible operations, e.g.
 - Read & write memory regions; Read from a keyboard input
 - Execute on a CPU
 - Create, delete, open, close, read, write, append files
- Protection mechanism operates based on domains
 - Application or user has permission to perform certain operations on certain objects
 - Again, can change dynamically

ECE 650 – Fall 2016

4

Protection Domains (2)

- Processes operate within a protection domain
 - Specifies objects (resources) the process has permission to access
- Access Right
 - A permission for a process to perform an operation on an object
 - Domain just a collection access rights
 - Access right can exist in more than one domain at a time
 - Example...
- Access rights can be static or dynamic for a process
 - Dynamic domains achieved via either:
 - A mechanism to change domain access rights
 - A mechanism for domain switching
 - Create new domain with proper new access rights; then switch to it
 - For example, user – supervisor mode we discussed for interrupts

ECE 650 – Fall 2016

5

UNIX Example of Domains

- In UNIX, each user belongs to a domain
- When the domain is changed, the user ID is changed
- Happens via a protection mechanism using the file system
 - Remember, in UNIX, "everything" is a file
 - Every file has an owner ID and a domain bit (setuid bit)
 - "set user ID upon execution"
 - Set just like file read/write/execute permissions
 - When a user executes a file:
 - If setuid bit is on, user ID is changed to the owner of the file
 - If setuid bit is off, user ID does not change
 - Temporary user ID change ends after process exits
- Allows a privileged component to be used by general users
 - E.g. an application that accesses the network or change user password
- What if a user creates a file with user ID of root & setuid on?

ECE 650 – Fall 2016

6

Access Matrix

- Domain protection model maps nicely to a matrix
- Rows = domains; columns = objects
- A matrix entry lists the access rights
- Provides a general mechanism for specifying policies
 - Enforce specific access rights for a user or process in a domain
- Example

Access Matrix – Additional Functions

- Base access matrix allows
 - Defining and enforcing strict access control policy
- How can we provide dynamic rights?
 - Domain switching
 - Add access matrix entries to enforce “switch” operation rights
 - Example
 - Allow controlled changing of the access matrix entries
 - Encode this permission in the access matrix as well!
 - New operations: copy, owner, control

Copy, Owner, Control

- Copy
 - Copy an access right from one domain to another for an object
 - Variants:
 - Move right instead of copy (transfer)
 - Copy with and without propagation of the copy right (limited copy)
- Owner
 - Allows creation of new rights in the access matrix
 - Process in a domain can add & remove any right for that object
- Control
 - Applies only to domain objects
 - Allows removal of rights in the access matrix
 - Process in a domain can remove any right from that row

Access Matrix Implementation

- In a real system, the matrix will be very sparse
 - But the way it is accessed & used cause special considerations
- Global Table
 - List of <domain, object, rights> tuples
 - Search for “right” when a process in “domain” accessed “object”
 - Drawbacks:
 - Table is huge
 - Objects may have “global” rights, but still listed in every domain
- Access Lists
 - Maintain a list per object with <domain, rights> (column-based)
 - Can extend with a “default” set of rights for each object

Security

- Protection is a mechanism for internal problem
 - Controlled access to programs, data
- Security deals with the external environment
 - Protection can be thwarted if security is compromised
 - Protection works well only if users behave as intended
 - E.g. if a user password is stolen or cracked
- In a secure system...
 - All resources (objects) are used only according to policy
 - Cannot be achieved in reality, but strive to limit violations

Security and the OS

- Why is security important for systems programming?
- Many attacks target an “escalation of privilege”
 - This often involves attempting to gain “root” privilege in a system
 - For purposes of reading or tampering with data
 - Data not otherwise accessible via protection mechanism
 - Systems programmers should be aware of:
 - Types of attacks
 - Mechanisms by which attacks are attempted and operate

Some Definitions

- Threat: potential for a security violation
 - E.g. a **vulnerability** in a program
- Attack: attempt to break security
 - E.g. an **exploit** is an attempt to utilize a program vulnerability
- Categories of security violations
 - Confidentiality breach: data theft
 - E.g. credit card, account information; very common goal
 - Integrity breach: unauthorized modification of code or data
 - Availability breach: unauthorized destruction of data
 - E.g. deleting customer account info or defacing a website
 - Theft of service: unauthorized use of resources
 - Denial of service: prevent legitimate use of a system

Domains of Security Measures

- Physical
 - Need to secure the sites where computer systems reside
 - Only authorized administrators / users have physical access
- Human
 - Social engineering may trick authorized users into performing an inadvertent breach of security
 - Phishing obtain information
 - Executing malicious code
- OS
 - Mechanisms to protect from accidental or purposeful breaches
- Network
 - Protect network-transmitted data from interception or tampering

Types of Threats

- Threats to Running Programs
 - Trojan Horse
 - Trap Door
 - Logic Bomb
 - Stack (or Buffer) Overflow
 - Viruses
- Threats to System and Network Resources
 - Worms
 - Port Scanning
 - Denial of Service

Trojan Horse

- Malware code that misuses its environment
 - Often disguised as legitimate software
 - User unknowingly is tricked into executing the malware code
 - Often takes advantage of access rights of the executing user
- Example: Take advantage of search paths on UNIX OS
 - PATH environment variable specifies order of locations to search for executable files (e.g. 'ls' command)
 - PATH usually has things like: /bin:/usr/bin:/usr/local/bin
 - Sometimes also has things like "." (current directory)
 - Malicious user creates a Trojan program
 - With a common command name (e.g. 'cd')
 - Unknowing user goes into the directory with this program and executes what they think is the "normal" cd command
 - Executes Trojan code

Trojan Horse (2)

- Emulate a login prompt
 - User enters a login ID and password
 - Trojan code captures user ID and password
 - Trojan code prints a login failure message & exits
 - Returning to the real login prompt
 - User thinks they have mistyped password; suspects nothing
 - Reason behind the <ctrl>+<alt>+<delete> Windows
 - Non-trappable key sequence
 - Trojan code cannot intercept this signal and ignore it
- Spyware
 - Code contained along with software to display ads
 - Or capture information and send it somewhere for mining
 - This is called a covert channel (e.g. by opening a network daemon)
 - Fundamental violation of principle of least privilege

Trap Door

- A security hole purposely left in legitimate software
 - Can be exploited by those with knowledge of the vulnerability
- Financial code might include tiny rounding errors
 - Route rounded money to a specific bank account



Trap Door (2)

- Tweak authentication procedures for an application
 - E.g. obscure user name and ID password combo is always valid
- Extra sneaky scenario:
 - Embed the trap door for an application in compiler(s)
 - Compiler checks to see if it is compiling the specific application
 - If so, it inserts the trap door code
 - Inspection of the application source code reveals no issues!
 - Additionally, if the application is re-compiled, problem still exists!


Logic Bomb

- Malware that is triggered only under certain conditions
 - E.g. causes a security incident only at a particular time
 - Application may run normally for a long period of time
- What kind of conditions? Almost anything...
 - Certain year, month, day, time
 - If certain information is present on the system
 - Check a database to see if the programmer is still employed

Stack (Buffer) Overflow

- Previous threats involve:
 - A programmer that can create malicious programs
 - A way to install or place malicious code in the system
- What if this is not possible?
 - How could an attacker execute malicious code?
- By far the most common way is through stack overflow
 - Takes advantage of the stack mechanism to:
 - Allow injection of malicious code
 - Force a process to execute the code
- As we will see, specific to architecture, OS, & application

Stack Overflow Mechanism



- Stack organized as frames
 - Every function call creates & pushes a new frame on stack
 - Function return pops frame from stack
- Frames may contain
 - Return address (PC)
 - Address of previous frame
 - Space for function args
 - Space for function local vars

Stack Overflow Mechanism


```
void func(int a, int b, char *c) {
    char buff1[5];
    char buff2[10];
    strcpy(buff2, c);
}

int main(int argc, char *argv[]) {
    func(1, 2, argv[1]);
}
```

compile

```
main:
<snip>
0x4 push $1
0x8 push $2
0xC push $3
0x10 call func #pushes pc=0x14
<snip>
```

Stack After Calling func()




Stack Overflow Mechanism (2)

```
void func(int a, int b, char *c) {
    char buff1[5];
    char buff2[10];
    strcpy(buff2, c);
}

int main(int argc, char *argv[]) {
    func(1, 2, argv[1]);
}
```

- What if buffer c has more than 10 bytes?
- Lack of bounds checking means other stack addresses will be overwritten
- Including the Return PC!!
- When a function exits ('ret' instruction)
 - This "Return PC" is placed in the CPU program counter

Stack After Calling func()



Stack Overflow

- If a stack buffer is filled based on user input:
 - A lack of bounds checking means the return PC can be changed
- What can the PC be changed to?
 - Where have we learned about PCs that point to specific code?
 - Or the attacker can fill the buffer with code!
 - Return PC will point back to the written buffer
- For example, code to execute a shell:

```
int main(int argc, char *argv[]) {  
    execvp("/bin/sh", "/bin/sh", NULL);  
    return 0;  
}
```

- If this is interesting, read "Smashing the Stack for Fun and Profit": <http://insecure.org/stf/smashstack.html>

Prevention Mechanisms

- Historically this type of attack has dominated incidents
 - Lack of bounds checking is the vulnerability
 - Stack overflow is the exploit
- What can be done? Lots of R&D on this:
 - strncpy()!! (i.e. more careful programming)
 - Non-executable stacks
 - But attackers just become more sophisticated
 - Use chains of system calls to perform attacks
 - Address space randomization
 - Stack meta-data to record modifications to stack information
 - And on and on

Virus

- Malware embedded in a legitimate program
- Replicates itself and actively spreads
 - Key distinction from the threats we have thus far seen
- Typically spread via social engineering
 - Users execute programs via spam email or internet downloads
- One common source
 - Microsoft office files, as they can execute macros
 - Attacker can embed malicious macro code in a file
- A delivery program (usually a Trojan horse) contains a program called a virus dropper
 - Virus dropper executes and injects virus into the system
 - E.g. copies to memory & starts executing virus program

Some Common Virus Categories

- File
 - Virus appends to a file
 - Changes start of program to jump to virus
 - After executing, returns control to program so virus is not noticed
- Boot
 - Infect boot sector; execute every time system is booted
 - Also infects other bootable media such as USB
 - Viruses don't appear in the file system
- Macro (mentioned on last page)
- Source code
 - Modifies source of programs to include the virus & help spread

More Common Virus Categories

- Lots of variants geared to avoiding detection
- Polymorphic
 - Changes each time it is installed to change its "signature"
 - Helps defeat virus scanners that look for patterns in code
- Stealth
 - Virus modifies parts of a system that check for virus existence
 - E.g. modify the read() system call
- Tunneling
 - Virus bypasses detection by installing itself in interrupt handler or device drivers
- Encrypted
 - Virus includes decryption code to decrypt itself and then execute

Worms

- Essentially a virus that uses network resources to spread
- Does not attach to existing program like a virus
- Often consist of 2 pieces
 - "Grappling hook"
 - Initial code that is established and executes on a machine
 - It communicates with an established machine & requests the worm
 - Main worm malware

Morris Worm

- Robert Morris – 1998, graduate student at Cornell
- First internet worm
- Worm attempted 3 network attack methods
 - Try to exploit 'rsh' to execute a task remotely
 - Searched for host-login name pairs in special files
 - Exploit vulnerability in 'finger' program
 - Buffer overflow of a stack frame to execute /bin/sh
 - Exploit vulnerability in 'sendmail' program
 - Debugging option commonly left enabled by system admins
 - Worm mailed itself & executed grappling hook program

Port Scanning

- Attempt TCP/IP connections to host on range of ports
- Each attempt tries to connect to a vulnerable service
 - E.g. sendmail
- If successful, this indicates an exploit opportunity
 - E.g. to exploit a buffer overflow

Denial of Service

- Not gained at stealing or modifying information
- Goal is to disrupt legitimate use of a system or service
- Usually network-based attacks
 - E.g. initiate but do not complete many TCP/IP connections
 - Results in no new legitimate connections being serviced
- Distributed Denial of Service (DDOS)
 - Launched from multiple sites at a common target
 - Sometimes the multiple sites are infected machines (zombies)
- Sometimes security measures introduce DOS vulnerabilities
 - E.g. increasing delays between unsuccessful login attempts

Rootkits

- Malware that acquires privileged access to the OS
 - Also maintains that access
 - By hiding its presence from normal OS activity
- Goals of a rootkit
 - Run (without restriction) on a target system
 - Use social engineering or vulnerabilities in protection (e.g. ACLs)
 - Remain invisible to security software, OS, users
 - Perform malicious action (called the payload)
 - Steal information or access to resources; install other malware

How do Rootkits Hide?

- Processes (including security software) depend on the OS to provide information about the environment
 - E.g. through APIs that expose system calls
- Rootkit software can monitor these API questions to OS
 - Rootkit intercepts questions related to its existence
 - E.g. an 'ls' of the directory where the rootkit program exists
 - E.g. an 'ls' of the /proc/ directory containing info on all processes
 - E.g. a 'read' request on a file modified by the Rootkit

Rootkit Subversion Mechanisms

- "Hooking" OS APIs
 - Change address of OS APIs by pointing to own malware code
 - Can be done for user or supervisor mode
 - In response to system calls, the code at modified address is run
- Hide in unused disk space
 - Unused disk space is not visible to normal file system APIs
 - Modify device driver(s) to execute rootkit code when loaded
- Infect the Master Boot Record (MBR)
 - Control what is loaded into memory before OS is even booted