

Pattern Classification and Recognition:

Bayes Classifiers

ECE 681

Spring 2016

Stacy Tantum, Ph.D.

Bayesian Decision Theory



Quantifies trade-offs between decisions using

- Available data



- Probability

- *prior*

- *likelihood*

- *evidence*

- Costs of decision outcomes

		Truth	
Decision	 ()		
	 ()		

Bayes' Theorem



Rev. Thomas Bayes


Rev. Thomas Bayes presented a solution to the problem of *inverse probability*

- Given data, what can we say about the process that generated it?
- What is the probability of the underlying state from which observed data was generated?

$$P(\check{S}_j | x) = \frac{p(x | \check{S}_j)P(\check{S}_j)}{p(x)}$$

$$\textit{posterior} = \frac{\textit{likelihood} \times \textit{prior}}{\textit{evidence}}$$

Making Decisions




Making Decisions

Bayes Decision Rule (alternate form)

Decide \check{S}_0 if $p(x | \check{S}_0)P(\check{S}_0) > p(x | \check{S}_1)P(\check{S}_1)$

Decide \check{S}_1 otherwise



Making Cost-Aware Decisions







Every action (decision) has a cost, which depends on the “truth”

$$c_{ij} = \text{cost}(r_i | \check{S}_j)$$


Risk of a decision = expected cost

$$R(r_0 | x) = c_{00} P(\check{S}_0 | x) + c_{01} P(\check{S}_1 | x)$$

$$R(r_1 | x) = c_{10} P(\check{S}_0 | x) + c_{11} P(\check{S}_1 | x)$$

		Truth	
			
Decision	 ()		
	 ()		

Likelihood Ratio



Discriminant Functions

Generalized function of the data, $g_i(x)$, that supports *discriminating* class i from other candidate classes

Decide \check{S}_0 if $g_0(x) > g_1(x)$

Decide \check{S}_1 if $g_1(x) > g_0(x)$

For the general case with risks

$$g_i(x) = -R(r_i | x)$$

Two-Class Discriminant Functions

Combine $g_0(x)$ and $g_1(x)$ into a single function

$$g(x) \equiv g_1(x) - g_0(x)$$


Decide \check{S}_1 if $g(x) > 0$

Discriminant Function for Normal Densities

$$p(\mathbf{x} | \check{S}_i) \sim N(\boldsymbol{\mu}_i, \Sigma_i) = \frac{1}{(2f)^{d/2} |\Sigma_i|^{1/2}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_i)^T \Sigma_i^{-1} (\mathbf{x} - \boldsymbol{\mu}_i) \right]$$







$$g_i(\mathbf{x}) = -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_i)^T \Sigma_i^{-1} (\mathbf{x} - \boldsymbol{\mu}_i) - \frac{d}{2} \ln 2f - \frac{1}{2} \ln |\Sigma_i| + \ln P(\check{S}_i)$$

Decision Regions for Gaussian Data



$$\boldsymbol{\mu}_B = \begin{bmatrix} 3 \\ 6 \end{bmatrix} \quad \boldsymbol{B} = \begin{bmatrix} 1/2 & 0 \\ 0 & 2 \end{bmatrix}$$

$$\boldsymbol{\mu}_R = \begin{bmatrix} 3 \\ -2 \end{bmatrix} \quad \boldsymbol{R} = \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}$$



Discrete Data

Data \mathbf{x} takes on only 1 of m discrete values $\mathbf{v}_1, \mathbf{v}_2, \dots \mathbf{v}_m$


Binary Data

Data \mathbf{x} takes on only 1 of 2 discrete values, i.e., 0 or 1





Relation to KNN

Numerical estimate of pdfs

- Total of n samples
- What proportion of the k samples in a volume V belong to class i ?



KNN... Hmm...




Distance Likelihood Ratio Test (DLRT)

Use KNN density estimation to *estimate* the likelihood ratio


- Considers distance to neighbors

KNN and DLRT Comparison

KNN (k=7)



DLRT(k=7)



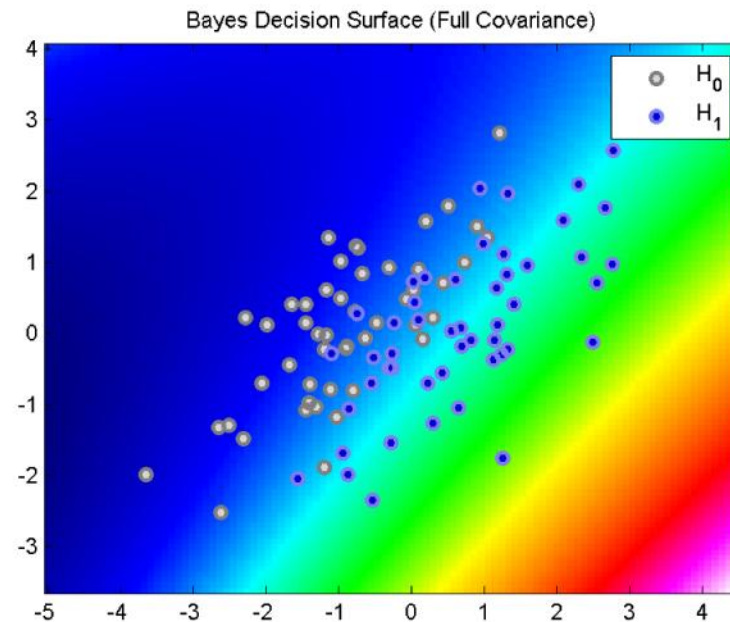
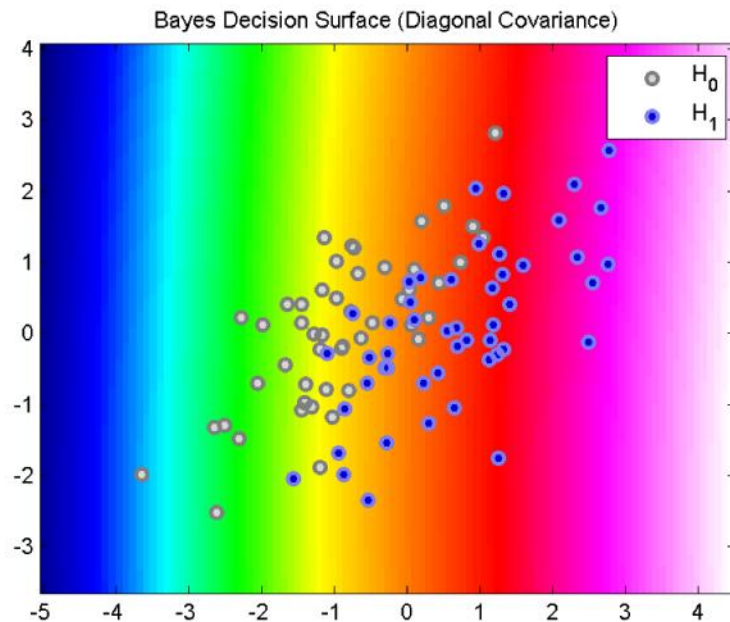
Decision Statistic: $g_1(x) - g_0(x)$

Bayes (Gaussian pdfs)

(Generalized) Likelihood Ratio Test

$$= \left[-\frac{1}{2}(x - \tilde{\mu}_1)^T \Sigma^{-1}(x - \tilde{\mu}_1) - \frac{1}{2} \ln |\Sigma_1| + \ln P(\tilde{\Sigma}_1) \right] - \left[-\frac{1}{2}(x - \tilde{\mu}_0)^T \Sigma^{-1}(x - \tilde{\mu}_0) - \frac{1}{2} \ln |\Sigma_0| + \ln P(\tilde{\Sigma}_0) \right]$$

See Matlab functions mean, cov, and diag



Bayes (Gaussian pdfs)

Decision Statistic: $g_1(x) - g_0(x)$

$$= \left[-\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) - \frac{1}{2} \ln |\Sigma_1| + \ln P(\check{S}_1) \right] \\ - \left[-\frac{1}{2}(x - \mu_0)^T \Sigma^{-1}(x - \mu_0) - \frac{1}{2} \ln |\Sigma_0| + \ln P(\check{S}_0) \right]$$

Training a Bayes classifier:

- What do we need to run it?

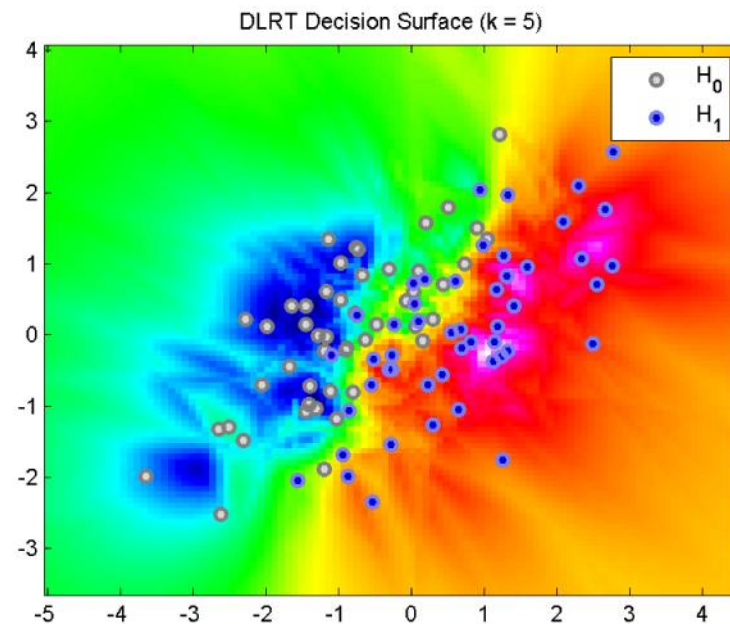
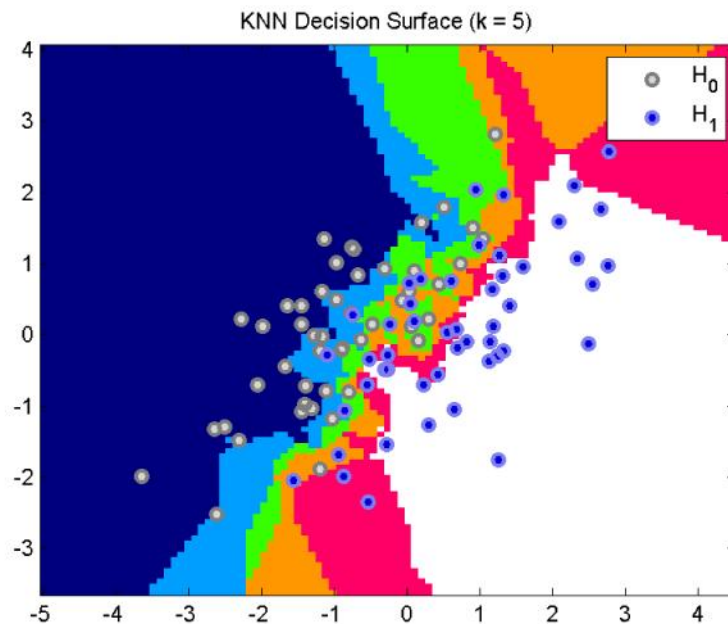
Running a Bayes classifier:

- How do we get a decision statistic?

DLRT

$$\text{Decision Statistic: } \ln\left(\frac{n_0}{n_1}\right) + D[\ln \Delta_{k0} - \ln \Delta_{k1}]$$

Like KNN, but takes into account the distance to the k^{th} neighbor



DLRT

$$\text{Decision Statistic: } \ln\left(\frac{n_0}{n_1}\right) + D[\ln \Delta_{k0} - \ln \Delta_{k1}]$$

See Matlab function `log` (for `ln`, `log10` is for `log10`)

Training a DLRT:

- What do we need to run it?

Running a DLRT:

- How do we get a decision statistic?